

O'REILLY®

3rd Edition  
Covers Scala 3.0

# Programming Scala

Scalability = Functional Programming + Objects

Early  
Release

RAW &  
UNEDITED



Dean Wampler

# **Programming Scala**

THIRD EDITION

**Dean Wampler**

# **Programming Scala**

by Dean Wampler

Copyright © 2021 Kevin Dean Wampler. All rights reserved.

Printed in the United States of America.

Published by O'Reilly Media, Inc., 1005 Gravenstein Highway North, Sebastopol, CA 95472.

O'Reilly books may be purchased for educational, business, or sales promotional use. Online editions are also available for most titles (<http://oreilly.com/safari>). For more information, contact our corporate/institutional sales department: 800-998-9938 or *corporate@oreilly.com*.

Acquisitions Editor: Suzanne McQuade

Production Editor: Caitlin Ghegan

Copyeditor: TK

Proofreader: TK

Indexer: TK

Interior Designer: TK

Cover Designer: Susan Thompson

Illustrator: Catherine Dullea

## **Revision History for the Early Release**

- 2020-10-09: First release
- 2020-10-29: Second release

- 2020-11-26: Third release
- 2021-02-09: Fourth release

See <http://oreilly.com/catalog/errata.csp?isbn=9781491949856> for release details.

The O'Reilly logo is a registered trademark of O'Reilly Media, Inc. *Programming Scala*, the cover image, and related trade dress are trademarks of O'Reilly Media, Inc.

While the publisher and the authors have used good faith efforts to ensure that the information and instructions contained in this work are accurate, the publisher and the authors disclaim all responsibility for errors or omissions, including without limitation responsibility for damages resulting from the use of or reliance on this work. Use of the information and instructions contained in this work is at your own risk. If any code samples or other technology this work contains or describes is subject to open source licenses or the intellectual property rights of others, it is your responsibility to ensure that your use thereof complies with such licenses and/or rights.

978-1-492-07782-4

[LSI]

## **Dedication**

*To Peggy Williams Hammond, September 10, 1933 - May 11, 2018.*

—Dean

## Foreword

If there has been a common theme throughout my career as a programmer, it has been the quest for better abstractions and better tools to support the craft of writing software. Over the years, I have come to value one trait more than any other: composability. If one can write code with good composability, it usually means that other traits we software developers value—such as orthogonality, loose coupling, and high cohesion—are already present. It is all connected.

When I discovered Scala some years ago, the thing that made the biggest impression on me was its composability.

Through some very elegant design choices and simple yet powerful abstractions that were taken from the object-oriented and functional programming worlds, Martin Odersky has managed to create a language with high cohesion and orthogonal, deep abstractions that invites composability in all dimensions of software design. Scala is truly a SCAlable LAnguage that scales with usage, from scripting all the way up to large-scale enterprise applications and middleware.

Scala was born out of academia, but it has grown into a pragmatic and practical language that is very much ready for real-world production use.

What excites me most about this book is that it's so practical. Dean has done a fantastic job, not only by explaining the language through interesting discussions and samples, but also by putting it in the context of the real world. It's written for the programmer who wants to get things done.

I had the pleasure of getting to know Dean some years ago when we were both part of the aspect-oriented programming community. Dean holds a rare mix of deep analytical academic thinking and a pragmatic, get-things-done kind of mentality.

You are about to learn how to write reusable components using mixin and function composition; how to write Reactive applications using Akka; how to make effective use of advanced features in Scala such as macros and higher kinded types; how to utilize Scala's rich, flexible, and expressive syntax to build domain-specific languages; how to effectively test your Scala code; how to let Scala simplify your Big Data problems; and much, much more.

Enjoy the ride. I sure did.

***Jonas Bonér***  
***CTO & cofounder Typesafe***  
***August 2014***

# Preface

---

## A NOTE FOR EARLY RELEASE READERS

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

This will be the preface of the final book. Please note that the GitHub repo will be made active later on.

If you have comments about how we might improve the content and/or examples in this book, or if you notice missing material within this chapter, please reach out to the author at [mcronin@oreilly.com](mailto:mcronin@oreilly.com).

*Programming Scala* introduces an exciting and powerful language that offers all the benefits of a modern object model, *functional programming* (FP), and an advanced type system, while leveraging the industry’s investment in the Java Virtual Machine (JVM). Packed with code examples, this comprehensive book teaches you how to be productive with Scala quickly, and explains what makes this language ideal for today’s scalable, distributed, component-based applications that support concurrency and distribution. You’ll also learn how Scala takes advantage of the advanced JVM as a platform for programming languages.

Learn more at <http://programming-scala.org> or at the book’s [catalog page](#).

## Welcome to Programming Scala, Third Edition



*Programming Scala, Second Edition* was published six years ago, in the fall of 2014. At that time, interest in Scala was surging, driven by two factors.

First, alternative languages for the JVM instead of Java were very appealing. Java's evolution had slowed, in part because its steward, Sun Microsystems was sold to Oracle Corporation a few years previously. Developers wanted improvements like more concise syntax for some constructs and features they saw in other languages, like support for functional programming.

Second, *big data* was a hot sector of the software industry and some of the most popular tools in that sector, especially **Apache Spark** and **Apache Kafka**, were written in Scala and offered elegant Scala APIs, along with APIs in other languages.

A lot has changed in six years. Oracle deserves a lot of credit for reinvigorating Java after the Sun Microsystems acquisition. The pace of innovation has improved considerably. Java 8 was a ground-breaking release, as it introduced two of the most important improvements needed to address limitations compared to Scala. One was support for anonymous functions, called *lambdas*, which addressed the biggest missing feature needed for functional programming. The second feature was support for “default” implementations of the methods declared in interfaces, which made Java interfaces more useful as composable “mixins”.

Also, the Kotlin language was created by the tool vendor Jet Brains, as a “better Java” that isn't as sophisticated as Scala. Kotlin received a big boost when Google endorsed it as the preferred language for Android apps. Around the same time, Apple introduced a language called Swift primarily for iOS development that has a very Scala-like syntax, although it does not target the JVM.

Big data drove the emergence of data science as a profession. Actually, this was just a rebranding and refinement of what data analysts and statisticians had been doing for years. The specialties

of data science called deep learning (DL - i.e., using neural networks), reinforcement learning (RL), and artificial intelligence (AI) are currently the hottest topics in the data world. All fit under the umbrella of machine learning (ML). A large percentage of the popular tools for data science, especially ML, are written in Python or expose Python APIs on top of C++ “kernels”. As a result, interest in Python is growing strongly again, while Scala’s growth in the data world has slowed.

But Scala hasn’t been sitting still. This edition introduces you to version 3 of the language, with significant changes to improve the expressiveness and correctness of Scala, and to remove deprecated and less useful features.

Also, Scala is now a viable language for targeting JavaScript applications through [Scala.js](#). Early support for Scala as a native language is now available through [Scala Native](#), although the version of Scala supported tends to lag the JVM version.

I currently split my time between the Python-based ML world and the Scala-based JVM world. When using Python, I miss the concision, power, and correctness of Scala, but when using Scala, I miss the wealth of data-centric libraries available in the Python world. So, I think we’re entering a period of consolidation for Scala, where developers who want that power and elegance will keep the Scala community vibrant and growing, especially in larger enterprises that are JVM-centered. Scala will remain a preferred choice for hosted services, even while the data and mobile communities tend to use other languages. Who knows what the next five or six years will bring, when it’s time for the fourth edition of *Programming Scala*?

With each addition of this book, I have attempted to provide a comprehensive introduction to Scala features and core libraries, illustrated with plenty of pragmatic examples, tips, and tricks. However, each edition has shifted more towards pragmatism and away from comprehensive surveying of features.

I think this makes the book more useful to you, in an age when so much of our information is gathered in small, ad hoc snippets through Google searches. Libraries come and go. For example, when you need the best way to parse JSON, an Internet search for *Scala JSON libraries* is your best bet. Second, what doesn't change so quickly and what's harder to find on [Stack Overflow](#), is the wisdom of how best to leverage Scala for real-world development. Hence, my goal in this edition is to teach you how to use Scala effectively for a wide class of pragmatic problems, without getting bogged down in corner cases, obscure features, or advanced capabilities that you won't need to know until you become an advanced Scala developer.

Hence, I also won't discuss how to use Scala.js or Scala Native, as the best sources of information for targeting those platforms are their respective websites.

Finally, I wrote this book for professional programmers. I'll err on the side of tackling deeply technical topics, rather than keeping the material "light". Other books provide less thorough, but more gentle introductions, if that's what you prefer. This is a book if you are serious about using Scala professionally.

## Welcome to Programming Scala, Second Edition

*Programming Scala, First Edition* was published five years ago, in the fall of 2009. At the time, it was only the third book dedicated to Scala, and it just missed being the second by a few months. Scala version 2.7.5 was the official release, with version 2.8.0 nearing completion.

A lot has changed since then. At the time of this writing, the Scala version is 2.11.2. Martin Odersky, the creator of Scala, and Jonas Bonér, the creator of Akka, an actor-based concurrency framework, cofounded [Typesafe](#) (now [Lightbend](#)) to promote the language and tools built on it.

There are also a lot more books about Scala. So, do we really need a second edition of this book? Many excellent beginner's guides to Scala are now available. A few advanced books have emerged. The encyclopedic reference remains *Programming in Scala*, Second Edition, by Odersky et al. (Artima Press).

Yet, I believe *Programming Scala, Second Edition* remains unique because it is a *comprehensive* guide to the Scala language and ecosystem, a guide for beginners to advanced users, and it retains the focus on the pragmatic concerns of working professionals. These characteristics made the first edition popular.

Scala is now used by many more organizations than in 2009 and most Java developers have now heard of Scala. Several persistent questions have emerged. Isn't Scala complex? Since Java 8 added significant new features found in Scala, why should I switch to Scala?

I'll tackle these and other, real-world concerns. I have often said that I was *seduced by Scala*, warts and all. I hope you'll feel the same way after reading *Programming Scala, Second Edition*.

## Welcome to Programming Scala, First Edition

Programming languages become popular for many reasons. Sometimes, programmers on a given platform prefer a particular language, or one is institutionalized by a vendor. Most Mac OS programmers use Objective-C. Most Windows programmers use C++ and .NET languages. Most embedded-systems developers use C and C++.

Sometimes, popularity derived from technical merit gives way to fashion and fanaticism. C++, Java, and Ruby have been the objects of fanatical devotion among programmers.

Sometimes, a language becomes popular because it fits the needs of its era. Java was initially seen as a perfect fit for browser-based,

rich client applications. Smalltalk captured the essence of object-oriented programming as that model of programming entered the mainstream.

Today, concurrency, heterogeneity, always-on services, and ever-shrinking development schedules are driving interest in functional programming. It appears that the dominance of object-oriented programming may be over. Mixing paradigms is becoming popular, even necessary.

We gravitated to Scala from other languages because Scala embodies many of the optimal qualities we want in a general-purpose programming language for the kinds of applications we build today: reliable, high-performance, highly concurrent Internet and enterprise applications.

Scala is a multiparadigm language, supporting both object-oriented and functional programming approaches. Scala is scalable, suitable for everything from short scripts up to large-scale, component-based applications. Scala is sophisticated, incorporating state-of-the-art ideas from the halls of computer science departments worldwide. Yet Scala is practical. Its creator, Martin Odersky, participated in the development of Java for years and understands the needs of professional developers.

Both of us were seduced by Scala, by its concise, elegant, and expressive syntax and by the breadth of tools it put at our disposal. In this book, we strive to demonstrate why all these qualities make Scala a compelling and indispensable programming language.

If you are an experienced developer who wants a fast, thorough introduction to Scala, this book is for you. You may be evaluating Scala as a replacement for or complement to your current languages. Maybe you have already decided to use Scala, and you need to learn its features and how to use it well. Either way, we hope to illuminate this powerful language for you in an accessible way.

We assume that you are well versed in object-oriented programming, but we don't assume that you have prior exposure to functional programming. We assume that you are experienced in one or more other programming languages. We draw parallels to features in Java, C#, Ruby, and other languages. If you know any of these languages, we'll point out similar features in Scala, as well as many features that are new.

Whether you come from an object-oriented or functional programming background, you will see how Scala elegantly combines both paradigms, demonstrating their complementary nature. Based on many examples, you will understand how and when to apply OOP and FP techniques to many different design problems.

In the end, we hope that you too will be seduced by Scala. Even if Scala does not end up becoming your day-to-day language, we hope you will gain insights that you can apply regardless of which language you are using.

## Conventions Used in This Book

The following typographical conventions are used in this book:

### *Italic*

Indicates new terms, URLs, email addresses, filenames, and file extensions.

### *Constant width*

Used for program listings, as well as within paragraphs to refer to program elements such as variable or function names, databases, data types, environment variables, statements, and keywords.

### *Constant width bold*

Shows commands or other text that should be typed literally by the user.

*Constant width italic*

Shows text that should be replaced with user-supplied values or by values determined by context.

#### **TIP**

This element signifies a tip or suggestion.

#### **NOTE**

This element signifies a general note.

#### **WARNING**

This element indicates a warning or caution.

## **Using Code Examples**

This book is here to help you get your job done. In general, you may use the code in this book in your programs and documentation. You do not need to contact us for permission unless you're reproducing a significant portion of the code. For example, writing a program that uses several chunks of code from this book does not require permission. Selling or distributing a CD-ROM of examples from O'Reilly books does require permission. Answering a question by citing this book and quoting example code does not require permission. Incorporating a significant amount of example code from this book into your product's documentation does require permission.

We appreciate, but do not require, attribution. An attribution usually includes the title, author, publisher, and ISBN. For example:  
“*Programming Scala, Third Edition* by Dean Wampler. Copyright 2020 Kevin Dean Wampler, 978-1-491-94985-6.”

If you feel your use of code examples falls outside fair use or the permission given above, feel free to contact us at [permissions@oreilly.com](mailto:permissions@oreilly.com).

## Getting the Code Examples

### NOTE

TODO: At this time, the CLI tools are actually called `dotr` and `dotc`, not `scala` and `scalac`. The latter names are used in the expectation that when Scala 3 release candidates begin, the tools will be renamed. I'll refine this text depending on what final names are used for Scala 3 tools. Also, some features shown for `scala` are actually not yet implemented in `dotr`!

You can download the code examples from [GitHub](#). Unzip the files to a convenient location. See the *README* file in the distribution for instructions on building and using the examples. I'll summarize those instructions in the first chapter.

Some of the example files can be run as scripts using the `scala` command. Others must be compiled into class files. A few files are only compatible with Scala 2 and a few files are additional examples that aren't built by *SBT*, the build tool. To keep these groups separate, I have adopted the following directory structure conventions:

*src/main/scala/.../\*.scala*

All Scala 3 source files built with SBT. The standard Scala file extension is *.scala*.



*src/main/scala-2/.../\*.scala*

All Scala 2 source files, some of which won't compile with Scala 3. They are not built with SBT.

*src/test/.../\*.scala*

All Scala 3 test source files built and executed with SBT.

*src/script/.../\*.scala*

“Script” files that won't compile with `scalac`, but can be interpreted with the `scala` interpreter.

## O'Reilly Safari

*Safari* (formerly Safari Books Online) is a membership-based training and reference platform for enterprise, government, educators, and individuals.

Members have access to thousands of books, training videos, Learning Paths, interactive tutorials, and curated playlists from over 250 publishers, including O'Reilly Media, Harvard Business Review, Prentice Hall Professional, Addison-Wesley Professional, Microsoft Press, Sams, Que, Peachpit Press, Adobe, Focal Press, Cisco Press, John Wiley & Sons, Syngress, Morgan Kaufmann, IBM Redbooks, Packt, Adobe Press, FT Press, Apress, Manning, New Riders, McGraw-Hill, Jones & Bartlett, and Course Technology, among others.

For more information, please visit <http://oreilly.com/safari>.

## How to Contact Us

Please address comments and questions concerning this book to the publisher:

O'Reilly Media, Inc.

1005 Gravenstein Highway North

Sebastopol, CA 95472

800-998-9938 (in the United States or Canada)

707-829-0515 (international or local)

707-829-0104 (fax)

We have a web page for this book, where we list errata, examples, and any additional information. You can access this page at

[http://bit.ly/programmingScala\\_2E](http://bit.ly/programmingScala_2E).

To comment or ask technical questions about this book, send email to [bookquestions@oreilly.com](mailto:bookquestions@oreilly.com).

For more information about our books, courses, conferences, and news, see our website at <http://www.oreilly.com>.

Find us on Facebook: <http://facebook.com/oreilly>

Follow us on Twitter: <http://twitter.com/oreillymedia>

Watch us on YouTube: <http://www.youtube.com/oreillymedia>

## Acknowledgments for the Third Edition

Working with early builds of Scala 3, I often ran into unimplemented features and incomplete documentation. The members of the Scala community have provided valuable help while I learned what's new. The EPFL documentation for Dotty, <https://dotty.epfl.ch/docs/>, provided essential information.

## Acknowledgments for the Second Edition

As I, Dean Wampler, worked on this edition of the book, I continued to enjoy the mentoring and feedback from many of my Typesafe colleagues, plus the valuable feedback from people who reviewed the early-access releases. I'm especially grateful to Ramnivas Laddad, Kevin Kilroy, Lutz Huehnken, and Thomas Lockney, who reviewed drafts of the manuscript. Thanks to my long-time colleague and friend, Jonas Bonér, for writing an updated **Foreword** for the book.

And special thanks to Ann who allowed me to consume so much of our personal time with this project. I love you!

## Acknowledgments for the First Edition

As we developed this book, many people read early drafts and suggested numerous improvements to the text, for which we are eternally grateful. We are especially grateful to Steve Jensen, Ramnivas Laddad, Marcel Molina, Bill Venners, and Jonas Bonér for their extensive feedback.

Much of the feedback we received came through the Safari Rough Cuts releases and the online edition available at <http://programmingscala.com>. We are grateful for the feedback provided by (in no particular order) Iulian Dragos, Nikolaj Lindberg, Matt Hellige, David Vydra, Ricky Clarkson, Alex Cruise, Josh Cronemeyer, Tyler Jennings, Alan Supynuk, Tony Hillerson, Roger Vaughn, Arbi Sookazian, Bruce Leidl, Daniel Sobral, Eder Andres Avila, Marek Kubica, Henrik Huttunen, Bhaskar Maddala, Ged Byrne, Derek Mahar, Geoffrey Wiseman, Peter Rawsthorne, Geoffrey Wiseman, Joe Bowbeer, Alexander Battisti, Rob Dickens, Tim MacEachern, Jason Harris, Steven Grady, Bob Follek, Ariel Ortiz, Parth Malwankar, Reid Hochstedler, Jason Zaugg, Jon Hanson, Mario Gleichmann, David Gates, Zef Hemel, Michael Yee,

Marius Kreis, Martin Süsskraut, Javier Vegas, Tobias Hauth, Francesco Boichicchio, Stephen Duncan Jr., Patrik Dudits, Jan Niehusmann, Bill Burdick, David Holbrook, Shalom Deitch, Jesper Nordenberg, Esa Laine, Gleb Frank, Simon Andersson, Patrik Dudits, Chris Lewis, Julian Howarth, Dirk Kuzemczak, Henri Gerrits, John Heintz, Stuart Roebuck, and Jungho Kim. Many other readers for whom we only have usernames also provided feedback. We wish to thank Zack, JoshG, ewilligers, abcoates, brad, teto, pjcj, mkleint, dandoyon, Arek, rue, acangiano, vkelman, bryanl, Jeff, mbaxter, pjb3, kxen, hipertracker, ctran, Ram R., cody, Nolan, Joshua, Ajay, Joe, and anonymous contributors. We apologize if we have overlooked anyone!

Our editor, Mike Loukides, knows how to push and prod gently. He's been a great help throughout this crazy process. Many other people at O'Reilly were always there to answer our questions and help us move forward.

We thank Jonas Bonér for writing the **Foreword** for the book. Jonas is a longtime friend and collaborator from the aspect-oriented programming (AOP) community. For years, he has done pioneering work in the Java community. Now he is applying his energies to promoting Scala and growing that community.

Bill Venners graciously provided the quote on the back cover. The first published book on Scala, *Programming in Scala* (Artima), that he cowrote with Martin Odersky and Lex Spoon, is indispensable for the Scala developer. Bill has also created the wonderful ScalaTest library.

We have learned a lot from fellow developers around the world. Besides Jonas and Bill, Debasish Ghosh, James Iry, Daniel Spiewak, David Pollack, Paul Snively, Ola Bini, Daniel Sobral, Josh Suereth, Robey Pointer, Nathan Hamblen, Jorge Ortiz, and others have illuminated dark corners with their blog entries, forum discussions, and personal conversations.

Dean thanks his colleagues at Object Mentor and several developers at client sites for many stimulating discussions on languages, software design, and the pragmatic issues facing developers in industry. The members of the Chicago Area Scala Enthusiasts (CASE) group have also been a source of valuable feedback and inspiration.

Alex thanks his colleagues at Twitter for their encouragement and superb work in demonstrating Scala's effectiveness as a language. He also thanks the Bay Area Scala Enthusiasts (BASE) for their motivation and community.

Most of all, we thank Martin Odersky and his team for creating Scala.

# Chapter 1. Zero to Sixty: Introducing Scala

---

## A NOTE FOR EARLY RELEASE READERS

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

This will be the 1st chapter of the final book. Please note that the GitHub repo will be made active later on.

If you have comments about how we might improve the content and/or examples in this book, or if you notice missing material within this chapter, please reach out to the author at [mcronin@oreilly.com](mailto:mcronin@oreilly.com).

Let’s start with a brief look at why you should investigate Scala. Then we’ll dive in and write some code.

## Why Scala?

*Scala* is a language that addresses the needs of the modern software developer. It is a statically typed, object-oriented and functional, mixed-platform language with a succinct, elegant, and flexible syntax, a sophisticated type system, and idioms that promote scalability from small, interpreted scripts to large, sophisticated applications. So let’s consider each of those ideas in more detail:

*A JVM, JavaScript, and native language*

Scala started as a JVM language that exploits the performance and optimizations of the JVM, as well as the rich ecosystem of tools and libraries built around Java. More recently, *Scala.js* brings Scala to JavaScript and *Scala Native* is an experimental Scala that compiles to native machine code, bypassing the JVM and JavaScript runtimes.

### *Statically typed*

Scala embraces *static typing* as a tool for creating robust applications. It fixes many of the flaws of Java's type system and it uses type inference to eliminate much of the typing boilerplate.

### *Object-oriented programming*

Scala fully supports *object-oriented programming* (OOP). Scala improves Java's object model with the addition of *traits*, providing a clean way to implement types using *mixin composition*. In Scala, everything *really* is an object, even numeric types, providing much more consistent handling, especially in collections.

### *Functional programming*

Scala fully supports *functional programming* (FP). FP has emerged as the best tool for thinking about problems of concurrency, *Big Data*, and general code correctness. Immutable values, first-class functions, functions without side effects, "higher-order" functions, and functional collections all contribute to concise, powerful, and correct code.

### *A sophisticated type system*

Scala extends the type system of Java with more flexible generics and other enhancements to improve code correctness. With type inference, Scala code is often as concise as code in dynamically typed languages, yet inherently safer.

### *A succinct, elegant, and flexible syntax*

Verbose expressions in Java become concise idioms in Scala. Scala provides several facilities for building *domain-specific languages* (DSLs), APIs that feel “native” to users.<sup>1</sup>

### *Scalable—architectures*

You can write small, interpreted scripts to large, distributed applications in Scala.

The name *Scala* is a contraction of the words *scalable language*. It is pronounced *scah-lah*, like the Italian word for “staircase.” Hence, the two “a”s are pronounced the same.

Scala was started by Martin Odersky in 2001. The first public release was January 20th, 2004. Martin is a professor in the School of Computer and Communication Sciences at the Ecole Polytechnique Fédérale de Lausanne (EPFL). He spent his graduate years working in the group headed by Niklaus Wirth, of Pascal fame. Martin worked on Pizza, an early functional language on the JVM. He later worked on GJ, a prototype of what later became Generics in Java, along with Philip Wadler, one of the designers of Haskell. Martin was hired by Sun Microsystems to produce the reference implementation of *javac*, the descendant of which is the Java compiler that ships with the Java Developer Kit (JDK) today.

## **The Seductions of Scala**

The growth of Scala users since it was introduced over fifteen years ago confirms my view that Scala is a language for our time. You can leverage the maturity of the JVM and JavaScript ecosystems while enjoying state-of-the-art language features with a concise, yet expressive syntax for addressing today’s development challenges.



In any field of endeavor, the professionals need sophisticated, powerful tools and techniques. It may take a while to master them, but you make the effort because mastery is the key to your success.

I believe Scala is a language for *professional* developers. Not all users are professionals, of course, but Scala is the kind of language a professional in our field needs, rich in features, highly performant, expressive for a wide class of problems. It will take you a while to master Scala, but once you do, you won't feel constrained by your programming language.

## Why Scala 3?

If you used Scala before, you used Scala 2, the major version since March 2006! Scala 3 aims to improve Scala in several ways.

First, Scala 3 strengthens Scala's foundations, especially in the type system. Martin Odersky and collaborators have been developing the *dependent object typing* (DOT) calculus, which provides a more sound foundation for Scala's type system. Scala 3 integrates DOT.

Second, Scala 2 has many powerful features, but sometimes they can be hard to use. Scala 3 improves the usability and safety of these features, especially *implicit*s. Other language warts and "puzzlers" are removed.

Third, Scala 3 improves the consistency and expressiveness of Scala's language constructs and it removes unimportant constructs to make the language smaller and more regular. The previous experimental approach to macros is replaced with a principled approach to meta-programming.

We'll call out these changes as we explore the corresponding language features.

## Migrating to Scala 3

The Scala team has worked hard to make migration to Scala 3 from Scala 2 as painless as possible, while still allowing the language to make improvements that require breaking changes. Scala 3 uses the same collections library as Scala 2.13. Hence, if you are using a Scala 2 version earlier than 2.13, I recommend upgrading to Scala 2.13 first, to update uses of the library, then upgrade to Scala 3.

However, to make this transition as painless as possible, there are several ways to compile Scala code that allows or disallows deprecated Scala 2 constructs. There are even compiler flags that will do some rewrites for you. See *later chapter [x]* and *the section on the command line* in *Chapter 21* for details.

## Installing Scala

Let's learn how to install the command-line tools that you need to work with the book's code examples.<sup>2</sup> The examples used in this book were written and compiled using Scala version 3.0, the latest release at the time of this writing.

All the examples will use Scala on the JVM. See the [Scala.js](#) and [Scala Native](#) websites for information on targeting those platforms.

At a minimum, you need to install a recent [Java JDK](#) and the de facto build tool for Scala, [SBT](#). Then you can use the `sbt` command to bootstrap everything else for the examples. However, the following sections go into more details about tools you might want to install and how to install them.

The Scala website [Getting Started](#) page discusses even more options to get started with Scala.

## Coursier

Coursier ([get-coursier.io](https://get-coursier.io)) is a new dependency resolver and tool manager. It replaces Maven and Ivy, the traditional dependency

resolvers for Java and Scala projects. Written in Scala, it is fast and easy to embed in other applications.

Installing Coursier is not required, but it is recommended as it will install all the tools you need and more, in addition to providing many other convenient features. For example, the Coursier CLI is handy for managing dependency metadata and artifacts from Maven and Ivy repositories, although we'll do this indirectly through the SBT build for the examples. Coursier also has convenient commands for working with applications embedded in libraries and it can even manages installations of different Java JDK versions.

See the installation instructions at [get-coursier.io/docs/cli-installation](https://get-coursier.io/docs/cli-installation) for details. After installing Coursier, see [get-coursier.io/docs/cli-install](https://get-coursier.io/docs/cli-install) for a description of the `coursier install` command for installing other tools, like SBT (`sbt-launcher`) and Scala. For example, this command shows the default set of available applications (at the time of this writing in early 2020):

```
$ unzip -l "$(cs fetch io.get-coursier:apps:0.0.8)" | grep json
188  02-09-2020 17:48  ammonite.json
175  02-09-2020 17:48  coursier.json
332  02-09-2020 17:48  cs.json
241  02-09-2020 17:48  dotty-repl.json
150  02-09-2020 17:48  echo-graalvm.json
108  02-09-2020 17:48  echo-java.json
172  02-09-2020 17:48  echo-native.json
335  02-09-2020 17:48  giter8.json
135  02-09-2020 17:48  mdoc.json
323  02-09-2020 17:48  mill-interactive.json
321  02-09-2020 17:48  mill.json
524  02-09-2020 17:48  sbt-launcher.json
222  02-09-2020 17:48  scala.json
209  02-09-2020 17:48  scalac.json
213  02-09-2020 17:48  scaladoc.json
180  02-09-2020 17:48  scalafix.json
184  02-09-2020 17:48  scalafmt.json
204  02-09-2020 17:48  scalap.json
```

(Many of the tools shown here are discussed in *Chapter 21*.) For now, run the following command to install the sbt command. Note that the `.json` suffix is removed from the name. I recommend installing all the `scala*` commands, but they aren't strictly necessary, as discussed below.

```
coursier install --install-dir path sbt-launcher
```

Note that the `--install-dir path` arguments are optional. Depending on your platform, the default installation location will vary. Whatever location you use, make sure you add it to your `PATH`.

## Java JDK

Use a recent Java JDK release. Version 11 or newer is recommended, although Java 8 should work. To install the JDK, go to the [Oracle Java website](#) and follow the instructions to install the full Java Development Kit (JDK).

## SBT

The most popular build tool for Scala, *SBT*, version 1.3.8 or newer, is used for the code examples. Install using Coursier or follow the instructions at [scala-sbt.org](#).

When you are finished, you will have an `sbt` command that you can run from a Linux or OS X terminal or Windows command window.

## Scala

You actually don't need to install the Scala command-line tools separately, as SBT will install the basics that you need as dependencies. However, if you want to install these tools, use Coursier or see the instructions at [scala-lang.org/download](#).

## Building the Code Examples

Now that you have the tools you need, you can download and build the code examples.

### *Get the Code*

Download the code examples as described in “[Getting the Code Examples](#)”.

### *Start SBT*

Open a terminal and change to the root directory for the code examples. Type the command `sbt test`. It will download all the library dependencies you need, including the Scala compiler. This will take a while and you’ll need an Internet connection. Then `sbt` will compile the code and run the unit tests. You’ll see lots of output, ending with a “success” message. If you run the command again, it should finish very quickly because it won’t need to do anything again.

Congratulations! You are ready to get started.

#### TIP

For most of the book, we’ll use the Scala tools indirectly through SBT, which downloads the Scala compiler version we want, the Scala interpreter, Scala’s standard library, and the required third-party dependencies automatically.

## More Tips

In your browser, it’s useful to bookmark the [URL](#) for the Scala library’s *Scaladocs*, the analog of *Javadocs* for Scala. For your convenience, most times when I mention a type in the Scala library, I’ll include a link to the corresponding Scaladocs entry.

Use the search field at the top of the page to quickly find anything in the docs. The documentation page for each type has a link to view the corresponding source code in Scala's [GitHub repository](#), which is a good way to learn how the library was implemented. Look for the link on the line labeled "Source."

Any text editor or IDE (integrated development environment) will suffice for working with the examples. Scala plug-ins exist for all the popular editors and IDEs. For more details, see *later chapters*. In general, the community for your favorite editor is your best source of up-to-date information on Scala support.

## Using SBT

Let's cover the basics of using SBT, which you'll need to work with the code examples.

When you start the `sbt` command, if you don't specify a task to run, SBT starts an interactive REPL (*Read, Eval, Print, Loop*). Let's try that now and see a few of the available "tasks."

In the listing that follows, the `$` is the shell command prompt (e.g., `bash`), where you start the `sbt` command, the `>` is the default SBT interactive prompt, and the `#` starts an `sbt` comment. You can type most of these commands in any order:

```
$ sbt
> help      # Describe commands.
> launchIDE # Open the project in Visual Studio Code.
> tasks     # Show the most commonly-used, available tasks.
> tasks     # Show ALL the available tasks.
> compile   # Incrementally compile the code.
> test      # Incrementally compile the code and run the tests.
> clean     # Delete all build artifacts.
> console   # Start the Scala REPL.
> run       # Run one of the "main" routines in the project.
> show      # Show the definition of variable "x".
> exit      # Quit the REPL (also control-d works).
```

## TIP

The *SBT* project for the code examples is actually configured to show the following as the SBT prompt:

```
sbt:Programming Scala, Third Edition - Code examples>
```

We'll use the more concise prompt, `>`, the default for SBT, to save space.

The `launchIDE` task is convenient for those of you who prefer IDEs, although currently only Visual Studio Code is supported. The details can be found at [dotty.epfl.ch/docs/usage/ide-support.html](http://dotty.epfl.ch/docs/usage/ide-support.html).

Both *IntelliJ IDEA* and *Visual Studio Code* can open SBT project, once you install a Scala plug-in.

A handy SBT technique is to add `~` at the front of any command. Whenever file changes are saved to disk, the command will be rerun. For example, I use `\~test` all the time to keep compiling and running my code and tests. SBT uses an incremental compiler, so you don't have to wait for a full rebuild every time. Break out of this loop by hitting the return key.

Scala has its own REPL. Invoke it using the `console` command in SBT. You will use this a lot to try examples in the book. The Scala REPL prompt is `scala>`.

Before starting the REPL, SBT will build your project and set up the `CLASSPATH` with your built artifacts and dependent libraries. This convenience means it's rare to use the `scala` command-line tool outside of SBT.

You can exit both the SBT REPL and the Scala REPL with *Ctrl-D*.

## TIP

Using the Scala REPL is a very effective way to experiment with code idioms and to learn an API, even Java APIs. Invoking it from SBT using the `console` task conveniently adds project dependencies and the compiled project code to the classpath for the REPL.

## Running the Scala Command-Line Tools

If you installed the Scala command-line tools separately, the Scala compiler is called `scalac`, analogous to the Java compiler `javac`. We will let SBT run it for us, but the command syntax is straightforward if you've ever run `javac`. Use `scalac -help` to see the options.

Similarly, the `scala` command, which is similar to `java`, is used to run programs, but it also supports the interactive REPL mode we just discussed *and* the ability to run Scala “scripts”. Consider this example script from the code examples:

```
// src/script/scala/progscala3/introscala/Upper1.scala

class Upper1:
  def convert(strings: Seq[String]): Seq[String] =
    strings.map((s: String) => s.toUpperCase())

val up = new Upper1()
println(up.convert(List("Hello", "World!")))
```

Let's run it with the `scala` command. Change your current working directory to the root of the code examples. For Windows, use backslashes in the next command:

```
$ scala src/script/scala/progscala3/introscala/Upper1.scala
List(HELLO, WORLD!)
...
```

And thus we have satisfied the *Prime Directive* of the Programming Book Authors Guild, which states that our first program must print



“Hello World!”

### TIP

As you can see from the listing above for `Upper1.scala`, each of these files is listed starting with a comment that contains the file path in the code examples. That makes it easy to find the file.

If you invoke `scala` without a compiled main class to run or a script file, `scala` enters the REPL mode. Here is a REPL session illustrating some useful commands. If you didn't install Scala separately, just start `console` in `sbt`. The REPL prompt is `scala>` (some output elided):

```
$ scala
...
scala> :help
The REPL has several commands available:

:help                print this summary
:load <path>         interpret lines in a file
:quit                exit the interpreter
:type <expression>  evaluate the type of the given expression
:doc <expression>   print the documentation for the given
expression
:imports              show import history
:reset               reset the repl to its initial state, ...

scala> val s = "Hello, World!"
val s: String = Hello, World!

scala> println("Hello, World!")
Hello, World!

scala> 1 + 2
val res0: Int = 3

scala> s.con<tab>
concat  contains  containsSlice  contentEquals
```

```
scala> s.contains("el")
val res1: Boolean = true

scala> :quit
$    # back at the terminal prompt.
```

We assigned a string, "Hello, World!", to a variable named `s`, which we declared as an immutable value using the `val` keyword. The `println` function prints a string to the console, followed by a line feed.

This `println` is effectively the same thing as Java's `System.out.println`. Also, Scala Strings are Java `Strings`.

When we added two numbers, we didn't assign the result to a variable, so the REPL made up a name for us, `res0`, which we could use in subsequent expressions.

The REPL supports tab completion. The input shown is used to indicate that a tab was typed after `s.con`. The REPL responded with a list of methods on `String` that could be called. The expression was completed with a call to the `contains` method.

We didn't always explicitly specify type information. When type information is shown, either when it is inferred or explicit type information is added to declarations, these *type annotations*, as they are called, follow a colon after the item name. The output of REPL shows several examples.

Why doesn't Scala follow Java conventions? When type annotations aren't explicitly in the code, then the type is *inferred*. Compared to Java's `item` type convention, the `item: type` convention is easier for the compiler to analyze unambiguously when you omit the type annotation and just write `item`.

As a general rule, Scala follows Java conventions, departing from them for specific reasons, like supporting a new feature that would be difficult using Java syntax.

### TIP

Showing the types in the REPL is very handy for learning the types that Scala infers for particular expressions. It's one example of exploration that the REPL enables.

Finally, we used `:quit` to exit the REPL. Ctrl-D can also be used.

We'll see additional REPL commands as we go and we'll explore the REPL commands in depth in *a later chapter*.

## A Taste of Scala

We've already seen a bit of Scala as we discussed tools, including how to print "Hello World!". The rest of this chapter and the two chapters that follow provide a rapid tour of Scala features. As we go, we'll discuss just enough of the details to understand what's going on, but many of the deeper background details will have to wait for later chapters. Think of this tour as a primer on Scala syntax and a taste of what programming in Scala is like day to day.

### TIP

When we mention a type in the Scala library, you might find it useful to read more in the Scaladocs about it. The Scaladocs for the current release of Scala can be found [here](#). Note that Scala 3 uses the Scala 2.13 collections library, while other parts of the library have changed in Scala 3.

All examples shown in the book start with a comment line like this:

```
// src/script/scala/progscala3/introscala/Upper1.scala
```

Scala follows the same comment conventions as Java, C#, C, etc. A `// comment` goes to the end of a line, while a `/* comment */` can

cross line boundaries. *Scaladoc* comments follow Java conventions, `/** comment */`.

When the path starts with `src/script`, use `scala` to run the script, as follows:

```
$ scala src/script/scala/progscala3/introscala/Upper1.scala
```

However, this may not work if the script uses libraries. Instead, run the SBT console, then use the `:load` command:

```
scala> :load src/script/scala/progscala3/introscala/Upper1.scala
```

Finally, you can also copy code and paste it at the `scala>` prompt.

Files named `src/test/scala/.../*Suite.scala` are tests written using **MUnit** (see *later chapters*). To run all the tests, use the `sbt` command `test`. To run just one particular test, use `testOnly path`, where *path* is the fully-qualified type name for the test.

```
> testOnly progscala3.objectsystem.equality.EqualitySuite
[info] Compiling 1 Scala source to ...
progscala3.objectsystem.equality.EqualitySuite:
+ The == operator is implemented with the equals method 0.01s
+ The != operator is implemented with the equals method 0.001s
...
[info] Passed: Total 14, Failed 0, Errors 0, Passed 14
[success] Total time: 1 s, completed Feb 29, 2020, 5:00:41 PM
>
```

The corresponding source file is

`src/test/scala/progscala3/objectsystem/equality/EqualitySuite.scala`. SBT follows Maven conventions that directories for compiled source code go under `src/main/scala` and tests go under `src/test/scala`. So, in this example, the package definition for this test is `progscala3.objectsystem.equality`. The compiled class name is `EqualitySuite`.

## NOTE

Java requires that package and file names must match the declared package and public class declared within the file. Scala doesn't require this practice. However, I follow these conventions most of the time for compiled code (less often for scripts) and I recommend you do this, too, for your production code.

Finally, many of the files under the `src/main/scala` define a `main` method that you can execute in one of several ways.

First, use SBT's `run` command. It will find all the classes with `main` methods and prompt you to pick which one. Note that SBT will only search `src/main/scala` and `src/main/java` directories, ignoring the other directories under `src`, including `src/script`.

Let's use another example we'll study later in the chapter, `src/main/scala/progscala3/introscala/UpperMain1.scala`. Invoke `run hello world`, then enter the number shown for `progscala3.introscala.UpperMain1`. Note we are passing arguments to `run`, `hello world`, which will be passed to the program to convert to upper case:

```
> run hello world
...
```

```
Multiple main classes detected, select one to run:
```

```
...
[20] progscala3.introscala.UpperMain1
...
```

```
20
```

```
[info] running progscala3.introscala.UpperMain1 hello world
HELLO WORLD
[success] Total time: 2 s, completed Feb 29, 2020, 5:08:18 PM
```

The second way to run this program is to use `runMain` and specify the specific program class name. This skips the prompt:

```
> runMain progscale3.introscala.UpperMain1 hello world
[warn] Multiple main classes detected. Run 'show
discoveredMainClasses' ...
[info] running progscale3.introscala.UpperMain1
HELLO WORLD
[success] Total time: 0 s, completed Feb 29, 2020, 5:18:05 PM
>
```

Finally, once your program is ready for production runs, you'll use the `scala` command, similar to how `java` is used. Now the correct `classpath` must be defined, including all dependencies. This example is relatively easy; we just point to the output directory for the compiled code:

```
$ scala -cp target/scala-3.0.0/classes/
progscale3.introscala.UpperMain1 words
```

Let's explore other differences between scripts, like the `Upper1` script we've used, and compiled code, like the `UpperMain1` example we just executed.

Here is the script again:

```
// src/script/scala/progscale3/introscala/Upper1.scala

class Upper1:
  def convert(strings: Seq[String]): Seq[String] =
    strings.map((s: String) => s.toUpperCase())

val up = new Upper1()
println(up.convert(List("Hello", "World!")))
```

We declare a class, `Upper1`, using the `class` keyword. The entire class body is indented on the subsequent lines (or inside curly braces `{...}` if you use that syntax instead).

Upper1 contains a method called `convert`. Method definitions start with the `def` keyword, followed by the method name and an optional parameter list. The method signature ends with an optional return type. The return type can be inferred in many cases, but adding the return type explicitly, as shown, provides useful documentation and also avoids occasional surprises from the type inference process.

### NOTE

I'll use *parameters* to refer to the things a method or function is defined as accepting when you call it. I'll use *arguments* to refer to values you actually pass to it when making the call.

Type *annotations* are specified using `name: type`. This is used here for both the parameter list and the return type of the method, the last `Seq[String]` before the equals sign.

An equals sign (`=`) separates the signature from the method body. Why an equals sign?

One reason is to reduce ambiguity. Scala infers the return type if the colon and type are omitted. If the method takes no parameters, you can omit the parentheses, too. So, the equal sign makes parsing unambiguous when either or both of these features are omitted. It's clear where the signature ends and the method body begins.

The equals sign also reminds us of the *functional programming* principle that variables and functions are treated uniformly. As we saw in the invocation of `map`, functions can be passed as arguments to other functions, just like values. They can also be returned from functions, and assigned to variables. In fact, it's correct to say that *functions are values*.

This method takes a *sequence* (`Seq`) of zero or more input strings and returns a new sequence, where each of the input strings is converted to uppercase. `Seq` is an abstraction for collections that you

can iterate through. The actual type returned by this method will be the same concrete type that was passed into it as an argument, like `Vector` or `List` (both of which are immutable collections).

*Collection* types like `Seq` are *parameterized types*, very similar to *generic* types in Java. They are a “collection of something,” in this example a sequence of strings. Scala uses square brackets (`[...]`) for parameterized types, whereas Java uses angle brackets (`<...>`).

### NOTE

Scala allows angle brackets to be used in *identifiers*, like method and variable names. For example, defining a “less than” method and naming it `<` is common and allowed by Scala, whereas Java doesn’t allow characters like that in identifiers. So, to avoid ambiguity, Scala uses square brackets instead for parameterized types and disallows them in identifiers.

Inside the body, we use one of the powerful methods available for most collections, `map`, which iterates through the collection, calls the provided method on each element, and returns a new collection with the transformed elements.

The function passed to `map` is an unnamed *function literal* (`parameters`) `=>` `body`, similar to Java’s *lambda* syntax:

```
(s: String) => s.toUpperCase()
```

It takes a parameter list with a single `String` named `s`. The body of the function literal is after the “arrow,” `=>`. (The UTF8 `⇒` characters was also allowed in Scala 2, but is now deprecated.) The body calls `toUpperCase()` on `s`. The result of this call is automatically returned by the function literal. In Scala, the last *expression* in a function or method is the return value. The `return` keyword exists in Scala, but it can only be used in methods, not in anonymous functions like this one. In fact, it is rarely used in methods.



## METHODS VERSUS FUNCTIONS

Following the convention in most *object-oriented programming* languages, the term *method* is used to refer to a function defined within a class. Methods have an implied `this` reference to the object as an additional argument when they are called. Like most OO languages, the syntax used is `this.method_name(other_args)`. We'll use the term *method* this way. We'll use the term *function* to refer to non-methods, but also sometimes use it generically to include methods. The context will indicate the distinction.

The expression `(s: String) => s.toUpperCase()` in `Upper1.scala` is an example of a *function* that is not a method.

On the JVM, functions are implemented using JVM *lambdas*:

```
scala> (s: String) => s.toUpperCase()  
val res0: String => String = Lambda$7775/0x00000008035fc040@7673711e=
```

The last two lines create an instance of `Upper1`, named `up`, and use it to convert two strings to uppercase and finally print the resulting `Seq`. As in Java, the syntax `new Upper1()` creates a new instance. The `up` variable is declared as a read-only “value” using the `val` keyword. It behaves like a `final` variable in Java.

Now let's look at the compiled example, where I added `Main` to the name. Note the path to the source file now contains `src/main`, instead of `src/script`:

```
// src/main/scala/progscala3/introscala/UpperMain1.scala  
package progscala3.introscala ❶  
  
object UpperMain1:  
  def main(params: Array[String]): Unit = ❷  
    params.map(s => s.toUpperCase()).foreach(s => printf("%s ",s))  
    println("")
```

```

    end main ❸

    @main def hello(params: String*) = main(params.toArray) ❹
end UpperMain1 ❺

```

- ❶ Declare the package location.
- ❷ Declare a `main` method, the program entry point.
- ❸ For long methods (unlike this one), you can use `end name`, but this is optional.
- ❹ An alternative way to define an entry point method.
- ❺ Optional end to the object definition.

Packages work much like they do in Java. Packages provide a “namespace” for scoping. Here we specify that this class exists in the `progscale3.introscala` package.

To declare a `main` method in Java, you would put it in a `class` and declare it `static`, meaning not tied to any one instance. You can then call it with the syntax `MyClass.main`. This pattern is so pervasive, that Scala builds it into the language. We instead declare an object, named `UpperMain1`, using the `object` keyword, then declare `main` as we would any other method. At the JVM level, it looks like a `static` method. When running this program like we did above, this `main` method is invoked.

Note the parameter list is an `Array[String]` and it returns `Unit`, which is analogous to `void` in programs like Java, where nothing useful is returned.

You use `end name` with all of the constructs like method and type definitions, `if`, `while` expressions, etc. that support the braceless, indentation-based structure. It is optional, intended for long definitions where it can be hard to see the beginning and end at the

same time. Hence, for this short definition of `main`, you wouldn't use it normally. Here are other examples for control constructs:

```
if sequence.length > 0 then
  println(sequence)
end if      // "end if" is optional

while i < 10 do
  i += 1
end while  // "end while" is optional
```

The `end ...` lines shown are not required, but they are useful for readability if the indented blocks are many lines long.

The `@main def hello` method is another way to declare an entry point, especially useful when you don't need to process arguments, so you don't need to declare the parameter list. Note that in the list printed by the `sbt run` command, this entry point is shown as `progscala3.introscala.hello`, while the `main` method was shown by the type name, `progscala3.introscala.UpperMain1`. Try run and invoke `hello`.

Declaring `UpperMain1` as an object makes it a *singleton*, meaning there will always only be one instance of it, controlled by the Scala runtime library. You can't create your own instances with `new`.

Scala makes the *Singleton Design Pattern* a first-class member of the language. In most ways, these object declarations are just like `class` declarations. Scala uses these objects to replace "class-level" members, like `statics` in Java. You use the declared object like an instance created from a regular class, reference the member variables and functions as required, e.g., `UpperMain1.hello`.

When you define the `main` method for your program, you *must* declare it inside an object, but in general, `UpperMain1` is a good candidate for a singleton, because we don't need more than one *instance* and it carries no state information.

The *Singleton Design Pattern* has drawbacks. It's hard to replace a singleton instance with a *test double* in unit tests and forcing all computation through a single instance raises concerns about thread safety and performance. However, there are times when singletons make sense, as in this example where these concerns don't apply.

### NOTE

Why doesn't Scala support statics? Compared to languages that allow static members (or equivalent constructs), Scala is more true to the vision that *everything* should be an object. The object construct keeps this policy more consistent than in languages that mix static and *instance* class members.

UpperMain1.main takes the user arguments in the params array, maps over them to convert to upper case, then uses another common method, foreach, to print them out.

The function we passed to map drops the type annotation for s:

```
s => s.toUpperCase()
```

Most of the time, Scala can infer the types of parameters for function literals, because the context provided by map tells the compiler what type to expect.

The foreach method is used when we want to process each element and do something with *complete side effects*, without returning a new value, unlike map. Here we print a string to *standard out*, without a newline after each one. The last println call prints the newline before the program exits.

The notion of *side effects* means that the function we pass to foreach does something to affect state outside the local context. We could write to a database or to a file system, or print to the console, as we do here.

Look again at the first line inside `main`, how concise it is, where we compose operations together. Sequencing transformations together lets us create concise, powerful programs, as we'll see over and over again.

We haven't needed to "import" any library items yet, but Scala imports work the same way as they do in Java. Scala automatically imports many commonly used types and features, like the `Seq` and `List` types above and methods for I/O, like `println`, a method on the `scala.Console` object.

In Java, to use `println`, you have to write `System.out.println` or import `System.out` and then write `out.println`. In Scala, you can import objects and even individual methods from them. This is done automatically for you for `Console.println`, so we can just use `println` by itself. This method is one of many methods and types imported automatically that are defined in a library object called `Predef`.

To run this code, you must first compile to a JVM-compatible `.class` file using `scalac`. (Sometimes multiple class files are generated.) SBT will do this for you, but let's see how to do it yourself. If you installed the Scala command-line tools separately, open a terminal window and change the working directory to the root of the project. Then run the following command (ignoring the `$` prompt):

```
$ scalac src/main/scala/progscala3/introscala/UpperMain1.scala
```

You should now have a new directory named `progscala3/introscala` that contains several `.class` and `.tasty` files, including a file named `UpperMain1.class`. ("Tasty" files are an intermediate representation generated by the compiler.) Scala must generate valid JVM byte code and files. For example, the directory structure must match the package structure.

Now, you can execute this program to process a list of strings. Here is an example:

```
$ scala -cp . progscale3.introscala.UpperMain1 Hello World!  
HELLO WORLD!
```

The `-cp .` option adds the current directory to the search classpath, although this is actually the default behavior.

Allowing SBT to compile it for us instead, we can run it at the SBT prompt using this command:

```
> runMain progscale3.introscala.UpperMain1 Hello World!
```

For completeness, if you compile with SBT, but run it using the `scala` command outside SBT, then set the classpath to point to the correct directory where SBT writes class files:

```
$ scala -cp target/scala-3.0.0/classes  
progscale3.introscala.UpperMain1 hello  
HELLO
```

## INTERPRETING VERSUS COMPILING AND RUNNING SCALA CODE

To summarize, if you type `scala` on the command line without a file argument, the REPL is started. You type in commands, expressions, and statements that are evaluated on the fly. If you give the `scala` command a Scala source file argument, it compiles and runs the file as a script. Otherwise, if you provide a JAR file or the name of class with a `main` routine, `scala` executes it just like the `java` command.

Returning to the script version, we can actually simplify it even more. Consider this simplified version:

```
// src/script/scala/progscale3/introscala/Upper2.scala  
  
object Upper2:  
  def convert(strings: Seq[String]) = strings.map(_.toUpperCase())
```

```
println(Upper2.convert(List("Hello", "World!")))
```

This code does exactly the same thing, but it's more concise.

I omitted the return type for the method declaration, because it's "obvious" what's returned. However, we can't omit the type annotations for parameters. Technically, the type inference algorithm does *local type inference*, which means it doesn't work globally over your whole program, but locally within certain scopes. It can't infer caller's expectations for the parameter types, but it is able to infer the type of the method's returned value in most cases, because it sees the whole function body. Recursive functions are one exception where the return type must be declared.

However, explicit type annotations in the parameter lists and explicit return type annotations provide useful documentation for the reader. Just because Scala can infer the return type of a function, should you let it? For simple functions, where the return type is obvious to the reader, perhaps it's not that important to show it explicitly. However, sometimes the inferred type won't be what's expected, perhaps due to a bug or some subtle behavior triggered by certain input arguments or expressions in the function body. Explicit return types express what you *think* should be returned and the compiler confirms it. Hence, I recommend adding return types rather than inferring them, especially in public APIs.

We have also exploited a shorthand for the function literal. Previously we wrote it in the following two, equivalent ways:

```
(s: String) => s.toUpperCase()  
s => s.toUpperCase()
```

We have now shortened it even further to the following expression:

```
_.toUpperCase()
```

The `map` method takes a single function parameter, where the function itself takes a single parameter. In this case, the function body only uses the parameter once, so we can use the anonymous variable `_` instead of a named parameter. The string parameter will be assigned to it before `toUpperCase` is called.

Finally, using an object instead of a `class` simplifies the invocation, because we don't need to first create an instance with `new`. We just call `Upper2.convert` directly.

Let's do one last version of the compiled code (under `src/main/scala`) to show another way of working with collections for this situation:

```
// src/main/scala/progscala3/introscala/UpperMain2.scala
package progscala3.introscala

object UpperMain2:
  def main(params: Array[String]): Unit =
    val output = params.map(_.toUpperCase()).mkString(" ")
    println(output)
```

Instead of using `foreach` to print each transformed word as before, we map the array to a new array of strings, then call a convenience method to concatenate the strings into a final string. There are two `mkString` methods. One takes a single parameter to specify the delimiter between the collection elements. The second version that takes three parameters, a leftmost prefix string, the delimiter, and a rightmost suffix string. Try changing the code to use `mkString("[", ", ", " ]")`.

As a little exercise, return to the script `Upper2.scala` and try simplifying it further. Eliminate the `Upper2` object completely and just call `map` on the list of words directly. You should have just one line of code when you're done! (See the code examples for one implementation.)



## A Sample Application

Let's finish this chapter by exploring several more seductive features of Scala using a sample application. We'll use a simplified hierarchy of geometric shapes, which we will "send" to another object for drawing on a display. Imagine a scenario where a game engine generates scenes. As the shapes in the scene are completed, they are sent to a display subsystem for drawing.

However, to keep it simple, this will be a single-threaded implementation. We won't actually do anything with concurrency right now.

To begin, we define a Shape class hierarchy:

```
// src/main/scala/progscala3/introscala/shapes/Shapes.scala
package progscala3.introscala.shapes

case class Point(x: Double = 0.0, y: Double = 0.0)
❶

abstract class Shape():
❷
  /**
   * Draw the shape.
   * @param f is a function to which the shape will pass a
   * string version of itself to be rendered.
   */
  def draw(f: String => Unit): Unit = f(s"draw: $this")
❸

case class Circle(center: Point, radius: Double) extends Shape
❹

case class Rectangle(lowerLeft: Point, height: Double, width: Double)
❺
  extends Shape

case class Triangle(point1: Point, point2: Point, point3: Point)
❻
  extends Shape
```

- ❶ Declare a class for two-dimensional points.
- ❷ Declare an abstract class for geometric shapes.
- ❸ Implement a draw method for “rendering” the shapes. The documentation comment uses the same conventions that Java uses.
- ❹ A circle with a center and radius.
- ❺ A rectangle with a lower-left point, height, and width. We assume for simplicity that the sides are parallel to the horizontal and vertical axes.
- ❻ A triangle defined by three points.

Let’s unpack what’s going on.

The parameter list after the `Point` class name is the list of constructor parameters. In Scala, the *whole* class body is the constructor, so you list the parameters for the *primary* constructor after the class name and before the class body. In this case, there is no class body, so we can omit the colon (or curly braces if you use those instead).

Because we put the `case` keyword before the class declaration, each constructor parameter is automatically converted to a read-only (immutable) field of `Point` instances. That is, if you instantiate a `Point` instance named `point`, you can read the fields using `point.x` and `point.y`, but you *can’t change their values*. Attempting to use `point.y = 3.0` triggers a compilation error.

You can also provide default values for method parameters, including constructors. The `= 0.0` after each parameter definition specifies `0.0` as the default. Hence, the user doesn’t have to provide them explicitly, but they are inferred left to right.

You can also construct instances without using `new`.

Let's use our SBT project to explore these points:

```
> console
[info] ...

scala> import progscale3.introscala.shapes._

scala> val p00 = Point()
val p00: progscale3.introscala.shapes.Point = Point(0.0,0.0)

scala> val p20 = Point(2.0)
val p20: progscale3.introscala.shapes.Point = Point(2.0,0.0)

scala> val p20b = Point(2.0)
val p20b: progscale3.introscala.shapes.Point = Point(2.0,0.0)

scala> val p02 = Point(y = 2.0)
val p02: progscale3.introscala.shapes.Point = Point(0.0,2.0)

scala> p20 == p20b
val res0: Boolean = true

scala> p20 == p02
val res1: Boolean = false
```

Running `console` will automatically compile the code first, so we don't need to run `compile` first.

The import statement uses `_` as a wildcard to import everything in the `progscale3.introscala.shapes` package. It behaves the same way as using `*` in Java. Scala uses `_` because you might want to use `*` as a method name for multiplication and since Scala lets you import methods into the local scope, using `*` as a wildcard would be ambiguous! This is the second use for `_` that we've seen. The first use was in function literals where `_` was an anonymous placeholder for parameters, used instead of naming them.

In the definition of `p00`, no arguments are specified, so Scala used `0.0` for both of them. (However, you must provide the empty

parentheses.) When one argument is specified, Scala applies it to the leftmost argument, `x`, and used the default value for the remaining argument, as shown for `p20` and `p20b`. We can even specify the arguments with the associated parameter name. The definition of `p02` uses the default value for `x`, but specifies the value for `y`, using `Point(y = 2.0)`.

### TIP

Using named arguments explicitly, even when it isn't required, like `Point(x = 0.0, y = 2.0)`, can make your code easier to understand.

While there is no class body for `Point`, another feature of the `case` keyword is the compiler automatically generates several methods for us, including the familiar `toString`, `equals`, and `hashCode` methods in Java. The output shown for each point, e.g., `Point(2.0,0.0)`, is the default `toString` output. The `equals` and `hashCode` methods are difficult for most developers to implement correctly, so autogeneration of these methods is a real benefit. However, you can provide your own definitions for any of these methods, if you prefer.

When we asked if `p20 == p20b` and `p20 == p02`, Scala invoked the generated `equals` method. This is in contrast with Java, where `==` just compares *references*. In Java, you have to call `equals` explicitly to do a *logical* comparison.

The last feature of case classes that we'll mention now is that the compiler also generates a *companion object*, a “singleton” of the same name, for each case class. In other words, we declared the class `Point` and the compiler also created an *object* `Point`.

## NOTE

You can define companions yourself. Any time an object and a class have the same name *and* they are defined in the same file, they are *companions*.

The compiler also adds several methods to the companion object automatically, one of which is named `apply`. It takes the same parameter list as the constructor.

For *any* instance you use, either a declared object or an instance of a class, if you put an argument list after it, Scala looks for a corresponding `apply` method to call. Therefore, the following two lines are equivalent:

```
val p1 = Point.apply(1.0, 2.0) // Point is the companion object
here!
val p2 = Point(1.0, 2.0)
```

It's a compilation error if no `apply` method exists for the instance. Also, the argument list supplied must conform to the expected parameter list.

The `Point.apply` method is effectively a *factory* for constructing `Points`. The behavior is simple here; it's just like calling the `Point` class constructor *without* the `new` keyword. The companion object generated is equivalent to this:

```
object Point {
  def apply(x: Double = 0.0, y: Double = 0.0) = new Point(x, y)
  ...
}
```

You can add methods to the companion object. A more sophisticated `apply` method might instantiate a different subclass with specialized behavior, depending on the argument supplied. For example, a data structure might have an implementation that is optimal for a small

number of elements and a different implementation that is optimal for a larger number of elements. The `apply` method can hide this logic, giving the user a single, simplified interface. Hence, putting an `apply` method on a companion object is a common idiom for defining a *factory* method for a class hierarchy, whether or not case classes are involved.

An *instance* `apply` method defined on any `class` has whatever meaning is appropriate for instances. For example, `Seq.apply(index: Int)` retrieves the element at position `index` (counting from zero).

### NOTE

To recap, when an argument list is put after an object or `class` instance, Scala looks for an `apply` method to call where the parameter list matches the argument list types. Syntactically, anything with an `apply` method behaves like a *function*, e.g., `Point(2.0, 3.0)`.

A companion object `apply` method is a factory method for the companion class instances. A class `apply` method has whatever meaning is appropriate for instances of the class, for example `Seq.apply(index: Int)` returns the item at position `index`.

`Shape` is an abstract class. We can't instantiate an abstract class, even if none of the members is abstract. `Shape.draw` is defined, but we only want to instantiate concrete shapes: `Circle`, `Rectangle`, and `Triangle`.

The parameter `f` for `draw` is a function of type `String => Unit`. We saw `Unit` above. It is a real type, but it behaves roughly like `void` in other languages.<sup>3</sup>

The idea is that callers of `draw` will pass a function that does the actual drawing when given a string representation of the shape. For simplicity, we just use the string returned by `toString`, but a

structured format like JSON would make more sense in a real application.

### TIP

When a function returns `Unit` it is *totally side-effecting*. There's nothing useful returned from the function, so it can only perform side effects on some state, like performing input or output (I/O).

Normally in functional programming, we prefer *pure* functions that have no side effects and return all their work as their return value. These functions are far easier to reason about, test, compose, and reuse. Side effects are a common source of bugs, so they should be used carefully.

### TIP

Use side effects only when necessary and in well-defined places. Keep the rest of the code *pure*.

`Shape.draw` is another example that functions are *first-class values*, just like instances of `Strings`, `Ints`, `Points`, etc. We saw this previously with `map` and `foreach`. Like other values, we can assign functions to variables, pass them to other functions as arguments, and return them from functions. This is a powerful tool for building composable, yet flexible software.

When a function accepts other functions as parameters or returns functions as values, it is called a *higher-order function* (HOF).

You could say that `draw` defines a *protocol* that all shapes have to support, but users can customize. It's up to each shape to serialize its state to a string representation through its `toString` method. The

`f` method is called by `draw` and it constructs the final string using an *interpolated string*.

An interpolated string starts with `s` before the opening double quote: `s"draw: ${this.toString}"`. It builds the final string by substituting the result of the expression `this.toString` into the larger string.

Actually, we don't need to call `toString`; it will be inferred, so we can use just `${this}`. However, now we're just referring to a variable, so we can drop the curly braces and just write `$this`. Hence, the expression becomes `s"draw: $this"`.

### WARNING

If you forget the `s` before the interpolated string, you'll get the literal output `draw: $this`, with no interpolation.

Back to the code listing, `Circle`, `Rectangle`, and `Triangle` are concrete subclasses of `Shape`. They have no class bodies, because the `case` keyword defines all the methods we need, such as the `toString` methods required by `Shape.draw`.

In our simple program, the `f` we will pass to `draw` will just write the string to the console, but you could build a real graphics application that uses an `f` to parse the string and render the shape to a display, write JSON to a web service, etc.

Even though this will be a single-threaded application, let's anticipate what we might do in a concurrent implementation by defining a set of possible Messages that can be exchanged between modules.

```
// src/main/scala/progscala3/introscala/shapes/Messages.scala
package progscala3.introscala.shapes

sealed trait Message
❶
case class Draw(shape: Shape) extends Message
❷
```



```
case class Response(message: String) extends Message
```

③

```
case object Exit extends Message
```

④

- ① Declare a trait called Message. A trait defines an interface for behavior and can be used as an abstract base class, which is how we use it here. All messages exchanged are subclasses of Message. The sealed keyword is explained below.
- ② A message to draw the enclosed Shape.
- ③ Response is used to return an arbitrary string message to a caller in response to a message received from the caller.
- ④ Exit has no state or behavior of its own, so it is declared a case object, since we only need one instance of it. It functions as a “signal” to trigger a state change, termination in this case.

The sealed keyword means that we can only define subclasses of Message in the same file. This prevents bugs where users define their own Message subtypes that could break the code in the next file! Note that Shape was not declared sealed previously, because we intend for people to create their own subclasses of it!

Now that we have defined our shapes and messages types, let's define an object for processing messages:

```
// src/main/scala/progscala3/introscala/shapes/ProcessMessages.scala  
package progscala3.introscala.shapes
```

```
object ProcessMessages:
```

①

```
  def apply(message: Message): Message =
```

②

```
    message match
```

```
      case Exit =>
```

```
        println(s"ProcessMessage: exiting...")
```

③

```

Exit
case Draw(shape) =>
  shape.draw(str => println(s"ProcessMessage: $str"))
  Response(s"ProcessMessage: $shape drawn")
case Response(unexpected) =>
  val response = Response(s"ERROR: Unexpected Response:
$unexpected")
  println(s"ProcessMessage: $response")
  response

```

- ❶ If we only need one instance, we can declare it an object, but it would be easy enough to make this a class and instantiate as many as we need, for scalability, etc.
- ❷ Define the apply method that takes a Message, processes it, then returns a new Message.
- ❸ Match on the incoming message to determine what to do.

The apply method introduces a powerful feature call *pattern matching*:

```

def apply(message: Message): Message =
  message match:
    case Exit =>
      expressions
    case Draw(shape) =>
      expressions
    case Response(unexpected) =>
      expressions

```

The message match: ... expression consists only of case clauses, which do *pattern matching* on the message passed into the function. This is an expression that returns a value, which you can assign to a variable or use as the return value, as we do here for apply.

Match expressions work a lot like “if/else” expressions. When one of the patterns matches, the expressions are evaluated after the arrow

(=>) up to the next case keyword (or the end of the definition). Matching is eager; the first match wins.

If the case clauses don't cover all possible values that can be passed to the match expression, a `MatchError` is thrown at runtime. Fortunately, the compiler can detect many cases where the match clauses don't handle all possible inputs. Note that our sealed hierarchy of messages is crucial here. If a user created a new subtype of `Message`, our match expression would no longer cover all possibilities. Hence, a bug would be introduced in this code!

A powerful feature of pattern matching is the ability to extract data from the object matched, sometimes called *deconstruction* (the inverse of construction). Here, when the input message is a `Draw`, we extract the enclosed `Shape` and assign it to the variable `shape`. Similarly, if `Response` is detected, we extract the message as `unexpected`, because `ProcessMessages` doesn't expect to receive a `Response`!

Now let's look at the expressions invoked for each case match:

```
def apply(message: Message): Message =  
  message match  
    case Exit => ❶  
      println(s"ProcessMessage: exiting...")  
      Exit  
    case Draw(shape) => ❷  
      shape.draw(str => println(s"ProcessMessage: $str"))  
      Response(s"ProcessMessage: $shape drawn")  
    case Response(unexpected) => ❸  
      val response = Response(s"ERROR: Unexpected Response:  
$unexpected")  
      println(s"ProcessMessage: $response")  
      response
```

- ❶ We're done, so print a message that we're exiting and return `Exit` to the caller.

- ❷ Call `draw` on `shape`, passing it an anonymous function that knows what to do with the string generated by `draw`. In this case, it just prints the string to the console and sends a `Response` to the caller.
- ❸ `ProcessMessages` doesn't expect to receive a `Response` message from the caller, so it treats it as an error. It returns a new `Response` to the caller.

One of the commonly taught tenets of object-oriented programming is that you should never use case statements that match on instance type, because inheritance hierarchies evolve, which breaks these case statements. Instead, polymorphic functions should be used. So, is the pattern-matching code just discussed an *antipattern*?

Recall that we defined `Shape.draw` to call the `toString` method on the `Shape`, which is automatically generated by the compiler for each concrete subclass because they are case classes. Hence, the code in the first case statement invokes a *polymorphic* `toString` operation and we don't match on specific subtypes of `Shape`. This means our code won't break if a user adds a new shape to the class hierarchy by subclassing `Shape`, which we encourage.

The case clauses match on subtypes of `Message`, but we protected ourselves from unexpected change by making `Message` a sealed hierarchy. If we add a new message type here, we can modify the match expression accordingly.

Hence, we have combined polymorphic dispatch from object-oriented programming with pattern matching, a workhorse of functional programming. This is one way that Scala elegantly integrates these two programming paradigms.

## PATTERN MATCHING VERSUS SUBTYPE POLYMORPHISM

*Pattern matching* plays a central role in functional programming just as *subtype polymorphism* (i.e., overriding methods in subtypes) plays a central role in object-oriented programming. Functional pattern matching is much more important and sophisticated than the corresponding switch/case statements found in most languages. The combination of functional-style pattern matching with polymorphic dispatch, as used here, is a powerful combination that is a benefit of a mixed paradigm language like Scala.

Finally, here is the `ProcessShapesDriver` that runs the example:

```
//  
src/main/scala/progscala3/introscala/shapes/ProcessShapesDriver.scala  
package progscala3.introscala.shapes  
  
@main def ProcessShapesDriver =  
  ❶ val messages = Seq(  
    Draw(Circle(Point(0.0,0.0), 1.0)),  
    Draw(Rectangle(Point(0.0,0.0), 2, 5)),  
    Response(s"Say hello to pi: 3.14159"),  
    Draw(Triangle(Point(0.0,0.0), Point(2.0,0.0), Point(1.0,2.0))),  
    Exit)  
    ❷  
  
    messages.foreach { message =>  
      val response = ProcessMessages(message)  
      println(response)  
    }  
    ❸
```

- ❶ An entry point for the application.
- ❷ A sequence of messages to send, including a `Response` in the middle that will be considered an error in `ProcessMessages`. The sequence ends with `Exit`.

- ③ Iterate through the sequence of messages, call `ProcessMessages.apply()` with each one, then print the response.

Let's try it! At the sbt prompt, type `run`, which will compile the code if necessary and then present you with a list of all the code examples that have a main method:

```
> run
[info] Compiling ...

Multiple main classes detected, select one to run:

...
[28] progscale3.introscala.shapes.ProcessShapesDriver
...

Enter number:
```

Enter 28 (or whatever number is shown for you) and the following output is written to the console (wrapped to fit):

```
Enter number: 28

[info] running progscale3.introscala.shapes.ProcessShapesDriver
ProcessMessage: draw: Circle(Point(0.0,0.0),1.0)
Response(ProcessMessage: Circle(Point(0.0,0.0),1.0) drawn)
ProcessMessage: draw: Rectangle(Point(0.0,0.0),2.0,5.0)
Response(ProcessMessage: Rectangle(Point(0.0,0.0),2.0,5.0) drawn)
ProcessMessage: Response(ERROR: Unexpected Response: Say hello to
pi: 3.14159)
Response(ERROR: Unexpected Response: Say hello to pi: 3.14159)
ProcessMessage: draw:
Triangle(Point(0.0,0.0),Point(2.0,0.0),Point(1.0,2.0))
Response(ProcessMessage: Triangle(Point(0.0,0.0),Point(2.0,0.0),...)
drawn)
ProcessMessage: exiting...
Exit
[success] ...
```

Make sure you understand how each message was processed and where each line of output came from.

## Recap and What's Next

We introduced many of the powerful and concise features of Scala. As you explore Scala, you will find other useful resources that are available on <http://scala-lang.org>. You will find links for libraries, tutorials, and various papers that describe features of the language.

Next we'll continue our introduction to Scala features, emphasizing the various concise and efficient ways of getting lots of work done.

- 
- 1 Lately, the acronym API has been used to refer to the interfaces exposed by services. I'll use API in its original sense, the types, methods, and values exposed by a code module.
  - 2 For more details on these and other tools, see *Chapter 21*.
  - 3 In fact, there is a single instance of `Unit` named `()`. If you are curious about that name, see *later chapter content*.

# Chapter 2. Type Less, Do More

---

## A NOTE FOR EARLY RELEASE READERS

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

This will be the 2nd chapter of the final book. Please note that the GitHub repo will be made active later on.

If you have comments about how we might improve the content and/or examples in this book, or if you notice missing material within this chapter, please reach out to the author at [mcronin@oreilly.com](mailto:mcronin@oreilly.com).

This chapter continues our tour of Scala features that promote succinct, flexible code. We’ll discuss organization of files and packages, importing other types, variable and method declarations, a few particularly useful types, and miscellaneous syntax conventions.

## New Scala 3 Syntax

If you have prior Scala experience, Scala 3 introduces a new *optional braces* syntax that makes it look a lot more like Python or Haskell, where Java-style curly braces (`{...}`) are replaced with indentation that becomes significant. The examples in the previous chapter and throughout the book use it.

This syntax has several benefits. It is more concise. For Python developers who decide to learn Scala, it will feel more familiar to



them (and vice versa).

There is also a new syntax for control structures like `for` loops and `if` expressions. For example, `if condition then ...` instead of the older `if (condition) ....` Also, `for ... do println(...)` instead of `for {...} println(...)`.

The disadvantage of these changes is that they are strictly not necessary. Some breaking changes in Scala 3 are necessary to move the language forward, but you could argue these syntax changes aren't required. You also have to be careful to use tabs and spaces consistently for indentation.

The new constructs are the defaults supported by the compiler, but using the compiler flags `-old-syntax` and `-noindent` will enforce the old syntax constructs. Another flag, `-new-syntax` makes the new keywords `then` and `do` required. Finally, the compiler can now rewrite your code to use whichever style you prefer. Add the `-rewrite` compiler flag, for example `-rewrite -new-syntax`.

I opposed these changes initially, because they aren't strictly necessary. However, now that I have worked with them, I believe the advantages outweigh the disadvantages. I also suspect that the syntax changes are the future for Scala and eventually the old syntax could be deprecated. We'll see. Hence, I have chosen to use the new conventions throughout this edition of the book. I will mention other pros and cons of these changes as we explore examples.

*Appendix A* provides examples of the old and new syntax.

### WARNING

Even if you continue using the brace syntax, the compiler now requires indentation to use tabs and spaces consistently, unless you use the `-noindent` compiler flag.

## Semicolons

Semicolons are expression delimiters and they are inferred. Scala treats the end of a line as the end of an expression, except when it can infer that the expression continues to the next line, even for the following example:

```
scala> val s = "hello"  
      |   + "world"  
      |   + "!"  
val s: String = helloworld!
```

The Scala 2 REPL was more aggressive at interpreting the end of a line as the end of an expression, so the previous example would infer “hello” for the definition of `s` and then throw an error for the next two lines.

### TIP

When using the Scala 2 REPL, use the `:paste` mode when multiple lines need to be parsed as a whole. Enter `:paste`, followed by the code you want to enter, then finish with `Ctrl-D`.

Conversely, you can put multiple expressions on the same line, separated by semicolons.

## Variable Declarations

Scala allows you to decide whether a variable is immutable (read-only) or not (read-write) when you declare it. We’ve already seen that an immutable “variable” is declared with the keyword `val`:

```
val array: Array[String] = new Array(5)
```

Scala is like Java in that most variables are actually references to heap-allocated objects. Hence, the array reference cannot be changed to point to a different Array, *but the array elements themselves are mutable*, so the elements can be modified:

```
scala> val array: Array[String] = new Array(5)
val array: Array[String] = Array(null, null, null, null, null)

scala> array = new Array(2)
1 | array = new Array(2)
  | ^^^^^^^^^^^^^^^^^^^
  | Reassignment to val array

scala> array(0) = "Hello"

scala> array
val res1: Array[String] = Array(Hello, null, null, null, null)
```

A `val` must be initialized when it is declared, except in certain contexts like abstract fields in type declarations.

Similarly, a mutable variable is declared with the keyword `var` and it must also be initialized immediately (in most cases), even though it can be changed later:

```
scala> var stockPrice: Double = 100.0
var stockPrice: Double = 100.0

scala> stockPrice = 200.0
stockPrice: Double = 200.0
```

To be clear, we changed the *value* of `stockPrice` itself. However, the “object” that `stockPrice` refers to can’t be changed, because `Doubles` in Scala are immutable.

In Java, so-called *primitive* types, `char`, `byte`, `short`, `int`, `long`, `float`, `double`, and `boolean`, are fundamentally different than reference objects. Indeed, there is no object and no reference, just the “raw” value. Scala tries to be consistently object-oriented, so

these types *are* actually objects with methods, like reference types (see “**Reference Versus Value Types**”). However, Scala compiles to primitives where possible, giving you the performance benefit they provide (see *later chapters* for details).

Consider the following REPL session, where we define a Person class with immutable first and last names, but a mutable age (because people age, I guess). The parameters are declared with `val` and `var`, respectively, making them both fields in Person:

```
// src/script/scala/progscala3/typelessdomore/Human.scala
scala> class Human(val name: String, var age: Int)
// defined class Human

scala> val p = new Human("Dean Wampler", 29)
val p: Human = Human@165a128d

scala> p.name
val res0: String = Dean Wampler

scala> p.name = "Buck Trends"
1 | p.name = "Buck Trends"
  | ^^^^^^^^^^^^^^^^^^^
  | Reassignment to val name

scala> p.name
val res1: String = Dean Wampler

scala> p.age
val res2: Int = 29

scala> p.age = 30

scala> p.age
val res3: Int = 30

scala> p.age = 31; p.age // Use semicolon to join two expressions...
val res4: Int = 31
```

## NOTE

The `var` and `val` keywords only specify whether the reference can be changed to refer to a different object (`var`) or not (`val`). They don't specify whether or not the object they reference is mutable.

Use immutable values whenever possible to eliminate a class of bugs caused by mutability. For example, a mutable object is dangerous as a key in hash-based maps. If the object is mutated, the output of the `hashCode` method will change, so the corresponding value won't be found at the original location.

More common is unexpected behavior when an object you are using is being changed by another thread. Borrowing a phrase from Quantum Physics, these bugs are *spooky action at a distance*. Nothing you are doing locally accounts for the unexpected behavior; it's coming from somewhere else.

These are the most pernicious bugs in multithreaded programs, where synchronized access to shared, mutable state is required, but difficult to get right. Using immutable values eliminates these issues.

## Ranges

Sometimes we need a sequence of numbers from some start to finish. A `Range` literal is just what we need. The following examples show how to create ranges for the types that support them, `Int`, `Long`, `Char`, `BigInt`, which represents integers of arbitrary size, and `BigDecimal`, which represents floating-point numbers of arbitrary size. `Float` and `Double` were supported in Scala 2, but floating point arithmetic makes range calculations error prone. Hence, Scala 3 drops ranges for `Float` and `Double`.

You can create ranges with an inclusive or exclusive upper bound, and you can specify an interval not equal to one (some output elided

to fit):

```
scala> 1 to 10 // Int range inclusive, interval of 1,
(1 to 10)
val res0: scala.collection.immutable.Range.Inclusive = Range 0 to 10

scala> 1 until 10 // Int range exclusive, interval of 1,
(1 to 9)
val res1: Range = Range 0 until 10

scala> 1 to 10 by 3 // Int range inclusive, every third.
val res2: Range = inexact Range 0 to 10 by 3

scala> 10 to 1 by -3 // Int range inclusive, every third,
counting down.
val res3: Range = Range 10 to 1 by -3

scala> 1L to 10L by 3 // Long
val res4: ...immutable.NumericRange[Long] = NumericRange 1 to 10 by 3

scala> 'a' to 'g' by 3 // Char
val res5: ...immutable.NumericRange[Char] = NumericRange a to g by 3

scala> BigInt(1) to BigInt(10) by 3
val res6: ...immutable.NumericRange[BigInt] = NumericRange 1 to 10
by 3

scala> BigDecimal(1.1) to BigDecimal(10.3) by 3.1
val res7: ...immutable.NumericRange.Inclusive[BigDecimal] =
  NumericRange 1.1 to 10.3 by 3.1
```

## Partial Functions

A `PartialFunction[A,B]` is a special kind of function with its own literal syntax. A is the type of the single parameter the function accepts and B is the return type.

The literal syntax for a `PartialFunction` consists only of case clauses, which we saw in “A Sample Application”, that do *pattern matching* on the input to the function. There is no function parameter

shown explicitly, but when each input is processed, it is passed to the body of the partial function.

For comparison, is a regular function that does pattern matching and similar partial function, adapted from the example we explored in “[A Sample Application](#)”:

```
//
src/script/scala/progscala3/typelessdomore/FunctionVsPartialFunction.s
cala
scala> import progscala3.introscala.shapes._

scala> val func: Message => String = message => message match
    | case Exit => "Got Exit"
    | case Draw(shape) => s"Got Draw($shape)"
    | case Response(str) => s"Got Response($str)"

scala> val pfunc: PartialFunction[Message, String] =
    | case Exit => "Got Exit"
    | case Draw(shape) => s"Got Draw($shape)"
    | case Response(str) => s"Got Response($str)"

scala> func(Draw(Circle(Point(0.0,0.0), 1.0)))
    | pfunc(Draw(Circle(Point(0.0,0.0), 1.0)))
    | func(Response(s"Say hello to pi: 3.14159"))
    | pfunc(Response(s"Say hello to pi: 3.14159"))
val res0: String = Got Draw(Circle(Point(0.0,0.0),1.0))
val res1: String = Got Draw(Circle(Point(0.0,0.0),1.0))
val res2: String = Got Response(Say hello to pi: 3.14159)
val res3: String = Got Response(Say hello to pi: 3.14159)
```

Function definitions can be a little harder to read than method definition. The function `func` is a named function of type `Message => String`. The equal sign starts the body, `message => message match` ....

The partial function, `pfunc`, is simpler. It's type is `PartialFunction[Message, String]`. There is no argument list, just a set of case match clauses, which happen to be identical to the clauses in `func`.

The concept of a *partial function* is simpler than it might appear. In essence, a partial function will only handle certain inputs, so don't send it something it doesn't know how to handle. A classic example from mathematics is division,  $x/y$ , which is undefined when the denominator  $y$  is 0. Hence, division is a *partial function*.

If a partial function is called with an input that doesn't match one of the case clauses, a `MatchError` is thrown at runtime. Both `func` and `pfunc` are actually *total*, because they handle all possible `Message` arguments. Try commenting out the case `Exit` clauses in both `func` and `pfunc`. You'll get a compiler warning for `func`, because it can determine that the match clauses don't handle all possible inputs. It won't complain about `pfunc`, because that situation is *by design*.

You can test if a `PartialFunction` will match an input using the `isDefinedAt` method. This function avoids the risk of throwing a `MatchError` exception.

You can also “chain” `PartialFunctions` together:

`pf1.orElse(pf2).orElse(pf3) ....` If `pf1` doesn't match, then `pf2` is tried, then `pf3`, etc. A `MatchError` is only thrown if none of them matches.

Let's explore these points with the following example:

```
// src/script/scala/progscala3/typelessdomore/PartialFunctions.scala
```

```
val pfs: PartialFunction[Any,String] =  
  ❶ case s:String => "YES"  
val pfd: PartialFunction[Any,String] = {  
  ❷ case d:Double => "YES"  
}  
  
val pfsd = pfs.orElse(pfd) ❸
```

- ❶ A partial function that only matches on strings, using the braceless syntax.



- ② A partial function that only matches on doubles, using braces.
- ③ Combine the two functions to construct a new partial function that matches on strings *and* doubles.

The next block of code in the script tries different values with the three partial functions to confirm expected behavior. Note that integers are not handled by any combination. A helper function `tryPF` is used to try the partial function and catch possible `MatchError` exceptions. So, a string is returned for both success and failure:

```
def tryPF(x: Any, f: PartialFunction[Any,String]): String =  
  try f(x)  
  catch case _: MatchError => "ERROR!"  
  
assert(tryPF("str", pfs) == "YES")  
assert(tryPF("str", pfd) == "ERROR!")  
assert(tryPF("str", pfsd) == "YES")  
assert(tryPF(3.142, pfs) == "ERROR!")  
assert(tryPF(3.142, pfd) == "YES")  
assert(tryPF(3.142, pfsd) == "YES")  
assert(tryPF(2, pfs) == "ERROR!")  
assert(tryPF(2, pfd) == "ERROR!")  
assert(tryPF(2, pfsd) == "ERROR!")  
  
assert(pfs.isDefinedAt("str") == true)  
assert(pfd.isDefinedAt("str") == false)  
assert(pfsd.isDefinedAt("str") == true)  
assert(pfs.isDefinedAt(3.142) == false)  
assert(pfd.isDefinedAt(3.142) == true)  
assert(pfsd.isDefinedAt(3.142) == true)  
assert(pfs.isDefinedAt(2) == false)  
assert(pfd.isDefinedAt(2) == false)  
assert(pfsd.isDefinedAt(2) == false)
```

Finally, we can *lift* a partial function into a regular (“total”) function that returns an option, `Some(value)`, when the partial function is defined for the input argument and `None` when it isn’t. We can also *unlift* a single-parameter function. Here is a session that uses `pfs`:

```

scala> val fs = pfs.lift
val fs: Any => Option[String] = <function1>

scala> fs("str")
val res0: Option[String] = Some(YES)

scala> fs(3.142)
val res1: Option[String] = None

scala> val pfs2 = fs.unlift
val pfs2: PartialFunction[Any, String] = <function1>

scala> pfs2("str")
val res3: String = YES

scala> tryPF(3.142, pfs2) // Use tryPF we defined above
val res4: String = ERROR!

```

## Infix Operator Notation

In the previous example, we combined two partial functions using `orElse`. This can be written equivalently in two ways:

```

val pfsd1 = pfs.orElse(pfd)
val pfsd2 = pfs orElse pfd

```

When a method takes a single parameter, you can drop the period after the object and drop the parentheses around the supplied argument. In this case, `pfs orElse pfd` has a cleaner appearance than `pfs.orElse(pfd)`, which is why this syntax is popular. This notation is called *infix notation*, because `orElse` is between the object and argument. This syntax is also called *operator notation*, because it is especially popular when writing libraries where algebraic notation is convenient. For example, you can write your own libraries for matrices and define a method named `*` for matrix multiplication using the `*` “operator”. Then you can write expressions like `val matrix3 = matrix1 * matrix2`.

## TIP

Scala method names can use most non-alphanumeric characters. When methods are called that take a single parameter, *infix operator notation* can be used where the period after the object and the parentheses around the supplied argument can be dropped.

## Method Declarations

Let's explore method definitions, using a modified version of our Shapes hierarchy from before.

## Method Default and Named Parameters

Here is an updated Point case class:

```
// src/main/scala/progscala3/typelessdomore/shapes/Shapes.scala
package progscala3.typelessdomore.shapes

case class Point(x: Double = 0.0, y: Double = 0.0):
  ❶ def shift(deltax: Double = 0.0, deltay: Double = 0.0) =
    ❷   copy(x + deltax, y + deltay)
```

- ❶ Define Point with default initialization values (as before). For case classes, both x and y are automatically immutable (val) fields.
- ❷ A new shift method for creating a new Point instance, offset from the existing Point. It uses the copy method that is also created automatically for case classes.

The copy method allows you to construct new instances of a case class while specifying just the fields that are changing. This is very useful for larger case classes:

```
scala> val p1 = new Point(x = 3.3, y = 4.4)    // Used named
arguments explicitly.
val p1: Point = Point(3.3,4.4)

scala> val p2 = p1.copy(y = 6.6)    // Copied with a new y value.
val p2: Point = Point(3.3,6.6)
```

Named arguments make client code more readable. They also help avoid bugs when a parameter list has several fields of the same type or it has a lot of parameters. It's easy to pass values in the wrong order. Of course, it's better to avoid such parameter lists in the first place.

## Methods with Multiple Parameter Lists

Next, consider the following changes to `Shape.draw()`:

```
abstract class Shape():
  def draw(offset: Point = Point(0.0, 0.0))(f: String => Unit): Unit
  =
    f(s"draw($offset, ${this})")
```

`Circle`, `Rectangle`, and `+Triangle` are unchanged and not shown.

Now `draw` has *two* parameter *lists*, each of which has a single parameter, rather than a single parameter list with two parameters. The first parameter list lets you specify an offset point where the shape will be drawn. It has a default value of `Point(0.0, 0.0)`, meaning no offset. The second parameter list is the same as in the original version of `draw`, a function that does the drawing.

You can have as many parameter lists as you want, but it's rare to use more than two.

So, why allow more than one parameter list? Multiple lists promote a very nice block-structure syntax when the last parameter list takes a single function. Here's how we might invoke this new `draw` method to draw a `Circle` at an offset:

```
val s = Circle(Point(0.0, 0.0), 1.0)
s.draw(Point(1.0, 2.0))(str => println(str))
```

Scala lets us replace parentheses with curly braces around a supplied argument (like a function literal) for a parameter list that has a single parameter. So, this line can also be written this way:

```
s.draw(Point(1.0, 2.0)){str => println(str)}
```

Suppose the function literal is too long for one line or it has multiple statements or expressions? We can rewrite it this way:

```
s.draw(Point(1.0, 2.0)) { str =>
  println(str)
}
```

Or equivalently:

```
s.draw(Point(1.0, 2.0)) {
  str => println(str)
}
```

If you use the traditional curly-brace syntax for Scala, it looks like a typical block of code we use with constructs like `if` and `for` expressions, method bodies, etc. However, the `{...}` block is still a function literal we are passing to `draw`.

So, this “syntactic sugar” of using `{...}` instead of `(...)` looks better with longer function literals; they look more like the block structure syntax we know.

Unfortunately, the new optional braces syntax doesn’t work here:

```
scala> s.draw(Point(1.0, 2.0)):
      |   str => println(str)
2 |   str => println(str)
      |   ^
      |   parentheses are required around the parameter of a lambda
      |   This construct can be rewritten automatically under -rewrite.
```

```

1 | s.draw(Point(1.0, 2.0)):
   | ^
   | not a legal formal parameter
2 |   str => println(str)

```

However, there is an experimental compiler flag `-Yindent-colons` that enables this capability, but it remains experimental (at the time of this writing), because it “is more contentious and less stable than the rest of the significant indentation scheme.” (Quote from this [Dotty documentation](#).)

Back to using parentheses or braces, if we use the default value for offset, the first set of parentheses is still required. Otherwise, the function would be parsed as the offset, triggering an error.

```

s.draw() {
  str => println(str)
}

```

To be clear, `draw` could just have a single parameter list with two values, like Java methods. If so, the client code would look like this:

```

s.draw(Point(1.0, 2.0), str => println(str))

```

That’s not nearly as clear and elegant. It would also prevent us from using the default value for the offset.

By the way, we can simplify our expressions even more. `str => println(str)` is an anonymous function that takes a single string argument and passes it to `println`. Although, `println` is implemented as a method in the Scala library, it can also be used as a function that takes a single string argument! Hence, the following two lines behave the same:

```

s.draw(Point(1.0, 2.0))(str => println(str))
s.draw(Point(1.0, 2.0))(println)

```

To be clear, they are *not* identical, they just do the same thing. In the first example, we pass an anonymous function that calls `println`. In the second example, we use `println` as a *named* function directly. Scala handles converting methods to functions in situations like this.

Another advantage of allowing two or more parameter lists is that we can use one or more lists for normal parameters and other lists for *using clauses*, formerly known as *implicit parameter lists*. These are parameter lists declared with the `using` or `implicit` keyword. When the methods are called, we can either explicitly specify arguments for these parameters, or we can let the compiler fill them in using a suitable value that's in scope. Using clauses provides a more flexible alternative to parameters with default values. Let's explore an example from the Scala library that uses this mechanism, Futures.

## A Taste of Futures

The `scala.concurrent.Future` API is another tool for concurrency. Akka uses Futures, but you can use them separately when you don't need the full capabilities of actors.

When you wrap some work in a Future, the work is executed asynchronously and the Future API provides various ways to process the results, such as providing callbacks that will be invoked when the result is ready. Let's use callbacks here and defer discussion of the rest of the API until *Chapter 18*.

The following example fires off five work items concurrently and handles the results as they finish:

```
// src/script/scala/progscala3/typelessdomore/Futures.scala
import scala.concurrent.Future
import scala.concurrent.ExecutionContext.Implicits.global
❶ import scala.util.{Failure, Success}

def sleep(millis: Long) = Thread.sleep(millis)
❷
```

```

(1 to 5) foreach { i =>
  val future = Future {
    val duration = (math.random * 1000).toLong
    sleep(duration)
    if i == 3 then throw new RuntimeException(s"$i -> $duration")
    duration
  }
  future onComplete {
    case Success(result)    => println(s"Success! # $i -> $result")
    case Failure(throwable) => println(s"FAILURE! # $i -> $throwable")
  }
}
sleep(1000) // Wait long enough for the "work" to finish.
println("Finished!")

```

- ❶ We'll discuss this import below.
- ❷ A `sleep` method to simulate staying busy for a period amount of time.
- ❸ Pass a block of work to the `scala.concurrent.Future.apply` method. It calls `sleep` with a duration, a randomly generated number of milliseconds between 0 and 1000, which it will also return. However, if `i` equals 3, we throw an exception.
- ❹ Use `onComplete` to assign a partial function to handle the computation result. Notice that the expected output is either `scala.util.Success` wrapping a value or `scala.util.Failure` wrapping an exception.

Success and Failure are subclasses of `scala.util.Try`, which encapsulates `try {...} catch {...}` clauses with less boilerplate. We can handle successful code *and* possible exceptions more uniformly. See “[Try: When There Is No Do](#)” for further discussion.

When we iterate through a `Range` of integers from 1 to 5, inclusive, we construct a `Future` with a block of work to do. `Future.apply` returns a new `Future` instance *immediately*. The body is executed



asynchronously on another thread. The `onComplete` callback we register will be invoked when the body completes.

A final `sleep` call waits before exiting to allow the futures to finish.

A sample run might go like this, where the order of the results and the numbers on the right-hand side are arbitrary, as expected:

```
Success! #2 -> 178
Success! #1 -> 207
FAILURE! #3 -> java.lang.RuntimeException: 3 -> 617
Success! #5 -> 738
Success! #4 -> 938
Finished!
```

You might wonder about the “body of work” we’re passing to `Future.apply`. Is it a function or something else? Here is part of the declaration of `Future.apply`

```
apply[T](body: => T)(/* explained below */): Future[T]
```

Note how the type of `body` is declared, `=> T`. This is called a *by-name parameter*. We are passing something that will return a `T` instance, but we want to evaluate `body` *lazily*. Go back to the example body we passed to `Future.apply` above. We *did not* want that code evaluated before it was passed to `Future.apply`. We wanted it evaluated inside the `Future` *after* construction. This is what by-name parameters do for us. We can pass a block of code that will be evaluated only when needed. The implementation of `Future.apply` evaluates this code.

Okay, let’s finally get back to implicit parameters. Note the second import statement:

```
import scala.concurrent.ExecutionContext.Implicits.global
```

`Future` methods use an `ExecutionContext` to run code in separate threads, providing concurrency. These methods use a given value of

an `ExecutionContext`. For example, here's the whole `Future.apply` declaration (using Scala 3 syntax):

```
apply[T](body: => T)(using executor: ExecutionContext): Future[T]
```

In the Scala 2 library, the `implicit` keyword is used instead of `using`. The second parameter list is called a *using clause*.

Because this parameter is in its own parameter list starting with `using` or `implicit`, users of `Future.apply` don't have to pass a value explicitly. This reduces code boilerplate. We imported the default `ExecutionContext.global` value that is declared as given (or `implicit` in Scala 2). It uses a thread pool with a *work-stealing* algorithm to balance the load and optimize performance.

We can tailor how threads are used by passing our own `ExecutionContext` explicitly:

```
Future(work)(using someExecutionContext)
```

Alternatively, we can declare our own given value that will be used implicitly when `Future.apply` is called:

```
given myEC as MyCustomExecutionContext(arguments)
...
val future = Future(work)
```

The `global` value is declared in a similar way, but our given value will take precedence.

The `Future.onComplete` method we used also has a `using` clause:

```
abstract def onComplete[U](
  f: (Try[T]) => U)(using executor: ExecutionContext): Unit
```

So, when `global` is imported into the current scope, the compiler will use it when methods are called that have a `using` clause with an `ExecutionContext` parameter, unless we specify a value explicitly.

For this to work, only given instances that are type compatible with the parameter will be considered.

The details for the new idioms and the reasons for their existence are explained in [Chapter 5](#).

## Nesting Method Definitions and Recursion

Method definitions can also be nested. This is useful when you want to refactor a lengthy method body into smaller methods, but the “helper” methods aren’t needed outside the original method. Nesting them inside the original method means they are invisible to the rest of the code base, including other methods in the type.

Here is an example for a factorial calculator:

```
// src/script/scala/progscala3/typelessdomore/Factorial.scala

def factorial(i: Int): Long =
  def fact(i: Int, accumulator: Long): Long =
    if (i <= 1) accumulator
    else fact(i - 1, i * accumulator)

  fact(i, 1L)

(0 to 5).foreach(i => println(s"$i: ${factorial(i)}"))
```

The last line prints the following:

```
0: 1
1: 1
2: 2
3: 6
4: 24
5: 120
```

The fact method calls itself recursively, passing an accumulator parameter, where the result of the calculation is “accumulated.” Note that we return the accumulated value when the counter *i* reaches 1.

(We're ignoring negative integer arguments, which would be invalid. The function just returns 1 for  $i \leq 1$ .) After the definition of the nested method, `factorial` calls it with the passed-in value `i` and the initial accumulator "seed" value of 1.

Notice that we use `i` as a parameter name twice, first in the `factorial` method and again in the nested `fact` method. The use of `i` as a parameter name for `fact` "shadows" the outer use of `i` as a parameter name for `factorial`. This is fine, because we don't need the outer value of `i` inside `fact`. We only use it the first time we call `fact`, at the end of `factorial`.

Like a local variable declaration in a method, a nested method is also only visible inside the enclosing method.

Look at the return types for the two functions. We used `Long` because factorials grow in size quickly. So, we didn't want Scala to infer `Int` as the return type. Otherwise, we don't need the type annotation on `factorial`.

However, we *must* declare the return type for `fact`, because it is recursive and Scala's local-scope type inference can't infer the return type of recursive functions.

You might be a little nervous about a recursive function. Aren't we at risk of blowing up the stack? The JVM and many other language environments don't do *tail-call optimizations*, which would convert a tail-recursive function into a loop. This prevents stack overflow and also makes execution faster by eliminating the additional function invocations.

The term *tail-recursive* means that the recursive call is the *last thing done* in an expression and only *one* recursive call is made. If we made the recursive call, *then* added something to the result, for example, that would not be a *tail* call. This doesn't mean that non-tail-call recursion is disallowed, just that we can't optimize it into a loop.

Recursion is a hallmark of functional programming and a powerful tool for writing elegant implementations of many algorithms. Hence, the Scala compiler does limited tail-call optimizations itself. It will handle functions that call themselves, but not so-called “trampoline” calls, i.e., “a calls b calls a calls b,” etc.

Still, you might want to know if you got it right and the compiler did in fact perform the optimization. No one wants a blown stack in production. Fortunately, the compiler can tell you if you got it wrong, if you add an annotation, `tailrec`, as shown in this refined version of factorial:

```
// src/script/scala/progscala3/typelessdomore/FactorialTailrec.scala
import scala.annotation.tailrec

def factorial(i: Int): Long =
  @tailrec
  def fact(i: Int, accumulator: Long): Long =
    if i <= 1 then accumulator
    else fact(i - 1, i * accumulator)

  fact(i, 1)

(0 to 5).foreach(i => println(s"$i: ${factorial(i)}"))
```

If fact is not actually tail recursive, the compiler will throw an error. Consider this attempt to write a naïve recursive implementation of Fibonacci sequences:

```
// src/script/scala/progscala3/typelessdomore/FibonacciTailrec.scala
scala> import scala.annotation.tailrec

scala> @tailrec
      | def fibonacci(i: Int): Long =
      |   if (i <= 1) 1L
      |   else fibonacci(i - 2) + fibonacci(i - 1)
4 |   else fibonacci(i - 2) + fibonacci(i - 1)
      |                                     ^^^^^^^^^^^^^^^^^^^^^
      |                                     Cannot rewrite recursive call: it is not in tail
position
```

```

4 | else fibonacci(i - 2) + fibonacci(i - 1)
  |           ^^^^^^^^^^^^^^^^^^^^^
  |           Cannot rewrite recursive call: it is not in tail position

```

We are attempting to make *two* recursive calls, not one, and *then* do something with the returned values, add them. So, this function is not tail recursive. (It's naïve because it is possible to write a tail recursive implementation.)

Finally, the nested function can see anything in scope, including arguments passed to the outer function. Note the use of *n* in *count* in the next example:

```

// src/script/scala/progscala3/typelessdomore/CountTo.scala
import scala.annotation.tailrec

def countTo(n: Int): Unit =
  @tailrec
  def count(i: Int): Unit =
    if (i <= n) then
      println(i)
      count(i + 1)
  count(1)

countTo(5)

```

## Inferring Type Information

Statically typed languages provide wonderful compile-time safety, but they can be very verbose if all the type information has to be explicitly provided. Scala's *type inference* removes most of this explicit detail, but where it is still required, it can provide an additional benefit of documentation for the reader.

Some functional programming languages, like Haskell, can infer almost all types, because they do global type inference. Scala can't do this, in part because Scala has to support *subtype polymorphism*, for object-oriented inheritance, which makes type inference harder.

We've already seen examples of Scala's type inference. Here are two more examples, showing different ways to declare a `Map`:

```
scala> val map1: Map[Int, String] = Map.empty
val map1: Map[Int, String] = Map()

scala> val map2 = Map.empty[Int, String]
val map1: Map[Int, String] = Map()
```

The second form is more idiomatic most of the time. However, `Map` is actually a trait with concrete subclasses, so you'll sometimes make declarations like this one for a `TreeMap`:

```
scala> import scala.collection.immutable.TreeMap
scala> val map3: Map[Int, String] = TreeMap.empty
val map3: Map[Int, String] = Map()
```

Here is a summary of the rules for when explicit type annotations are required in Scala.

## WHEN EXPLICIT TYPE ANNOTATIONS ARE REQUIRED

In practical terms, you have to provide explicit type annotations for the following situations:

- Abstract `var` or `val` declarations in an abstract class or trait.
- All method parameters (e.g., `def deposit(amount: Money) = ...`).
- Method return types in the following cases:
  - When you explicitly call `return` in a method (even at the end).
  - When a method is recursive.
  - When two or more methods are overloaded (have the same name) and one of them calls another. The *calling* method needs a return type annotation.
  - When the inferred return type would be more general than you intended, e.g., `Any`.

The last case is somewhat rare, fortunately.

### NOTE

The `Any` type is the root of the Scala type hierarchy. If a block of code is inferred to return a value of type `Any` unexpectedly, chances are good that the code is more general than you intended so that `Any` is the only common super type of all possible values.

We'll explore Scala's types in [Chapter 13](#).



Let's look at a few examples of cases we haven't seen yet where explicit types are required. First, look at overloaded methods:

```
//
src/script/scala/progscala3/typelessdomore/MethodOverloadedReturn.scala

case class Money(value: Double)
case object Money {
  def apply(s: String): Money    = apply(s.toDouble)
  def apply(d: BigDecimal): Money = apply(d.toDouble)
}
```

While the Money constructor expects a Double, we want the user to have the convenience of passing a String or a BigDecimal (ignoring possible errors). So, we *add* two more apply methods to the companion object. Both call the apply(d: Double) method the compiler automatically generates for the companion object, corresponding to the primary constructor Money(value: Double).

The two methods have explicit return types. If you try removing them, you'll get a compiler error:

```
scala> case class Money(value: Double)
      | case object Money {
      |   def apply(s: String)      = apply(s.toDouble) // no return
type  |   def apply(d: BigDecimal) = apply(d.toDouble) // no return
type  | }
      |
4 | def apply(d: BigDecimal) = apply(d.toDouble)
      |                               ^
      |                               Overloaded or recursive method apply needs
return type
3 | def apply(s: String)      = apply(s.toDouble)
      |                               ^
      |                               Overloaded or recursive method apply needs
return type
```

## Variadic Argument Lists

Scala supports methods that take *variable argument lists* (sometimes just called *variadic methods*). Consider this contrived example that computes the mean of Doubles:

```
// src/script/scala/progscala3/typelessdomore/VariadicArguments.scala
```

```
scala> object Mean1 {  
  |   def calc1(ds: Double*): Double = calc2(ds :_*)  
  |   def calc2(ds: Seq[Double]): Double = ds.sum/ds.size  
  | }
```

```
scala> Mean1.calc1(1.0, 2.0)  
val res0: Double = 1.5
```

```
scala> Mean1.calc2(Seq(1.0, 2.0))  
val res1: Double = 1.5
```

The syntax `ds: Double*` means zero or more Doubles, a *variable argument list*. Since `calc1` calls `calc2`, which expects a `Seq[Double]` argument, the unusual syntax `(ds :_*)` is how you take the variable argument list and convert to a sequence when needed.

Why have both functions? The examples show that both can be convenient for the user. In particular, using `calc1` doesn't require you to wrap values in a `Seq` first.

There are downsides. The API "footprint" is larger with two methods instead of one and the maintenance burden is larger.

Assuming you want both methods, why not use the same name, in particular, `apply`?

```
scala> object Mean2 {  
  |   def apply(ds: Double*): Double = apply(ds :_*)  
  |   def apply(ds: Seq[Double]): Double = ds.sum/ds.size  
  | }  
3 | def apply(ds: Seq[Double]): Double = ds.sum/ds.size  
  |   ^  
  |   Double definition:  
  |   def apply(ds: Double*): Double in object Mean2 at line 2  
and
```

```
|      def apply(ds: Seq[Double]): Double in object Mean2 at line 3
|      have the same type after erasure.
```

In fact, `ds: Double*$$` is converted to a kind of sequence internally, so effectively the methods look identical at the byte code level.

There's a common idiom to break the ambiguity, add a first `Double` parameter in the first `apply`, then use a variable list for the rest of the supplied arguments:

```
scala> object Mean {
|     def apply(d: Double, ds: Double*): Double = apply(d +: ds)
|     def apply(ds: Seq[Double]): Double = ds.sum/ds.size
| }
// defined object Mean

scala> Mean(1.0,2.0)
val res10: Double = 1.5

scala> Mean(Seq(1.0,2.0))
val res11: Double = 1.5

scala> Mean()
1 | Mean()
  | ^^^^
  | None of the overloaded alternatives of method apply in object Mean
  | with types
  |   (ds: Seq[Double]): Double
  |   (d: Double, ds: Double*): Double
  | match arguments ()

scala> Mean(Nil)
val res12: Double = NaN
```

When calling the second `apply`, the first one constructs a new sequence prepending `d` to `ds` using `d +: ds`. We'll explain this syntax in ["Matching on Sequences"](#).

Finally, `Nil` is an object representing an empty sequence with any type of elements.

## Reserved Words

Scala reserves some words for defining constructs like conditionals, declaring variables, etc. Some of the reserved words are marked with (*soft*), which means they can be used as regular identifiers for method and variable names, for example, but they are treated as keywords when used in particular contexts. All of the soft words are new reserved words in Scala 3. The reason for treating them as soft is to avoid breaking older code that happens to use them as identifiers.

**Table 2-1** lists the reserved keywords in Scala. Many are found in Java and they usually have the same meanings in both languages.

*Table 2-1. Reserved words*

Word	Description	See ...
abstract	Makes a declaration abstract.	“Class and Object Basics: Review”
as	(soft) Used with given.	“Context Bounds”
case	Start a case clause in a match expression. Define a “case class.”	Chapter 4
catch	Start a clause for catching thrown exceptions.	“Using try, catch, and finally Clauses”
class	Start a class declaration.	Chapter 9
def	Start a method declaration.	“Method Declarations”
do	New syntax for while and for loops without braces. Old Scala 2 do...while loop.	“Scala while Loops”
else	Start an else clause for an if clause.	“Scala Conditional Expressions”
extends	Indicates that the class or trait that follows is the parent type of the class or trait being declared.	

---

**W  
o  
r  
d**

**Description**

**See ...**

extension (soft) Marks one or more extension methods for a type.

[“Type Classes”](#)

false Boolean *false*.

[“Boolean Literals”](#)

final Applied to a class or trait to prohibit deriving child types from it. Applied to a member to prohibit overriding it in a derived class or trait.

[“Consider Making Concrete Members Final”](#)

finally Start a clause that is executed after the corresponding try clause, whether or not an exception is thrown by the try clause.

[“Using try, catch, and finally Clauses”](#)

for Start a for comprehension (loop).

[“Scala for Comprehensions”](#)

for Used in Scala 2 for *existential type* declarations to constrain the allowed concrete types that can be used. Dropped in Scala 3.

given Marks implicit definitions.

[Chapter 5](#)

if Start an if clause.

[“Scala Conditional Expressions”](#)

---

---

**W  
o  
r  
d**

Description	See ...
<div>im</div> <div>pl</div> <div>ic</div> <div>it</div> Marks a method or value as eligible to be used as an <i>implicit</i> type converter or value. Marks a method parameter as optional, as long as a type-compatible substitute object is in the scope where the method is called.	Chapter 5
<div>im</div> <div>po</div> <div>rt</div> Import one or more types or members of types into the current scope.	“Importing Types and Their Members”
<div>la</div> <div>zy</div> Defer evaluation of a val.	“lazy val”
<div>ma</div> <div>tc</div> <div>h</div> Start a pattern-matching clause.	Chapter 4
<div>ne</div> <div>w</div> Create a new instance of a class.	“Class and Object Basics: Review”
<div>nu</div> <div>ll</div> Value of a reference variable that has not been assigned a value.	Chapter 13
<div>ob</div> <div>je</div> <div>ct</div> Start a <i>singleton</i> declaration: a class with only one instance.	“A Taste of Scala”
<div>op</div> <div>aq</div> <div>ue</div> (soft) Declares a special type alias with zero runtime overhead.	“Opaque Types and Value Classes”
<div>op</div> <div>en</div> (soft) Declares a concrete class is open for subclassing.	“Open Versus Closed Types”

---

**W  
o  
r  
d**

	Description	See ...
override	Override a <i>concrete</i> member of a type, as long as the original is not marked <code>final</code> .	<a href="#">“Overriding Members of Classes and Traits”</a>
package	Start a package scope declaration.	<a href="#">“Organizing Code in Files and Namespaces”</a>
private	Restrict visibility of a declaration.	<i>Chapter 14</i>
protected	Restrict visibility of a declaration.	<i>Chapter 14</i>
deprecated	Deprecated. Was used for <i>self-typing</i> .	<i>Chapter 2</i>
return	Return from a function.	<a href="#">“A Taste of Scala”</a>
sealed	Applied to a parent type. Requires all derived types to be declared in the same source file.	<a href="#">“Sealed Class Hierarchies and Enumerations”</a>

---



---

**W  
o  
r  
d**

	<b>Description</b>	<b>See ...</b>
super	Analogous to <code>this</code> , but binds to the parent type.	“Defining Abstract Methods and Overriding Concrete Methods”
then	New syntax for <code>if</code> expressions	“Scala Conditional Expressions”
this	How an object refers to itself. The method name for <i>auxiliary constructors</i> .	“Constructors in Scala”
throw	Throw an exception.	“Using try, catch, and finally Clauses”
trait	A <i>mixin module</i> that adds additional state and behavior to an instance of a class. Can also be used to just declare methods, but not define them, like a Java interface.	Chapter 10
try	Start a block that may throw an exception.	“Using try, catch, and finally Clauses”
true	Boolean <i>true</i> .	“Boolean Literals”
type	Start a <i>type</i> declaration.	“Abstract Types Versus Parameterized Types”

---

## W o r d

Description	See ...
us (soft) Scala 3 alternative to <code>implicit</code> for <i>implicit arguments</i> . in g	Chapter 5
va Start a read-only “variable” declaration. l	“Variable Declarations”
va Start a read-write variable declaration. r	“Variable Declarations”
wh Start a <code>while</code> loop. il e	Later Chapter
wi Include the trait that follows in the class being declared th or the object being instantiated.	Chapter 10
yi Return an element in a <code>for</code> comprehension that el becomes part of a sequence. d	“Yielding New Values”
_ A placeholder, used in imports, function literals, etc.	Many
: Separator between identifiers and type annotations.	“A Taste of Scala”
= Assignment.	“A Taste of Scala”
=> Used in <i>function literals</i> to separate the parameter list from the function body.	“Anonymous Functions, Lambdas, and Closures”
<- Used in <code>for</code> comprehensions in <i>generator</i> expressions.	“Scala for Comprehensions”

Word	Description	See ...
<:	Used in <i>parameterized</i> and <i>abstract type</i> declarations to constrain the allowed types.	<i>later chapter</i>
<%	Used in <i>parameterized</i> and <i>abstract type</i> “view bounds” declarations.	<i>later chapter</i>
>:	Used in <i>parameterized</i> and <i>abstract type</i> declarations to constrain the allowed types.	<i>later chapter</i>
#	Used in <i>type projections</i> .	
@	Marks use of an <i>annotation</i> .	

Some Java methods use names that are reserved words by Scala, for example, `java.util.Scanner.match`. To avoid a compilation error, surround the name with single back quotes (“back ticks”), e.g., `java.util.Scanner.`match``.

## Literal Values

We’ve seen a few *literal values* already, such as `val book = "Programming Scala"`, where we initialized a `val book` with a `String` literal, and `(s: String) => s.toUpperCase`, an example of a function literal. Let’s discuss all the literals supported by Scala.

## Numeric Literals

Scala 3 expands the ways that numeric literals can be written and used as initializers. Consider these examples:

```
val i: Int = 123           // decimal
val x: Long = 0x123L      // hexadecimal (291 decimal)
val f: Float = 123_456.789F // 123456.789
```

```
val d: Double = 123_456_789.0123 // 123456789.0123
val y: BigInt = 0x123_a4b_c5d_e6f_789 // 82090347159025545
val z: BigDecimal = 123_456_789.0123 // 123456789.0123
```

Scala 3 allows underscores to make long numbers easier to read. They can appear anywhere in the literal (except between 0x), not just between every third character.

Hexadecimal numbers start with 0x followed by one or more digits and the letters a through f and A through F.

Indicate a negative number by prefixing the literal with a - sign.

The ability to use numeric literals for library and user-defined types like `BigInt` and `BigDecimal` is implemented with a new trait called `FromDigits`.

For Long literals, it is necessary to append the L or l character at the end of the literal, unless you are assigning the value to a variable declared to be Long. Otherwise, Int is inferred. The valid values for an integer literal are bounded by the type of the variable to which the value will be assigned. [Table 2-2](#) defines the limits, which are inclusive.

*Table 2-2. Ranges of allowed values for integer literals (boundaries are inclusive)*

Target type	Minimum (inclusive)	Maximum (inclusive)
Long	$-2^{63}$	$2^{63} - 1$
Int	$-2^{31}$	$2^{31} - 1$
Short	$-2^{15}$	$2^{15} - 1$
Char	0	$2^{16} - 1$
Byte	$-2^7$	$2^7 - 1$

A compile-time error occurs if an integer literal number is specified that is outside these ranges.

Floating-point literals are expressions with an optional minus sign, zero or more digits and underscores, followed by a period (.), followed by *one* or more digits. For `Float` literals, append the `F` or `f` character at the end of the literal. Otherwise, a `Double` is assumed. You can optionally append a `D` or `d` for a `Double`.

Floating-point literals can be expressed with or without exponentials. The format of the exponential part is `e` or `E`, followed by an optional `+` or `-`, followed by one or more digits.

Here are some example floating-point literals, where `Double` is inferred unless the declared variable is `Float` or an `f` or `F` suffix is used:

```
0.14
3.14
3.14f
3.14F
3.14d
3.14D
3e5
3E5
3.14e+5
3.14e-5
3.14e-5f
3.14e-5F
3.14e-5d
3.14e-5D
```

At least one digit must appear after the period, `3.` and `3.e5` are disallowed. Use `3.0` and `3.0e5` instead. Otherwise it would be ambiguous; do you mean method `e5` on the `Int` value of `3` or do you mean floating point literal `3.0e5`?

`Float` consists of all IEEE 754 32-bit, single-precision binary floating-point values. `Double` consists of all IEEE 754 64-bit, double-precision binary floating-point values.

## Boolean Literals

The Boolean literals are `true` and `false`. The type of the variable to which they are assigned will be inferred to be `Boolean`:

```
scala> val b1 = true
b1: Boolean = true

scala> val b2 = false
b2: Boolean = false
```

## Character Literals

A character literal is either a *printable* Unicode character or an escape sequence, written between single quotes. A character with a Unicode value between 0 and 255 may also be represented by an octal escape, i.e., a backslash (`\`) followed by a sequence of up to three octal characters. It is a compile-time error if a backslash character in a character or string literal does not start a valid escape sequence.

Here are some examples:

```
'A'
'\u0041' // 'A' in Unicode
'\n'
'\012'   // '\n' in octal
'\t'
```

The valid escape sequences are shown in [Table 2-3](#).

*Table 2-3. Character escape sequences*

Sequence	Meaning
\b	Backspace (BS)
\t	Horizontal tab (HT)
\n	Line feed (LF)
\f	Form feed (FF)
\r	Carriage return (CR)
\"	Double quote (")
\'	Single quote (')
\\	Backslash (\)

Note that *nonprintable* Unicode characters like \u0009 (tab) are not allowed. Use the equivalents like \t. Recall that three Unicode characters were mentioned in [Table 2-1](#) as valid replacements for corresponding ASCII sequences,  $\Rightarrow$  for  $=>$ ,  $\rightarrow$  for  $->$ , and  $\leftarrow$  for  $<-$ .

## String Literals

A string literal is a sequence of characters enclosed in double quotes or *triples* of double quotes, i.e., `"""..."""`.

For string literals in double quotes, the allowed characters are the same as the character literals. However, if a double quote (") character appears in the string, it must be “escaped” with a \ character. Here are some examples:

```
"Programming\nScala"  
"He exclaimed, \"Scala is great!\""  
"First\tSecond"
```

The string literals bounded by triples of double quotes are also called *multiline* string literals. These strings can cover several lines; the line feeds will be part of the string. They can include any characters, including one or two double quotes together, but not three together. They are useful for strings with `\` characters that don't form valid Unicode or escape sequences, like the valid sequences listed in [Table 2-3](#). Regular expressions are a good example, which use lots of escaped characters with special meanings. Conversely, if escape sequences appear, they aren't interpreted.

Here are three example strings:

```
"""Programming\nScala"""  
"""He exclaimed, "Scala is great!" """  
"""First line\nSecond line\t  
  
Fourth line"""
```

Note that we had to add a space before the trailing `"""` in the second example to prevent a parse error. Trying to escape the second `"` that ends the `"Scala is great!"` quote, i.e., `"Scala is great!\\"`, doesn't work.

When using multiline strings in code, you'll want to indent the substrings for proper code formatting, yet you probably don't want that extra whitespace in the actual string output. `String.stripMargin` solves this problem. It removes all whitespace in the substrings up to and including the first occurrence of a vertical bar, `|`. If you want some whitespace indentation, put the whitespace you want after the `|`. Consider this example:

```
// src/script/scala/progscala3/typelessdomore/MultilineStrings.scala  
  
def hello(name: String) = s"""Welcome!  
  Hello, $name!  
  * (Gratuitous Star!!)  
  |We're glad you're here.
```



```
| Have some extra whitespace.""".stripMargin  
  
val hi = hello("Programming Scala")
```

The last line prints the following:

```
val hi: String = Welcome!  
  Hello, Programming Scala!  
  * (Gratuitous Star!!)  
  We're glad you're here.  
  Have some extra whitespace.
```

Note where leading whitespace is removed and where it isn't.

If you want to use a different leading character than |, use the overloaded version of `stripMargin` that takes a `Char` (character) parameter. If the whole string has a prefix or suffix you want to remove (but not on individual lines), there are corresponding `stripPrefix` and `stripSuffix` methods, too:

```
scala> "<hello> <world>".stripPrefix("<").stripSuffix(">")  
val res0: String = hello> <world
```

The `<` and `>` inside the string are not removed.

## Symbol Literals

Scala 2 supported *symbols*, which are *interned* strings, meaning that two symbols with the same “name” (i.e., the same character sequence) will actually refer to the same object in memory. They are deprecated, but still exist in Scala 3. However, the literal syntax has been removed:

```
scala> val sym1 = 'name           // Scala 2 only; single "tick"  
scala> val sym2 = Symbol("name") // Scala 2 and 3
```

You might see symbols in older code, but don't use them yourself.

## Function Literals

As we've seen already, `(i: Int, d: Double) => (i+d).toString` is a function literal of type `Function2[Int,Double,String]`, where the last type is the return type.

You can even use the literal syntax for a type declaration. The following declarations are equivalent:

```
val f1: (Int,String) => String      = (i, s) => s+i
val f2: Function2[Int,String,String] = (i, s) => s+i
```

## Tuple Literals

Often, declaring a class to hold instances with two or more values is more than you need. You could put those values in a collection, but then you lose their specific type information. Scala implements tuples of values, where the individual types are retained. A literal syntax for tuples uses a comma-separated list of values surrounded by parentheses.

Here is an example of a tuple declaration and how we can access elements and use *pattern matching* to extract them:

```
// src/script/scala/progscala3/typelessdomore/TupleExample.scala
```

```
scala> val tup = ("Hello", 1, 2.3)           ❶
val tup: (String, Int, Double) = (Hello,1,2.3)
```

```
scala> val tup2: (String, Int, Double) = ("World", 4, 5.6)
val tup2: (String, Int, Double) = (World,4,5.6)
```

```
scala> tup._1                                ❷
val res0: String = Hello
```

```
scala> tup._2
val res1: Int = 1
```

```
scala> tup._3
val res2: Double = 2.3
```

```
scala> val (s, i, d) = tup
val s: String = Hello
val i: Int = 1
val d: Double = 2.3

scala> println(s"s = $s, i = $i, d = $d")
s = Hello, i = 1, d = 2.3
```

- ❶ Use the literal syntax to construct a three-element tuple. Note the literal syntax is used for the type, too.
- ❷ Extract the first element of the tuple. Tuple indexing is one-based, by historical convention, not zero-based. The next two lines extract the second and third elements.
- ❸ Declare three values, *s*, *i*, and *d*, that are assigned the three corresponding fields from the tuple using pattern matching.

Two-element tuples, sometimes called *pairs* for short, are so commonly used there is a special syntax for constructing them:

```
scala> (1, "one")
val res3: (Int, String) = (1,one)

scala> 1 -> "one"
val res4: (Int, String) = (1,one)

scala> Tuple2(1, "one") // Rarely used.
val res5: (Int, String) = (1,one)
```

For example, maps are often constructed with key-value pairs as follows:

```
//
src/script/scala/progscala3/typelessdomore/StateCapitalsSubset.scala

scala> val stateCapitals = Map(
  |   "Alabama" -> "Montgomery",
```

```

|     "Alaska"  -> "Juneau",
|     // ...
|     "Wyoming" -> "Cheyenne")
val stateCapitals: Map[String, String] =
  Map(Alabama -> Montgomery, Alaska -> Juneau, Wyoming -> Cheyenne)

```

## Option, Some, and None: Avoiding nulls

Let's discuss three useful types that express a very useful concept, when we may or may not have a value.

Most languages have a special keyword or type instance that's assigned to reference variables when there's nothing else for them to refer to. In Scala and Java, it's called `null`.

Using `null` is a giant source of nasty bugs. What `null` really signals is that we don't have a value in a given situation. If the value is not `null`, we do have a value. Why not express this situation explicitly with the type system and exploit type checking to avoid

`NullPointerException`?

`Option` lets us express this situation explicitly without using `null`. `Option` is an abstract class and its two concrete subclasses are `Some`, for when we have a value, and `None`, when we don't.

You can see `Option`, `Some`, and `None` in action using the map of state capitals in the United States that we declared in the previous section:

```

scala> stateCapitals.get("Alabama")
| stateCapitals.get("Wyoming")
| stateCapitals.get("Unknown")
val res6: Option[String] = Some(Montgomery)
val res7: Option[String] = Some(Cheyenne)
val res8: Option[String] = None

scala> stateCapitals.getOrElse("Alabama", "Oops1")
| stateCapitals.getOrElse("Wyoming", "Oops2")
| stateCapitals.getOrElse("Unknown", "Oops3")
val res9: String = Montgomery
val res10: String = Cheyenne
val res11: String = Oops3

```

`Map.get` returns an `Option[T]`, where `T` is `String` in this case. Either a `Some` wrapping the value is returned or a `None` when no value for the specified key was found.

In contrast, similar methods in other languages just return a `T` value, when found, or `null`.

By returning an `Option`, we can't "forget" that we have to verify that something was returned. In other words, the fact that a value may not exist for a given key is enshrined in the return type for the method declaration. This also provides clear documentation for the user of `Map.get` about what can be returned.

The second group uses `Map.getOrElse`. This method returns either the value found for the key or it returns the second argument passed in, which functions as the default value to return.

So, `getOrElse` is more convenient, as you don't need to process the `Option`, if a suitable default value exists.

To reiterate, because the `Map.get` method returns an `Option`, it automatically documents for the reader that there may not be an item matching the specified key. The map handles this situation by returning a `None`.

Also, thanks to Scala's static typing, you can't make the mistake of "forgetting" that an `Option` is returned and attempting to call a method supported by the type of the value inside the `Option`. You must extract the value first or handle the `None` case. Without an option return type, when a method just returns a value, it's easy to forget to check for `null` before calling methods on the returned object.

## When You Can't Avoid Nulls

Because Scala runs on the JVM, JavaScript, and native environments, it must interoperate with other libraries, which means Scala has to support `null`.

Scala 3 introduces a new way to indicate a possible `null` through the type `Scala.Null`, which is a subtype of all other types. Suppose you have a **Java HashMap** to access:

```
// src/script/scala/progscala3/typelessdomore/Null.scala

import java.util.{HashMap => JHashMap} ❶

val jhm = new JHashMap[String,String]()
jhm.put("one", "1")

val one1: String = jhm.get("one") ❷
val one2: String | Null = jhm.get("one") ❸

val two1: String = jhm.get("two") ❹
val two2: String | Null = jhm.get("two")
```

- ❶ Import the Java `HashMap`, but give it an alias so it doesn't shadow Scala's `HashMap`.
- ❷ Return the string "1".
- ❸ Declare explicitly that `one2` is of type `String` or `Null`. The value will still be "1" in this case.
- ❹ These two values will equal `null`.

Scala 3 introduces *union types*, which we use for `one2` and `two2`. They tell the reader that the value could be either a `String` or a `null`.

There is an optional feature to enable aggressive null checking. If you add the flag `-Yexplicit-nulls`, then the declarations of `one1` and `two1` will be disallowed, because the compiler knows you are referring to a Java library where `null` could be returned. I have not enabled this option in the code examples build, because it forces lots of changes to otherwise safe code and it has other implications that are described in more detail in the [documentation](#). However, if you use this same code in a REPL with this flag, you'll see the following:

```

$ scala -Yexplicit-nulls
...
scala> val one1: String = jhm.get("one")
1 | val one1: String = jhm.get("one")
    |                      ^^^^^^^^^^^^^^^^^^
    |                      Found:    String | UncheckedNull
    |                      Required: String
...

```

Tony Hoare, who invented the null reference in 1965 while working on a language called ALGOL W, called its invention his “billion dollar” mistake. Use `Option` instead. Consider enabling explicit nulls.

## Sealed Class Hierarchies and Enumerations

While we’re discussing `Option`, let’s discuss a useful design feature it uses. A key point about `Option` is that there are really only two valid subtypes. Either we have a value, the `Some` case, or we don’t, the `None` case. There are *no* other subtypes of `Option` that would be valid. So, we would really like to prevent users from creating their own.

Scala 2 and 3 have a keyword `sealed` for this purpose. `Option` could be declared as follows:

```

sealed abstract class Option[+A] {...}
final case class Some[+A](a: A) extends Option[A] {...}
final case object None extends Option[Nothing] {...}

```

The `sealed` keyword tells the compiler that all subclasses must be declared *in the same source file*. `Some` and `None` are declared in the same file with `Option` in the Scala library. This technique effectively prevents additional subtypes of `Option`.

`None` has an interesting declaration. It is a case class with only one instance, so it is declared `case object`. The `Nothing` type along with

the `Null` type are *subtypes* of all other types in Scala. We'll explain Nothing in more detail in “Sequences”, if you get what I mean...

You can also declare a type `final` if you want to prevent users from subtyping it. If you want to encourage subtyping of a concrete class, you can declare it `open`. (This is redundant for abstract types.)

This same constraint on subclassing can now be achieved more concise in Scala 3 with the new enum syntax:

```
enum Option[+A] {  
  case Some(a: A) {...}  
  case None {...}  
  ...  
}
```

## Organizing Code in Files and Namespaces

Scala adopts the package concept that Java uses for namespaces, but Scala offers more flexibility. Filenames don't have to match the type names and the package structure does not have to match the directory structure. So, you can define packages in files independent of their “physical” location.

The following example defines a class `MyClass` in a package `com.example.mypkg` using the conventional Java syntax:

```
// src/main/scala/progscala3/typelessdomore/PackageExample1.scala  
package com.example.mypkg  
  
class MyClass:  
  def mymethod(s: String): String = s
```

Scala also supports a block-structured syntax for declaring package scope:

```
// src/main/scala/progscala3/typelessdomore/PackageExample2.scala  
package com:  
  package example:
```



```

package pkg1:
  class Class11:
    def m = "m11"

  class Class12:
    def m = "m12"

package pkg2:
  class Class21:
    def m = "m21"
    def makeClass11 = new pkg1.Class11

    def makeClass12 = new pkg1.Class12

package pkg3.pkg31.pkg311:
  class Class311:
    def m = "m21"

```

Two packages, pkg1 and pkg2, are defined under the com.example package. A total of three classes are defined between the two packages. The makeClass11 and makeClass12 methods in Class21 illustrate how to reference a type in the “sibling” package, pkg1. You can also reference these classes by their full paths, com.example.pkg1.Class11 and com.example.pkg1.Class12, respectively.

The package pkg3.pkg31.pkg311 shows that you can “chain” several packages together in one statement. It is not necessary to use a separate package statement for each package.

If you have package-level declarations, like types, in each of several parent packages that you want to bring into scope, use separate package statements as shown for each level of the package hierarchy with these declarations. Each subsequent package statement is interpreted as a subpackage of the previously specified package, as if we used the block-structure syntax shown previously. The first statement is interpreted as an absolute path.

Following the convention used by Java, the root package for Scala’s library classes is named scala.

Although the package declaration syntax is flexible, one limitation is that packages cannot be defined within classes and objects, which wouldn't make much sense anyway.

### WARNING

Scala does not allow package declarations in scripts, which are implicitly wrapped in an object, where package declarations are not permitted.

## Importing Types and Their Members

To use declarations in packages, you have to import them. However, Scala offers flexible options how items are imported:

```
import java.awt._           ❶  
import java.io.File         ❷  
import java.io.File._       ❸  
import java.util.{Map, HashMap} ❹
```

- ❶ Import everything in a package, using underscore ( `_` ) as a wildcard.
- ❷ Import an individual type.
- ❸ Import all static member fields and methods in `File`.
- ❹ Selectively import two types from `java.util`.

Java uses the asterisk character ( `*` ) as the wildcard for imports. In Scala, this character is allowed as a method name (e.g., for multiplication), so `_` is used instead to avoid ambiguity.

The third line imports all the static methods and fields in `java.io.File`. The equivalent Java import statement would be `import static java.io.File.*`; Scala doesn't have an `import`

static construct because it treats object types uniformly like other types.

You can put import statements almost anywhere, so you can scope their visibility to just where they are needed, you can rename types as you import them, and you can suppress the visibility of unwanted types:

```
def stuffWithBigInteger() = {  
  
  import java.math.BigInteger.{  
    ONE => _,           ❶  
    TEN,                ❷  
    ZERO => JAVAZERO }  ❸  
  
  // println( "ONE: "+ONE )    // ONE is effectively undefined  
  println( "TEN: "+TEN )  
  println( "ZERO: "+JAVAZERO )  
}
```

- ❶ Alias to `_` to make it invisible. Use this technique when you want to import everything except a few items.
- ❷ Import `TEN` from `BigDecimal`. It can be referenced simply as `TEN`.
- ❸ Import `ZERO` but give it an alias. Use this technique to avoid shadowing other items with the same name. This is used a lot when mixing Java and Scala types, such as Java's `List` and Scala's `List`.

Because this import statement is inside `stuffWithBigInteger`, the imported items are not visible outside the function.

Finally, Scala 3 adds new ways to control how implicit definitions are imported. We'll discuss these details in "[Givens and Imports](#)", once we understand the new syntax and behaviors for *given instances*.

## Package Imports and Package Objects

Sometimes it's nice to give the user one import statement for a public API that brings in all types, as well as constants and methods not attached to a type. For example:

```
import progscale3.typelessdomore.api._
```

This is simple to do; just define anything you need under the package:

```
// src/main/scala/progscale3/typelessdomore/PackageObjects.scala
package progscale3.typelessdomore.api

val DEFAULT_COUNT = 5
def countTo(limit: Int = DEFAULT_COUNT) = (0 to
limit).foreach(println)

class Class1:
  def m = "cm1"

object Object1:
  def m = "om1"
```

In Scala 2, non-type definitions had to be declared inside a *package object*, like this:

```
// src/main/scala-2/progscale3/typelessdomore/PackageObjects.scala
package progscale3.typelessdomore // Notice, no ".api"

package object api {
  val DEFAULT_COUNT = 5
  def countTo(limit: Int = DEFAULT_COUNT) = (0 to
limit).foreach(println)

  class Class1 {
    def m = "cm1"
  }

  object Object1 {
    def m = "om1"
  }
}
```

```
}  
}
```

Package objects are still supported in Scala 3, but they are deprecated.

## Abstract Types Versus Parameterized Types

We mentioned in “[A Taste of Scala](#)” that Scala supports *parameterized types*, which are very similar to *generics* in Java. (The terms are somewhat interchangeable, but the Scala community uses parameterized types.) Java uses angle brackets (<...>), while Scala uses square brackets ([...]), because < and > are often used for method names.

Because we can plug in almost any type for a type parameter A in a collection like `List[A]`, this feature is called *parametric polymorphism*, because generic implementations of the `List` methods can be used with instances of any type A.

Consider the declaration of `Map`, which is written as follows, where K is the keys type and V is the values type.

```
trait Map[K, +V] extends Iterable[(K, V)] with ...
```

The + in front of the V means that `Map[K, V2]` is a *subtype* of `Map[K, V1]` for any V2 that is a *subtype* of V1. This is called *covariant typing*. It is a reasonably intuitive idea. If we have a function `f(map: Map[String, Any])`, it makes sense that passing a `Map[String, Double]` to it should work fine, because the function has to assume values of `Any`, a parent type of `Double`.

In contrast, the key K is *invariant*. We can't pass `Map[Any, Any]` to `f` nor any `Map[S, Any]` for some subtype or supertype S of `String`.

If there is a - in front of a type parameter, the relationship goes the other way; `Foo[B]` would be a *supertype* of `Foo[A]`, if B is a *subtype*

of A and the declaration is `Foo[-A]` (called *contravariant typing*). This is less intuitive, but also not as important to understand now. We'll see how it is important for function types in

Scala supports another type abstraction mechanism called *abstract types*, which can be applied to many of the same design problems for which parameterized types are used. However, while the two mechanisms overlap, they are not redundant. Each has strengths and weaknesses for certain design problems.

These types are declared as members of other types, just like methods and fields. Here is an example that uses an abstract type in a parent class, then makes the type member concrete in child classes:

```
// src/main/scala/progscala3/typelessdomore/AbstractTypes.scala
package progscala3.typelessdomore
import scala.io.Source

abstract class BulkReader:
  type In
  val source: In
  /** Read source and return a sequence of Strings */
  def read: Seq[String]

case class StringBulkReader(source: String) extends BulkReader:
  type In = String
  def read: Seq[String] = Seq(source)

case class FileBulkReader(source: Source) extends BulkReader:
  type In = Source
  def read: Seq[String] = source.getLines.toVector
```

- ❶ Abstract type, really just like any abstract field or method.
- ❷ Concrete subtype of `BulkReader` where `In` is defined to be `String`. Note that the type of the `source` parameter must match.

- ③ Concrete subtype of `BulkReader` where `In` is defined to be `Source`, the Scala library class for reading sources, like files. `Source.getLines` returns an iterator, which we can easily read into a `Vector` with `toVector`.

Strictly speaking, we don't need to declare the `source` field in the parent class, but I put it there to show you that the concrete case classes can make it a constructor parameter, where the specific type is specified.

Using these readers:

```
scala> import progscala3.typelessdomore.{StringBulkReader,
FileBulkReader}

scala> new StringBulkReader("Hello Scala!").read
val res1: Seq[String] = List(Hello Scala!)

scala> val lines = FileBulkReader(Source.fromFile("README.md")).read
val lines: Seq[String] = Vector(# Programming Scala, 3rd Edition,
...)

scala> lines(0)    // look at two lines...
      | lines(2)
val res2: String = # Programming Scala, 3rd Edition
val res3: String = ## README for the Code Examples
```

The type field is used like a type parameter in a parameterized type. In fact, as an exercise, try rewriting the example to use type parameters, e.g., `BulkReader[In]`.

So what's the advantage here of using type members instead of parameterized types? The latter are best for the case where the type parameter has no relationship with the parameterized type, like `List[A]` when `A` is `Int`, `String`, `Person`, etc. A type member works best when it "evolves" in parallel with the enclosing type, as in our `BulkReader` example, where the type member needed to match the "behaviors" expressed by the enclosing type, specifically the `read`

method. Sometimes this characteristic is called *family polymorphism* or *covariant specialization*.

For completeness, another use for type members is to provide a convenient alias for a more complicated type. For example, suppose you use (String, Double) tuples a lot in some code. You could either declare a class for it or use a type alias for a simple alternative:

```
scala> object Foo:
  |   type Record = (String, Double)
  |   def transform(record: Record): Record =
  |     (record._1.toUpperCase, 2*record._2)
  // defined object Foo

scala> Foo.transform("hello", 10)
val res0: Foo.Record = (HELLO,20.0)
```

Notice the type shown for res0. In fact, concrete type members are always type aliases.

## Recap and What's Next

We covered a lot of practical ground, such as literals, keywords, file organization, and imports. We learned how to declare variables, methods, and classes. We learned about Option as a better tool than null, plus other useful techniques. In the next chapter, we will finish our fast tour of the Scala “basics” before we dive into more detailed explanations of Scala’s features.



# Chapter 3. Rounding Out the Basics

---

## A NOTE FOR EARLY RELEASE READERS

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

This will be the 3rd chapter of the final book. Please note that the GitHub repo will be made active later on.

If you have comments about how we might improve the content and/or examples in this book, or if you notice missing material within this chapter, please reach out to the author at [mcronin@oreilly.com](mailto:mcronin@oreilly.com).

Let’s finish our survey of essential “basics” in Scala.

## Operator Overloading?

Almost all “operators” are actually methods. Consider this most basic of examples:

`1 + 2`

The plus sign between the numbers is a method on the `Int` type.

Scala doesn’t have special “primitives” for numbers and booleans that are distinct from types you define. They are regular types: `Float`, `Double`, `Int`, `Long`, `Short`, `Byte`, `Char`, and `Boolean`. Hence, they can

have methods. However, Scala will compile to native platform primitives when possible for greater efficiency.

Also we've already seen that Scala identifiers can have most nonalphanumeric characters and single-parameter methods can be invoked with *infix operator notation*. So, `1 + 2` is the same as `1.+(2)`.

### CAUTION

Actually, they don't always behave identically, due to operator precedence rules. While `1 + 2 * 3 = 7`, `1.+(2)*3 = 9`. When present, the period binds before the star.

In fact, you aren't limited to "operator-like" method names when using infix notation. It is common to write code like the following, where the dot and parentheses around `println` are omitted:

```
scala> Seq("one", "two") foreach println
one
two
```

Use this convention cautiously, as these expressions can sometimes be confusing to read or parse.

To bring a little more discipline to this practice, Scala 3 now issue a deprecation warning if a one-parameter method is declared *without* the annotation `@infix`, but is used with infix notation. However, to support traditional practice, most of the familiar collection "operators", like `map`, `foreach`, etc. are declared with `@infix`, so our `foreach` example still works in Scala 3 without a warning.

If you still want to use infix notation and you want to avoid the deprecation warning for a method that isn't annotated with `@infix`, put the name in "back ticks":

```
// src/script/scala/progscala3/rounding/InfixMethod.scala
```

```
case class Foo(str: String):
  def append(s: String): Foo = copy(str + s)
```

```
Foo("one").append("two") ❶
```

```
Foo("one") append "two" ❷
```

```
Foo("one") `append` "two" ❸
```

- ❶ Normal usage.
- ❷ Triggers a deprecation warning.
- ❸ Accepted usage, but odd looking.

You can also define your own operator methods with symbolic names. Suppose you want to allow users to create `java.io.File` objects by appending strings using `/`, the file separator for UNIX-derived systems. Consider the following implementation:

```
// src/main/scala/progscala3/rounding/Path.scala
package progscala3.rounding

import scala.annotation.alpha
import java.io.File

case class Path(
  value: String, separator: String = Path.defaultSeparator): ❶
  val file = new File(value)
  override def toString: String = file.getPath ❷

  @alpha("concat") def / (node: String): Path = ❸
    copy(value + separator + node) ❹

object Path:
  val defaultSeparator = sys.props("file.separator")
```

- ❶ Use the operating systems default path separator string as the default separator when constructing a `File` object.

❷

How to override the default toString method. Here, I use the path string from File.

- ③ I'll explain the @alpha in a moment.
- ④ Use the case class copy method to create a new instance, changing only the value.

Now users can create new File objects as follows:

```
scala> import progscale3.rounding.Path

scala> val one = Path("one")
val one: progscale3.rounding.Path = one

scala> val three = one / "two" / "three"
val three: progscale3.rounding.Path = one/two/three

scala> three.file
val res0: java.io.File = one/two/three

scala> val threeb = one./("two")./("three")
val threeb: progscale3.rounding.Path = one/two/three

scala> three == threeb
val res1: Boolean = true

scala> one `concat` "two"
1 |one `concat` "two"
  |^^^^^^^^^^^^^^
  |value concat is not a member of progscale3.rounding.Path
```

On Windows, the character \ would be used as the default separator. This method is designed to be used with infix notation. It looks odd to use normal invocation syntax.

In Scala 3, the @alpha annotation is recommended for methods with symbolic names. It will be required in a future release of Scala 3. In this example, concat is the name the compiler will use internally when it generates byte code.

This is the name you would use if you wanted to call the method from Java code, which doesn't support invoking methods with symbolic names. However, the name `concat` can't be used in Scala code, as shown at the end of the example. It only affects the byte code produced by the compiler.

The `@infix` annotation is *not* required for methods that use “operator” characters, like `*` and `/`, because support for symbolic operators has always existed for the particular purpose of allowing intuitive, infix expressions, like `a * b` and `path1 / path2`.

Types can also be written with infix notation, in many contexts. The same rules using the `@infix` annotation apply:

```
// src/script/scala/progscala3/rounding/InfixType.scala
import scala.annotation.{alpha, infix}
```

```
@alpha("BangBang") case class !![A,B](a: A, b: B)           ❶
val ab1: Int !! String = 1 !! "one"                         ❷
val ab2: Int !! String = !(1, "one")                        ❸

@infix case class bangbang[A,B](a: A, b: B)                 ❹
val ab1: Int bangbang String = 1 bangbang "one"
val ab2: Int bangbang String = bangbang(1, "one")
```

- ❶ Some type declaration with two type parameters.
- ❷ An attempt to use infix notation on both sides, but we get an error that `!!` is not a method on `Int`. We'll solve this problem in [Chapter 5](#).
- ❸ This declaration works, with the non-infix notation on the right-hand side.
- ❹ These three lines behave the same, but note the use of `@infix` now.

To recap:

- Annotate symbolic operator definitions with `@alpha("some_name")`.
- Annotate alphanumeric types and methods with `@infix` if you want to allow their use with infix notation.

Related to infix notation is postfix notation, where a method with no parameters can be invoked without the period. However, postfix invocations can sometimes be even more ambiguous and confusing, so Scala requires that you explicitly enable this feature first, in one of several ways:

1. Add the import statement `import scala.language.postfixOps` to the source code.
2. Invoke the compiler with the option `-language:postfixOps` “feature flag”.

There is no `@postfix` annotation.

### TIP

The SBT build for the examples is configured to use the `-feature` flag that enables warnings when features are used, like postfix expressions, that should be enabled explicitly.

Dropping the punctuation by using infix or even postfix expresses can make your code cleaner and help create elegant programs that read more naturally, but avoid cases that are actually harder to understand.

## Allowed Characters in Identifiers

Here is a summary of the rules for characters in identifiers for method and type names, variables, etc:

## Characters

Scala allows all the printable ASCII characters, such as letters, digits, the underscore ( `_` ), and the dollar sign ( `$` ), with the exceptions of the “parenthetical” characters, ( `,` ), [ `,` ], { `,` }, and the “delimiter” characters, ```, `'`, `'`, `"`, `.`, `;`, and `,`. Scala allows the characters between `\u0020` and `\u007F` that are not in the sets just shown, such as mathematical symbols, the so-called *operator characters* like `/` and `<`, and some other symbols. This includes white space characters, discussed below.

### *Reserved words can't be used*

We listed the reserved words in “**Reserved Words**”. Recall that some of them are combinations of operator and punctuation characters. For example, a single underscore ( `_` ) is a reserved word!

### *Plain identifiers—combinations of letters, digits, \$, \_, and operators*

A *plain identifier* can begin with a letter or underscore, followed by more letters, digits, underscores, and dollar signs. Unicode-equivalent characters are also allowed. Scala reserves the dollar sign for internal use, so you shouldn't use it in your own identifiers, although this isn't prevented by the compiler. After an underscore, you can have either letters and digits, *or* a sequence of operator characters. The underscore is important. It tells the compiler to treat all the characters up to the next whitespace as part of the identifier. For example, `val xyz_++= 1` assigns the variable `xyz_++=` the value 1, while the expression `val xyz++= 1` won't compile because the “identifier” could also be interpreted as `xyz ++=`, which looks like an attempt to append something to `xyz`. Similarly, if you have operator characters after the underscore, you can't mix them with letters and digits. This restriction prevents ambiguous expressions like this: `abc_-123`. Is

that an identifier `abc_-123` or an attempt to subtract 123 from `abc_`?

### *Plain identifiers—operators*

If an identifier begins with an operator character, the rest of the characters must be operator characters.

### *“Back-quote” literals*

An identifier can also be an arbitrary string between two back quote characters, e.g., `def `test that addition works` = assert(1 + 1 == 2)`. (Using this trick for literate test names is the one use I can think of for this otherwise-questionable technique for using white space in identifiers.) Also use back quotes to invoke a method or variable in a Java class when the name is identical to a Scala reserved word, e.g., `java.net.Proxy.`type`()`.

### *Pattern-matching identifiers*

In pattern-matching expressions (for example, “**A Sample Application**”), tokens that begin with a lowercase letter are parsed as *variable identifiers*, while tokens that begin with an uppercase letter are parsed as *constant identifiers* (such as class names). This restriction prevents some ambiguities because of the very succinct variable syntax that is used, e.g., no `val` keyword is present.

## **Syntactic Sugar**

Once you know that all operators are methods, it’s easier to reason about unfamiliar Scala code. You don’t have to worry about special cases when you see new operators. We’ve seen several examples where infix expressions like `matrix1 * matrix2` were used, which are actually just ordinary method invocations.



This flexible method naming gives you the power to write libraries that feel like a natural extension of Scala itself. You can write a new math library with numeric types that accept all the usual mathematical operators. The possibilities are constrained by just a few limitations for method names.

### CAUTION

Avoid making up operator symbols when an established ASCII name exists, because the latter is easier to understand and remember, especially for beginners reading your code.

## Methods with Empty Parameter Lists

Scala is flexible about the use of parentheses in methods with no parameters.

If a method takes no parameters, you can define it without parentheses. Callers must invoke the method without parentheses. Conversely, if you add empty parentheses to your definition, callers must add the parentheses.

For example, `List.size` has no parentheses, so you write `List(1, 2, 3).size`. If you try `List(1, 2, 3).size()`, you'll get an error.

However, exceptions are made for no-parameter methods in Java libraries. For example, the `length` method for `java.lang.String` does have parentheses in its definition, because Java requires them, but Scala lets you write both `"hello".length()` and `"hello".length`. This flexibility with Scala-defined methods was also allowed in earlier releases of Scala, but Scala 3 now requires that usage match the definition.

A convention in the Scala community is to omit parentheses for no-parameter methods that have no side effects, like returning the size of a collection, which could actually be a precomputed, immutable

field in the object. So, many no-parameter methods could be interpreted as simply reading a field. However, when the method performs side effects or does extensive work, parentheses are added by convention, providing a hint to the reader of nontrivial activity, for example `myFileReader.readLine()`.

Why bother with optional parentheses in the first place? They make some method call chains read better as expressive, self-explanatory “sentences” of code:

```
// src/script/scala/progscala3/rounding/NoDotBetter.scala  
  
def isEven(n: Int) = (n % 2) == 0  
  
Seq(1, 2, 3, 4) filter isEven foreach println
```

The second line is “clean”, but on the edge of non-obvious to read. Here is the same code repeated four times with progressively less explicit detail filled in. The last line is the original:

```
Seq(1, 2, 3, 4).filter((i: Int) => isEven(i)).foreach((i: Int) =>  
println(i))  
Seq(1, 2, 3, 4).filter(i => isEven(i)).foreach(i => println(i))  
Seq(1, 2, 3, 4).filter(isEven).foreach(println)  
Seq(1, 2, 3, 4) filter isEven foreach println
```

The first three versions are more explicit and hence better for the beginning reader to understand. However, once you’re familiar with the fact that `filter` is a method on collections that takes a single argument, `foreach` loops over a collection, and so forth, the last, “Spartan” implementation is much faster to read and understand. The other two versions have more visual noise that just get in the way, once you’re more experienced. Keep that in mind as you learn to read Scala code.

To be clear, this expression works because each method we used took a single parameter. If you tried to use a method in the chain that

takes zero or more than one parameter, it would confuse the compiler. In those cases, put some or all of the punctuation back in.

The previous explanation glossed over an important detail. The four versions are *not exactly* the same code with different details inferred by the compiler, although all four behave the same way. Consider the function arguments passed to `filter` in the second and third examples. They are different as follows:

- `filter(i => isEven(i))` passes an *anonymous* function that *calls* `isEven`.
- `filter(isEven)` passes `isEven` itself as the function. Note that `isEven` is a *named* function.

The same distinction applies to the functions passed to `foreach` in the second and third examples.

## Precedence Rules

So, if an expression like `2.0 * 4.0 / 3.0 * 5.0` is actually a series of method calls on `Doubles`, what are the *operator precedence* rules? Here they are in order from lowest to highest precedence:

1. *All letters*
2. `|`
3. `^`
4. `&`
5. `<` `>`
6. `=` `!`
7. `:`
8. `+` `-`

9. \* / %

10. *All other special characters*

Characters on the same line have the same precedence. An exception is = when it's used for assignment, in which case it has the lowest precedence.

Because \* and / have the same precedence, the two lines in the following scala session behave the same:

```
scala> 2.0 * 4.0 / 3.0 * 5.0
res0: Double = 13.333333333333332
```

```
scala> (((2.0 * 4.0) / 3.0) * 5.0)
res1: Double = 13.333333333333332
```

## Left vs. Right Associative Methods

Usually, method invocation using infix operator notation simply bind in left-to-right order, i.e., they are *left-associative*. Don't all methods work this way? No. In Scala, any method with a name that ends with a colon (:) binds to the *right* when used in infix notation, while all other methods bind to the *left*. For example, you can prepend an element to a Seq using the +: method (sometimes called “cons,” which is short for “constructor,” a term introduced by Lisp):

```
scala> val seq = Seq('b', 'c', 'd')
val seq: Seq[Char] = List(b, c, d)
```

```
scala> val seq2 = 'a' +: seq
val seq2: Seq[Char] = List(a, b, c, d)
```

```
scala> val seq3 = 'z'.+: (seq2)
1 | val seq3 = 'z'.+: (seq2)
  |           ^^^^^^
  |           value +: is not a member of Char
```

```
scala> val seq3 = seq2.+: ('z')
val seq3: Seq[Char] = List(z, a, b, c, d)
```

Note that if we don't use infix notation, we have to put `seq2` on the *left*.

### TIP

Any method whose name ends with a `:` binds to the *right*, not the *left*, in infix operator notation.

## Enumerations and Algebraic Data Types

While it's common to declare a type hierarchy to represent all the possible “kinds” of some abstraction, sometimes we know the list of them is fixed and mostly what we need are unique “flags” to indicate each one.

Take for example the days of the week, where we have seven fixed values. An enumeration for the English days of the week can be declared as follows:

```
// src/script/scala/progscala3/rounding/DaysEnumeration.scala

// tag::weekday[]
enum WeekDay(val fullName: String):                                ❶
  case Sun extends WeekDay("Sunday")                               ❷
  case Mon extends WeekDay("Monday")
  case Tue extends WeekDay("Tuesday")
  case Wed extends WeekDay("Wednesday")
  case Thu extends WeekDay("Thursday")
  case Fri extends WeekDay("Friday")
  case Sat extends WeekDay("Saturday")

  def isWorkingDay: Boolean = ! (this == Sat || this == Sun)      ❸

import WeekDay._

val sorted = WeekDay.values.sortBy(_.ordinal).toSeq               ❹
assert(sorted == List(Sun, Mon, Tue, Wed, Thu, Fri, Sat))

assert(Sun.fullName == "Sunday")
```

```

assert(Sun.ordinal == 0)
assert(Sun.isWorkingDay == false)
assert(WeekDay.valueOf("Sun") == WeekDay.Sun)
// end::weekday[]

```

5

6

```

assert(Mon.fullName == "Monday")
assert(Mon.ordinal == 1)
assert(Mon.isWorkingDay == true)
assert(WeekDay.valueOf("Mon") == Mon)

```

```

assert(Tue.fullName == "Tuesday")
assert(Tue.ordinal == 2)
assert(Tue.isWorkingDay == true)
assert(WeekDay.valueOf("Tue") == Tue)

```

```

assert(Wed.fullName == "Wednesday")
assert(Wed.ordinal == 3)
assert(Wed.isWorkingDay == true)
assert(WeekDay.valueOf("Wed") == Wed)

```

```

assert(Thu.fullName == "Thursday")
assert(Thu.ordinal == 4)
assert(Thu.isWorkingDay == true)
assert(WeekDay.valueOf("Thu") == Thu)

```

```

assert(Fri.fullName == "Friday")
assert(Fri.ordinal == 5)
assert(Fri.isWorkingDay == true)
assert(WeekDay.valueOf("Fri") == Fri)

```

```

assert(Sat.fullName == "Saturday")
assert(Sat.ordinal == 6)
assert(Sat.isWorkingDay == false)
assert(WeekDay.valueOf("Sat") == Sat)

```

- 1 Declare an enumeration, similar to declaring a class. You can have optional fields as shown. Declare them with `val` if you want them to be accessible, e.g., `WeekDay.Sun.fullName`. The `derives Eq` clause lets us do comparisons with `==` and `!=`. We'll explain this construct in “[Type Class Derivation](#)” and “[Multiversal Equality](#)”.

- ② The values are declared using the case keyword.
- ③ You can define methods.
- ④ The `WeekDay.values` order does not match the declaration order, so we sort by the ordinal.
- ⑤ The ordinal value matches the declaration order.
- ⑥ You can lookup an enumeration value by its name.

Scala 3 introduced a new syntax for enumerations, which we saw briefly in “**Sealed Class Hierarchies and Enumerations**”.<sup>1</sup> The new syntax lends itself to a more concise definition of *algebraic data types* (ADTs - not to be confused with *abstract data types*). An ADT is “algebraic” in the sense that transformations obey well defined properties (think of addition with integers as an example), such as transforming an element or combining two of them with an operation can only yield another element in the set.

Consider the following example:

```
// src/script/scala/progscala3/rounding/TreeADT.scala
```

```
object Scala2ADT:
  sealed trait Tree[T]
  final case class Branch[T](
    left: Tree[T], right: Tree[T]) extends Tree[T]
  final case class Leaf[T](elem: T) extends Tree[T]

  val tree = Branch(
    Branch(
      Leaf(1),
      Leaf(2)),
    Branch(
      Leaf(3),
      Branch(Leaf(4), Leaf(5))))

object Scala3ADT:
```

```

enum Tree[T]:
  case Branch(left: Tree[T], right: Tree[T])
  case Leaf(elem: T)

import Tree._
val tree = Branch(
  Branch(
    Leaf(1),
    Leaf(2)),
  Branch(
    Leaf(3),
    Branch(Leaf(4), Leaf(5))))

Scala2ADT.tree
Scala3ADT.tree

```

- ❶ Use a sealed type hierarchy. Valid for Scala 2 and 3.
- ❷ One subtype, a branch with left and right children.
- ❸ The other subtype, a leaf node.
- ❹ Scala 3 syntax using the new enum construct. It is much more concise.
- ❺ The elements of the enum need to be imported.
- ❻ Is the output the same for these two lines?

We saw sealed type hierarchies before in “[Sealed Class Hierarchies and Enumerations](#)”. The enum syntax provides the same benefits as sealed type hierarchies, but with much less code.

The types of the tree values are slightly different:

```

scala> Scala2ADT.tree
val res1: Scala2ADT.Branch[Int] = Branch(...)

scala> Scala3ADT.tree
val res2: Scala3ADT.Tree[Int] = Branch(...)

```



One last point; you may have noticed that `Branch` and `Leaf` don't extend `Tree`, while in `WeekDay` above, each day extends `WeekDay`. For `Branch` and `Leaf`, extending `Tree` is inferred by the compiler, although we could add this explicitly. For `WeekDay`, each day must extend `WeekDay` to provide a value for the `val fullName: String` field declared by `WeekDay`.

## Interpolated Strings

We introduced *interpolated* strings in “A Sample Application”. Let's explore them further.

A `String` of the form `s"foo ${bar}"` will have the value of expression `bar`, converted to a `String` and inserted in place of `${bar}`. If the expression `bar` returns an instance of a type other than `String`, a `toString` method will be invoked, if one exists. It is an error if it can't be converted to a `String`.

If `bar` is just a variable reference, the curly braces can be omitted. For example:

```
val name = "Buck Trends"
println(s"Hello, $name")
```

When using interpolated strings, to write a literal dollar sign `$`, use two of them, `$$`.

There are two other kinds of interpolated strings. The first kind provides *printf* formatting and uses the prefix `f`. The second kind is called “raw” interpolated strings. It doesn't expand escape characters, like `\n`.

Suppose we're generating financial reports and we want to show floating-point numbers to two decimal places. Here's an example:

```
val gross    = 100000F
val net      = 64000F
```

```
val percent = (net / gross) * 100
println(f"$$$${gross}%.2f vs. $$$${net}%.2f or ${percent}%.1f%%")
```

The output of the last line is the following:

```
$100000.00 vs. $64000.00 or 64.0%
```

Scala uses Java's `Formatter` class for `printf` formatting. The embedded references to expressions use the same `${...}` syntax as before, but `printf` formatting directives trail them with no spaces.

In this example, we use two dollar signs, `$$`, to print a literal US dollar sign, and two percent signs, `%%`, to print a literal percent sign. The expression `${gross}%.2f` formats the value of `gross` as a floating-point number with two digits after the decimal point.

The types of the variables used must match the format expressions, but some implicit conversions are performed. An `Int` expression in a floating point context is allowed. It just pads with zeros. However, attempting to use `Double` or `Float` in an `Int` context causes a compilation error, due to the truncation that would be required.

While Scala uses Java strings, in certain contexts the Scala compiler will wrap a Java `String` with extra methods defined in `scala.collection.StringOps`. One of those extra methods is an *instance* method called `format`. You call it on the format string itself, then pass as arguments the values to be incorporated into the string. For example:

```
scala> val s = "%02d: name = %s".format(5, "Dean Wampler")
val s: String = "05: name = Dean Wampler"
```

In this example, we asked for a two-digit integer, padded with leading zeros.

The final version of the built-in string interpolation capabilities is the “raw” format that doesn’t expand control characters. Consider these

examples:

```
scala> val name = "Dean Wampler"
val name: String = "Dean Wampler"

scala> val multiLine = s"123\n$name\n456"
val multiLine: String = 123
Dean Wampler
456

scala> val multiLineRaw = raw"123\n$name\n456"
val multiLineRaw: String = 123\nDean Wampler\n456
```

Finally, we can actually define our own string interpolators, but we'll need to learn more about *context abstractions* first. See [“Build Your Own String Interpolator”](#) for details.

## Scala Conditional Expressions

Scala conditionals start with the `if` keyword. They are *expressions*, meaning they return a value that you can assign to a variable. In many languages, `if` conditionals are *statements*, which can only perform side effect operations.

Scala 2 `if` expressions used braces just like Java's `if` statements, but in Scala 3, `if` expressions can also use the new conventions discussed in the previous section. A simple example:

```
// src/script/scala/progscala3/rounding/If.scala

(0 until 6) foreach { n =>
  if n%2 == 0 then
    println(s"$n is even")
  else if n%3 == 0 then
    println(s"$n is divisible by 3")
  else
    println(n)
}
```

Recall from “**New Scala 3 Syntax**” that the `then` keyword is required only if you pass the `-new-syntax` flag to the compiler or REPL, which we use in the code examples SBT file. However, if you don’t use that flag, you must wrap the predicate expressions, like `n%2 == 0`, in parentheses.

The bodies of each clause are so concise, we can write them on the same line as the `if` or `else` expressions:

```
// src/script/scala/progscala3/rounding/If2.scala

(0 until 6) foreach { n =>
  if n%2 == 0 then println(s"$n is even")
  else if n%3 == 0 then println(s"$n is divisible by 3")
  else println(n)
}
```

Here are the same examples using the curly brace syntax required by Scala 2 and optional for Scala 3:

```
// src/script/scala-2/progscala3/rounding/If.scala

(0 until 6) foreach { n =>
  if (n%2 == 0) {
    println(s"$n is even")
  } else if (n%3 == 0) {
    println(s"$n is divisible by 3")
  } else {
    println(n)
  }
}

// src/script/scala-2/progscala3/rounding/If2.scala

(0 until 6) foreach { n =>
  if (n%2 == 0) println(s"$n is even")
  else if (n%3 == 0) println(s"$n is divisible by 3")
  else println(n)
}
```

The older syntax can still be used, even with compiler flags for the new syntax, but you can also use flags to require the older syntax, as discussed in “[New Scala 3 Syntax](#)”.

What is the type of the returned value if objects of different types are returned by different branches? The type of the returned value will be the so-called *least upper bound* (LUB) of all the branches, the closest parent type that matches all the potential values from each clause.

In the following example, the LUB is `Option[String]`, because the three branches return either `Some[String]` or `None`. The returned sequence is of type `IndexedSeq[Option[String]]`:

```
// src/script/scala/progscala3/rounding/IfTyped.scala

scala> val seq = (0 until 6) map { n =>
  |   if n%2 == 0 then Some(n.toString)
  |   else None
  | }
val seq: IndexedSeq[Option[String]] = Vector(Some(0), None, Some(2),
...)
```

## Conditional Operators

Scala uses the same conditional operators as Java. [Table 3-1](#) provides the details.

Table 3-1. Conditional operators

Operator	Operation	Description
&&	and	The values on the left and right of the operator are true. The righthand side is <i>only</i> evaluated if the lefthand side is <i>true</i> .
	or	At least one of the values on the left or right is true. The righthand side is <i>only</i> evaluated if the lefthand side is <i>false</i> .
>	greater than	The value on the left is greater than the value on the right.
>=	greater than or equals	The value on the left is greater than or equal to the value on the right.
<	less than	The value on the left is less than the value on the right.
<=	less than or equals	The value on the left is less than or equal to the value on the right.
==	equals	The value on the left is the same as the value on the right.
!=	not equals	The value on the left is not the same as the value on the right.

The && and || operators are “short-circuiting”. They stop evaluating expressions as soon as the answer is known. This is handy when you must work with `null` values:

```
scala> val s: String|Null = null
val s: String | Null = null

scala> val okay = s != null && s.length > 5
val okay: Boolean = false
```

Calling `s.length` would throw a `NullPointerException` without the `s != null` test first. What happens if you use `||` instead? Also, we don't use `if` here, because we just want to know the `Boolean` value of the conditional.

Most of the operators behave as they do in Java and other languages. An exception is the behavior of `==` and its negation, `!=`. In Java, `==` compares instance references only. It doesn't check logical equality, i.e., comparing field values. You must call the `equals` method for that purpose. Instead in Scala, `==` and `!=` call the `equals` method for the left-hand instance. You actually don't implement `equals` yourself very often in Scala, because you mostly only compare case class instances and the compiler generates `equals` automatically for case classes! You can override it when you need to, however.

If you need to determine if two instances are identical references, use the `eq` method.

See “[Equality of Instances](#)” for more details.

## Scala for Comprehensions

Another familiar control structure that's particularly feature-rich in Scala is the `for` loop, called the *for comprehension*. They are expressions, not statements as in Java.

The term *comprehension* comes from functional programming. It expresses the idea that we are traversing one or more collections of some kind, “comprehending” something new from it, such as another collection. Python *list comprehensions* are a similar concept.

## for Loops

Let's start with a basic `for` expression. As for `if` expressions, I use the new format options consistently in the code examples, except

where noted.

```
// src/script/scala/progscala3/rounding/BasicFor.scala

for
  i <- 0 until 10
do println(i)
```

Since there is one expression inside the `for ... do`, you can put the expression on the same line after the `for` and you can even put everything on one line:

```
for i <- 0 until 10
do println(i)

for i <- 0 until 10 do println(i)
```

As you might guess, this code says, “For every integer between 0 inclusive and 10 exclusive, print it on a new line.”

Because this form doesn’t return anything, it only performs side effects. These kinds of `for` comprehensions are sometimes called *for loops*, analogous to Java `for` loops.

In the older Scala 2 syntax, this example would be written as follows:

```
// src/script/scala-2/progscala3/rounding/BasicFor.scala

for (i <- 0 until 10)
  println(i)

for (i <- 0 until 10) println(i)
```

### TIP

From now on, in the examples that follow, I’ll only show Scala 3 syntax, but you can find Scala 2 versions of some examples, in the code examples under the directory `src/*/scala-2/progscala3/...` and a table of differences in *Appendix A*.



## Generator Expressions

The expression `i <- 0 until 10` is called a *generator expression*, so named because it *generates* individual values in some way. The left arrow operator (`<-`) is used to iterate through any collection or iterator instance that supports sequential access, such as `Seq` and `Vector`.

## Guards: Filtering Values

We can add `if` expressions, called *guards*, to filter for just elements we want to keep:

```
// src/script/scala/progscala3/rounding/GuardFor.scala

for
  n <- 0 to 6
  if n%2 == 0
do println(n)
```

The output is the number 0, 2, 4, and 6 (because we use `to` to make the 6 inclusive). Note the sense of filtering; the guards express what to *keep*, not *remove*.

## Yielding New Values

So far our `for` loops have only performed side effects, writing to output. Usually, we want to return a new collection, making our `for` expressions *comprehensions* rather than loops. We use the `yield` keyword to express this intent:

```
// src/script/scala/progscala3/rounding/YieldingFor.scala

val evens = for
  n <- 0 to 10 // Note: 0 to 10, inclusive
  if n%2 == 0
yield n
```

```
assert(evens == Seq(0, 2, 4, 6, 8, 10))
```

Each iteration through the for expression “yields” a new value, named *n*. These are accumulated into a new collection that is assigned to the variable *evens*.

The type of the collection resulting from a comprehension expression is inferred from the type of the collection being iterated over. *Seq* is a *trait* and the actual concrete instance returned is of type *IndexedSeq*.

In the following example, a *Vector[Int]* is converted to a *Vector[String]*:

```
// src/script/scala/progscala3/rounding/YieldingForVector.scala

val odds = for
  number <- Vector(1,2,3,4,5)
  if number % 2 == 1
yield number.toString

assert(odds == Vector("1", "3", "5"))
```

## Expanded Scope and Value Definitions

You can define immutable values inside the for expressions without using the *val* keyword, like *fn* in the following example that uses the *WeekDay* enumeration we defined earlier in this chapter:

```
// src/script/scala/progscala3/rounding/ScopedFor.scala

import progscala3.rounding.WeekDay

val days = for
  day <- WeekDay.values
  if day.isWorkingDay
  fn = day.fullName
yield fn

assert(days.toSeq ==
  Seq("Friday", "Monday", "Tuesday", "Wednesday", "Thursday"))
```

In this case, the for comprehension now returns an `Array[String]`, because `WeekDay.values` returns an `Array[WeekDay]`. Because Arrays are Java Arrays and Java doesn't define a useful `equals` method, we convert to a `Seq` with `toSeq` and perform the assertion check.

Now let's consider a powerful use of `Option` with for comprehensions. Recall we discussed `Option` as a better alternative to using `null`. It's also useful to recognize that `Option` is a special kind of collection, limited to zero or one elements. In fact, we can "comprehend" it too:

```
// src/script/scala/progscala3/rounding/ScopedOptionFor.scala

import progscala3.rounding.WeekDay
import progscala3.rounding.WeekDay._

val dayOptions = Seq(
  Some(Mon), None, Some(Tue), Some(Wed), None,
  Some(Thu), Some(Fri), Some(Sat), Some(Sun), None)

val goodDays1 = for           // First pass
  dayOpt <- dayOptions
  day <- dayOpt
  fn = day.fullName
yield fn
assert(goodDays1 == Seq("Monday", "Tuesday", "Wednesday", ...))

val goodDays2 = for          // second, more concise pass
  case Some(day) <- dayOptions
  fn = day.fullName
yield fn
assert(goodDays1 == Seq("Monday", "Tuesday", "Wednesday", ...))
```

Imagine that we called some services to return days of the week. The services returned `Options`, because some of the services couldn't return anything, so they returned `None`. Now we want to remove and ignore the `None` values.

In the first expression of the first for comprehension, each element extracted is an `Option`, assigned to `dayOpt`. The next line uses the arrow to extract the value in the option and assign it to `day`.

But wait! Doesn't `None` throw an exception if you try to extract a value from it? Yes, but the comprehension effectively checks for this case and skips the `Nones`. It's as if we added an explicit `if dayOpt != None` before the second line.

Hence, we construct a collection with only values from `Some` instances.

The second for comprehension makes this filtering even cleaner and more concise, using *pattern matching*. The expression `case Some(day) <- dayOptions` only succeeds when the instance is a `Some`, also skipping the `None` values, and it extracts the value into `day`, all in one step.

To recap the difference between using the left arrow (`<-`) versus the equals sign (`=`), use the arrow when you are iterating through a collection and extracting values. Use the equals sign when you're assigning a value from another value that doesn't involve iteration. A limitation is that the first expression in a for comprehension has to be an extraction/iteration using the left arrow. If you really need to define a value first, put it before the for comprehension.

When working with loops in many languages, they provide `break` and `continue` keywords for breaking out of a loop completely or continuing to the next iteration, respectively. Scala doesn't have either of these keywords, but when writing idiomatic Scala code, they aren't missed. Use conditional expressions to test if a loop should continue, or make use of recursion. Better yet, filter your collections ahead of time to eliminate complex conditions within your loops.

## Scala while Loops

The while loop is seldom used. It executes a block of code as long as a condition is true:

```
// src/script/scala/progscala3/rounding/While.scala

var count = 0
while count < 10
  count += 1
  println(count)

assert(count == 10)
```

## Scala do-while Loops

Scala 3 dropped the do-while construct in Scala 2, because it was rarely used. It can be rewritten using while, although awkwardly:

```
// src/script/scala/progscala3/rounding/DowhileAlternative.scala

var count = 0
while
  count += 1
  println(count)
  count < 10
do ()
assert(count == 10)
```

## Using try, catch, and finally Clauses

Through its use of functional constructs and strong typing, Scala encourages a coding style that lessens the need for exceptions and exception handling. But exceptions are still supported. In particular, they are needed when using Java libraries.

Unlike Java, Scala does not have checked exceptions. Java's checked exceptions are treated as unchecked by Scala. There is also no throws clause on method declarations. However, there is a @throws annotation that is useful for Java interoperability. See *later chapter*

Scala uses pattern matching to specify the exceptions to be caught. The following example implements a common application scenario, resource management. We want to open files and process them in some way. In this case, we'll just count the lines. However, we must handle a few error scenarios. The file might not exist, perhaps because the user misspelled the filenames. Also, something might go wrong while processing the file. (We'll trigger an arbitrary failure to test what happens.) We need to ensure that we close all open file handles, whether or not we process the files successfully:

```
// src/main/scala/progscala3/rounding/TryCatch.scala
package progscala3.rounding
import scala.io.Source
❶
import scala.util.control.NonFatal

/** Usage: scala rounding.TryCatch filename1 filename2 ... */
@main def TryCatch(fileNames: String*) =
❷
  fileNames foreach { fileName =>
    var source: Option[Source] = None
❸
    try
      source = Some(Source.fromFile(fileName))
      val size = source.get.getLines.size
      println(s"file $fileName has $size lines")
    catch
      case NonFatal(ex) => println(s"Non fatal exception! $ex")
❹
❺
❻
    finally
      for s <- source do
        println(s"Closing $fileName...")
        s.close
      }
    }
    }
  }
```

- ❶ Import `scala.io.Source` for reading input and `scala.util.control.NonFatal` for matching on “nonfatal” exceptions, i.e., those where it’s reasonable to attempt recovery.

- ❷ In Scala 3, we don't need an object to wrap the "main" method. By using the `@main` annotation, we can name the method whatever we want and we can specify the number and types of the arguments expected. Here, we just expect zero or more strings.
- ❸ Declare the source to be an `Option`, so we can tell in the `finally` clause if we have an actual instance or not. We use a mutable variable, but it's hidden inside the implementation and thread safety isn't a concern here.
- ❹ Start of `try` clause.
- ❺ `Source.fromFile` will throw a `java.io.FileNotFoundException` if the file doesn't exist. Otherwise, wrap the returned `Source` instance in a `Some`. Calling `get` on the next line is safe, because if we're here, we know we have a `Some`.
- ❻ Catch nonfatal errors. For example, out of memory would be fatal.
- ❼ Use a `for` comprehension to extract the `Source` instance from the `Some` and close it. If `source` is `None`, then nothing happens.

Note the catch clause. Scala uses pattern matches to pick the exceptions you want to catch. This is more compact and more flexible than Java's use of separate catch clauses for each exception. In this case, the clause `case NonFatal(ex) => ...` uses `scala.util.control.NonFatal` to match any exception that isn't considered fatal.

The `finally` clause is used to ensure proper resource cleanup in one place. Without it, we would have to repeat the logic at the end of the `try` clause and the catch clause, to ensure our file handles are closed. Here we use a `for` comprehension to extract the `Source` from

the option. If the option is actually a `None`, nothing happens; the block with the `close` call is not invoked. Note that since this is “main” method, the handles would be cleaned up anyway on exit, but you’ll want to close resources in other contexts.

### TIP

When resources need to be cleaned up, whether or not the resource is used successfully, put the cleanup logic in a `finally` clause.

This program is already compiled by `sbt` and we can run it from the `sbt` prompt using the `runMain` task, which lets us pass arguments. I have elided some output:

```
> runMain progscale3.rounding.TryCatch README.md foo/bar
file README.md has 116 lines
Closing README.md...
Non fatal exception! java.io.FileNotFoundException: foo/bar (...)
```

You throw an exception by writing `throw new MyBadException(...)`. If your custom exception is a case class, you can omit the `new`.

Automatic resource management is a common pattern. Let’s use a Scala library facility, `scala.util.Using` for this purpose.<sup>2</sup> Then we’ll actually implement our own version to illustrate some powerful capabilities in Scala and better understand how the library version works.

```
// src/main/scala/progscale3/rounding/FileSizes.scala
package progscale3.rounding

import scala.util.Using
import scala.io.Source

/** Usage: scala rounding.FileSizes filename1 filename2 ... */
@main def FileSizes(fileNames: String*) =
  val sizes = fileNames map { fileName =>
```



```

    Using.resource(Source.fromFile(fileName)) { source =>
        source.getLines.size
    }
}
println(s"Returned sizes: ${sizes.mkString(", ")}")
println(s"Total size: ${sizes.sum}")

```

This simple program also counts the number of lines in the files specified on the command line. However, if a file is not readable or doesn't exist, an exception is thrown and processing stops. No other results are produced, unlike the previous TryCatch example, which continues processing the arguments specified.

See the `scala.util.Using` documentation for a few other ways this utility can be used. For more sophisticated approaches to error handling, see “Nedelcu2020,” Alexandru Nedelcu’s blog.

## Call by Name, Call by Value

Now let’s implement our own *application resource manager* to learn a few powerful techniques that Scala provides for us. This implementation will build on the TryCatch example:

```

// src/main/scala/progscala3/rounding/TryCatchArm.scala
package progscala3.rounding
import scala.language.reflectiveCalls
import reflect.Selectable.reflectiveSelectable
import scala.util.control.NonFatal
import scala.io.Source

object manage:
  def apply[R <: { def close():Unit }, T](resource: => R)(f: R => T):
  T =
    var res: Option[R] = None
    try
      res = Some(resource)           // Only reference "resource" once!!
      f(res.get)                     // Return the T instance
    catch
      case NonFatal(ex) =>
        println(s"manage.apply(): Non fatal exception! $ex")

```

```

        throw ex
    finally
        res match
            case Some(resource) =>
                println(s"Closing resource...")
                res.get.close()
            case None => // do nothing

/** Usage: scala rounding.TryCatchARM filename1 filename2 ... */
@main def TryCatchARM(fileNames: String*) =
    val sizes = fileNames map { fileName =>
        try
            val size = manage(Source.fromFile(fileName)) { source =>
                source.getLines.size
            }
            println(s"file $fileName has $size lines")
            size
        catch
            case NonFatal(ex) =>
                println(s"caught $ex")
                0
    }
    println("Returned sizes: " + (sizes.mkString(", ")))

```

The output will be similar what we saw for TryCatch.

This is a lovely little bit of *separation of concerns*, but to implement it, we used a few new power tools.

First, we named our object `manage` rather than `Manage`. Normally, you follow the convention of using a leading uppercase letter for type names, but in this case we will use `manage` like a function. We want client code to look like we're using a built-in operator, similar to a `while` loop. This is another example of Scala's tools for building little DSLs.

That `manage.apply` method declaration is hairy looking. Let's deconstruct it. Here is the signature again, spread over several lines and annotated:

```

def apply[
    R <: { def close():Unit },    ❶

```

```

T ]                                ❷
(resource: => R)                   ❸
(f: R => T) = {...}               ❹

```

- ❶ Two new things are shown here. `R` is the type of the resource we'll manage. The `<:` means `R` is a subclass of something else. In this case, any type used for `R` must contain a `close():Unit` method. We declare this using a *structural type* defined with the braces. What would be more intuitive, especially if you are new to structural types, would be for all resources to implement a `Closable` interface that defines a `close():Unit` method. Then we could say `R <: Closable`. Instead, structural types let us use reflection and plug in any type that has a `close():Unit` method (like `Source`). Reflection has a lot of overhead and structural types are a bit scary, so reflection is another *optional feature*, like postfix expressions, which we saw earlier. So we add the import statements to tell the compiler we know what we're doing.
- ❷ `T` will be the type returned by the anonymous function passed in to do work with the resource.
- ❸ It looks like `resource` is a function with an unusual declaration. Actually, `resource` is a *by-name* parameter, which we first encountered in ["A Taste of Futures"](#).
- ❹ Finally we have a second parameter list containing a function for the work to do with the resource. This function will take the resource as an argument and return a result of type `T`.

Recapping point 1, here is how the `apply` method declaration would look if we could assume that all resources implement a `Closable` abstraction:

```

object manage {
  def apply[R <: Closable, T](resource: => R)(f: R => T) =

```

```
} ...
```

The line, `res = Some(resource)`, is the *only* place `resource` is evaluated, which is important, because it is a *by-name* parameter. We learned in “[A Taste of Futures](#)” that they are lazily evaluated, only when used, but they are evaluated every time they are referenced, just like a function call would be. The thing we pass as `resource` inside `TryCatchARM`, `Source.fromFile(fileName)`, should only be evaluated *once* inside `apply` to construct the `Source` for a file. The code correctly evaluates it once.

So, you have to use *by-name* parameters carefully, but their virtue is the ability to control when and even if a block of code is evaluated. We’ll see another example shortly where we will evaluate a *by-name* parameter repeatedly for a good reason.

To recap, it’s as if the `res = ...` line is actually this:

```
res = Some(Source.fromFile(fileName))
```

After constructing `res`, it is passed to the work function `f`.

See how `manage` is used in `TryCatchARM`. It looks like a built-in control structure with one parameter list that creates the `Source` and a second parameter list that is a block of code that works with the `Source`. So, using `manage` looks something like a conventional `while` statement.

Like most languages, Scala normally uses *call-by-value* semantics. If we write `val source = Source.fromFile(fileName)`, it is evaluated immediately.

Supporting idiomatic code like our use of `manage` is the reason that Scala offers *by-name* parameters, without which we would have to pass an anonymous function that looks ugly:

```
manage(() => Source.fromFile(fileName)) { source =>
```

Then, within `manage.apply`, our reference to `resource` would now be a function call:

```
val res = Some(resource())
```

Okay, that's not a terrible burden, but *call by name* enables a syntax for building our own control structures, like our `manage` utility.

Here is another example using a call by name, this time repeatedly evaluating *two* by-name parameters; an implementation of a while-like loop construct, called `continue`:

```
// src/script/scala/progscala3/rounding/CallByName.scala
import scala.annotation.tailrec
```

```
@tailrec
```

```
def continue(conditional: => Boolean)(body: => Unit): Unit =
```

```
  ②
```

```
    if conditional then
```

```
      body
```

```
      continue(conditional)(body)
```

```
var count = 0
```

```
continue (count < 5) {
```

```
  println(s"at $count")
```

```
  count += 1
```

```
}
```

```
assert(count == 5)
```

- ❶ Ensure the implementation is tail recursive.
- ❷ Define a `continue` function that accepts two argument lists. The first list takes a single, by-name parameter that is the conditional. The second list takes a single, by-name value that is the body to be evaluated for each iteration.
- ❸ Evaluate the condition. If true, evaluate the body and call `continue` recursively.

#### ④ Try it with traditional brace syntax.

It's important to note that the by-name parameters are evaluated every time they are referenced. So, by-name parameters are in a sense *lazy*, because evaluation is deferred, but possibly repeated over and over again. Scala also provides lazy values. By the way, this implementation shows how “loop” constructs can be replaced with recursion.

Unfortunately, this ability to define our own control structures only works with the old Scala 2 syntax using parentheses and braces. If `continue` really behaved like `while` or similar built-in constructs, we would be able to write the last two examples. The syntax uniformity with user-defined constructs is a nice feature we must give up if we use the new syntax... or do we??

```
count = 0
continue (count < 5)
  println(s"at $count")
  count += 1
assert(count == 5)
```

This actually parses, but it executes in an unexpected way. Here it is again, annotated to explain what actually happens:

```
continue (count < 5)    // Returns a "partially-applied function"
  println(s"at $count") // A separate statement that prints "at 0"
once
  count += 1           // Increments count once
```

A *partially-applied function* is one where we provide some, but not all arguments, returning a new function that will accept the remaining arguments (see “**Partially Applied Functions Versus Partial Functions**”). Here’s what the REPL will print for that line:

```
scala> continue (count < 5)
val res1: (=> Unit) => Unit =
```

`Lambda$11103/0x00000008048c0040@4d0f9ec0`

This is an anonymous function, implemented with a JDK *lambda* on the JVM, that takes a by-name parameter returning `Unit` (recall the second parameter for `continue`) and then returns `Unit`.

The second and third lines are not treated as part of the body that should be passed to `continue`. They are treated as single statements that are evaluated once.

So, the optional braces don't work.

## lazy val

By-name parameters show us that lazy evaluation is useful, but they are evaluated every time they are referenced.

There are times when you want to evaluate an expression *once* to initialize a field in an object, but you want to defer that invocation until the value is actually needed. In other words, on-demand evaluation. This is useful when:

- The expression is expensive (e.g., opening a database connection) and you want to avoid the overhead until the value is actually needed, which could be never.
- You want to improve startup times for modules by deferring work that isn't needed immediately.
- A field in an object needs to be initialized lazily so that other initializations can happen first.

We'll explore the last scenario when we discuss "**Initializing Abstract Fields**".

Here is a "sketch" of an example using a `lazy val`:

```
// src/script/scala/progscala3/rounding/LazyInitVal.scala
```

```

case class DBConnection():
  println("In constructor")
  type MySQLConnection = String
  lazy val connection: MySQLConnection =
    // Connect to the database
    println("Connected!")
    "DB"

```

The `lazy` keyword indicates that evaluation will be deferred until the value is accessed.

Let's try it. Notice when the `println` statements are executed:

```

scala> val dbc = DBConnection()
In constructor
val dbc: DBConnection = DBConnection()

scala> dbc.connection
Connected!
val res4: dbc.MySQLConnection = DB

scala> dbc.connection
val res5: dbc.MySQLConnection = DB

```

So, how is a `lazy val` different from a method call? We see that “Connected!” was only printed once, whereas if `connection` were a method, the body would be executed *every* time and we would see “Connected!” printed each time. Furthermore, we didn't see that message until when we referenced `connection` the first time.

One-time evaluation makes little sense for a mutable field. Therefore, the `lazy` keyword is not allowed on vars.

Lazy values are implemented with the equivalent of a *guard*. When client code references a lazy value, the reference is intercepted by the guard to check if initialization is required. This guard step is really only essential the *first* time the value is referenced, so that the value is initialized first before the access is allowed to proceed.

Unfortunately, there is no easy way to eliminate these checks for subsequent calls. So, lazy values incur overhead that “eager” values



don't. Therefore, you should only use lazy values when initialization is expensive, especially if the value may not actually be used. There are also some circumstances where careful ordering of initialization dependencies is most easily implemented by making some values lazy (see [“Initializing Abstract Fields”](#)).

There is a `@threadUnsafe` annotation you can add to a lazy val. It causes the initialization to use a faster mechanism which is not thread-safe, so use with caution.

## Traits: Interfaces and “Mixins” in Scala

Until now, I have emphasized the power of functional programming in Scala. I waited until now to discuss Scala's features for object-oriented programming, such as how abstractions and concrete implementations are defined, and how inheritance is supported. We've seen some details in passing, like abstract and case classes and objects, but now it's time to cover these concepts.

Scala uses *traits* to define abstractions. We'll explore most details in [Chapter 10](#), but for now, think of them as interfaces for declaring abstract member fields, methods, and types, with the option of defining any or all of them, too.

Traits enable true *separation of concerns* and composition of behaviors (“mixins”).

Here is a typical enterprise developer task, adding logging. Let's start with a service:

```
// tag::service[]  
// src/script/scala/progscala3/rounding/Traits.scala  
import util.Random  
  
class Service(name: String):  
  def work(i: Int): (Int, Int) =  
    (i, Random.between(0, 1000))
```

```

val service1 = new Service("one")
(1 to 3) foreach (i => println(s"Result: ${service1.work(i)}"))
// end::service[]

// tag::logging[]
trait Logging:
  def info (message: String): Unit
  def warning(message: String): Unit
  def error (message: String): Unit

trait StdoutLogging extends Logging:
  def info (message: String) = println(s"INFO: $message")
  def warning(message: String) = println(s"WARNING: $message")
  def error (message: String) = println(s"ERROR: $message")
// end::logging[]

// tag::example[]
val service2 = new Service("two") with StdoutLogging:
  override def work(i: Int): (Int, Int) =
    info(s"Starting work: i = $i")
    val result = super.work(i)
    info(s"Ending work: result = $result")
    result

(1 to 3) foreach (i => println(s"Result: ${service2.work(i)}"))
// end::example[]

```

We ask the service to do some (random) work and get this output:

```

Result: (1,975)
Result: (2,286)
Result: (3,453)

```

Now we want to mix in a standard logging library. For simplicity, we'll just use `println`.

Here are two traits, one that defines the abstraction with no concrete members and the other that implements the abstraction for “logging” to standard output:

```

// tag::service[]
// src/script/scala/progscala3/rounding/Traits.scala

```

```

import util.Random

class Service(name: String):
  def work(i: Int): (Int, Int) =
    (i, Random.between(0, 1000))

val service1 = new Service("one")
(1 to 3) foreach (i => println(s"Result: ${service1.work(i)}"))
// end::service[]

// tag::logging[]
trait Logging:
  def info (message: String): Unit
  def warning(message: String): Unit
  def error (message: String): Unit

trait StdoutLogging extends Logging:
  def info (message: String) = println(s"INFO: $message")
  def warning(message: String) = println(s"WARNING: $message")
  def error (message: String) = println(s"ERROR: $message")
// end::logging[]

// tag::example[]
val service2 = new Service("two") with StdoutLogging:
  override def work(i: Int): (Int, Int) =
    info(s"Starting work: i = $i")
    val result = super.work(i)
    info(s"Ending work: result = $result")
    result

(1 to 3) foreach (i => println(s"Result: ${service2.work(i)}"))
// end::example[]

```

Note that Logging is pure abstract. It works *exactly* like a Java interface. It is even implemented the same way in JVM byte code.

Finally, let's declare a service that "mixes in" logging and use it:

```

// tag::service[]
// src/script/scala/progscala3/rounding/Traits.scala
import util.Random

class Service(name: String):
  def work(i: Int): (Int, Int) =

```

```

        (i, Random.between(0, 1000))

val service1 = new Service("one")
(1 to 3) foreach (i => println(s"Result: ${service1.work(i)}"))
// end::service[]

// tag::logging[]
trait Logging:
  def info (message: String): Unit
  def warning(message: String): Unit
  def error (message: String): Unit

trait StdoutLogging extends Logging:
  def info (message: String) = println(s"INFO: $message")
  def warning(message: String) = println(s"WARNING: $message")
  def error (message: String) = println(s"ERROR: $message")
// end::logging[]

// tag::example[]
val service2 = new Service("two") with StdoutLogging:
  override def work(i: Int): (Int, Int) =
    info(s"Starting work: i = $i")
    val result = super.work(i)
    info(s"Ending work: result = $result")
    result

(1 to 3) foreach (i => println(s"Result: ${service2.work(i)}"))
// end::example[]

```

We override the `work` method to log when we enter and before we leave the method. Scala requires the `override` keyword when you override a concrete method in a parent class. This prevents mistakes when you didn't know you were overriding a method, for example from a library parent class and it catches misspelled method names that aren't actually overrides! Note how we access the parent class `work` method, using `super.work`.

Here is the output:

```

INFO:    Starting work: i = 1
INFO:    Ending work: result = (1,737)
Result:  (1,737)

```

```
INFO:    Starting work: i = 2
INFO:    Ending work: result = (2,310)
Result:  (2,310)
INFO:    Starting work: i = 3
INFO:    Ending work: result = (3,273)
Result:  (3,273)
```

## WARNING

Be very careful about overriding concrete methods! In this case, we don't change the behavior of the the parent-class method. We just log activity, then call the parent method, then log again. We are careful to return the result unchanged that was returned by the parent method.

To mix in traits while constructing an instance as shown, we use the `with` keyword. We can mix in as many as we want. Some traits might not modify existing behavior at all, and just add new useful, but independent methods.

In this example, we're actually *modifying* the behavior of `work`, in order to inject logging, but we are not changing its "contract" with clients, that is, its external behavior.<sup>3</sup>

If we needed multiple instances of `Service` with `StdoutLogging`, we should declare a class:

```
class LoggedService(name: String)
  extends Service(name) with StdoutLogging:
  ...
```

Note how we pass the `name` argument to the parent class `Service`. To create instances, `new LoggedService("three")` works as you would expect it to work.

There is a lot more to discuss about traits and *mixin composition*, as we'll see.

## Recap and What's Next

We've covered a lot of ground in these first chapters. We learned how flexible and concise Scala code can be. In this chapter, we learned some powerful constructs for defining DSLs and for manipulating data, such as for comprehensions. Finally, we learned more about enumerations and the basic capabilities of traits.

You should now be able to read quite a lot of Scala code, but there's plenty more about the language to learn. Next we'll begin a deeper dive into Scala features.

- 
- 1 You can find a Scala 2 version of `WeekDay` in the code examples, `src/script/scala-2/progscala3/rounding/WeekDay.scala`.
  - 2 Not to be confused with the keyword `using` that we discussed in “[A Taste of Futures](#)”.
  - 3 That's not strictly true, in the sense that the extra I/O has changed the code's interaction with the “world.”

# Chapter 4. Pattern Matching

---

## A NOTE FOR EARLY RELEASE READERS

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

This will be the 4th chapter of the final book. Please note that the GitHub repo will be made active later on.

If you have comments about how we might improve the content and/or examples in this book, or if you notice missing material within this chapter, please reach out to the author at [mcronin@oreilly.com](mailto:mcronin@oreilly.com).

Scala’s pattern matching provides deep inspection and decomposition of objects in a variety of ways. It’s one of my favorite features in Scala. For your own types, you can follow a protocol that allows you to control the visibility of internal state and how to expose it to users. The terms “extraction” and “destructuring” are sometimes used for this capability.

Pattern matching can be used in several code contexts, as we’ve already seen in “[A Sample Application](#)” and “[Partial Functions](#)”. We’ll start with a quick tour of common and straightforward usage examples, then explore more advanced scenarios.

## Values, Variables, and Types in Matches

Let’s cover several kinds of matches. The following example matches on specific values, all values of specific types, and it shows

one way of writing a “default” clause that matches anything:

```
// src/script/scala/progscala3/patternmatching/MatchVariable.scala

val seq = Seq(1, 2, 3.14, 5.5F, "one", "four", true, (6, 7)) ❶
val result = seq map {
  case 1                => "int 1"                               ❷
  case i: Int           => s"other int: $i"
  case d: (Double | Float) => s"a double or float: $d"          ❸
  case "one"            => "string one"                         ❹
  case s: String        => s"other string: $s"
  case (x, y)           => s"tuple: ($x, $y)"                   ❺
  case unexpected       => s"unexpected value: $unexpected"     ❻
}
assert(result == Seq(
  "int 1", "other int: 2",
  "a double or float: 3.14", "a double or float: 5.5",
  "string one", "other string: four",
  "unexpected value: true",
  "tuple: (6, 7)"))
```

- ❶ Because of the mix of values, seq is of type Seq[Any].
- ❷ If one or more case clauses specify particular values of a type, they need to occur before more general clauses that just match on the type.
- ❸ When two cases are handled the same, how to match on either case.
- ❹ Here’s another example with strings.
- ❺ Match on a two-element tuple where the elements are of any type, extract the elements into the variables x and y.
- ❻ Match all other inputs, where unexpected is the variable to which the value of x is assigned. The variable name is arbitrary. Because no type annotation is given, Any is inferred. This functions as the “default” clause.



I pass a partial function to `Seq.map()`. It's actually *total* because the last clause matches anything.

I didn't include a clause with a floating-point literal, because matching on floating-point literals is a bad idea, as rounding errors mean two values that appear to be the same often differ in the least significant digits.

Matches are eager, so more specific clauses must appear before less specific clauses. Otherwise, the more specific clauses will never get the chance to match. So, the clauses matching on particular values of types must come before clauses matching on the type (i.e., on any value of the type). The default clause shown must be the last one.

Pattern matching is an expression, so it returns a value. In this case, all clauses return strings, so the return type of the whole thing is `List[String]`. The compiler infers the closest supertype (also called the *least upper bound*) for types of values returned by all the case clauses.

We passed a *partial function literal* to `map`, rather than a typical anonymous function. Recall that the literal syntax requires case statements.

Here is a similar example that passes an anonymous function to `map`, rather than a partial function, plus some other changes:

```
// src/script/scala/progscala3/patternmatching/MatchVariable2.scala

val seq2 = Seq(1, 2, 3.14, "one", (6, 7))
val result2 = seq2 map {
  x => x match
    case _: Int           => s"int: $x"           ❶
    case _                => s"unexpected value: $x" ❷
}
assert(result2 == Seq(
  "int: 1", "int: 2", "unexpected value: 3.14",
  "unexpected value: one", "unexpected value: (6,7)"))
```

- ❶ Use `_` for the variable name, meaning we don't capture it. We don't actually need to, because we already have `x`.
- ❷ Catch-all clause that also uses `x` instead of capturing to a new variable.

The first case clause doesn't need to capture the variable because it doesn't exploit the fact that the value is an `Int`. Otherwise, just using `x` wouldn't be sufficient, as it has type `Any`.

We used the braceless syntax inside the function, because now we have a match expression, whereas before we had a partial function literal where the braces were necessary to mark the whole anonymous function. In general, using a partial function is more concise, because we eliminate the need for `x => x match`.

### TIP

When you use any of the collection methods like `map` and `foreach`, and you will pattern match, use a partial function.

There are a few rules and gotchas to keep in mind when writing case clauses. The compiler assumes that a term that starts with a capital letter is a type name, while a term that begins with a lowercase letter is assumed to be the name of a variable that will hold an extracted or matched value.

This rule can cause surprises, as shown in the following example, where we want to match on a particular value that's passed into a method:

```
// src/script/scala/progscala3/patternmatching/MatchSurprise.scala  
  
def checkY(y: Int): Seq[String] =  
  for
```

```

    x <- Seq(99, 100, 101)
  yield
    x match
      case y => "found y!"
      case i: Int => "int: "+i // <2> ERROR: Unreachable code

  checkY(100)

```

We want the ability to pass in a specific value for the first case clause, rather than hard-code it. So, we might expect the first case clause to match when `x` equals `y`, which holds the value of 100, but this is what we actually get:

```

def checkY(y: Int): Seq[String]
10 |      case i: Int => "int: "+i // <2> ERROR: Unreachable code
    |          ^^^^^^
    |          Unreachable case

```

Calling `checkY(100)` returns `found y!` for all three numbers.

The case `y` actually means, “match anything (because there is no type annotation) and assign it to this *new* variable named `y`.” The `y` here is not interpreted as a reference to the method parameter `y`. So, we actually wrote a default, match-all clause first, triggering the “Unreachable case” warning, so we will never reach the second case expression.

The solution is to use “back ticks” to indicate we really want to match against the value held by `y`:

```

// src/script/scala/progscala3/patternmatching/MatchSurpriseFix.scala

def checkY(y: Int): Seq[String] =
  for
    x <- Seq(99, 100, 101)
  yield
    x match
      case `y` => "found y!" // Note the backticks
      case i: Int => "int: "+i

```

## WARNING

In case clauses, a term that begins with a lowercase letter is assumed to be the name of a new variable that will hold an extracted value. To refer to a previously defined variable, enclose it in back ticks. Conversely, a term that begins with an uppercase letter is assumed to be a type name.

Finally, most match expressions should be exhaustive:

```
// src/script/scala/progscala3/patternmatching/MatchExhaustive.scala
```

```
scala> val seq3 = Seq(Some(1), None, Some(2), None)
val seq3: Seq[Option[Int]] = List(Some(1), None, Some(2), None)
```

```
scala> val result3 = seq3 map {
  |   case Some(i) => "int 1"
  | }
scala.MatchError: None (of class scala.None$)
...
```

```
2 | case Some(i) => "int 1"
  | ^
  | match may not be exhaustive.
  |
  | It would fail on pattern case: None
```

The compiler knows that the elements of `seq3` are of type `Option[Int]`, but the partial function isn't exhaustive, so we get a warning that `None` isn't covered and a `MatchError` when a `None` is encountered. The fix is straightforward:

```
//
src/script/scala/progscala3/patternmatching/MatchExhaustiveFix.scala
```

```
scala> val result3 = seq3 map {
  |   case Some(i) => "int 1"
  |   case None    => ""
  | }
val result3: Seq[String] = List(int 1, "", int 1, "")
```

## Matching on Sequences

Let's examine the classic idiom for iterating through a Seq using pattern matching and recursion, and along the way, learn some useful fundamentals about sequences:

```
// src/script/scala/progscala3/patternmatching/MatchSeq.scala
```

```
val nonEmptySeq    = Seq(1, 2, 3)           ❶
val emptySeq       = Seq.empty[Int]
val nonEmptyVector = Vector(1, 2, 3)        ❷
val emptyVector    = Vector.empty[Int]
val nonEmptyMap    = Map("one" -> 1, "two" -> 2, "three" -> 3)  ❸
val emptyMap       = Map.empty[String,Int]

def seqToString[T](seq: Seq[T]): String = seq match
4  case head +: tail => s"($head +: ${seqToString(tail)})"  ❺
  case Nil           => "Nil"                                ❻
val results = Seq(                                           ❼
  nonEmptySeq, emptySeq, nonEmptyVector, emptyVector,
  nonEmptyMap.toSeq, emptyMap.toSeq) map {
  seq => seqToString(seq)
}

assert(results == Seq(
  "(1 +: (2 +: (3 +: Nil)))",
  "Nil",
  "(1 +: (2 +: (3 +: Nil)))",
  "Nil",
  "((one,1) +: ((two,2) +: ((three,3) +: Nil)))",
  "Nil"))
```

- ❶ Construct a nonempty scala-  
[lang.org/api/current/scala/collection/immutable/Seq.html](http://lang.org/api/current/scala/collection/immutable/Seq.html)[Seq[Int]]. A **List** is actually returned, which uses a linked-list implementation. Accessing the head of a list is  $O(1)$ , while accessing an arbitrary element is  $O(N)$ , where  $N$  is the length of the list. The next line shows the idiomatic way of constructing an empty Seq[Int].

- ② Construct a nonempty `Vectors[Int]`, a subtype of `Seq` with  $O(1)$  access patterns, followed by an empty `Vector[Int]`.
- ③ Construct a nonempty `Map[String,Int]`, which *isn't* a subtype of `Seq`; we'll use `toSeq` to return a sequence of key-value tuples.
- ④ Define a recursive method that constructs a `String` from a `Seq[T]` for some type `T`, which will be inferred from the sequence passed in. The body is a single match expression. Note this method is not tail recursive.
- ⑤ There are two match clauses and they are exhaustive. The first matches on any nonempty `Seq`, extracting the head, the first element, and the tail, which is the rest of the `Seq`. (`Seq` has `head` and `tail` methods, but here these terms are interpreted as variable names as usual for case clauses.) The body of the clause constructs a `String` with the head followed by `+`: followed by the result of calling `seqToString` on the tail.
- ⑥ The only other possible case is an empty `Seq`. We can use the special case object for an empty `List`, `Nil`, to match all the empty cases. This clause terminates the recursion. Note that any `Seq` can always be interpreted as terminating with an empty instance of the same type, although only some types, like `List`, are actually implemented that way.
- ⑦ Put the `Seqs` in another `Seq`, calling `Map.toSeq` as required, then iterate through it and call `seqToString` on each one.

`Map` is not a subtype of `Seq`, because it doesn't guarantee a particular order when you iterate over it. Calling `Map.toSeq` creates a sequence of key-value tuples that happen to be in insertion order, which is a side effect of the implementation for small `Maps` and not the general case.

There are two new kinds of case clauses. The first, `head +: tail`, matches the head element and the tail `Seq` (the remainder) of a sequence. The operator `+:` is the “cons” (construction) operator for sequences. Recall that methods that end with a colon (`:`) bind to the *right*, toward the `Seq` tail.

I’m calling them “operators” and “methods,” but actually that’s not quite right in this context; we’ll come back to this expression shortly and examine what’s going on.

The Scala library has an object called `Nil` for empty lists, but we can use it for matching any empty sequence. The types don’t have to match exactly.

Because `Seq` behaves conceptually like a linked list, where each head node holds an element and it points to the tail (the rest of the sequence), creating a hierarchical structure that is indicated by parentheses inserted, like this example:

```
(1 +: (2 +: (3 +: Nil)))
```

So, we process sequences with just two case clauses and recursion. Note that this implies something fundamental about all sequences; they are either empty or not. That sounds trite, but once you recognize simple patterns like this, it gives you a surprisingly general tool for “divide and conquer.” The idiom used by `processSeq` is widely reusable.

We can copy and paste the output of the previous examples to reconstruct the original objects, which is not accidental:

```
scala> val is = (1 +: (2 +: (3 +: Nil)))
val is: List[Int] = List(1, 2, 3)

scala> val kvs = (("one",1) +: (("two",2) +: (("three",3) +: Nil)))
val kvs: List[(String, Int)] = List((one,1), (two,2), (three,3))

scala> val map = Map(kvs :_*)
val map: Map[String, Int] = Map(one -> 1, two -> 2, three -> 3)
```

The `Map.apply` method expects a variable argument list of two-element tuples. So, in order to use the sequence `kvs`, we had to use the `:_*` idiom for the compiler to convert it to a variable-argument list.

So, there's an elegant symmetry between construction and pattern matching (“deconstruction”) when using `+:.`

## Matching on Tuples

Tuples are also easy to match on, using their literal syntax:

```
// src/script/scala/progscala3/patternmatching/MatchTuple.scala

val langs = Seq(
  ("Scala", "Martin", "Odersky"),
  ("Clojure", "Rich", "Hickey"),
  ("Lisp", "John", "McCarthy"))

val results = langs map {
  case ("Scala", _, _) => "Scala"                                ❶
  case (lang, first, last) => s"$lang, creator $first $last"    ❷
}
assert(results == Seq(
  "Scala",
  "Clojure, creator Rich Hickey",
  "Lisp, creator John McCarthy"))
```

- ❶ Match a three-element tuple where the first element is the string “Scala” and we ignore the second and third arguments.
- ❷ Match any three-element tuple, where the elements could be any type, but they are inferred to be Strings due to the input `langs`. Extract the elements into variables `lang`, `first`, and `last`.

A tuple can be taken apart into its constituent elements. We can match on literal values within the tuple, at any positions we want, and we can ignore elements we don't care about.



Just as we can construct pairs (two-element tuples) with `->`, we can deconstruct them that way, too:

```
// src/script/scala/progscala3/patternmatching/MatchPair.scala

val langs2 = Seq("Scala" -> "Odersky", "Clojure" -> "Hickey")

val results = langs2 map {
  case "Scala" -> _ => "Scala"           ❶
  case lang -> last => s"$lang: $last"   ❷
}
assert(results == Seq("Scala", "Clojure: Hickey"))
```

## Parameter Untupling

Consider this example of tuples with no nesting:

```
//
src/script/scala/progscala3/patternmatching/ParameterUntupling.scala

val tuples = Seq((1,2,3), (4,5,6), (7,8,9))
val counts1 = tuples.map { // result: List(6, 15, 24)
  case (x, y, z) => x + y + z
}
```

A disadvantage of the tuple syntax here is the implication that it's not exhaustive, when we know it is for three-element tuples. It is also a bit inconvenient. Scala 3 introduces *parameter untupling* that simplifies special cases like this. We can drop the `case` keyword:

```
val counts2 = tuples.map {
  (x, y, z) => x + y + z
}
```

We can even use anonymous variables:

```
val counts3 = tuples.map(_+_+_)
```

However, this untupling only happens to one level:

```
scala> val tuples2 = Seq((1,(2,3)), (4,(5,6)), (7,(8,9)))
      | val counts2b = tuples2.map {
      |   (x, (y, z)) => x + y + z
      | }
      |
3 | (x, (y, z)) => x + y + z
  |      ^^^^^^
  |      not a legal formal parameter
```

## Guards in Case Clauses

Matching on literal values is very useful, but sometimes you need a little additional logic:

```
// src/script/scala/progscala3/patternmatching/MatchGuard.scala

val results = Seq(1,2,3,4) map {
  case e if e%2 == 0 => s"even: $e"           ❶
  case o              => s"odd:  $o"         ❷
}
assert(results == Seq("odd:  1", "even: 2", "odd:  3", "even: 4"))
```

- ❶ Match only if `i` is even.
- ❷ Match the only other possibility, that `i` is odd.

Note that we didn't need parentheses around the condition in the `if` expression, just as we don't need them in `for` comprehensions. This was true in Scala 2's conditional syntax, too.

## Matching on Case Classes and Enums

It's no coincidence that the same `case` keyword is used for declaring "special" classes and for case expressions in `match` expressions. The features of case classes were designed to enable convenient pattern matching, which also works for enums. The compiler implements pattern matching and extraction for us. We can use it with nested

objects and we can bind variables at any level of the extraction, which we are seeing for the first time now:

```
// src/script/scala/progscala3/patternmatching/MatchDeep.scala

case class Address(street: String, city: String)
case class Person(name: String, age: Int, address: Address)

val alice    = Person("Alice",    25, Address("1 Scala Lane",
"Chicago"))
val bob      = Person("Bob",      29, Address("2 Java Ave.",  "Miami"))
val charlie  = Person("Charlie",  32, Address("3 Python Ct.",  "Boston"))

val results = Seq(alice, bob, charlie) map {
  case p @ Person("Alice", age, a @ Address(_, "Chicago")) => ❶
    s"Hi Alice! $p"
  case p @ Person("Bob", 29, a @ Address(street, city)) =>    ❷
    s"Hi ${p.name}! age ${p.age}, in ${a}"
  case p @ Person(name, age, Address(street, city)) =>       ❸
    s"Who are you, $name (age: $age, city = $city)?"
}
assert(results == Seq(
  "Hi Alice! Person(Alice,25,Address(1 Scala Lane,Chicago))",
  "Hi Bob! age 29, in Address(2 Java Ave.,Miami)",
  "Who are you, Charlie (age: 32, city = Boston)?"))
```

- ❶ Match on any person named “Alice”, of any age at any street address in Chicago. Use `p @` to bind variable `p` to `Person`, while also extracting fields inside the instance, in this case `age`.
- ❷ Match on any person named “Bob”, age 29 at any street and city. Bind `p` the whole `Person` instance and `a` to the nested `Address` instance.
- ❸ Match on any person, binding `p` to the `Person` instance and `name`, `age`, `street`, and `city` to the nested fields.

If you aren't extracting fields from the `Person` instance, we can just write `p: Person => ...`

This nested matching can go arbitrarily deep. Consider this example that revisits the enum `Tree[T]` algebraic data type from “Enumerations and Algebraic Data Types”:

```
// src/script/scala/progscala3/patternmatching/MatchTreeADT.scala
```

```
enum Tree[T]:  
  case Branch(left: Tree[T], right: Tree[T])  
  case Leaf(elem: T)  
  
import Tree._  
val tree1 = Branch(  
  Branch(Leaf(1), Leaf(2)),  
  Branch(Leaf(3), Branch(Leaf(4), Leaf(5)))  
)  
val tree2 = Branch(Leaf(6), Leaf(7))  
  
for t <- Seq(tree1, tree2, Leaf(8))  
yield t match  
  case Branch(  
    l @ Branch(_,_),  
    r @ Branch(rl @ Leaf(rli), rr @ Branch(_,_)) =>  
      s"l=$l, r=$r, rl=$rl, rli=$rli, rr=$rr"  
  case Branch(l, r) => s"Other Branch($l, $r)"  
  case Leaf(x) => s"Other Leaf($x)"
```

The same extraction could be done for the alternative version using a sealed class hierarchy in the original example.

The last two case clauses are relatively easy to understand. The first one is highly “tuned” to match `tree1`, although it uses `_` to ignore some parts of the tree. In particular, note that it isn’t sufficient to write `l @ Branch`. We need to write `l @ Branch(_,_)`. Try removing the `(, )` here and you’ll notice the first case no longer matches `tree1`, without any obvious explanation.

## WARNING

If a nested pattern match expression doesn't match when you think it should, make sure that you capture the full structure, like `l @ Branch(_,_)` instead of `l @ Branch`.

It's worth experimenting with this example to capture different parts of the trees, so you develop an intuition about what works, what doesn't, and how to debug match expressions.

Here's an example using tuples. Imagine we have a sequence of `(String, Double)` tuples for the names and prices of items in a store and we want to print them with their index. The `Seq.zipWithIndex` method is handy here:

```
// src/script/scala/progscala3/patternmatching/MatchDeepTuple.scala

val itemsCosts = Seq(("Pencil", 0.52), ("Paper", 1.35), ("Notebook",
2.43))

val results = itemsCosts.zipWithIndex map {
  case ((item, cost), index) => s"$index: $item costs $cost each"
}
assert(results == Seq(
  "0: Pencil costs 0.52 each",
  "1: Paper costs 1.35 each",
  "2: Notebook costs 2.43 each"))
```

Note that `zipWithIndex` returns a sequence of tuples of the form `(element, index)`, or `+((name, cost), index)` in this case. We matched on this form to extract the three elements and construct a string with them. I write code like this *a lot*.

## Matching on Regular Expressions

Regular expressions (or *regexes*) are convenient for extracting data from strings that have a particular structure.

Scala wraps Java's regular expressions.<sup>1</sup> Here is an example:

```
// src/script/scala/progscala3/patternmatching/MatchRegex.scala

val BookExtractorRE = """"Book: title=([^,]+),\s+author=(.+)""".r
❶
val MagazineExtractorRE = """"Magazine: title=([^,]+),\s+issue=
(.+)""".r

val catalog = Seq(
  "Book: title=Programming Scala Third Edition, author=Dean Wampler",
  "Magazine: title=The New Yorker, issue=January 2020",
  "Unknown: text=Who put this here??"
)

val results = catalog map {
  case BookExtractorRE(title, author) =>
❷
    s""""Book "$title", written by $author""""
  case MagazineExtractorRE(title, issue) =>
    s""""Magazine "$title", issue $issue""""
  case entry => s"Unrecognized entry: $entry"
}
assert(results == Seq(
  """"Book "Programming Scala Third Edition", written by Dean
Wampler""",
  """"Magazine "The New Yorker", issue January 2020""",
  "Unrecognized entry: Unknown: text=Who put this here??"))
```

- ❶ Match a book string, with two *capture groups* (note the parentheses), one for the title and one for the author. Calling the `r` method on a string creates a regex from it. Also match a magazine string, with *capture groups* for the title and issue (date).
- ❷ Use the regular expressions much like using case classes, where the string matched by each capture group is assigned to a variable.

Normally you'll want to use triple-quoted strings for the regexes. Otherwise, you must escape the regex "backslash" constructs, like

`\\s` instead of `\s`. You can also define regular expressions by creating new instances of the `Regex` class, as in `new Regex("""\W+""")`, but this isn't as common.

### WARNING

Using interpolation in triple-quoted strings doesn't work cleanly for the regex escape sequences. You still need to escape these sequences, e.g., `s"""$first\\s+$second""".r` instead of `s"""$first\s+$second""".r`. If you aren't using interpolation, escaping isn't necessary.

`scala.util.matching.Regex` defines several methods for other manipulations, such as finding and replacing matches.

## More on Type Matching

Consider the following example, where we attempt to discriminate between `List[Double]` and `List[String]` inputs:

```
// src/script/scala/progscala3/patternmatching/MatchTypes.scala

scala> val results = Seq(Seq(5.5,5.6,5.7), Seq("a", "b")) map {
  |   case seqd: Seq[Double] => ("seq double", seqd)
  |   case seqs: Seq[String] => ("seq string", seqs)
  |   case other              => ("unknown!", other)
  | }
val results: Seq[(String, Seq[Double | String])] =
  List((seq double,List(5.5, 5.6, 5.7)), (seq double,List(a, b)))
2 |   case seqd: Seq[Double] => ("seq double", seqd)
  |       ^^^^^^^^^^^^^^^^^
  |       the type test for Seq[Double] cannot be checked at runtime
3 |   case seqs: Seq[String] => ("seq string", seqs)
  |       ^^^^^^^^^^^^^^^^^
  |       the type test for Seq[String] cannot be checked at runtime
```

When Scala code runs on the JVM these warnings result from the JVM's *type erasure*, a historical legacy of Java's introduction of

*generics* in Java 5, long ago. In order to avoid breaking older code, the JVM byte code doesn't retain information about the actual type parameters that were used for instances of generic (parameterized) types, like `Seq[T]`.

So, the compiler is warning us that, while it can check that a given object is a `Seq`, it can't check at *runtime* that it's a `Seq[Double]` or a `Seq[String]`.

In fact, the second case clause for `Seq[String]` is effectively unreachable. Note the output for results. It shows that "seq double" was written for both inputs, even the `Seq[String]`!

One ugly, but effective workaround is to match on the collection first, then use a nested match on the head element to determine the type. We now have to handle an empty sequence, too:

```
// src/script/scala/progscala3/patternmatching/MatchTypesFix.scala

def doSeqMatch[T](seq: Seq[T]): String = seq match
  case Nil => "Nothing"
  case head +: _ => head match
    case _ : Double => "Double"
    case _ : String => "String"
    case _ => "Unmatched seq element"

val results = Seq(Seq(5.5, 5.6, 5.7), Seq("a", "b"), Nil) map {
  case seq: Seq[_] => (s"seq ${doSeqMatch(seq)}", seq)
}

assert(results == Seq(
  ("seq Double", Seq(5.5, 5.6, 5.7)),
  ("seq String", Seq("a", "b")),
  ("seq Nothing", Seq())))
```

## Sealed Hierarchies and Exhaustive Matches

Let's revisit the need for exhaustive matches and consider the situation where we have an enum or the equivalent sealed class



hierarchy. As an example, suppose we define the following code to represent the allowed “methods” for HTTP:

```
// src/script/scala/progscala3/patternmatching/HTTPMethods.scala

enum HttpMethod:                                     ❶
  def body: String                                   ❷

  case Connect(body: String)                         ❸
  case Delete (body: String)
  case Get     (body: String)
  case Head    (body: String)
  case Options (body: String)
  case Post    (body: String)
  case Put     (body: String)
  case Trace   (body: String)

import HttpMethod._
def handle (method: HttpMethod) = method match        ❹
  case Connect (body) => s"Connect: length = ${body.length}, body =
$body"
  case Delete   (body) => s"Delete:  length = ${body.length}, body =
$body"
  case Get      (body) => s"Get:      length = ${body.length}, body =
$body"
  case Head     (body) => s"Head:     length = ${body.length}, body =
$body"
  case Options  (body) => s"Options:  length = ${body.length}, body =
$body"
  case Post     (body) => s"Post:     length = ${body.length}, body =
$body"
  case Put      (body) => s"Put:      length = ${body.length}, body =
$body"
  case Trace    (body) => s"Trace:    length = ${body.length}, body =
$body"

assert(handle(Connect("CONNECT")) == "Connect: length = 7, body =
CONNECT")
assert(handle(Delete ("DELETE"))  == "Delete:  length = 6, body =
DELETE")
assert(handle(Get      ("GET"))    == "Get:      length = 3, body =
GET")
assert(handle(Head     ("HEAD"))   == "Head:     length = 4, body =
HEAD")
```

```
assert(handle(Options("OPTIONS")) == "Options: length = 7, body =  
OPTIONS")  
assert(handle(Post  ("POST"))    == "Post:    length = 4, body =  
POST")  
assert(handle(Put   ("PUT"))      == "Put:     length = 3, body =  
PUT")  
assert(handle(Trace ("TRACE"))    == "Trace:   length = 5, body =  
TRACE")
```

- ❶ Define an enum `HttpMethod`, the equivalent of a sealed class hierarchy. Only the values defined in this file are allowed `HttpMethod` members. Similarly, for a sealed base type, all the derived types must be defined in the same file.
- ❷ Define a method for the body of the HTTP message.
- ❸ Define eight methods. Note that each declares a constructor parameter `body: String`, which is a `val` because each of these types is effectively a case class. This `val` *implements* the abstract `def` method in `HttpMethod`.
- ❹ An *exhaustive* pattern-match expression, even though we don't have a default clause, because the method argument can only be an instance of one of the eight cases we defined.

### TIP

When pattern matching on an instance of an enum or a sealed base class, the match is exhaustive if the case clauses cover all the derived types defined in the same source file. Because no user-defined derived types are allowed, the match can never become non-exhaustive as the project evolves, since users are prevented from defining new types.

A corollary is to avoid using `sealed` if the type hierarchy is at all likely to change (see *later chapter*). Instead, rely on your traditional object-oriented inheritance principles, including polymorphic methods. What

if you added a new derived type, either in this file or in another file, and you removed the `sealed` keyword on `HttpMethod`? You would have to find and fix all pattern-match clauses in your code base *and* all your users would have to fix their code.

As a side note, we are exploiting a useful feature for implementing certain methods. An abstract, no-parameter method declaration in a parent type can be implemented by a `val` in a subtype. This is because a `val` has a single, fixed value (of course), whereas a no-parameter method returning the same type can return any value of the type. Hence, the `val` implementation is more restrictive in the return type, which means using it where the method is “called” is always just as safe as calling a method. Why? Because clients of the base-type method have to account for any value being returned, but any one concrete implementation only returns one value. In fact, this is an application of *referential transparency*, where we are substituting a value for an expression that *should* always return the same value!

### TIP

When declaring an abstract field in a parent type, use a no-parameter method declaration instead. This allows concrete implementations to choose whether to use a `val` or a method to implement it.

## Chaining Match Expressions

Scala 3 changed the parsing rules for match expressions to allow chaining, as in this contrived example:

```
// src/script/scala/progscala3/patternmatching/MatchChaining.scala
```

```
scala> for opt <- Seq(Some(1), None)
      | yield opt match {
      |   case None => ""
      |   case Some(i) => i.toString
```

```

| } match {
|   case "" => false
|   case _ => true
| }
val res10: Seq[Boolean] = List(true, false)

```

## Pattern Matching in Other Contexts

Fortunately, this powerful feature is not limited to match expressions. You can use pattern matching in assignment statements, called *pattern bindings*:

```

// src/script/scala/progscala3/patternmatching/OtherUses1.scala

scala> case class Address(street: String, city: String, country:
String)
scala> case class Person(name: String, age: Int, address: Address)

scala> val addr = Address("1 Scala Way", "CA", "USA")
scala> val pers = Person("Dean", 29, addr)
val addr: Address = Address(1 Scala Way,CA,USA)
val pers: Person = Person(Dean,29,Address(1 Scala Way,CA,USA))

scala> val Person(name, age, Address(_, state, _)) = pers
val name: String = Dean
val age: Int = 29
val state: String = CA

```

This works with sequences, too:

```

scala> val seq = 0 to 4
val seq: scala.collection.immutable.Range.Inclusive = Range 0 to 4
scala> val head1 +: head2 +: tail = seq
val head1: Int = 0
val head2: Int = 1
val tail: IndexedSeq[Int] = Range 2 to 4

```

They work in for comprehensions:

```

scala> val people = seq.map {
|   i => Person(s"Name$i", 10+i, Address(s"$i Main Street", "CA",

```

```

"USA"))
  | }
val people: IndexedSeq[Person] =
Vector(Person(Name0,10,Address(...)), ...)

scala> val na = for
  |   Person(name, age, address) <- people
  | yield (name, age)
val na: IndexedSeq[(String, Int)] =
Vector((Name0,10), (Name1,11), (Name2,12), (Name3,13), (Name4,14))

```

Suppose we have a function that takes a sequence of doubles and returns the count, sum, average, minimum value, and maximum value in a tuple:

```

// src/script/scala/progscala3/patternmatching/OtherUsesTuples.scala

/** Return the count, sum, average, minimum value, and maximum value.
 */
def stats(seq: Seq[Double]): (Int, Double, Double, Double, Double) =
  assert(seq.size > 0)
  val sum = seq.sum
  (seq.size, sum, sum/seq.size, seq.min, seq.max)

val (count, sum, avg, min, max) = stats((0 until 100).map(_.toDouble))

```

Finally, we can use pattern matching on a regular expression to decompose a string. Here's an example extracted from tests I once wrote for parsing (simple!) SQL strings:

```

// src/script/scala/progscala3/patternmatching/RegexAssignments.scala

scala> val c = """\*|[\w, ]+"" // cols
      | val t = """\w+"" // table
      | val o = """.*"" // other substrings
      | val selectRE =
      |   s""SELECT\s*(DISTINCT)?\s+(\$c)\s*FROM\s+(\$t)\s*
($o)?;""
val distinct: String = DISTINCT

```

```
val cols: String = "col1, col2 "  
val table: String = atable  
val otherClauses: String = WHERE col1 = 'foo'
```

Note that I had to add extra backslashes, e.g., `\\s` instead of `\s`, in the regular expression string, because I used string interpolation. See the source file for other examples.

Obviously, using regular expressions to parse complex text, like XML or programming languages, has its limits. Beyond simple cases, consider a parser library, like the ones we'll discuss in *Chapter 20*

## Problems in Pattern Bindings

In general, keep in mind that pattern matching will throw `MatchError` exceptions when the match fails. This can make your code fragile. For example:

```
//  
src/script/scala/progscala3/patternmatching/FragileAssignments.scala  
  
scala> val h4a :+: h4b :+: t4 = Seq(1,2,3,4)  
val h4a: Int = 1  
val h4b: Int = 2  
val t4: Seq[Int] = List(3, 4)  
  
scala> val h2a :+: h2b :+: t2 = Seq(1,2)  
val h2a: Int = 1  
val h2b: Int = 2  
val t2: Seq[Int] = List()  
  
scala> val h1a :+: h1b :+: t1 = Seq(1)  
scala.MatchError: List(1) (of class  
scala.collection.immutable.$colon$colon)  
...
```

In fact, if you use the `-strict` flag, expressions like the last one will fail to compile, even though Scala 2 allowed them. In Scala 3.1, this expression or any expression of the following form will fail to compile:

```
val head :+: tail = someSequence // Error in Scala 3.1, but not 3.0
```

Hence, while Scala 3.0 still allows some expressions like this, subsequent Scala 3 versions will tighten the type-checking rules for pattern bindings, closing some loopholes in Scala 2 that would parse successfully, but fail with type errors at runtime, like the last several examples.

If you *know* that `someSequence` has at least one element, you can use the `@unchecked` annotation to allow the binding:

```
val head :+: tail : @unchecked = someSequence // Compiles without error
```

Note how the annotation is used in the type location, because it controls how Scala types the binding.

Consider this final example of a `for` comprehension:

```
scala> val elems: Seq[Any] = Seq((1, 2), "hello", (3, 4))
scala> val what = for ((x, y) <- elems) yield (y, x)
val what: Seq[(Any, Any)] = List((2,1), (4,3))
```

In Scala 2 and 3.0, the `elems` sequence is filtered to retain only the elements of tuple type that match the pattern `(x, y)`. In Scala 3.1, this code will trigger a compile-time error. If filtering is the behavior you want in Scala 3.1 and later, use a `case` expression:

```
scala> val what = for case (x, y) <- elems yield (y, x)
val what: Seq[(Any, Any)] = List((2,1), (4,3))
```

## Extractors

So, how does pattern matching and destructuring or extraction work? Scala defines a pair of object methods that are implemented automatically for case classes and also for many types in the Scala library. You can implement these extractors yourself to customize the

behavior for your types. I'll start with how this process typical works in Scala 2, then discuss how it has been generalized in recent releases of Scala, including Scala 3.

Since you rarely need to implement your own extractors, you can safely skip to “[Concluding Remarks on Pattern Matching](#)” near the end of the chapter and return to this discussion later, as needed.

## unapply Method

Recall that the companion object for a case class has at least one factory method named `apply`, which is used for construction. Using “symmetry” arguments, we might infer that there must be another method generated called `unapply`, which is used for extraction. Indeed there is such an *extractor* method and it is invoked in pattern-match expressions for most types.

Consider again `Person` and `Address` from before:

```
person match {  
  case Person(name, age, Address(street, city)) => ...  
  ...  
}
```

Scala looks for `Person.unapply(...)` and `Address.unapply(...)` and calls them. All Scala 2-compatible `unapply` methods return an `Option[(...)]`, where the tuple type corresponds to the number of values and their types that can be extracted from the object, three for `Person` (of types `String`, `Int`, and `Address`) and two for `Address` (both of types `String`). So, the `Person` companion object that the Scala 2 compiler generates looks like this:

```
object Person {  
  def apply(name: String, age: Int, address: Address) =  
    new Person(name, age, address)  
  def unapply(p: Person): Option[(String, Int, Address)] =  
    Some((p.name, p.age, p.address))  
}
```



Why is an `Option` used, if the compiler already knows that the object is a `Person`? Scala allows an implementation of `unapply` to “veto” the match for some reason and return `None`, in which case Scala will attempt to use the next case clause. Also, we don’t have to expose all fields of the instance if we don’t want to. We could suppress our age, if we’re embarrassed by it. We’ll explore the details in “[unapplySeq Method](#)”, but for now, just note that the extracted fields are returned in a `Some` wrapping a three-element tuple. The compiler then extracts those tuple elements for comparison with literal values, when used, assignment to variables, or they are dropped for `_` placeholders.

The `unapply` methods are invoked recursively when necessary. Here we process the nested `Address` object first, then `Person`.

Recall the head `+: tail` expression we used above. Now let’s understand how it actually works. We’ve seen that the `+: (cons)` operator can be used to construct a new sequence by prepending an element to an existing sequence, and we can construct an entire sequence from scratch this way:

```
val list = 1 +: 2 +: 3 +: 4 +: Nil
```

Because `+:` is a method that binds to the right, we first prepend 4 to `Nil`, then prepend 3 to that list, and so forth.

If the construction of sequences is done with a method named `+:`, how can extraction be done with the same syntax, so that we have uniform syntax for *construction* and *destruction/extraction*?

To do that, the Scala library defines a special singleton object named `+:.`. Yes, that’s the name. Like methods, types can have names with a wide variety of characters.

It has just one method, the `unapply` method the compiler needs for our extraction case statement. The declaration of `unapply` is schematically as follows:<sup>2</sup>

```
def unapply[T, Coll](collection: Coll): Option[(T, Coll)]
```

The head is of type T, which is inferred, and some collection type Coll, which represents the type of the input collection and the output tail collection. So, an Option of a two-element tuple with the head and tail is returned.

How can the compiler see the expression `case head +: tail => ...` and use a method `+:.unapply(collection)`? We might expect that the case clause would have to be written `case +:(head, tail) => ...` to work consistently with the behavior we just examined for pattern matching with Person, Address, and tuples.

As a matter of fact, we can write it that way:

```
scala> def seqToString2[T](seq: Seq[T]): String = seq match
      |   case +:(head, tail) => s"($head +: ${seqToString2(tail)})"
      |   case Nil => "Nil"
def seqToString2[T](seq: Seq[T]): String

scala> seqToString2(Seq(1,2,3,4))
val res0: String = (1 +: (2 +: (3 +: (4 +: Nil))))
```

But we can also use *infix* notation, `head +: tail`, because the compiler exploits another bit of syntactic sugar. Types with two type parameters can be written with infix notation and so can case clauses:

```
// src/script/scala/progscala3/patternmatching/Infix.scala

case class With[A,B](a: A, b: B)

val with1: With[String,Int] = With("Foo", 1)
val with2: String With Int = With("Bar", 2)
// val with3: String With Int = "Baz" With 3 // ERROR

val results = Seq(with1, with2).map {
  case s With i => s"$s with $i"
}
assert(results == Seq("Foo with 1", "Bar with 2"))
```

For completeness, there is an analog of `+`: that can be used to process the sequence elements in reverse, `:+`:

```
// src/script/scala/progscala3/patternmatching/MatchReverseSeq.scala
// Compare to match-seq.sc

val nonEmptySeq = Seq(1, 2, 3, 4, 5)

def reverseSeqToString[T](l: Seq[T]): String = l match
  case prefix :+ end => s"(${reverseSeqToString(prefix)} :+ $end)"
  case Nil => "Nil"

assert(reverseSeqToString(nonEmptySeq) ==
  "((((Nil :+ 1) :+ 2) :+ 3) :+ 4) :+ 5)")
```

Note that `Nil` comes first this time.

As before, you could use this output to reconstruct collections (skipping the duplicate second line of the previous output):

```
scala> val revList = (((((Nil :+ 1) :+ 2) :+ 3) :+ 4) :+ 5)
val revList: List[Int] = List(1, 2, 3, 4, 5)
```

The parentheses are unnecessary. Try it!

We mentioned above that you can pattern match pairs with `->`. This feature is implemented with a `val` defined in `Predef`, `->`. This is an alias for `Tuple2`, which subclasses `Product2`, which defines an `unapply` method that is used for these pattern matching expressions.

## Alternatives to Option Return Values

Back to implementation requirements for `unapply`, recent Scala releases relaxed the requirement to return an `Option`. As of Scala 2.11, any type with this signature is allowed (which `Option` also implements):

```
def isEmpty: Boolean
def get: T
```

Scala 3 also allows a Boolean to be returned or a Product type, which is a parent class of Tuples, but more broad than that. Here's an example using Boolean, where we want to discriminate between two kinds of strings:

```
// src/script/scala/progscala3/patternmatching/UnapplyBoolean.scala

object ScalaSearch:
  ❶ def unapply(s: String): Boolean = s.toLowerCase.contains("scala")

  val books = Seq(
    "Programming Scala",
    "JavaScript: The Good Parts",
    "Scala Cookbook").zipWithIndex // add an "index"

  val result = for s <- books yield s match
    ❷ case (_ @ ScalaSearch(), index) => s"$index: found Scala"
      case (_, index) => s"$index: no Scala"

  assert(result == Seq("0: found Scala", "1: no Scala", "2: found
Scala"))
```

- ❶ Define an object with an unapply method that takes a string, converts to lower case, and returns the result of a predicate; does it contain “scala”?
- ❷ Try it on a list of strings, where the first case match succeeds only when the string contains “scala”.

Other single values can be returned. Here is an example that converts a Scala Map to a Java **HashMap**:

```
//
src/script/scala/progscala3/patternmatching/UnapplySingleValue.scala

import java.util.{HashMap => JHashMap}
```

```

case class JHashMapWrapper[K,V](jmap: JHashMap[K,V])
object JHashMapWrapper:
  def unapply[K,V](map: Map[K,V]): JHashMapWrapper[K,V] =
    val jmap = new JHashMap[K,V]()
    for (k,v) <- map do jmap.put(k, v)
    new JHashMapWrapper(jmap)

```

In action:

```

scala> val map = Map("one" -> 1, "two" -> 2)
      | map match
      |   case JHashMapWrapper(jmap) => jmap
val map: Map[String, Int] = Map(one -> 1, two -> 2)
val res2: java.util.HashMap[String, Int] = {one=1, two=2}

```

However, it's not possible to implement a similar extractor for Java's **HashSet** and combine them into one match expression (because there are two, not one possible return values):

```

//
src/script/scala/progscala3/patternmatching/UnapplySingleValue2.scala
scala> ...
scala> val map = Map("one" -> 1, "two" -> 2)
scala> val set = map.keySet
scala> for x <- Seq(map, set) yield x match
      |   case JHashMapWrapper(jmap) => jmap
      |   case JHashSetWrapper(jset) => jset // Not shown; see the
file
... errors ...

```

Anyway, the Scala collections already have tools for converting between Scala and Java collections (see *later chapter*).

Another option for `unapply` is to return a **Product**, or more specifically an object that mixes in this trait, which is an abstraction for types when it is useful to treat the member fields uniformly, such as retrieving by an index or iterating over them. Tuples implement `Product`. We can use it as a way to provide several return values extracted by `unapply`:

```
// src/script/scala/progscala3/patternmatching/UnapplyProduct.scala
```

```
class Words(words: Seq[String], index: Int) extends Product:
  ❶
    def _1 = words
    def _2 = index
    ❷

    def canEqual(that: Any): Boolean = ???
    ❸
    def productArity: Int = ???
    def productElement(n: Int): Any = ???

object Words:
    def unapply(si: (String, Int)): Words =
      ❹
        val words = si._1.split("""\W+""").toSeq
        new Words(words, si._2)
        ❺

val books = Seq(
  "Programming Scala",
  "JavaScript: The Good Parts",
  "Scala Cookbook").zipWithIndex // add an "index"

val result = books.map {
  case Words(words, index) => s"$index: count = ${words.size}"
}
assert(result == Seq("0: count = 2", "1: count = 4", "2: count = 2"))
```

- ❶ Now we need a class to instantiate with the results, when a match succeeds. It implements Product
- ❷ Define two methods for retrieving the first and second items. Note the method names are the same as for two-element tuples.
- ❸ The Product trait declares these methods, too, so we have to provide definitions, but we don't need "working" implementations. This is because Product is actually a *marker trait* for our purposes. All we really need is for Words to mixin this type. So we simply invoke The `Predef.???`, which always throws `NotImplementedError`.

- ④ Matches on a tuple of String and Int.
- ⑤ Split the string on runs of whitespace.

## unapplySeq Method

When you want to return a sequence of extracted items, rather than a fixed number of them, use `unapplySeq`. It turns out the `Seq` companion object implements `apply` and `unapplySeq`, but not `unapply`:

```
def apply[A](elems: A*): Seq[A]
final def unapplySeq[A](x: Seq[A]): UnapplySeqWrapper[A]
```

`UnapplySeqWrapper` is a helper class. Recall that `A*` means that `elems` is a variable argument list.

Matching with `unapplySeq` is invoked in this variation of our previous example for `+:,`, where we examine a “sliding window” of pairs of elements at a time:

```
// src/script/scala/progscala3/patternmatching/MatchUnapplySeq.scala
```

```
val nonEmptyList = List(1, 2, 3, 4, 5)
val emptyList    = Nil
val nonEmptyMap   = Map("one" -> 1, "two" -> 2, "three" -> 3)
```

```
// Process pairs
```

```
def windows[T](seq: Seq[T]): String = seq match
```

```
  case Seq(head1, head2, _) =>
```

```
    s"($head1, $head2), " + windows(seq.tail)
```

```
  case Seq(head, t: _) =>
```

```
③    s"($head, _), " + windows(t)
```

```
  case Nil => "Nil"
```

```
val results = Seq(nonEmptyList, emptyList, nonEmptyMap.toSeq) map {
  seq => windows(seq)
}
```

```
assert(results == Seq(
```

①

②

④

```
"(1, 2), (2, 3), (3, 4), (4, 5), (5, _), Nil",
"Nil",
"((one,1), (two,2)), ((two,2), (three,3)), ((three,3), _), Nil"))
```

- ❶ It looks like we're calling `Seq.apply(...)`, but in a match clause, we're actually calling `Seq.unapplySeq`. We grab the first two elements and ignore the rest of the variable argument list with `_*`. Think of the `*` as matching zero to many, like in regular expressions.
- ❷ Format a string with the first two elements, then move the "window" by one (not two) calling `seq.tail`.
- ❸ We also need a match for a one-element sequence, such as near the end, or we won't have exhaustive matching. This time we capture the tail as `t` and use it in the recursive call, although we actually know that this call to `windows(t)` will simply return `Nil`
- ❹ The `Nil` case with temninates the recursion.

You could rewrite the second case statement to skip the final invocation of `windows(t)`, but I left it in to show capturing of the tail using `t: +*`.

We could still use the `+:` matching we saw before, which is more elegant and what I would do:

```
//
src/script/scala/progscala3/patternmatching/MatchWithoutUnapplySeq.scala

val nonEmptyList = List(1, 2, 3, 4, 5)
val emptyList    = Nil
val nonEmptyMap  = Map("one" -> 1, "two" -> 2, "three" -> 3)

// Process pairs
def windows2[T](seq: Seq[T]): String = seq match
  case head1 +: head2 +: _ => s"($head1, $head2), " +
```



```

windows2(seq.tail)
  case head +: tail => s"($head, _), " + windows2(tail)
  case Nil => "Nil"

val results = Seq(nonEmptyList, emptyList, nonEmptyMap.toSeq) map {
  seq => windows2(seq)
}
assert(results == Seq(
  "(1, 2), (2, 3), (3, 4), (4, 5), (5, _), Nil",
  "Nil",
  "((one,1), (two,2)), ((two,2), (three,3)), ((three,3), _), Nil"))

```

Working with sliding windows is actually so useful that Seq gives us two methods to create them:

```

scala> val seq = 0 to 5
val seq: scala.collection.immutable.Range.Inclusive = Range 0 to 5

scala> seq.sliding(2).foreach(println)
ArraySeq(0, 1)
ArraySeq(1, 2)
ArraySeq(2, 3)
ArraySeq(3, 4)

scala> seq.sliding(3,2).foreach(println)
ArraySeq(0, 1, 2)
ArraySeq(2, 3, 4)

```

Both `sliding` methods return an iterator, meaning they are “lazy” and don’t immediately make a copy of the collection, which is desirable for large collections. The second method takes a `stride` argument, which is how many steps to go for the next sliding window. The default is one step. Note that none of sliding windows contain our last element, 5.

## Implementing `unapplySeq`

Let’s implement an `unapplySeq` method adapted from our Words example above. We’ll tokenize the words as before, but also remove all words shorter than a specified value.

```
// src/script/scala/progscala3/patternmatching/UnapplySeq.scala

object Tokenize:
  // def unapplySeq(s: String): Option[Seq[String]] =
  Some(tokenize(s)) ❶
  def unapplySeq(lim_s: (Int,String)): Option[Seq[String]] =
  ❷
    val (limit, s) = lim_s
    if limit > s.length then None
    else
      val seq = tokenize(s).filter(_.length >= limit)
      Some(seq)

  def tokenize(s: String): Seq[String] = s.split("\\W+").toSeq
  ❸

val message = "This is Programming Scala v3"
val limits = Seq(1, 3, 20, 100)

val results = for limit <- limits yield (limit, message) match
  case Tokenize() => s"No words of length >= $limit!"
  case Tokenize(a, b, c, d: _) => s"limit: $limit => $a, $b, $c,
d=$d" ❹
  case x => s"limit: $limit => Tokenize refused! x=$x"

assert(results == Seq(
  "limit: 1 => This, is, Programming, d=ArraySeq(Scala, v3)",
  "limit: 3 => This, Programming, Scala, d=ArraySeq()",
  "No words of length >= 20!",
  "limit: 100 => Tokenize refused! x=(100,This is Programming Scala
v3)")))
```

- ❶ If we didn't match on the `limit` value, this is what the declaration would be.
- ❷ We match on a tuple with the limit for word size and the string of words. If successful, we return `Some(Seq(words))`, where the words are filtered for those of length at least `limit`. We consider it “unsuccessful” and return `None` when the input `limit` is greater than the length of the input string.
- ❸ Split on whitespace.

- ④ Capture the first three words returned and the rest of them as a variable argument list (d).

Try simplifying this example to not do length filtering. Uncomment the line for comment 1 and work from there.

## Concluding Remarks on Pattern Matching

Along with `for` comprehensions, pattern matching makes idiomatic Scala code concise, yet powerful. It provides a “protocol” for extracting data inside data structures in a principled way, one you can control by implementing custom `unapply` and `unapplySeq` methods (“**Extractors**”). These methods let you extract that information while hiding the implementation details. In fact, the information returned by `unapply` might be a transformation of the actual fields in the instance.

When designing pattern-matching statements, be wary of relying on a default case clause. Under what circumstances would “none of the above” be the correct answer? It may indicate that the design should be refined so you know more precisely all the possible matches that might occur. Often the right design uses sealed class hierarchies or enums, especially those that implement algebraic data types.

## Recap and What’s Next

Pattern matching is a hallmark of many functional languages. It is a flexible and concise technique for extracting data from data structures. We saw examples of pattern matching in case clauses and how to use pattern matching in other expressions.

The next chapter discusses a unique, powerful, but controversial feature in Scala, *context abstractions*, formerly known as *implicit*s, which are a set of tools for building intuitive DSLs, reducing

boilerplate, and making APIs both easier to use and more amenable to customization.

- 
- 1 See *The Java Tutorials. Lesson: Regular Expressions*.
  - 2 I have simplified the actual declaration, because we haven't yet covered details about the type system that we need to understand the actual signature.

# Chapter 5. Abstracting over Context: Type Classes and Extension Methods

---

## A NOTE FOR EARLY RELEASE READERS

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

This will be the 5th chapter of the final book. Please note that the GitHub repo will be made active later on.

If you have comments about how we might improve the content and/or examples in this book, or if you notice missing material within this chapter, please reach out to the author at [mcronin@oreilly.com](mailto:mcronin@oreilly.com).

In previous editions of this book, this chapter was titled *Implicits* after the mechanism used to implement many powerful idioms in Scala. Scala 3 begins the migration to new language constructs that emphasize purpose over mechanism, both to make learning and using these idioms easier and to address some shortcomings of the prior implementations. The transition will happen over several 3.X releases of Scala to make it easier, especially for existing code bases. Therefore, I will cover both the Scala 2 and 3 techniques, while emphasizing the latter.

All of these idioms fall under the umbrella *abstracting over context*. We saw a few examples already, such as the `ExecutionContext`

parameters needed in many Future methods, discussed in “**A Taste of Futures**”. We’ll see many more idioms now in this chapter and the next. In all cases, the idea of “context” will be some situation where an extension to a type, a transformation to a new type, or an insertion of values automatically is desired for easier programming. Frankly, in all cases, it would be possible to live without the tools described here, but it would require more work on the user’s part. This raises an important point, though. Make sure you use these tools judiciously; all constructs have pros and cons.

The most sweeping changes introduced in Scala 3 are to the type system and to how we define and use context abstractions. The changes to the latter are designed to make the purpose and application of these abstractions more clear. The underlying implicit mechanism is still there, but it’s now easier to use for specific purposes. Not only do the changes make intention more clear, they also eliminate some boilerplate previously required when using implicits and fix other drawbacks of the Scala 2 idioms.

## Four Changes

If you know Scala 2 implicits, the changes in Scala 3 can be summarized as follows:<sup>1</sup>

### *Given Instances*

Instead of declaring implicit *terms* (i.e. `vals` or methods) to be used to resolve implicit parameters, the new `given` clause specifies how to synthesize the required term from a type. The change deemphasizes the previous distinction where you had to know when to declare an instance vs. a class. Most of the time you will specify that a particular class should be used to satisfy the need for an implicit value and the compiler does the rest.

### *Using Clauses*

Instead of using the keyword `implicit` to declare an *implicit parameter list* for a method, you use the keyword `using`. The different keyword eliminates several ambiguities and it allows a method definition to have more than one implicit parameter list, which are now called *using clauses*.

### *Given Imports*

When you use wild cards in import statements, they no longer import given instances along with everything else. Instead, you use the `given` keyword to explicitly ask for givens to be imported.

### *Implicit Conversions*

For the special case of given terms that are used for implicit conversions between types, they are now declared as given instances of a standard **Conversion** class. All other forms of implicit conversions will be phased out.

This chapter explores the context abstractions for extending types with additional state and behavior using *extension methods* and *type classes*, which are defined using given instances. We'll also cover given imports and implicit conversions. In **Chapter 6**, we'll explore using clauses and the specific idioms they support.

## **Extension Methods**

In Scala 2, if we wanted to simulate adding new methods to existing types, we had to do an implicit conversion to a wrapper type that implements the method. Scala 3 adds extension methods that allow us to extend a type with new methods without conversion. By themselves, extension methods only allow us to add one or more methods, but not new fields for additional state, nor is there a mechanism for implementing a common abstraction. We'll address those limitations when we discuss *type classes*.

But first, why not just modify the original source code? You may not have that option, for example if it's a third-party library. Also, adding too many methods and fields to classes makes them very difficult to maintain. Keep in mind that every modification to an existing type forces users to recompile their code, at least. This is especially annoying if the changes involve functionality they don't even use.

Context abstractions help us avoid the temptation to create types that contain lots of utility methods that are used only occasionally. Our types can avoid *mixing concerns*. For example, if some users want toJSON methods on a hierarchy of types, like our Shapes in “[A Sample Application](#)”, then only those users are affected.

Hence, the goal is to enable ad hoc additions to types in a principled way. By *principled*, type implementations can remain focused on their core abstractions, while additional behaviors can be added only where needed, as opposed to making global modifications that affect all users. These tools also preserve type safety.

However, a drawback of this *separation of concerns* is that the separate toJSON functionality needs to track changes in the code for the type hierarchy. If a field is renamed, the compiler will catch it for us. If a new field is added, for example Shape.color, it will be easy to miss.

Let's explore an example. Recall that we used the pair construction idiom, `a -> b`, to create tuples `(a, b)`, which is popular for creating Map instances:

```
val map = Map("one" -> 1, "two" -> 2)
```

In Scala 2, this is done using an *implicit conversion* to a library type ArrowAssoc in `Predef` (some details omitted for simplicity):

```
implicit final class ArrowAssoc[A](private val self: A) {  
  @infix def -> [B](y: B): (A, B) = (self, y)  
}
```



Here is how implicit conversion works in Scala 2. When the compiler sees the expression "one" -> 1, it sees that String does not have the -> method. However, ArrowAssoc[T] is in scope, it has this method, *and* the class is declared implicit. So, the compiler can emit code to create an instance of ArrowAssoc[String], with the string "one" passed as the self argument, followed by code to call ->(1) to construct and return the tuple ("one", 1).

If ArrowAssoc were not declared implicit, the compiler would not attempt to use it for this purpose.

Let's re-implement this using a Scala 3 extension method. To avoid ambiguity with ->, let's use ~> instead, but it works identically:

```
// src/script/scala/progscala3/contexts/ArrowAssocExtension.scala
```

```
scala> import scala.annotation.{alpha, infix} ❶
```

```
scala> extension [A,B] (a: A): ❷  
  |   @infix @alpha("arrow2") def ~>(b: B): (A, B) = (a, b)  
def extension_~>[A, B](a: A)(b: B): (A, B)
```

```
scala> "one" ~> 1  
val res0: (String, Int) = (one,1)
```

- ❶ Use the @infix annotation to allow *infix operator notation*, i.e., "one" ~> 1. Use @alpha to define an alphanumeric name in byte code for this method.
- ❷ The syntax for defining an extension method. Any type parameters used by the methods that follow must go after the keyword extension. (The whitespace is arbitrary.) Next we define the method ~>, like we would if this were a “regular” type.

Note the signature the compiler reports for the generated method. It is named extension\_~> and it takes two parameter lists. The first one is for the target instance of type A being extended. The second list is

the same one specified when we declared the `~>` method. We'll see this naming convention again for anonymous given instances below.

Now, when the compiler sees `"one" ~> 1`, it will look for a corresponding `extension_~>` that is in scope and type compatible in the left-hand and right-hand types. Our definition satisfies this requirement for all types. Then the compiler will emit code to call the extension method. No wrapping in a new instance is required. Hence, implicit conversion is eliminated.

Let's complete an example we started in “Operator Overloading?”, where we showed that parameterized types with two parameters can be written with infix notation, but at the time, we didn't know how to support using the same type name as an *operator* for constructing instances. Specifically, we defined a type `!!` allowing declarations like `Int !! String`, but we couldn't define a value of this type using the same literal syntax, for example, `2 !! "two"`. Now we can do this by defining an *extension method* `!!` as follows:

```
// src/script/scala/progscala3/contexts/InfixTypeRevisited.scala
```

```
scala> import scala.annotation.{alpha, infix}
```

```
scala> @alpha("BangBang") case class !![A,B](a: A, b: B) ❶
```

```
scala> extension [A,B] (a: A) def !!(b: B): A !! B = !!(a, b) ❷  
def extension_!![A, B](a: A)(b: B): A !! B
```

```
scala> val ab1: Int !! String = 1 !! "one" ❸  
      | val ab2: Int !! String = !!(1, "one")  
val ab1: Int !! String = !!(1,one)  
val ab2: Int !! String = !!(1,one)
```

- ❶ The same case class defined in “Operator Overloading?”.
- ❷ The extension method definition. When only one method is defined, you can omit the colon (or curly braces) and even define it on the same line as shown.

- ③ This line failed to compile before, but now the extension method is applied to `Int` and invoked with the `String` argument `"one"`.

### TIP

When defining just one extension method for a type, the colon at the end of the opening line or curly braces can be omitted. For consistency and easier reading, consider always using the colon or braces.

Both of our examples did not need to add additional fields to the target type nor was there an interface that made sense to implement. When those are required, we'll use *type classes*, discussed next.

There was a “context” for both extensions. Users only need `->` or `!!` in certain, limited circumstances. These would not be good methods to add to the source code for *all* types! With extension methods, we get the best of both worlds, calling “methods” like `->` when we need them, while keeping types as focused as possible.

So far we have extended classes. What about extension methods on objects? An object can be thought of as a *singleton*. To get its type, use `Foo.type`:

```
scala> object Foo:
      |   def one: Int = 1

scala> extension (foo: Foo.type) def two: String = "two"
def extension_two(foo: Foo.type): String

scala> Foo.one
      |   Foo.two
val res0: Int = 1
val res1: String = two
```

## Type Classes

The next step beyond extension methods is to add not only methods to types, but also state (fields) and to implement an abstraction, so all type extensions are done uniformly. A term that is popular for these kinds of extensions is *type classes*, which comes from the Haskell language, where this idea was pioneered. The word *class* in this context is not the same as Scala's OOP concept of classes, which can be confusing.

As an example, suppose we have a collection of Shapes from “[A Sample Application](#)” and we want the ability to call a `toJSON` method on them that returns a JSON representation appropriate for each type? If we write `someShape.toJSON`, We want the Scala compiler to invoke some mechanism that implements this functionality.

## Scala 3 Type Classes

A type class defines an abstraction with optional state (fields) and behavior (methods). They provide a another way to implement *mixin composition* (“[Traits: Interfaces and “Mixins” in Scala](#)”). The abstraction is valuable for ensuring that all “instances” of the type class follow the same protocol uniformly.

First, we need a trait for the state and behavior we want to add. To keep things simple, we'll return JSON-formatted strings, not objects from some JSON library (of which there are many...):

```
// src/main/scala/progscala3/contexts/json/ToJSON.scala
package progscala3.contexts.json

trait ToJSON[T]:
  extension (t: T) def toJSON(name: String, level: Int): String

  protected val INDENTATION = "  "
  protected def indentation(level: Int): (String,String) =
    (INDENTATION * level, INDENTATION * (level+1))
```

This is the Scala 3 type class pattern. We define a trait with a type parameter and we define extension methods. Although we don't

actually use the type parameter in the body of this particular type class, the parameter will be essential for disambiguating one type class instance, such as the one for `Circle`, from another, such as the one for `Rectangle`.

The public method users care about is `toJSON`. The protected method, `indentation`, and immutable state, `INDENTATION`, are implementation details.

The type class pattern solves the limitation discussed above for extension methods alone. We can define and implement an abstraction and we can add arbitrary state as fields.

Now we create instances for our Shapes:

```
//
src/main/scala/progscala3/contexts/typeclass/new1/ToJSONTypeClasses.sc
ala
package progscala3.contexts.typeclass.new1

import progscala3.introscala.shapes.{Point, Shape, Circle,
Rectangle, Triangle}
import progscala3.contexts.json.ToJSON

given ToJSON[Point]:                                     ❶
  extension (point: Point) def toJSON(name: String, level: Int):
String =
  val (outdent, indent) = indentation(level)
  s""""$name": {
    |${indent}"x": "${point.x}",
    |${indent}"y": "${point.y}"
    |$outdent}""".stripMargin

given ToJSON[Circle]:                                    ❷
  extension (circle: Circle) def toJSON(name: String, level: Int):
String =
  val (outdent, indent) = indentation(level)
  s""""$name": {
    |${indent}${circle.center.toJSON("center", level + 1)},
    |${indent}"radius": ${circle.radius}
    |$outdent}""".stripMargin

// And similarly for Rectangle and Triangle
```

```
@main def TryJSONTypeClasses() =
  println(Circle(Point(1.0,2.0), 1.0).toJSON("circle", 0))
  println(Rectangle(Point(2.0,3.0), 2, 5).toJSON("rectangle", 0))
  println(Triangle(
    Point(0.0,0.0), Point(2.0,0.0), Point(1.0,2.0)).toJSON("triangle",
    0))
```

- ❶ The given keyword declares a conversion for ToJSON[Point]. The extension method for ToJSON is implemented as required for point
- ❷ A given for ToJSON[Circle].

Running TryJSONTypeClasses prints the following:

```
> runMain progscala3.contexts.typeclass.new1.TryJSONTypeClasses
...
"circle": {
  "center": {
    "x": "1.0",
    "y": "2.0"
  },
  "radius": 1.0
}
...
```

### NOTE

If you use braces, the colon on the given line is replaced with an opening curly brace and a corresponding closing brace is required. The same applies for the definition of the anonymous ToJSON subtype, of course.

There's a flaw with our implementation, though. If we put those shapes in a sequence, shapes, and try shapes.foreach(s => println(s.toJSON("shape", 0))), we get an error that Shape doesn't have a toJSON method. Polymorphic dispatch doesn't work here.

What if we add a given for Shape that delegates to the others?

```
given ToJSON[Shape]:
  extension (shape: Shape) def toJSON(name: String, level: Int):
String =
  shape.toJSON(name, level)
```

Seems legit, but the compiler says we have an “infinite recursion”. Again, we aren’t actually calling a polymorphic method toJSON defined in the old-fashioned way for the hierarchy. So, the call `shape.toJSON(level)` attempts to call the extension method recursively.

What about pattern matching on the type of Shape?

```
given ToJSON[Shape]:
  extension (shape: Shape) def toJSON(name: String, level: Int):
String =
  shape match
    case c: Circle    => c.toJSON(name, level)
    case r: Rectangle => r.toJSON(name, level)
    case t: Triangle  => t.toJSON(name, level)
```

We still get an infinite recursion at runtime, but the compiler can’t detect it! So, instead, let’s call the compiler generated toJSON implementations directly. A synthesized object is output for each specific given:

```
//
src/main/scala/progscala3/contexts/typeclass/new2/ToJSONTypeClasses.sc
ala
...
given ToJSON[Shape]:
  extension (shape: Shape) def toJSON(name: String, level: Int):
String =
  shape match
    case c: Circle    =>
      given_ToJSON_Circle.extension_toJSON(c)(name, level)
    case r: Rectangle =>
      given_ToJSON_Rectangle.extension_toJSON(r)(name, level)
    case t: Triangle  =>
      given_ToJSON_Triangle.extension_toJSON(t)(name, level)
  ...
```

```

@main def TryJSONTypeClasses() =
  val c = Circle(Point(1.0,2.0), 1.0)
  val r = Rectangle(Point(2.0,3.0), 2, 5)
  val t = Triangle(Point(0.0,0.0), Point(2.0,0.0), Point(1.0,2.0))
  println("=== Use shape.toJSON:")
  Seq(c, r, t).foreach(s => println(s.toJSON("shape", 0)))
  println("=== call toJSON on each shape explicitly:")
  println(c.toJSON("circle", 0))
  println(r.toJSON("rectangle", 0))
  println(t.toJSON("triangle", 0))

```

The output of `...typeclass.new2.TryJSONTypeClasses` (not shown) indicates that calling `shape.toJSON` and `circle.toJSON` (for example) now both work as desired, but we are relying on an obscure implementation detail that could change in a future release of the compiler. Note the naming conventions, `given_...` and `extension_...`, which we saw previously.

There is an easy fix. We can give names to the givens and then call them:

```

//
src/main/scala/progscala3/contexts/typeclass/new3/ToJSONTypeClasses.sc
ala
...
given ToJSON[Shape]:
  extension (shape: Shape) def toJSON(name: String, level: Int):
String =
  shape match
    case c: Circle    => circleToJSON.extension_toJSON(c)(name,
level)
    case r: Rectangle => rectangleToJSON.extension_toJSON(r)(name,
level)
    case t: Triangle  => triangleToJSON.extension_toJSON(t)(name,
level)

given circleToJSON as ToJSON[Circle]:
  ...

```

Note the `as` keyword after the name. Instead of the synthesized name `given_ToJSON_Circle`, it will be `circleToJSON`, and similarly for



the other shapes. We must still use the synthesized method name for the extension method, but at least we have more control over the object name.

### NOTE

When used as shown, `as` is a “soft” reserved word. You can still use it as an identifier elsewhere. However, `given` is always reserved.

At this point, if we still want to simulate polymorphic behavior, consider refactoring the code to move the implementations of the `toJSON` methods for each concrete `Shape` to a regular helper object. Then call its methods instead of using the compiler generated objects. See `...typeclass.new4.TryJSONTypeClasses` in the code examples for one approach.

### TIP

Keep the `given` instances as concise as possible. Consider moving some of the code to “regular” types that operate as helpers.

Finally, note that `Point` and the concrete subtypes of `Shape` are not related in the type hierarchy (well, except for `AnyRef` way at the top). Hence, this extension mechanism is *ad hoc polymorphism*, because the polymorphic behavior of `toJSON` is not tied to the type system, as it would be in *subtype polymorphism*. Subtype polymorphism is nice for allowing parent types to declare behaviors that can be defined in subtypes. We had to hack around this missing feature for `toJSON`! For completeness, recall that we discussed a third kind of polymorphism, *parametric polymorphism*, in “[Abstract Types Versus Parameterized Types](#)”, where containers like `Container[A]` behave uniformly for any type `A`.

Sometimes a type class will define members that make more sense as the analogs of *companion object* members, rather than *instance* members. To see this, let's look at type classes for *Semigroup* and *Monoid*. *Semigroup* generalizes the notion of addition or composition. You know how addition works for numbers, and even strings can be “added”. *Monoid* adds the idea of a “unit” value. If you add zero to a number, you get the number back. If you add a string to an empty string, you get the first string back.

Here are the definitions for these types:

```
//tag::definitions[]
// src/script/scala/progscala3/contexts/MonoidTypeClass.scala

trait SemiGroup[T]:
  extension (t: T):
    def combine(other: T): T
    ❶ def <+>(other: T): T = t.combine(other)

trait Monoid[T] extends SemiGroup[T]:
  def unit: T
  ❷

given StringMonoid as Monoid[String]:
  ❸ def unit: String = ""
  extension (s: String) def combine(other: String): String = s +
  other

given IntMonoid as Monoid[Int]:
  def unit: Int = 0
  extension (i: Int) def combine(other: Int): Int = i + other
//end::definitions[]

//tag::usage[]
"2" <+> ("3" <+> "4")           // "234"
("2" <+> "3") <+> "4"           // "234"
StringMonoid.unit <+> "2"       // "2"
"2" <+> StringMonoid.unit       // "2"

2 <+> (3 <+> 4)                 // 9
(2 <+> 3) <+> 4                 // 9
```

```

IntMonoid.unit <+> 2           // 2
2 <+> IntMonoid.unit          // 2
//end::usage[]

//tag::numericdefinition[]
given NumericMonoid[T](using num: Numeric[T]) as Monoid[T]:
  def unit: T = num.zero
  extension (t: T) def combine(other: T): T = num.plus(t, other)

2.2 <+> (3.3 <+> 4.4)           // 9.9
(2.2 <+> 3.3) <+> 4.4           // 9.9

BigDecimal(3.14) <+> NumericMonoid.unit
NumericMonoid[BigDecimal].unit <+> BigDecimal(3.14)
//tag::numericdefinition[]

```

- ❶ Define an *instance* extension method `combine` and an alternative operator methods `<+>` that calls `combine`. Note that combining one element with another of the same type returns a new element of the same type, like adding numbers. For given instances of the type class, we will only need to define `combine`.
- ❷ The definition for `unit`, i.e., zero for addition of numbers. It's not defined as an extension method, but an *object* method, because we only need one instance of the value for all `T`.
- ❸ Instances of the type class for `String` and `Int`. Note how `unit` and `combine` are defined.

The Monoid combine operation is associative, so here are examples of both instances in action.

Notice how each `unit` is referenced. This is why we gave these given instances explicit names, so it's easier to remember what to call them, instead of the default `given_Monoid_String`, etc.

Finally, we don't actually need to define separate instances for each numeric type. Here is how to implement it once for a `T` in

**Numeric**:left\_sb: [T:right\_sb: ]:

When accessing `NumericMonoid[BigDecimal].unit`, the type parameter could be inferred when it was used as the argument for the extension method `<+>`. However, when it was used as the *instance to be extended* with `<+>`, the type parameter was required.

## Scala 2 Type Classes

To implement the same type class and instances in Scala 2 syntax (which is still allowed), you write an implicit conversion that wraps the `Point` and `Shape` instances in a new instance of a type that has the `toJSON` method, then call the method.

First, we need a slightly different `ToJSON` trait, because the extension method code in `ToJSON` won't work with Scala 2:

```
//  
src/main/scala/progscala3/contexts/typeclass/old/ToJSONOldTypeClasses.  
scala  
package progscala3.contexts.typeclass.old  
  
import scala.language.implicitConversions  
  
trait ToJSONOld[T]:  
  def toJSON(level: Int): String  
  
  protected val INDENTATION = "  "  
  protected def indentation(level: Int): (String, String) =  
    (INDENTATION * level, INDENTATION * (level+1))
```

- ❶ We must enable implicit conversions.
- ❷ Now this is a regular method. In the previous `ToJSON` implementation, it was an extension method.

Indiscriminate use of implicit conversions can be confusing for code comprehension and it sometimes leads to unexpected behavior.

Therefore, implicit conversions are treated as an optional feature by Scala. This means you must enable the feature explicitly with the import statement used for `implicitConversions` or use the global - `language:implicitConversions` compiler flag.

Now here is an implementation of a `toJSON` type class instances for Scala 2, which also works in Scala 3. We'll only show the implementations for `Point` and `Circle`:

```
implicit final class PointToJSON(  
  point: Point) extends ToJSONOld[Point]:  
  def toJSON(name: String, level: Int): String =  
    val (outdent, indent) = indentation(level)  
    s""""$name": {  
      |${indent}"x": "${point.x}",  
      |${indent}"y": "${point.y}"  
      |$outdent}""".stripMargin  
  
implicit final class CircleToJSON(  
  circle: Circle) extends ToJSONOld[Circle]:  
  def toJSON(name: String, level: Int): String =  
    val (outdent, indent) = indentation(level)  
    s""""$name": {  
      |${indent}${circle.center.toJSON("center", level + 1)},  
      |${indent}"radius": ${circle.radius}  
      |$outdent}""".stripMargin  
...  
  
@main def TryJSONOldTypeClasses() =  
  val c = Circle(Point(1.0,2.0), 1.0)  
  val r = Rectangle(Point(2.0,3.0), 2, 5)  
  val t = Triangle(Point(0.0,0.0), Point(2.0,0.0), Point(1.0,2.0))  
  println(c.toJSON("circle", 0))  
  println(r.toJSON("rectangle", 0))  
  println(t.toJSON("triangle", 0))
```

- ❶ The type class “instance” that implements `ToJSONOld.toJSON` for `Point`. It is called a type class *instance* following Haskell conventions, but we actually declare a *class* for it, which can be confusing.

- ② The type class *instance* that implements `toJSON` for `Circle`.

Because these classes are declared as implicit, when the compiler sees `circle.toJSON()`, for example, it will look for an implicit conversion in scope that returns some wrapper type that has this method.

The output of `TryJSONOldTypeClasses` works as expected. However, we didn't solve the problem of iterating through some `Shapes` and calling `toJSON` polymorphically. You can try that yourself.

We didn't declare our implicit classes as case classes. In fact, Scala doesn't allow an implicit class to also be a case class. It wouldn't make much sense anyway, because the extra, auto-generated code for the case class would never be used. Implicit classes have a very narrow purpose. Similarly, declaring them `final` is recommended to eliminate some potential surprises when the compiler resolves which type classes to use.

If you need to support Scala 2 code for a while, then using the original type class pattern will work for a few versions of Scala 3. However, in most cases, it will be better to migrate to the new type class syntax, because it is more concise and purpose-built, and it doesn't require implicit conversions.

## Implicit Conversions

We saw that an implicit conversion called `ArrowAssoc` was used in the Scala 2 library to implement the `"one" -> 1` idiom, whereas we could use an extension method in Scala 3. We also saw implicit conversions used for type classes in Scala 2, while Scala 3 combines extension methods and `givens` to avoid doing conversions.

Hence, in Scala 3, the need to do implicit conversions is greatly reduced, but it hasn't disappeared completely. Sometimes you will want to convert between types for other reasons. Consider the

following example that defines types to represent Dollars, Percentages, and a person's Salary, where the gross salary and the percentage to deduct for taxes are encapsulated. When constructing a Salary instance, we want to allow users to enter Doubles, for convenience:

```
// tag::definitions[]
// src/main/scala/progscala3/contexts/NewImplicitConversions.scala
package progscala3.contexts
import scala.language.implicitConversions

case class Dollars(amount: Double):
  override def toString = f"$$$amount%.2f"

case class Percentage(amount: Double):
  override def toString = f"${(amount*100.0)}%.2f%"

case class Salary(gross: Dollars, taxes: Percentage):
  def net: Dollars = Dollars(gross.amount * (1.0 - taxes.amount))
// end::definitions[]

// tag::entrypoint[]
@main def TryImplicitConversions() =
  given Conversion[Double,Dollars] = d => Dollars(d)
  given Conversion[Double,Percentage] = d => Percentage(d)

  val salary = Salary(100_000.0, 0.20)
  println(s"salary: $salary. Net pay: ${salary.net}")
// end::entrypoint[]
```

Note that we import `scala.language.implicitConversions`. The Dollars class encapsulates a Double for the amount, with `toString` overridden to return the familiar “\$dollars.cents” output. Similarly, Percentage wraps a Double and overrides `toString`.

Let's try it:

```
// tag::definitions[]
// src/main/scala/progscala3/contexts/NewImplicitConversions.scala
package progscala3.contexts
import scala.language.implicitConversions
```

```

case class Dollars(amount: Double):
  override def toString = f"$$$amount%.2f"

case class Percentage(amount: Double):
  override def toString = f"${(amount*100.0)}%.2f%"

case class Salary(gross: Dollars, taxes: Percentage):
  def net: Dollars = Dollars(gross.amount * (1.0 - taxes.amount))
// end::definitions[]

// tag::entrypoint[]
@main def TryImplicitConversions() =
  given Conversion[Double,Dollars] = d => Dollars(d)
  given Conversion[Double,Percentage] = d => Percentage(d)

  val salary = Salary(100_000.0, 0.20)
  println(s"salary: $salary. Net pay: ${salary.net}")
// end::entrypoint[]

```

- ❶ The syntax for declaring a given conversion from Double to Dollars and a second conversion from Double to Percentage.

Running this example prints the following:

```
salary: Salary($100000.00,20.00%). Net pay: $80000.00
```

The declaration of the `Conversion[Double,Dollars]` given is shorthand for the following longer form:

```

given Conversion[Double,Dollars]:
  def apply(d: Double): Dollars = Dollars(d)

```

By the way, if you define the given for converting Doubles to Dollars in the REPL, observe what happens:

```

...
scala> given Conversion[Double,Dollars] = d => Dollars(d)
def given_Conversion_Double_Dollars: Conversion[Double, Dollars]

```



```
scala> given dd as Conversion[Double,Dollars] = d => Dollars(d)
def dd: Conversion[Double, Dollars]
```

In our type classes above, objects were generated. Here, methods are generated. As we saw previously, for an anonymous given, the generated name follows the convention given\_....

Scala 3 still supports the Scala 2 mechanism of using an implicit method for conversion:

```
implicit def toDollars(d: Double): Dollars = Dollars(d)
```

## Rules for Implicit Conversion Resolution

Here is a summary of the lookup rules used by the compiler to find and apply conversions. I'll use given and given instances to refer to both new and old style conversions:

1. No conversion will be attempted if the object and method combination type check successfully.
2. Only given instances for conversion are considered.
3. Only given instances in the current scope are considered, as well as givens defined in the *companion object* of the *target* type.
4. Given conversions aren't chained to get from the available type, through intermediate types, to the target type. Only one conversion will be considered.
5. No conversion is attempted if more than one possible conversion could be applied and have the same scope. There must be one and only one, unambiguous possibility.

## Type Class Derivation

TODO: Verify which mixins support derivation at the time Scala 3 is released. Also verify maturity of the implementation For example, derives Eql shouldn't be necessary for enum and case class definitions, but it appears to be necessary. What is the final Scala 3 implementation?

*Type class derivation* is the idea that we should be able to automatically generate type class given instances as long as they obey a minimum set of requirements. A class uses the new keyword derives, which works like extends or with, to trigger derivation.

For example, Scala 3 introduces `scala.Eql`, which restricts use of the comparison operators `==` and `!=` for instances of arbitrary types. Normally, it's allowed to do these comparisons, but when the compiler flag `-language:strictEquality` or the import statement `import scala.language.strictEquality` is used, then the comparison operators are only allowed in certain specific contexts. Here is an example:

```
// src/main/scala/progscala3/contexts/Derivation.scala

package progscala3.contexts
import scala.language.strictEquality

enum Tree[T] derives Eql:
  case Branch(left: Tree[T], right: Tree[T])
  case Leaf(elem: T)

@main def TryDerived() =
  import Tree._
  val l1 = Leaf("l1")
  val l2 = Leaf(2)
  val b = Branch(l1, Branch(Leaf("b1"), Leaf("b2")))
  assert(l1 == l1)
  // assert(l1 != l2)    // Compilation error!
  assert(l1 != b)
  assert(b == b)
  println(s"For String, String:
  ${implicitly[Eql[Tree[String], Tree[String]]}]")
  println(s"For Int, Int: ${implicitly[Eql[Tree[Int], Tree[Int]]]}")
  // Compilation error:
```

```
// println(s"For String, Int:
${implicitly[Eql[Tree[String], Tree[Int]]}]")
```

Because of the `derives Eql` clause in the `Tree` declaration, the equality checks in the assertions are allowed. The `derives Eql` clause has the effect of generating the following given instance:

```
given Eql[Tree[T], Tree[T]] = Eql.derived
```

`Eql.derived` is the following:

```
object Eql:
  object derived extends Eql[Any, Any]
  ...
```

Furthermore, `T` will be constrained to types with `given Eql[T, T] = Eql.derived`. What all this effectively means that we can only compare `Tree[T]` instances for the same `T` types.

The terminology used is `Tree` is the *deriving type* and the `Eql` instance is a *derived instance*.

In general, any type `T` defined with a companion object that has the derived instance or method can be used with `derives T` clauses. We'll discuss how methods are implemented in *later chapters*, after we have learned the metaprogramming details required. The reason `Eql` and the `strictEquality` language feature were introduced is discussed in "[Multiversal Equality](#)".

### TIP

If you want to enforce stricter use of comparison operators, use `language:strictEquality`, but expect to add `derives Eql` to many of your types.

## Givens and Imports

In “[A Taste of Futures](#)” we imported an implicit `ExecutionContext`, `scala.concurrent.ExecutionContext.Implicits.global`. The name of the enclosing object `Implicits` reflects a common convention in Scala 2 for making implicit definitions more explicit in code that uses them, at least if you pay attention to the import statements.

Scala 3 introduces a new way to control imports of givens and implicits, which provides an effective alternative form of visibility, as well as allowing developers to use wild-card imports frequently while restricting if and when givens and implicits are also imported.

Consider the following example adapted from the [Dotty documentation](#):

```
// src/script/scala/progscala3/contexts/GivenImports.scala

object 01:
  val name = "01"
  val m(s: String) = s"$s, hello from $name"
  class C1
  given c1 as C1
  class C2
  given c2 as C2
```

Now consider these import statements:

```
import 01._           // Imports everything EXCEPT the givens, c1
                        and c2
import 01.{given _}    // Imports ONLY the givens, c1 and c2
import 01.{given c1}   // Imports just c1 explicitly
import 01.{given _, _} // Imports everything in 01
```

A given import selector also brings old style implicits into scope.

What if the given instances are anonymous and you don't want to use the wild card?

```
object 02:
  class C1
  given C1
```

```
class C2
given C2
given intOrd as Ordering[Int]
given listOrd[T: Ordering] as Ordering[List[T]]
```

You can import *by type*. Note the ? wild card for the type parameter, which means both Ordering givens will be imported:

```
import 02.{given C1, given Ordering[?]}
```

Because this is a breaking change in how \_ wild cards work for imports, it is being implemented gradually:

- In Scala 3.0 an old-style implicit definition can be brought into scope either by a \_ or a given \_ wildcard selector.
- In Scala 3.1 an old-style implicit accessed through a \_ wildcard import will give a deprecation warning.
- In some version after 3.1, old-style implicits accessed through a \_ wildcard import will give a compiler error.

Finally, the older Scala 2 implicit conversions are still allowed, where an implicit def is used, for example:

```
implicit def doubleToDollars(d: Double): Dollars = Dollars(d)
```

Unlike the Scala 2 alternative to extension methods, we don't need an implicit class here, like ArrowAssoc above, because Dollars has all the methods we need. This method would be invoked to do the conversion exactly the same way as the given Conversion[Double,Dollars] above.

## Resolution Rules for Givens and Extension Methods

Extension methods and given definitions obey the same scoping rules as other declarations, i.e., they must be visible to be considered. The previous examples scoped the extension methods to packages, such as the `new1`, `new2`, etc. packages. They were not visible unless the package contents were imported or we were already in the scope of that package.

Within a particular scope, there could be several candidate givens or extension methods that the compiler might use for a type extension. The [Dotty documentation](#) has the details for Scala 3's resolution rules. I'll summarize the key points here. Givens are also used to resolve implicit parameters in method *using clauses*, which we'll explore in the next chapter. The same resolution rules apply.

I'll use the term "given" in the following discussion to include given instances, extension methods, and Scala 2 implicit methods, values, and classes, depending on the scenario. Resolving to a particular given happens in the following order:

## RULES FOR GIVEN RESOLUTION

1. Any type-compatible given that doesn't require a prefix path. In other words, it is defined in the same scope, such as within the same block of code, within the same type, within its companion object (if any), or within a parent type.
2. A given that was imported into the current scope. It also doesn't require a prefix path to use it.
3. Imported givens take precedence over the already-in-scope givens. That is, when an import brings in new givens.
4. In some cases, several possible matches are type compatible. The most specific match wins. For example, if a `Foo` is needed in a context, then a given in scope of type `Foo` will be chosen over a given of type `AnyRef`, if both are in scope.
5. If two or more candidate givens are ambiguous, for example they have the same exact type, it triggers a compiler error.

The compiler always puts some library givens in scope, while other library givens require an import statement. For example, `Predef` extends a type called `LowPriorityImplicits`, which makes the givens defined in `Predef` lower priority when potential conflicts arise with other givens in scope. The rationale is that the other givens are likely to be user defined or imported from special libraries, and hence more “important” to the user.

## Build Your Own String Interpolator

Let's look at a final example of extension, one that lets us define our own string interpolation capability. Recall from “**Interpolated Strings**” that Scala has several built-in ways to format strings through interpolation. For example:

```
val (first, last) = ("Buck", "Trends")
println(s"Hello, ${first} ${last}")
```

There are also `StringContext` methods named `f` and `raw`, where `f` supports `printf` format directives and `raw` doesn't interpret escape characters:

```
scala> val pi=3.14159

scala> f"$pi%5.3f or ${pi}%7.5f"
val res0: String = 3.142 or 3.14159

scala> raw"\t $pi \n $pi again"
val res1: String = \t 3.14159 \n 3.14159 again
```

We'll look at a simplistic implementation of a SQL query compiler named `sql`. When the compiler sees an expression like `sql"SELECT $column FROM $table;"`, it will be translated to the following:

```
StringContext("SELECT ", "FROM ", ";").sql(column, table)
```

Note how the embedded expressions become arguments to `sql`, while the other string tokens are arguments to `StringContext.apply`. However, `scala.StringContext` doesn't have a `sql` method, so an implicit conversion to another type or an extension method is required.

Let's define `sql` for `StringContext`. For simplicity, it will only handle SQL queries of the form `sql"SELECT columns FROM table;"` with the columns and table strings specified as part of the string or using embedded expressions. The extracted column and table names are returned in an instance of a simple case class `SQLQuery`. It would be



possible to use the same approach with a real SQL parser and return a query object for a library like JDBC. I won't show the whole implementation (which is somewhat of a "hack" for this simple case), but just the declarations:

```
// src/main/scala/progscala3/contexts/SQLStringInterpolator.scala
package progscala3.contexts

object SimpleSQL:
  case class SQLQuery(columns: Seq[String] = Nil, table: String = "")

  extension (sc: StringContext):
    def sql(values: String*): SQLQuery =
      // Extract the column names and table name.
      SQLQuery(columns, table)
```

See the SQLStringInterpolator source code for the full details. Here is how to use it:

```
scala> import progscala3.contexts.SimpleSQL._

scala> val query1 = sql"SELECT one, two, three FROM t1;"
val query1: ...SimpleSQL.SQLQuery = SQLQuery(Vector(one, two,
three),t1)

scala> val cols = Seq("four", "five").mkString(", ")
      | val table = "t2"
      | val query2 = sql"SELECT $cols FROM $table;"
val cols: Seq[String] = List(four, five)
val table: String = t2
val query2: ...SQLQuery = SQLQuery(Vector(four, five),t2)
```

As shown, custom string interpolators can return any type you want, not just a new String, like s, f, and raw return. Hence, they can function as instance factories that are driven by data encapsulated in strings.

## The Expression Problem

Let's step back for a moment and ponder what we just accomplished in the previous example. We added new functionality to an existing library type without editing the source code for it!

This desire to extend modules without modifying their source code is called the *Expression Problem*, a **term coined by Philip Wadler**.

Object-oriented programming solves this problem with subtyping, more precisely called *subtype polymorphism*. We program to abstractions and use derived classes when we need changed behavior. Bertrand Meyer coined the term *Open/Closed Principle* to describe a principled OOP approach, where base types declare the behaviors as abstract that should be open for extension or variation in subtypes, while keeping invariant behaviors closed to modification. The base types are never modified for extension.

Scala certainly supports this technique, but it has drawbacks. What if it's questionable that we should have that behavior defined in the type hierarchy in the first place? What if the behavior is only needed in a few contexts, while for most contexts, it's just a burden that the code carries around?

It can be a burden for several reasons. First, the extra, mostly-unused code is a maintenance burden. Developers have to understand it, even when working on other aspects of the code, so they don't break it inadvertently. Second, it's also inevitable that most defined behaviors will be refined over time. Every change to a feature that some clients aren't using forces unwanted updates on client code.

This problem led to the *Single Responsibility Principle*, a classic design principle that encourages us to define abstractions and implement types with just a single behavior.

Still, in realistic scenarios, it's sometimes necessary for an object to combine several behaviors. For example, a service often needs to "mix in" the ability to log messages. Scala makes these *mixin* features relatively easy to implement, as we saw in **"Traits: Interfaces**

and “Mixins” in Scala”. We can even declare instances that mix in traits without first defining a class.

In general, mixins, extension methods, and type classes provide robust and principled solutions to the Open/Closed Principle while allowing the core implementations of types to obey the Single Responsibility Principle.

## Wise Use of Type Extensions

Why not take an extreme approach and define types with very little state and behavior (sometimes called *anemic* types), then add most behaviors using mixin traits, type classes, extension methods, even implicit conversions?

First, when using a type class or implicit conversion, the resolution algorithm requires more work by the compiler than just finding the logic inside the type’s original definition. Also, there can be more boilerplate writing extensions compared to the alternative of defining constructs inside the type. Therefore, a project that over-uses these tools is a project that is slow to build, as well as potentially hard to comprehend when reading the code.

Another problem for the extension mechanisms explored in this chapter is that you effectively lose several benefits of object orientation!

First, code evolution can be challenging if the extensions depend on details of the types they extend, like our ToJSON example in the last chapter. The details have to be coordinated in more than one place, the locations of the extensions, as well as the type definitions themselves. Fortunately, coupling between type definitions and extensions are limited to the public interfaces, as an extension has no access to private or protected members of a type.

A second issue is the loss of object-oriented method dispatch. We had to do some hacking to support `Shape.toJSON` in a polymorphic

way.

If instead, `toJSON` were declared abstract in `Shape` and implemented in `Circle`, `Rectangle`, etc., then this code would work with the usual object-oriented dispatch rules.

Most of the time, the core domain logic belongs in the type definition. Ancillary behaviors, like serializing to JSON and logging, belong in mixins or type classes. However, if *your* applications use `toJSON` frequently for your domain classes, it might be a good idea to move this behavior into the type definitions, on balance.

When should use use type classes and extension methods vs. *mix-in* composition? For example, recall the Logging trait example we saw in “[Traits: Interfaces and “Mixins” in Scala](#)”. If the trait has *orthogonal* state and behavior, like logging, that can be mixed into many different objects, then a mixin trait is often best. If the behavior has to be defined carefully for each type, like `toJSON`, then type classes are best.

## Recap and What’s Next

We started our exploration of context abstractions in Scala 2 and 3, beginning with tools to extend types with additional state and behavior, such as type classes, extension methods, and implicit conversions.

Part 2 explores *using clauses*, which work with given instances to address particular design scenarios and to simplify user code.

---

<sup>1</sup> Adapted from this [Dotty documentation](#).

# Chapter 6. Abstracting over Context: Using Clauses

---

## A NOTE FOR EARLY RELEASE READERS

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

This will be the 6th chapter of the final book. Please note that the GitHub repo will be made active later on.

If you have comments about how we might improve the content and/or examples in this book, or if you notice missing material within this chapter, please reach out to the author at [mcronin@oreilly.com](mailto:mcronin@oreilly.com).

In [Chapter 5](#), we began our discussion of the powerful tools and idioms in Scala 2 and 3 for abstracting over *context*. In particular, we discussed type classes, extension methods, and implicit conversions as tools for extending the behaviors of existing types.

This chapter explores *using clauses*, which work with given instances to address particular design scenarios and to simplify user code.

## Using Clauses

The other major use of context abstractions is to provide method parameters implicitly rather than explicitly. When a method argument list begins with the keyword `using` (Scala 3) or `implicit` (Scala 2

and 3), the user does not have to provide values explicitly for the parameters, as long as *given instances* (explored in the previous chapter) are in scope that the compiler can use instead.

In Scala 2 terminology, those parameters were called *implicit parameters* and the whole list of parameters is an *implicit parameter list* or *implicit parameter clause*. In Scala 3, they are *context parameters* and the whole parameter list is a *using clause*.<sup>1</sup> Here is an example:

```
class BankAccount(...):  
  def debit(amount: Money)(using transaction: Transaction)  
  ...
```

Here, the *using clause* starts with the `using` keyword and contains the *context parameter* `transaction`.

The values in scope that can be used to fill in these parameters are called *implicit values* in Scala 2. In Scala 3 they are the *given instances* or *givens* for short that we studied last chapter.

I'll mostly use the Scala 3 terminology in this book, but when I use Scala 2 terminology, it will usually be when discussing a Scala 2 library that uses implicit definitions and parameters. Scala 3 more or less treats them interchangeably, although the Scala 2 implicits will be phased out eventually.

For each parameter in a *using clause*, a type-compatible given must exist in the enclosing scope. Using Scala 2-style implicits, an implicit value or an implicit function returning a compatible value must be in scope.

For comparison, recall you can also define default values for method parameters. While sufficient in many circumstances, they are statically scoped to the definition at compile time and they are defined by the implementer of the method. *Using clauses*, on the other hand, provide greater flexibility for users of a method.

As an example, suppose we implement a simple type that wraps sequences for convenient sorting (ignoring the fact this capability is already provided by Seq). One way to do this is for the user to supply an implementation of `math.Ordering`, which knows how to sort elements of the particular type used in the sequence. That object could be passed as an argument to the sort method, but the user might also like the ability to specify the value once, as an implicit, and then have all sequences of the same element type use it automatically.

This first implementation uses syntax valid for both Scala 2 and 3:

```
// src/script/scala-2/progscala3/contexts/ImplicitClauses.scala

case class SortableSeq[A](seq: Seq[A]) {
  ❶ def sortBy1[B](transform: A => B)(implicit o: Ordering[B]):
    SortableSeq[A] =
      new SortableSeq(seq.sortBy(transform)(o))

  def sortBy2[B : Ordering](transform: A => B): SortableSeq[A] =
    new SortableSeq(seq.sortBy(transform)(implicitly[Ordering[B]]))
}

val seq = SortableSeq(Seq(1,3,5,2,4))

def defaultOrdering() = {
  assert(seq.sortBy1(i => -i) == SortableSeq(Seq(5, 4, 3, 2, 1)))
  ❸ assert(seq.sortBy2(i => -i) == SortableSeq(Seq(5, 4, 3, 2, 1)))
}
defaultOrdering()

def oddEvenOrdering() = {
  implicit val oddEven: Ordering[Int] = new Ordering[Int]:
    ❹ def compare(i: Int, j: Int): Int = i%2 compare j%2 match
      case 0 => i compare j
      case c => c

  ❺ assert(seq.sortBy1(i => -i) == SortableSeq(Seq(5, 3, 1, 4, 2)))
}
```

```
    assert(seq.sortBy2(i => -i) == SortableSeq(Seq(5, 3, 1, 4, 2)))
  }
  oddEvenOrdering()
```

- ❶ Use braces, because this is also valid Scala 2 code.
- ❷ Wrap examples in methods to scope the use of implicits.
- ❸ Uses the default ordering provided by `math.Ordering` for `Ints`.
- ❹ Define a custom `oddEven` ordering, which will be the “closest” implicit value in scope for the following lines.
- ❺ Implicitly use the custom `oddEven` ordering.

Let's focus on `sortBy1` for now. All the implicit parameters must be declared in their own parameter list. Here we need two lists, because we have a regular parameter, the function transform. If we only had implicit parameters, we would need only one parameter list.

The implementation of `sortBy1` just uses the `Seq.sortBy` method in the collections library. It takes a function that transforms the values to affect the sorting, and an `Ordering` instance to sort the values after transformation.

There is already a default implicit implementation in scope for `math.Ordering[Int]`, so we don't need to supply one if we want the usual numeric ordering. The anonymous function `i => -i` transforms the integers to their negative values for the purposes of ordering, which effectively results in sorting from highest to lowest.

Next, let's discuss the other method, `sortBy2`, and also explore new Scala 3 syntax for this purpose.

## Context Bounds



If you think about it, while `SortableSeq` is declared to support any element type `A`, the two `sortBy*` methods “bound” the allowed types to those for which an `Ordering` exists. Hence, the term *context bound* is used for the implicit value in this situation.

In `SortableSeq.sortBy1`, the implicit parameter `o` is a context bound. A major clue is the fact that it has type `Ordering[B]`, meaning it is parameterized by the output element type, `B`. So, while it doesn’t bound `A` explicitly, the result of applying `transform` is to convert `A` to `B` and then `B` is context bound by `Ordering[B]`.

Context bounds are so common that Scala 2 defined a more concise way of declaring them in the types, as shown in `sortBy2`, where the syntax `B : Ordering` appears. (Note that it’s not `B : Ordering[B]`.)

In the generated byte code for Scala 2, this is just short hand for the same code we wrote explicitly for `sortBy1`, with one difference. In `sortBy1`, we defined a name for the `Ordering` parameter, `o`, in the second argument list. We don’t have a name for it in `sortBy2`, but we need it in the body of the method. The solution is to use the method `Predef.implicitly`, as shown in the method body. It “binds” the implicit `Ordering` that is in scope so it can be passed as an argument.

Let’s rewrite this code in Scala 3:

```
// src/script/scala/progscala3/contexts/UsingClauses.scala

case class SortableSeq[A](seq: Seq[A]):
  def sortBy1a[B](transform: A => B)(using o: Ordering[B]):
    SortableSeq[A] =
      new SortableSeq(seq.sortBy(transform)(o))

  def sortBy1b[B](transform: A => B)(using Ordering[B]):
    SortableSeq[A] =
      new SortableSeq(seq.sortBy(transform)(summon[Ordering[B]]))

  def sortBy2[B : Ordering](transform: A => B): SortableSeq[A] =
    new SortableSeq(seq.sortBy(transform)(summon[Ordering[B]]))
```

The `sortBy1a` method is identical to the previous `sortBy1` method with a `using` clause instead of an implicit parameter list. In `sortBy1b`, we see that the name can be omitted and a new `Predef` method, `summon` is used to bind the value, instead. (It is identical to `+implicitly`.) The `sortBy2` here is written identically to the previous one in `ImplicitClauses`, but in Scala 3 it is implemented with a `using` clause.

The previously defined test methods, `defaultOrdering` and `oddEvenOrdering`, are almost the same in this source file, but are not shown here. There is an additional test method in this file that uses a `given` instance instead of an implicit value:

```
def evenOddGivenOrdering() =  
  given evenOdd as Ordering[Int] = new Ordering[Int]:  
    def compare(i: Int, j: Int): Int = i%2 compare j%2 match  
      case 0 => i compare j  
      case c => -c  
  
  val expected = SortableSeq(Seq(4, 2, 5, 3, 1))  
  assert(seq.sortBy1a(i => -i) == expected) ❶  
  assert(seq.sortBy1b(i => -i) == expected)  
  assert(seq.sortBy2(i => -i) == expected)  
  
  assert(seq.sortBy1a(i => -i)(using evenOdd) == expected) ❷  
  assert(seq.sortBy1b(i => -i)(using evenOdd) == expected)  
  assert(seq.sortBy2(i => -i)(using evenOdd) == expected)  
  
evenOddGivenOrdering()
```

- ❶ Use the given `evenOdd` instance implicitly.
- ❷ Use the given `evenOdd` instance explicitly with `using`.

The syntax `given foo as Type[T]` is used instead of `implicit val foo: Type[T]`, essentially the same way we used `givens` when discussing type classes. Recall the use of `as`, too.

If the using clause is provided explicitly, as marked with comment 2, the using keyword is *required* in Scala 3, whereas Scala 2 didn't require the implicit keyword here. The reason using is now required is two fold. First, it's better documentation for the reader. Second, it removes an ambiguity that is illustrated in the following contrived Scala 2 example:

```
case class FastSeq[T](implicit storage: Storage[T]):           ❶
  def apply(i: Int): Option[T] = storage.get(i)                ❷

implicit val customStorage: Storage = ???
val opt = FastSeq[String](5)                                   ❸
```

- ❶ A “fast” sequence implementation with user-pluggable storage.
- ❷ Optionally return the item at index *i*.
- ❸ Does this line return a None, because the sequence is empty?

In Scala 2, the last line would cause a compiler error for the argument 5, saying that a Storage instance was expected as the argument. The actual user intention was for the instance to be constructed with the implicit value `customStorage`, then the `apply` method was to be called with 5. Instead, you would have to use the unintuitive expression `FastSeq[String].apply(5)`.

Now this ambiguity is removed by requiring `using` when the implicit is provided explicitly. As written, the compiler knows that you want to use the implicit for the storage *and then* call `apply(5)`.

### TIP

The intent of the new `given name as ...` syntax and the `using ...` syntax is to make their purpose more explicit, but it functions almost identically to Scala 2 implicit definitions and parameters.

## By-Name Context Parameters

Context parameters can be by-name parameters. Here is an example adapted from [this Dotty documentation](#).

```
// src/script/scala/progscala3/contexts/ByNameContextParameters.scala

trait Codec[T]:
  def write(x: T): Unit

given intCodec as Codec[Int]:
  def write(i: Int): Unit = print(i)

given optionCodec[T](using ev: => Codec[T]) as Codec[Option[T]]:
  def write(xo: Option[T]) = xo match
    case Some(x) => ev.write(x)
    case None =>

val s = summon[Codec[Option[Int]]]

s.write(Some(33))
s.write(None)
```

Note that `ev` for `optionCodec[T]` is a by-name parameter, which means its evaluation is delayed until used. Using a by-name parameter here can avoid certain cases where a *divergent expansion* can happen, as the compiler chases its tail trying to resolve all using clause parameters.

## Other Context Parameters

In “A Taste of Futures”, we saw that `Future.apply` has a second, implicit argument list that is used to pass an `ExecutionContext`:

```
object Future:
  apply[T](body: => T)(implicit executor: ExecutionContext):
    Future[T]
  ...
```

It is not a context bound, because the `ExecutionContext` is independent of `T`.

We didn't specify an `ExecutionContext` when we called these methods, but we imported a global default that the compiler used:

```
import scala.concurrent.ExecutionContext.Implicits.global
Future(...) // Use the implicit value
Future(...)(using customExecutionContext) // Explicit argument with
"using"
```

`Future` supports a lot of the operations like `filter`, `map`, etc. All take two argument lists, like `Future.apply`. Have a `using` clause for the `ExecutionContext` makes the code much cleaner:

```
given customExecutionContext: ExecutionContext = ???
val f1 = Future(...)(using customExecutionContext)
  .map(...)(using customExecutionContext)
  .filter(...)(using customExecutionContext)
// versus:
val f2 = Future(...).map(...).filter(...)
```

Other example of *using contexts* (my term) include transaction identifiers, database connections, and web sessions.

### TIP

The example shows that using contexts can make code more concise, but they can be overused in Scala code. When you see the same `using FooContext` all over a code base, it feels more like a global variable than pure functional programming.

## Passing Context Functions

**Context Functions** are functions with context parameters only. Scala 3 introduces a new *context function type* for them, indicated by `? => T`, with special handling depending on how they are used.

In essence, context functions abstract over context parameters.

Consider this alternative for handling The ExecutionContext passed to Future.apply(), using a wrapper FutureCF (for “context function”):

```
// src/script/scala/progscala3/contexts/ContextFunctions.scala

import scala.concurrent.{Await, ExecutionContext, Future}
import scala.concurrent.ExecutionContext.Implicits.global
import scala.concurrent.duration._

object FutureCF:
  type Executable[T] = ExecutionContext => T
  ❶

  def apply[T](body: => T): Executable[Future[T]] = Future(body)
  ❷

  def sleepN(dur: Duration): Duration =
    ❸
    val start = System.currentTimeMillis()
    Thread.sleep(dur.toMillis)
    Duration(System.currentTimeMillis - start, MILLISECONDS)

  val future1 = FutureCF(sleepN(1.second))
  ❹
  val future2 = FutureCF(sleepN(1.second))(using global)
  ❺
  val duration1 = Await.result(future1, 2.second)
  val duration2 = Await.result(future2, 2.second)
  ❻
```

- ❶ Type alias for a context function with an ExecutionContext.
- ❷ Compare this definition of FutureCF.apply() to Future.apply() above, which we are calling here. The implicit ExecutionContext is passed to Future.apply().
- ❸ Define some work that will be passed to futures; sleep for some Duration and return the actual elapsed time as a Duration.
- ❹

Create a future with the implicit value, like calling `Future.apply()`.

- ⑤ Create another future specifying the implicit argument explicitly.
- ⑥ Await the results of the futures. Wait no longer than two seconds.

The last two lines print the following (your actual numbers may vary slightly):

```
val duration1: concurrent.duration.Duration = 1004 milliseconds
val duration2: concurrent.duration.Duration = 1002 milliseconds
```

Let's look at a more extensive example inspired by the [context functions documentation page](#). We'll create a *domain-specific language* (DSL) for constructing JSON. For simplicity, we won't construct instances for some JSON library, just JSON-formatted strings. To motivate the example, let's begin with an entry point that shows the DSL in action:

```
// src/main/scala/progscala3/contexts/json/JSONBuilder.scala
package progscala3.contexts.json
```

```
@main def TryJSONBuilder(): Unit =
  val js = obj {
    "config" -> obj {
      "master" -> obj {
        "host" -> "192.168.1.1"
        "port" -> 8000
        "security" -> "null"
        // "foo" -> (1, 2.2, "three") // doesn't compile!
      }
      "nodes" -> array {
        aobj { // "array object"
          "name" -> "node1"
          "host" -> "192.168.1.10"
        }
        aobj {
          "name" -> "node2"
          "host" -> "192.168.1.20"
        }
      }
    }
  }
```

```

    }
    "otherThing" -> 2
  }
}
println(js)

```

Let's try it in SBT. I reformatted the output for better legibility:

```

> runMain proscala3.contexts.json.TryJSONBuilder
...
{
  "config": {
    "master": {
      "host": "192.168.1.1", "port": 8000, "security": null
    },
    "nodes": [
      {"name": "node1", "host": "192.168.1.10"},
      {"name": "node2", "host": "192.168.1.20"},
      "otherThing" -> 2
    ]
  }
}

```

Now let's work through the implementation (same source file):

```

object JSONElement:
  def valueString[T](t: T): String = t match
    ❶
    case "null" => "null"
    case s: String => "\"" + s + "\""
    case _ => t.toString

sealed trait JSONElement
    ❷
case class JSONKeyedElement[T](key: String, element: T) extends JSONElement:
  override def toString = "\"" + key + "\": " + JSONElement.valueString(element)
case class JSONArrayElement[T](element: T) extends JSONElement:
  override def toString = JSONElement.valueString(element)

```



- ❶ Return the correct string representation for a value. JSON allows `nulls`, for which we'll expect the user to use the string `"null"` (as shown in the example). Hence, `valueString` returns `null` without quotes, all other strings in double quotes, and for everything else, the output of `toString`.
- ❷ We can model everything as either a “keyed” element of the form `"key": value` or just a value, but the latter only appear as elements in arrays.

Continuing, we have types for JSON objects and arrays:

```
import scala.collection.mutable.ArrayBuffer

trait JSONContainer(open: String, close: String) extends
  JMLElement: ❶
  val elements = new ArrayBuffer[JMLElement]
  def add(e: JMLElement): Unit = elements += e
  override def toString = elements.mkString(open, ", ", close)

class JSONObject extends JSONContainer("{", "}")
class JSONArray extends JSONContainer("[", "]")
```

- ❶ For both JSON objects and arrays, we add elements to a mutable array buffer. There are two places the `add` method is called, discussed below.

Note that traits can define constructor parameters, like classes. For our purposes, only the opening and closing delimiter differ between objects and arrays. The concrete classes for them define the correct delimiters.

```
sealed trait ValidJSONValue[T]
❶
given ValidJSONValue[Int]
given ValidJSONValue[Double]
given ValidJSONValue[String]
```

```

given ValidJSONValue[Boolean]
given ValidJSONValue[JSONObject]
given ValidJSONValue[JSONArray]

extension [T : ValidJSONValue] (name: String)
  ❷
  def ->(element: T)(using jc: JSONContainer) =
    jc.add(JSONNamedElement(name, element))

```

- ❶ These given instances of `ValidJSONValue[T]` are *witnesses*, constraining the allowed types of JSON values (see “[Constraining Allowed Instances](#)”).
- ❷ This `String` extension method that is constrained by `ValidJSONValue[T]`. It constructs `JSONKeyedElements` using “key” `-> value`, just like tuple pairs, but constrained by the *context bound* `T : ValidJSONValue`. We use a `JSONContainer` because these key-value pairs only occur inside containers (objects or arrays) in the DSL. It is here that we add the key-value pairs to the container `jc`.

If you try a tuple value, it will fail to compile, as shown in a comment in `TryJSONBuilder`!

In Scala 2, you would need to declare derived classes of `ValidJSONValue[T]`, like this:

```

implicit object VJSONInt extends ValidJSONValue[Int]
...

```

Finally, we see the actual context functions in action:

```

def obj(init: JSONObject ?=> Unit) =
  ❶
  given jo as JSONObject
  init
  jo

```

```

def aobj(init: JSONObject ?=> Unit)(using jc: JSONContainer) =
  ❷
  given jo as JSONObject
  init
  jc.add(jo)

def array(init: JSONArray ?=> Unit) =
  ❸
  given ja as JSONArray
  init
  ja

```

- ❶ A whole JSON object, as well as nested objects, starts with `obj`. Refer to the example in `TryJSONBuilder`. Where does the `init` context function of type `JSONObject ?=> Unit` come from? It is constructed by the compiler from the expressions inside the braces passed as the argument to `obj`. Or, as it appears in the DSL, the braces after the `obj` “keyword”. Next, the `given` clause creates an instance of `JSONObject` named `jo`. Then, `init` is evaluated, where `jo` will be used to satisfy using clauses inside those nested expressions. Finally, we return `jo`.
- ❷ Use `aobj` to define objects as array elements. Note that this function has a `using` clause, unlike `obj`, which will always a `JSONArray`. Unfortunately, the name `obj` can’t be overloaded here, because the compiler would consider the two definitions ambiguous. The body of `aobj` is the second place where the `add` method is called. Recall that the other location is inside the `String` extension method `->`.
- ❸ Define an array. This body is very similar to `obj`.

So, if you find you need a small DSL for expressing structure, context functions is one tool at your disposal. We’ll explore more tools for DSLs in *Chapter 20*.

## Constraining Allowed Instances

The “given” instances of `ValidJSONValue[T]` in the previous example were used as context bounds that constrained the allowed types that could be used for type parameter `T` in the `String` extension method `->(element: T)`.

What was new is that we did no actual work with these instances. Only their existence mattered. They “witnessed” the allowed types for JSON elements. So, because we didn’t provide an instance for three-element tuples, for example, attempting to use a tuple value in the DSL, such as `"stuff" -> (1, "two", 3.3)`, causes a compilation error.

Sometimes a context bound is used in both ways, as a witness and to do work. Consider the following sketch of an API for data “records” with ad hoc schemas, like in some NoSQL databases. Each row is encapsulated in a `Map[String,Any]`, where the keys are the field names and the “column” values are unconstrained. However, the `add` and `get` methods, for adding column values to a row and retrieving them, do constrain the allowed instance types. Here is the example:

```
// src/main/scala/progscala3/contexts/NoSQLRecords.scala
package progscala3.contexts.scaladb

import scala.language.implicitConversions
import scala.util.Try

case class InvalidFieldName(name: String)
  extends RuntimeException(s"Invalid field name $name")

object Record:
  ❶
  def make: Record = new Record(Map.empty)
  type Conv[T] = Conversion[Any,T]

case class Record private (contents: Map[String,Any]):
  ❷
  import Record.Conv
  def add[T](nameValue: (String, T))(using Conv[T]): Record =
```

```

③    Record(contents + nameValue)
    def get[T](colName: String)(using toT: Conv[T]): Try[T] =
④    Try(toT(col(colName)))
    private def col(colName: String): Any =
        contents.getOrElse(colName, throw InvalidFieldName(colName))

@main def TryScalaDB =
    import Record.Conv
    given Conv[Int] = _.asInstanceOf[Int]
⑤    given Conv[Double] = _.asInstanceOf[Double]
    given Conv[String] = _.asInstanceOf[String]
    given ab[A : Conv, B : Conv] as Conv[(A, B)] =
        _.asInstanceOf[(A,B)]

    val rec = Record.make.add("one" -> 1).add("two" -> 2.2)
        .add("three" -> "THREE!").add("four" -> (4.4, "four"))
        .add("five" -> (5, ("five", 5.5)))

    val one    = rec.get[Int]("one")
    val two    = rec.get[Double]("two")
    val three  = rec.get[String]("three")
    val four   = rec.get[(Double, String)]("four")
    val five   = rec.get[(Int, (String, Double))]("five")
    val bad1   = rec.get[String]("two")
    val bad2   = rec.get[String]("five")
    val bad3   = rec.get[Double]("five")
    // val error = rec.get[Byte]("byte")

    println(s"one, two, three, four, five -> $one, $two, $three, $four, $five")
    println(s"bad1, bad2, bad3 -> $bad1, $bad2, $bad3")
⑥

```

- ❶ The companion object defines make to start “safe” construction of a Record. It also defines a type alias for **Conversion**, where we always use Any as the first type parameter. This alias is necessary when we define given ab below.
- ❷ Define Record with a single field Map[String,Any] to hold the user-defined fields and values. Use of private after the type

name declares the constructor private, forcing users to create records using `Record.make` followed by `add` calls. This prevents users from using an unconstrained `Map` to construct a `Record`!

- ③ A method to add a field with a particular type and value. The anonymous context parameter is used only to constrain the allowed values for `T`. It's `apply` method won't be used. Since `Records` are immutable, a new instance is returned.
- ④ A method to retrieve a field value with the desired type `T`. Here the context parameter both constrains the allowed `T` types and it handles conversion from `Any` to `T`. On failure, an exception is returned in the `Try`. Hence, this example can't catch all type errors at compile time, as shown below.
- ⑤ Only `Int`, `Double`, `String`, and pairs of them are supported. These definitions work as *witnesses* for the allowed types in both the `add` and `get` methods, as well as function as *implicit conversions* from `Any` to specific types when used in `get`. Note that given `ab` declares a `given` for pairs, but the `A` and `B` types are constrained to be other allowed types, including other pairs!
- ⑥ Attempting to retrieve columns with the wrong types. Attempting to retrieve an unsupported `Byte` column would cause a compilation error.

Running this example with `runMain`  
`progscale3.contexts.scaladb.TryScalaDB`, you get the following output (abbreviated):

```
one, two, three, four, five -> Success(1), Success(2.2),  
Success(THREE!),  
Success((4.4,four)), Success((5,(five,5.5)))  
bad1, bad2, bad3 ->  
Failure(... java.lang.Double cannot be cast to class  
java.lang.String ...),
```

```
Failure(... scala.Tuple2 cannot be cast to class java.lang.String
...),
Failure(... scala.Tuple2 cannot be cast to class java.lang.Double
...))
```

Hence, the only runtime failure we have we can't prevent at compile time is attempting to get a column with the wrong type.

The type alias `Conv[T]` not only made the code more concise than using `Conversion[Any,T]`, it is necessary for the context bounds on `A` and `B` in `ab`. This is because context bounds *a/ways* require one and only one type parameter, but `Conversion[A,B]` has two. Fortunately, the `A` is always `Any` in our case, so we were able to define the type alias `Conv[T]` and use it for the bounds in `ab`.

### TIP

Using a type alias to fill in some of the type parameters is a useful trick when the number of type parameters in a type don't match what your needs.

As a reminder, use of `given` provides a more concise syntax than the Scala 2 way of declaring an implicit value (which is still supported):

```
given Conv[Int] = _.asInstanceOf[Int]           // Scala 3
// vs.
implicit val toInt: Conv[Int] = new Conv[Int]:   // Scala 2
  def apply(any: Any): Int = any.asInstanceOf[Int]
```

To recap, we limited the allowed types that can be used for a parameterized method by passing an implicit parameter and only defining given values that match the types we want to allow.

This example was inspired by an API I once wrote to work with Cassandra.

## Implicit Evidence

In the previous example, the `Record.add` method showed one way to constrain the allowed types without doing anything else with the context bounds. Now we'll discuss another technique called *implicit evidence*.

A nice example of this technique is the `toMap` method available for all iterable collections. Recall that the `Map` constructor wants key-value pairs, i.e., two-element tuples, as arguments. If we have a sequence of pairs, wouldn't it be nice to create a `Map` out of them in one step? That's what `toMap` does, but we have a dilemma. We can't allow the user to call `toMap` if the sequence is *not* a sequence of pairs.

The `toMap` method is defined in `IterableOnceOps`:

```
trait IterableOnceOps[+A]:  
  def toMap[K, V](implicit ev: <::[A, (K, V)]): immutable.Map[K, V]  
  ...
```

The implicit parameter `ev` is the “evidence” we need to enforce our constraint. It uses a type defined in `Predef` called `<::`, named to resemble the type parameter constraint `<::`, e.g., `A <:: (K,V)`. In “**Call by Name, Call by Value**”, we learned that this notation means `A` is a subtype of `(K,V)`.

Recall we said that types with two type parameters can be written in “infix” notation. So, the following two expressions are equivalent:

```
<::[A, (T,U)]  
A <:: (T,U)
```

Now, when we have a traversable collection that we want to convert to a `Map`, the implicit evidence `ev` value we need will be synthesized by the compiler, but only if `A <:: (T,U)`; that is, if `A` is actually a pair of types. If true, then `toMap` can be called and it simply passes the elements of the traversable to the `Map` constructor. However, if `A` is not a pair type, the code fails to compile.



Hence, evidence only has to exist to enforce a type constraint, which the compiler generates for us. We don't have to define a given or implicit value ourselves.

There is also a related type in Predef for providing evidence that two types are equivalent, called `==`.

## Working Around Type Erasure with Context Bounds

Context bounds can also be used to work around limitations due to *type erasure* on the JVM.

For historical reasons, the JVM “forgets” the supplied type arguments for parameterized types. For example, consider the following definitions for an *overloaded* method with unique type signatures, at least to human readers:

```
scala> object O:
  |   def m(seq: Seq[Int]): Unit = println(s"Seq[Int]: $seq")
  |   def m(seq: Seq[String]): Unit = println(s"Seq[String]: $seq")
  |
3 |   def m(seq: Seq[String]): Unit = println(s"Seq[String]: $seq")
  |   ^
  |   Double definition:
  |   def m(seq: Seq[Int]): Unit in object O at line 2 and
  |   def m(seq: Seq[String]): Unit in object O at line 3
  |   have the same type after erasure.
```

So, the compiler disallows the definitions because they are effectively the same in byte code.

## WARNING

Be careful defining overloaded methods in the REPL without an enclosing type. Because the REPL lets you *redefine* anything, for convenience, if you enter these two method definitions without the enclosing object, you'll see no complaints. You will have one method, for `Seq[String]`, instead of two.

However, we can add an implicit parameter to disambiguate the methods:

```
//
src/script/scala/progscala3/contexts/UsingTypeErasureWorkaround.scala
object M:
  implicit object IntMarker           ❶
  implicit object StringMarker

  def m(seq: Seq[Int])(using IntMarker.type): String =           ❷
    s"Seq[Int]: $seq"

  def m(seq: Seq[String])(using StringMarker.type): String =
    s"Seq[String]: $seq"

import M._
```

The last three lines produce the following results:

```
scala> m(Seq(1,2,3))
      | m(Seq("one", "two", "three"))
val res0: String = Seq[Int]: List(1, 2, 3)
val res1: String = Seq[String]: List(one, two, three)

scala> m(Seq("one" -> 1, "two" -> 2, "three" -> 3)) // ERROR
1 | m(Seq("one" -> 1, "two" -> 2, "three" -> 3))
   | ^
   | None of the overloaded alternatives of method m in object M with
types
   | ...
```

Define two special-purpose implicit objects that will be used to disambiguate the methods affected by type erasure.

- ② Redefinition of the method that takes `Seq[Int]`. It now has a second parameter list expecting an implicit `IntMarker.type` (because `IntMarker` is an object). Then define a similar method for `Seq[String]`.

Now the compiler considers the two `m` methods to be distinct after type erasure.

You might wonder why I didn't use implicit `Int` and `String` values, rather than invent new types. Using implicit values for very common types is not recommended. It would be too easy for one or more implicit `String` values, for example, to show up in a particular scope. If you don't expect one to be there, you might be surprised when it gets used. If you do expect one to be in scope, but there are several of them, you'll get a compiler error because all of them are valid choices and the compiler can't decide which to use.

At least the second scenario triggers an immediate error rather than allowing unintentional behavior to occur.

The safer bet is to limit your use of implicit parameters and values to very specific, purpose-built types.

### WARNING

Avoid using context parameters for very common types like `Int` and `String`, as they are more likely to cause confusing behavior or compilation errors.

We'll discuss type erasure in more detail in *Chapter 15*.

## Rules for Using Clauses

The sidebar lists the general rules for using clauses.

## RULES FOR USING CLAUSES

1. Zero or more argument lists can be using clauses.
2. The `implicit` or `using` keyword must appear first and only once in the parameter list and all the parameters are *context parameters*.

Hence, any one parameter list can't mix context parameters with other parameters. Here are a few more examples, including what happens when a regular parameter list follows an using clause, which is allowed in Scala 3, but not in Scala 2:

```
// src/script/scala/progscala3/contexts/UsingClausesLists.scala

case class U1[T](t: T)
case class U2[T](t: T)

def f1[T1,T2](name: String)(using u1: U1[T1], u2: U2[T2]): String =
  ❶
  s"f1: $name: $u1, $u2"
def f2[T1,T2](name: String)(using u1: U1[T1])(using u2: U2[T2]):
String = ❷
  s"f2: $name: $u1, $u2"
def f3[T1,T2](name: String)(using u1: U1[T1])(u2: U2[T2]): String =
  ❸
  s"f3: $name: $u1, $u2"

given u1i as U1[Int](0)
given u2s as U2[String]("one")
```

- ❶ One using clause with two values.
- ❷ Two using clauses, each with one value.
- ❸

One using clause sandwiched between two regular parameter lists.

Now use them:

```
scala> f1("f1a")
| f1("f1b")(using u1i, u2s)
| f2("f2a")
| f2("f2b")(using u1i)(using u2s)
| f3("f3a")
| f3("f3b")(using u1i)
| f3("f3c")(using u1i)(u2s)
val res0: String = f1: f1a: U1(0), U2(one)
val res1: String = f1: f1b: U1(0), U2(one)
val res2: String = f2: f2a: U1(0), U2(one)
val res3: String = f2: f2b: U1(0), U2(one)
val res4: U2[Any] => String = Lambda$7814/0x000000080360d040@4aa25f5d
val res5: U2[Any] => String = Lambda$7815/0x000000080360c840@521dc499
val res6: String = f3: f3c: U1(0), U2(one)
```

The results for f1 and f2 should make sense; they are functionally equivalent. Recall that when passing values explicitly, the `using` keyword is required.

Now consider res4 through res6. First, res6 should be unsurprising, as we explicitly provided arguments for all three lists.

*Partial application* is the explanation for res4 and res5. For methods with more than one parameter list, if you invoke them with a subset of the *leading* parameter lists, a new function is returned expecting the *rest* of the parameter lists. For res4, the given is used for the second parameter list, the using clause, while a value is explicitly provided for the using clause in the res5 definition. Hence, both return the same thing.

For both res4 and res5, the third parameter list is not a using clause and it was not provided explicitly. Therefore, the expressions returned a *function* that expects the remaining parameter list, which takes an instance of U2, and returns a String:

```
scala> val u2a = U2[Any](1.1)    // Declare a U2[Any] we can use.
      | res4(u2a)                // Pass it to the res4 and res5
      | res5(u2a)                functions.
val u2a: U2[Any] = U2(1.1)
val res7: String = f3: f3a: U1(0), U2(1.1)
val res8: String = f3: f3b: U1(0), U2(1.1)
```

Because we didn't provide the third parameter list when we constructed `res4` and `res5`, the type parameter `T2` for `f3` was inferred to be the widest possible type, `Any`. Try calling `res4(m1)` or `res4(m2)` and you'll get a type error, as `m1` and `m2` are not type compatible with `U2[Any]`. It doesn't matter that `m1` and `m2` were declared as givens; we can still use them as regular parameters, as long as the types are compatible.

## Improving Error Messages

Let's finish the discussion of using clauses by discussing how to improve the errors reported when a value isn't found. The compiler's default messages are usually sufficiently descriptive, but you can customize them with the `implicitNotFound` annotation,<sup>2</sup> as follows:

```
// src/script/scala/progscala3/contexts/ImplicitNotFound.scala[]

import scala.annotation.implicitNotFound

@implicitNotFound("Stringer: No implicit found ${T} : Tagify[${T}]")
trait Tagify[T]:
  def toTag(t: T): String

case class Stringer[T : Tagify](t: T):
  override def toString: String =
    s"Stringer: ${implicitly[Tagify[T]].toTag(t)}"

object O:
  def makeXML[T](t: T)(
    implicit @implicitNotFound("No Tagify[${T}] implicit found")
```

```

    tagger: Tagify[T]): String =
      s"<xml>${tagger.toTag(t)}</xml>"

given Tagify[Int]:
  def toTag(i: Int): String = s"<int>$i</int>"
given Tagify[String]:
  def toTag(s: String): String = s"<string>$s</string>"

```

Let's try it:

```

scala> Stringer("Hello World!")
      | Stringer(100)
      | O.makeXML("Hello World!")
      | O.makeXML(100)
val res0: Stringer[String] = Stringer: <string>Hello World!</string>
val res1: Stringer[Int] = Stringer: <int>100</int>
val res2: String = <xml><string>Hello World!</string></xml>
val res3: String = <xml><int>100</int></xml>

scala> Stringer(3.14569)
      | O.makeXML(3.14569)
1 | Stringer(3.14569)
   |               ^
   |               Stringer: No implicit found Double :
   Tagify[Double]
2 | O.makeXML(3.14569)
   |               ^
   |               Stringer: No implicit found Double :
   Tagify[Double]

```

TODO: Still true in Scala 3 final? What about a new annotation??

Only the annotation on Tagify is used. The annotation on the parameter to O.makeXML is supposed to take precedence for the last output. This appears to be a current limitation in Scala 3.

You can only annotate types intended for use as givens. This is another reason for creating custom types for this purpose, rather than using more common types, like Int, String, or our Person type. You can't use this annotation with those types.

## Recap and What's Next

We completed our exploration into the details of abstracting over context in Scala 2 and 3. I hope you can appreciate their power and utility, but also the need to use them wisely. Unfortunately, because the old *implicit* idioms are still supported for backwards compatibility, at least for a while, it will be necessary to understand how to use both the old and new constructs, even though they are redundant.

Now we're ready to dive into the principles of functional programming. We'll start with a discussion of the core concepts and why they are important. Then we'll look at the powerful functions provided by most container types in the library. We'll see how we can use those functions to construct concise, yet powerful programs.

- 
- 1 A “regular” parameter list is also known as a *normal parameter clause*, but I have just used the more familiar *parameter list* in this book. *Using clause* is more of a formal term in Scala 3 documentation than *implicit parameter clause* was, which is why I emphasize it here.
  - 2 At the time of this writing, there is no `givenNotFound` or similar replacement annotation in Scala 3.



# Chapter 7. Functional Programming in Scala

---

## A NOTE FOR EARLY RELEASE READERS

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

This will be the 7th chapter of the final book. Please note that the GitHub repo will be made active later on.

If you have comments about how we might improve the content and/or examples in this book, or if you notice missing material within this chapter, please reach out to the author at [mcronin@oreilly.com](mailto:mcronin@oreilly.com).

*It is better to have 100 functions operate on one data structure than 10 functions on 10 data structures.*

—Alan J. Perlis

Every decade or two, a major computing idea goes mainstream. The idea may have lurked in the background of academic Computer Science research or in obscure corners of industry, perhaps for decades. The transition to mainstream acceptance comes in response to a perceived problem for which the idea is well suited. Object-oriented programming (OOP), which was invented in the 1960s, went mainstream in the 1980s, arguably in response to the emergence of graphical user interfaces, for which the OOP paradigm is a natural fit.

Functional programming (FP) experienced a similar breakout over the last 15 years or so. FP is actually much older than OOP, going back to the 1930s! FP offers effective techniques for three major challenges that became pressing in the 2000s and remain pressing today:

1. The need for pervasive concurrency, so we can scale our applications horizontally and improve their resiliency against service disruptions. Concurrent programming is now an essential skill for every developer to master.
2. The need to write data-centric (e.g., “Big Data”) applications. Of course, at some level all programs are about data, but the growth of Big Data highlighted the importance of effective techniques for working with large data sets.
3. The need to write bug-free applications. Of course, this is an old problem, but it has grown more pressing as software has become more pervasive and bugs have become more potentially disruptive to society. FP gives us new tools from mathematics that move us further in the direction of *provably bug-free* programs.

Immutability eliminates the hardest problem in concurrency, coordinating access to shared, mutable state. Hence, writing immutable code is an essential tool for writing robust, concurrent programs. We’ll explore examples in *Chapter 18*.

The benefits of FP for data-centric applications will become apparent as we master the functional operations discussed in this and subsequent chapters. We’ll explore the connection in depth in *later chapters*.

Finally, Embracing FP is the best way to write immutable code. Immutability and the mathematical rigor of functional thinking lead to programs with fewer logic flaws.

Scala is a mixed-paradigm language, supporting both FP and OOP. It encourages you to use both programming models to get the best of both of them.

## What Is Functional Programming?

All programming languages have functions of some sort. Functional programming is based on the rules of mathematics for the behavior of functions and values. This starting point has far-reaching implications for software.

### Functions in Mathematics

In mathematics, functions have no *side effects*. Consider the classic function  $y = \sin(x)$ .

No matter how much work  $\sin(x)$  does, all the results are returned and assigned to  $y$ . No global state of any kind is modified internally by  $\sin(x)$ . Also, all the data it needs to compute the value is passed in through  $x$ . Hence, we say that such a function is free of *side effects*, or *pure*.

Purity drastically simplifies the challenge of analyzing, testing, and debugging a function. You can do all these things without having to know anything about the context in which the function is invoked. It's easy to write unit tests for these functions.

This obliviousness to the surrounding context provides *referential transparency*, which has two implications. First, you can call such a function anywhere and be confident that it will always behave the same way, independent of the calling context. Because no global state is modified, concurrent invocation of the function is also straightforward and reliable. No tricky thread-safe coding is required.

The second sense of the term is that you can substitute the value computed by an expression for subsequent invocations of the

expression. Consider, for example, the equation  $\sin(\pi/2) = 1.0$ . A code analyzer could replace repeated calls to  $\sin(\pi/2)$  with 1.0 with no loss of correctness, as long as `sin` is truly pure.

### TIP

A function that returns `Unit` can only perform side effects. It can only modify mutable state somewhere. A simple example is a function that just does input or output, which modifies “the world.”

Note that there is a natural uniformity between values and functions, due to the way we can substitute one for the other. What about substituting functions for values, or treating functions as values?

In fact, functions are *first-class* values in functional programming, just like data values. You can compose functions from other functions (for example,  $\tan(x) = \sin(x) / \cos(x)$ ). You can assign functions to variables. You can pass functions to other functions as arguments. You can return functions as values from other functions.

A function that takes another function as a parameter or returns a function is called a *higher-order function*. In calculus, two examples of higher-order functions are derivation and integration. We pass an expression, like a function, to the derivation “operation,” which returns a new function, the derivative.

We’ve seen many examples of higher-order functions already, such as the `map` method on collections, which takes a single function parameter that is applied to each element.

## Variables That Aren’t

In most programming languages, variables are *mutable*. In functional programming, variables are *immutable*, as they are in mathematics.

This is another consequence of the mathematical orientation. In the expression  $y = \sin(x)$ , once you pick  $x$ , then  $y$  is fixed. Similarly, values are immutable; if you increment the integer 3 by 1, you don't "modify the 3 object," you create a new value to represent 4. We've been using the term "value" as a synonym for immutable instances.

Immutability is difficult to work with at first when you're not used to it. If you can't change a variable, then you can't have loop counters, you can't have objects that change their state when methods are called on them, and you can't do input and output, which changes the state of the world! Learning to think in immutable terms takes some effort.

Obviously you can't be pure always. Without input and output, your computer would just heat the room. Practical functional programming makes *strategic* decisions about where code should perform side effects, leaving the rest of it pure.

### WHY AREN'T INPUT AND OUTPUT PURE?

It's easy to grasp that idea that  $\sin(x)$  is a pure function without side effects. Why are input and output considered side effects and therefore not pure? They modify the state of the world around us, such as the contents of files or what we see on the screen. They aren't referentially transparent either. Every time I call `readLine` (defined in `Predef`), I get a different input. Every time I call `println` (also defined in `Predef`), I pass a different argument, but `Unit` is always returned.

This does not mean that functional programming is stateless. If so, it would also be useless. You can always represent state changes with new instances.

Recall this example from [Chapter 2](#):

```
// src/script/scala/progscala3/typelessdomore/Factorial.scala
```

```
def factorial(i: Int): Long =  
  def fact(i: Int, accumulator: Long): Long =  
    if (i <= 1) accumulator  
    else fact(i - 1, i * accumulator)  
  
  fact(i, 1L)  
  
(0 to 5).foreach(i => println(s"$i: ${factorial(i)}"))
```

We calculate factorials using recursion. Updates to the accumulator are pushed on the stack. We don't modify a running value in place. At the end of the example, we “mutate” the world by printing the results.

Immutability has enormous benefits for writing code that scales through concurrency. Almost all the difficulty of multithreaded programming lies in synchronizing access to shared, mutable state. If you remove mutability, the problems vanish. It is the combination of referentially transparent functions and immutable values that make functional programming compelling as a better way to write concurrent software.

Almost all the constructs we have invented in 80-odd years of computer programming have been attempts to manage complexity. Higher-order, pure functions are called *combinators*, because they compose together very well as flexible, fine-grained building blocks for constructing larger, more complex programs.

Pure functions and immutability drastically reduce the occurrence of bugs. Mutable state is the source of the most pernicious bugs, the ones that are hardest to detect with testing before deploying to production, and often the hardest bugs to fix.

Mutable state means that state changes made by one module are *unexpected* by other modules, causing a “spooky action at a distance” phenomenon. This can happen even without concurrency.

Purity simplifies designs by eliminating a lot of the defensive boilerplate required in object-oriented code, where it's common to encapsulate access to mutable data structures in objects, because we can't risk sharing them with clients unprotected. Such accessors increase code size and the ad-hoc quality of code. Therefore, they increase the testing and maintenance burden. They also broaden the footprint of APIs, which increases the learning burden for users.

In contrast, when we have immutable data structures, most of these problems simply vanish. We can make internal data public without fear of data loss or corruption. Encapsulation is still useful to minimize coupling and to expose coherent abstractions, but there is less fear about data access.

What about performance? If you can't mutate an object, then you must copy it when the state changes, right? Fortunately, *functional data structures* minimize the overhead of making copies by sharing the unmodified parts of the data structures between the two copies.

In contrast, the data structures often used in object-oriented languages don't support efficient copying. A defensive module might be forced to make expensive copies of mutable data structures when it's necessary to share them without risking corruption by uncontrolled mutation.

Another source of potential performance gains is the use of data structures with *lazy evaluation*, such as Scala's `LazyList`. Evaluation is delayed until an element is required. Instances can be infinite, like this definition of the natural numbers (the integers 0 to infinity):

```
// src/script/scala/progscala3/fp/datastructs/LazyListNaturals.scala[]
```

```
scala> val natNums = LazyList.from(0)
val natNums: LazyList[Int] = LazyList(0, 1, 2, ..., 99, <not
computed>)
```

```
scala> natNums.take(1000).toList
val res0: List[Int] = List(0, 1, 2, ..., 998, 999)
```

- ❶ Take the first  $n$  (1000) elements, returning another `LazyList`, then convert to a regular `List`.

Scala uses *eager* or *strict* evaluation by default, but lazy evaluation avoids work that isn't necessary now and may never be necessary. A `LazyList` can be used for processing a very large stream of incoming data, yielding results from early, rather than forcing the client to wait until all the data has been processed.

Note that sequential collections have `take(n)`, as shown, and also `drop(n)`, which drops the leading  $n$  values, and `splitAt(n)`, which returns a tuple of `take(n)` and `drop(n)`.

So why isn't Scala lazy by default? There are many scenarios where lazy evaluation is less efficient and it is harder to predict the performance of lazy evaluation. Hence, most functional languages use eager evaluation, but most also provide lazy data structures for when they are useful.

We're exploring Scala's functional programming support before its object-oriented support to encourage you to really understand its benefits and embrace FP.

This chapter covers what I consider to be the essentials that every new Scala programmer needs to know. Functional programming is a large and rich field. We'll cover some of the more advanced topics that are less essential for new programmers in *a later chapter*.

## Functional Programming in Scala

As a hybrid object-functional language, Scala does not require functions to be pure, nor does it require variables to be immutable. It does encourage you to write your code this way whenever possible.

Let's quickly recap a few things we've seen already.



Here are several higher-order functions that we compose together to iterate through a list of integers, filter for the even ones, map each one to its value multiplied by two, and finally multiply them together using `reduce`:

Recall that `_ % 2 == 0`, `_ * 2`, and `_ * _` are *function literals*. The first two functions take a single parameter assigned to the placeholder `_`. The last function, which is passed to `reduce`, takes two parameters.

The `reduce` method is new for us. It is used to multiply the elements together. That is, it “reduces” the collection of integers to a single value. Here is the signature for `reduce`:

```
trait Seq[+A]:  
  def reduce[B >: A](op: (B, B) => B): B  
  ...
```

For instances of type `A`, the result will be either of type `A` or a *supertype* of `A`, represented by `B`. You pass a two-parameter function that accepts two `B` values and “reduces” them to a single `B` value. (If it’s confusing to differentiate `A` and `B`, just think of the simplest case where `B` is `A`, for example when both are `Double`.)

As `reduce` works through the collection, it could process the values from the left, from the right, or as parallel “trees”. This choice is unspecified for `reduce`. An implication of this is the requirement that the function passed to `reduce` is *associative*, like multiplication or addition of integers  $(a + b) + c = a + (b + c)$ , because we are not guaranteed that the collection elements will be processed in a particular order! Similar functions, like `reduceLeft` and `reduceRight` define the traversal order and also how the first and second parameters are used. Hence, they can be used when associativity is not required.

Back to our example, we use the `_` placeholder for both parameters, `_ * _` is equivalent to `(x,y) => x * y` for the function passed to `reduce`.

So, with a single line of code, we successfully “looped” through the list without the use of a mutable counter to track iterations, nor did we require mutable accumulators for the reduction as it was performed.

## Anonymous Functions, Lambdas, and Closures

Consider the following modifications of the previous example:

```
// src/script/scala/progscala3/fp/basics/HOFsClosures.scala

scala> var factor = 2
var factor: Int = 2

scala> val multiplier = (i: Int) => i * factor
val multiplier: Int => Int = Lambda$...

scala> val result1 =
  |   (1 to 10).filter(_ % 2 == 0).map(multiplier).reduce(_ * _)
val result1: Int = 122880

scala> factor = 3
  |   val result2 =
  |   (1 to 10).filter(_ % 2 == 0).map(multiplier).reduce(_ * _)
factor: Int = 3
val result2: Int = 933120
```

We define a variable, `factor`, to use as the multiplication factor, and we extract the previous anonymous function `_ * 2` into a function *value* called `multiplier` that now uses `factor`. Because functions are first-class in Scala, we can define values that are functions.

Running the same expression with different values for `factor` changes the results. Even though `multiplier` was an immutable function value, its behavior changed when `factor` changed.

Of the two variables in `multiplier`, `i` and `factor`, `i`, is called a *formal parameter* to the function. It is *bound* to a new value each time `multiplier` is called.

However, `factor` is not a formal parameter, but a *free variable*, a reference to a variable in the enclosing scope. Hence, the compiler creates a *closure* that encompasses (or “closes over”) `multiplier` and the external context of the unbound variables that `multiplier` references, thereby binding those variables as well.

If a function has no external references, it is trivially closed over itself. No external context is required.

This works even if `factor` is a local variable in some scope, like a method, and we passed `multiplier` to another scope, like another method. The free variable `factor` would be carried along for the ride.

This is illustrated in the following refactoring of the example, where `mult` returns a function of type `Int => Int`. That function references the local variable `factor` value, which goes out of scope once `mult` returns. That’s okay, because the 2 is captured in the returned function.

```
scala> def mult: Int => Int =  
      |   val factor = 2  
      |   (i: Int) => i * factor  
def mult: Int => Int  
  
scala> val result3= (1 to 10).filter(_ % 2 == 0).map(mult).reduce(_ *  
_)  
val result3: Int = 122880
```

There are a few partially overlapping terms that are used a lot:

### *Function*

An operation that is named or anonymous. Its code is not evaluated until the function is called. It may or may not have free (unbound) variables in its definition.

## *Lambda*

An anonymous (unnamed) function. It may or may not have free (unbound) variables in its definition.

## *Closure*

A function, anonymous *or* named, that closes over its environment to bind variables in scope to free variables within the function.

### WHY THE TERM LAMBDA?

The term *lambda* for anonymous functions originated in *lambda calculus*, where the greek letter lambda ( $\lambda$ ) is used to represent anonymous functions. First studied by Alonzo Church as part of his research in the mathematics of computability theory, lambda calculus formalizes the properties of functions as abstractions of calculations, which can be evaluated or *applied* when we bind values or substitute other expressions for variables. (The term *applied* is the origin of the default method name `apply` we've already seen.) Lambda calculus also defines rules for simplifying function expressions, substituting values for variables, etc.

Different programming languages use these and other terms to mean slightly different things. In Scala, we typically just say *anonymous function* or *function literal* for lambdas, while Java uses the term *lambda*. Also, we don't distinguish closures from other functions unless it's important for the discussion.

## Methods as Functions

While discussing variable capture in the preceding section, we defined an anonymous function `multiplier` as a value:

```
val multiplier = (i: Int) => i * factor
```

However, you can also use a method:

```
// src/script/scala/progscala3/fp/basics/HOFsClosures2.scala

scala> object Multiplier2:
  |   var factor = 2
  |   def multiplier(i: Int) = i * factor ❶
// defined object Multiplier2

scala> val result3 =
  |   (1 to 10).filter(_ % 2 ==
0).map(Multiplier2.multiplier).reduce(_ * _)
val result3: Int = 122880

scala> Multiplier2.factor = 3
  |   val result4 =
  |   (1 to 10).filter(_ % 2 ==
0).map(Multiplier2.multiplier).reduce(_ * _)
val result4: Int = 933120
```

- ❶ Now `multiplier` is defined as a method in an object and used below with `Multiplier2.multiplier`.

Despite the fact `Multiplier2.multiplier` is now a method, we use it just like a function, because it doesn't reference `this`. When a method is used where a function is required, we say that Scala *lifts* the method to be a function.

## Purity Inside Versus Outside

If we called  $\sin(x)$  thousands of times with the same value of  $x$ , it would be wasteful if it calculated the same value every single time.<sup>1</sup> Even in “pure” functional libraries, it is acceptable to perform internal optimizations like caching previously computed values (also called *memoization*).

Caching *is* a side effect, as the cache has to be modified, of course. If the performance benefits are worth it, caching must be

implemented in a thread-safe way, fully encapsulated from the user, and referentially transparent.

## Recursion

Recursion plays a larger role in functional programming than in imperative programming. Recursion is the “pure” way to implement “looping” without mutable loop counters.

Calculating factorials provides a good example, which we saw in “**Nesting Method Definitions and Recursion**”:

```
// src/script/scala/progscala3/typelessdomore/FactorialTailrec.scala
import scala.annotation.tailrec

def factorial(i: Int): Long =
  @tailrec
  def fact(i: Int, accumulator: Long): Long =
    if i <= 1 then accumulator
    else fact(i - 1, i * accumulator)

  fact(i, 1)

(0 to 5).foreach(i => println(s"$i: ${factorial(i)}"))
```

There are no mutable variables and the implementation is tail recursive.

## Tail Calls and Tail-Call Optimization

“**Nesting Method Definitions and Recursion**” explained the importance of tail recursion. More precisely, we have used *tail-call* self-recursion, where a function calls itself and the call is the final (“tail”) operation performed.

Tail-call self-recursion is very important, because it is the easiest kind of recursion to optimize by conversion into a loop by the compiler. This eliminates the function call for each iteration, thereby

improving performance and eliminating the potential for a stack overflow. You should use the `@tailrec` annotation to cause the compiler to throw an error if the annotated method is not actually tail recursive.

### WARNING

The tail-call optimization *won't* be applied when a method that calls itself might be overridden in a derived type. Hence, the recursive method must be `private` or `final`, or it must be nested in another method.

## Trampoline for Tail Calls

A *trampoline* is a loop that works through a list of functions, calling each one in turn. The metaphor of bouncing the functions back and forth off a trampoline is the source of the name.

Consider a recursion where a function `f1` doesn't call itself recursively, but instead it calls another function `f2`, which then calls `f1`, which calls `f2`, etc. This is obviously not self recursion, but it can also be converted into a loop using a *trampoline* data structure.

The Scala library has a `scala.util.control.TailCalls` object for this purpose.

The following example defines an inefficient way of determining if a number is even (adapted from the `TailCalls` *scaladoc*:

```
// src/script/scala/progscala3/fp/recursion/Trampoline.scala
```

```
scala> import scala.util.control.TailCalls._
|
| def isEven(xs: Seq[Int]): TailRec[Boolean] =
|   if xs.isEmpty then done(true) else tailcall(isOdd(xs.tail))
|
| def isOdd(xs: Seq[Int]): TailRec[Boolean] =
|   if xs.isEmpty then done(false) else tailcall(isEven(xs.tail))
|
```

```
| val eo = (1 to 5).map(i => (i, isEven(1 to i).result))

...
val eo: IndexedSeq[(Int, Boolean)] =
  Vector((1,false), (2,true), (3,false), (4,true), (5,false))
```

The code “bounces” back and forth between `isOdd` and `isEven` for each list element until the end of the list. If it hits the end of the list while it’s in `isEven`, it returns `true`. If it’s in `isOdd`, it returns `false`.

## Partially Applied Functions Versus Partial Functions

Consider the following session where we define and use three different string concatenation methods:

```
scala> def cat1(s1: String)(s2: String) = s1 + s2
      | def cat2(s1: String) = (s2: String) => s1 + s2
def cat1(s1: String)(s2: String): String
def cat2(s1: String): String => String

scala> cat1("hello")("world")
      | cat2("hello")("world")
val res1: String = helloworld
val res2: String = helloworld
```

When used, it appears that `cat1` and `cat2` are the same, but while `cat1` has two parameter lists, `cat2` is one parameter list, but it *returns a function* that takes the equivalent of the second parameter list. `cat3` has a single parameter list with two arguments. As invoked here, all three return the same result, `helloworld`.

Let’s look at a function definition that is equivalent to `cat2`:

```
scala> val cat2F = (s1: String) => (s2: String) => s1+s2
val cat2F: String => String => String = Lambda$...
```



It takes a bit of practice to read signatures like this. Keep in mind that the the type signature (when shown) is separated from the body by the `=`. Here is the same function with the type signature, where now we don't need the types on the right-hand side:

```
scala> val cat2F: String => String => String = s1 => s2 => s1+s2
      | (s1: String) => (s2: String) => s1+s2
val cat2F: String => String => String = Lambda$...
```

The binding is right to left. This definition is equivalent:

```
scala> val cat2Fb: String => (String => String) = s1 => s2 => s1+s2
val cat2Fb: String => String => String = Lambda$...
```

This means if we apply one argument, a `String`, we get back a function `String => String`:

```
scala> val c2f = cat2F("hello")
val c2f: String => String = Lambda$...

scala> c2f("world")
val res4: String = helloworld
```

Note the signature for `c2f`. We get the same behavior for `cat2Fb`.

Applying some, but not all arguments for the parameters is called *partial application*. A new function is returned that can be called with the remaining parameters. In Scala all the arguments for a particular parameter list have to be provided, but there is no limit to the number of parameter lists you can have and you can partially apply any *prefix* of them you want. You can't "skip over" parameter lists:

```
scala> def cat3(s1: String, s2: String) = s1 + s2
def cat3(s1: String, s2: String): String

scala> cat3("hello")
1 |cat3("hello")
  |^^^^^^^^^^^^^^
  |missing argument for parameter s2 of method cat3...
```

```
scala> def f123 =
  |   (one:String) => (two:String) => (three:String) =>
  s"$one|$two|$three"
def f123: String => String => String => String

scala> val f23 = f123("first")
val f23: String => String => String = Lambda$...

scala> val f3 = f23("second")
val f3: String => String = Lambda$...

scala> val s123 = f3("third")
val s123: String = "first|second|third"
```

## NOTE

A *partially applied function* is an expression with some, but not all of a function's parameter lists applied, returning a new function that takes the remaining parameter lists. In contrast, a *partial function* is a single-parameter function that is not defined for all values of the type of its parameter. The literal syntax for a partial function is one or more case match clauses (see “[Partial Functions](#)”).

## Currying and Uncurrying Functions

Methods and functions with multiple parameter lists have a fundamental property called *currying*, which is named after the mathematician Haskell Curry (for whom the Haskell language is named). Actually, Curry's work was based on an original idea of Moses Schönfinkel, but *Schönfinkeling* or maybe *Schönfinkelization* never caught on...

Currying is the transformation of a function that takes multiple parameters into a chain of functions, each taking a single parameter, as we just saw with `cat2` and `f123`. Scala also provides ways to convert between curried and “uncurried” or more typical syntax, with one parameter list and multiple parameters.

Recall that `cat3` was defined in the previous section with a single parameter list taking two parameters.

```
scala> val cat3Curried = cat3.curried
val cat3Curried: String => String => String =
scala.Function2$$Lambda$...

scala> cat3Curried("hello")("world")
val res4: String = helloworld
```

Note that the returned type signature is effectively the same as `cat2`, but the actual type is now a subtype of a Scala trait `scala.Function2`.

We can also “uncurry” a function:

```
scala> val cat3b = Function.uncurried(cat3Curried)
val cat3b: (String, String) => String = scala.Function$$$Lambda$...

scala> cat3b("hello", "world")
val res6: String = helloworld
```

So, `cat3b` is equivalent of the original `cat3` method.

To summarize, curried functions are defined with multiple argument lists, each with a *single* argument. Or they can be generated from other functions or methods by calling `curried` on them.

A practical use for currying is to specialize functions for particular types of data. You can start with an extremely general case, and use the curried form of a function to narrow down to particular cases. This might be done to create an easy-to-use version of the more general version, for a particular scenario.

As a simple example of this approach, the following code provides specialized forms of a base function that handles multiplication:

```
scala> def multiplier(i: Int)(factor: Int) = i * factor
def multiplier(i: Int)(factor: Int): Int
```

```
scala> val byFive = multiplier(5)
val byFive: Int => Int = Lambda$...
```

```
scala> val byTen = multiplier(10)
val byTen: Int => Int = Lambda$...
```

```
scala> val ten = byFive(2)
val ten: Int = 10
```

```
scala> val thirty = byTen(3)
val thirty: Int = 30
```

## Tupled and Untupled Functions

One scenario you'll encounter occasionally is when you have data in a tuple, let's say a three-element tuple, and you need to call a three-parameter function:

```
scala> def mult(d1: Double, d2: Double) = d1 * d2
def mult(d1: Double, d2: Double): Double
```

```
scala> val d23 = (2.2, 3.3)
val d23: (Double, Double) = (2.2,3.3)
```

```
scala> val d = mult(d23._1, d23._2)
val d: Double = 7.26
```

It's tedious extracting the tuple elements like this.

Because of the literal syntax for tuples, like (2.2, 3.3), there seems to be a natural symmetry between tuples and function parameter lists. We would love to have a new version of `mult` that takes the tuple itself as a single parameter. Fortunately, the `scala.Function` object provides `tupled` and `untupled` methods for us:

```
scala> val multTupled1 = Function.tupled(mult)
      | val multTupled2 = mult.tupled
val multTupled1: ((Double, Double)) => Double =
scala.Function$$$Lambda$...
val multTupled2: ((Double, Double)) => Double =
```

```

scala.Function$$$Lambda$...

scala> val d21 = multTupled1(d23)
      | val d22 = multTupled2(d23)
val d21: Double = 7.26
val d22: Double = 7.26

scala> val mult2 = Function.untupled(multTupled1)
val mult2: (Double, Double) => Double = scala.Function$$$Lambda$...

scala> val d3 = mult2(d23._1, d23._2)
val d3: Double = 7.26

```

However, there isn't a corresponding `multTupled1.untupled` available. Also, `Function.tupled` and `Function.untupled` only work for arities between two and five, inclusive, an arbitrary limitation. Above arity five, you can call `myfunc.tupled` up to arity 22, which is another historical limitation. In Scala 2, functions and tuples were limited to 22 elements.

Fortunately, Scala 3 introduces several new types to eliminate these arbitrary arities:

- `scala.Tuple` - Tuples of arbitrary arity.
- `scala.Tuple` - A type class that defines methods `tupled`, for converting to a function that takes a tuple argument, and `untupled` for arbitrary arities.

We'll revisit `Tuple` in [Chapter 13](#). The code examples contain `src/script/scala/progscale3/fp/basics/Tupling.scala`, which explores the use of `TupledFunction`.

## Partial Functions vs. Functions Returning Options

In “[Partial Functions](#)”, we discussed the synergy between partial functions and total functions that return an `Option`, either

Some(value) or None, corresponding to the case where the partial function can return a value and when it can't, respectively. Scala provides transformations between partial functions and total functions returning options:

```
scala> val finicky: PartialFunction[String,String] =
  |   case "finicky" => "FINICKY"
val finicky: PartialFunction[String, String] = <function1>

scala> val fin = finicky("finicky")
val fin: String = FINICKY

scala> scala> finicky("other")
scala.MatchError: other (of class java.lang.String)
...
```

Now “lift” it to a total function returning Option[String]:

```
scala> val finickyOption = finicky.lift
val finickyOption: String => Option[String] = <function1>

scala> val fo1 = finickyOption("finicky")
val fo1: Option[String] = Some(FINICKY)

scala> val fo2 = finickyOption("other")
val fo2: Option[String] = None
```

We can go from a total function returning an option using unlift:

```
scala> val finicky2 = Function.unlift(finickyOption)
val finicky2: PartialFunction[String, String] = <function1>

scala> val fin2 = finicky2("finicky")
val fin2: String = FINICKY

scala> finicky2("other")
scala.MatchError: other (of class java.lang.String)
...
```

Note that unlift only works on single parameter functions returning an option.

Lifting a partial function is especially useful when we would prefer handling optional values instead of thrown `MatchErrors`. Conversely, unlifting is useful when we want to use a regular function returning an option in a context where we need a partial function.

## Functional Data Structures

Functional programming emphasizes the use of a core set of data structures and algorithms, where the algorithms, usually implemented with functions, are maintained separately from the data structures.

In contrast, object oriented programming emphasizes putting data and functions (methods) together, often hiding the actual structure of the data. OOP also puts greater emphasis on types (classes) that closely model the application domain concepts.

In practice, OOP languages led to the proliferation of ad hoc classes that undermined the promised goal of code reuse. Functional data structures and algorithms are much more reusable and the programs written with them tend to be more concise, too.

This doesn't mean that OOP is bad. Scala has demonstrated how to combine the two paradigms. In particular, Scala's object model provides an effective *modularity* system that's needed for large programs, which we'll explore in [Chapter 9](#).

Back to data structures and algorithms, the minimum set typically includes sequential collections like lists, vectors, and arrays, hash-based structures like maps and sets, and perhaps others, like tree structures. Each collection supports a subset of the same higher-order, side effect-free functions, called *combinators*, such as `map`, `filter`, `fold`, etc. Once you learn these combinators, you can pick the appropriate collection to meet your requirements for data access and performance, then apply the same familiar combinators to manipulate that data. These collections are the most successful tools

for code reuse and composition that we have in all of software development. Note that SQL, the immensely popular data query and manipulation tool, can be considered a limited functional language with a set-theory orientation.

## Sequences

Let's look at a few of the most common data structures used in Scala, focusing on their functional characteristics. We'll explore the Scala collections library in more thorough detail in *Chapter 13*.

Many data structures are *sequential*, where the elements can be traversed in a predictable order, which might be the order of insertion or sorted in some way. The `collection.Seq` trait is the abstraction for all mutable and immutable sequential types. Child traits `collection.mutable.Seq` and `collection.immutable.Seq` represent mutable and immutable sequences, respectively.

We have used the `Seq.apply()` method frequently in examples. It constructs a linked `List`.

By convention, when adding an element to a list, it is prepended to the existing list, becoming the “head” of the new list. The existing list remains unchanged and it becomes the “tail” of the new list.

**Figure 7-1** shows two lists, `List(1,2,3,4,5)` and `List(2,3,4,5)`, where the latter is the tail of the former.




Figure 7-1. Two linked lists



Note that we created a new list from an existing list using an  $O(1)$  operation. We shared the tail with the original list and just constructed a new link from the new head element to the old list. This is our first example of an important idea in functional data structures, sharing a structure to minimize the cost of making copies. To support immutability, we need the ability to make copies with minimal cost.

Lists are immutable, so the tail list is unaffected by prepending elements to construct a new list.

Any operation that requires list traversal, such as computing the size or accessing an arbitrary element (e.g., `mylist(5)`) is  $O(N)$ . Other types that implement `Seq`, like `Vector`, can be faster.

The following example demonstrates ways to construct Lists:

```
// src/script/scala/progscala3/fp/datastructs/Sequence.scala

scala> val seq1 = Seq("Programming", "Scala")           ❶
val seq1: Seq[String] = List(Programming, Scala)

scala> val seq2 = "Programming" +: "Scala" +: Nil       ❷
val seq2: Seq[String] = List(Programming, Scala)

scala> seq2.head                                         ❸
| seq2.tail
val res0: String = Programming
val res1: Seq[String @uncheckedVariance] = List(Scala)
```

- ❶ `Seq.apply()` takes a variable argument list and constructs a `List`.
- ❷ The `+:` method on `Seq` used in infix notation, with `Nil` as the empty list.
- ❸ Getting the head and tail of the sequence.<sup>2</sup>

We encountered the “cons” (for “construct”) method, `+:`, in “**Left vs. Right Associative Methods**”. It binds to the *right*, because the name

ends with `:`. If we used regular method syntax, we would write `list.+=(element)`.

The case object `Nil` is a subclass of `List` that is a convenient object when an empty list is required. Note that the chain of cons operators requires a list on the far right, empty or non-empty. `Nil` is equivalent to `List.empty[Nothing]`, where `Nothing` is the subtype of *all* other types in Scala.

The construction of `seq2` is parsed as follows, with parentheses added to make the ordering of construction explicit. Each set of parentheses encloses an immutable list:

```
val seq2b = ("Programming" +: ("Scala" +: (Nil)))
```

There are many more methods available on `collection.Seq` for concatenation of sequences, transforming them, etc.

What about other types that implement `Seq`? Let's consider `immutable.Vector`, which is important because some operations, like `head`, `tail`, `+: (prepend)`, and `:+ (append)` on vectors are  $O(1)$  and all other operations are  $O(\log(N))$ , worst case. The elements are stored in a tree structure, providing these performance characteristics.

```
// src/script/scala/progscala3/fp/datastructs/Vector.scala
```

```
scala> val vect1 = Vector("Programming", "Scala")
      | val vect2 = "People" +: "should" +: "read" +: Vector.empty ❶
      |
val vect1: Vector[String] = Vector(Programming, Scala)
val vect2: Vector[String] = Vector(People, should, read)

scala> vect2.head
      | vect2.tail
val res2: String = People
val res3: Vector[String] = Vector(should, read)

scala> val vect3 = seq1.toVector
val vect3: Vector[String] = Vector(Programming, Scala) ❷
```

- ❶ Use `Vector.empty` instead of `Nil` so that the whole sequence is constructed as a `Vector` from the beginning.
- ❷ An alternative; take an existing sequence and convert it to a `Vector`. This is less efficient as two collections with identical elements have to be constructed.

You'll note that we didn't import any of these collection types. `Predef` exposes several immutable collection types without requiring explicit import statements. Other examples include `immutable.Map` and `immutable.Set`

Until Scala 2.13, the default `Seq` trait in scope was the trait `collection.Seq`, the parent trait of both `collection.immutable.Seq` and `collection.mutable.Seq`. While this parent trait doesn't expose mutation methods, it opened the door for mutable sequences, especially `Arrays`, to enter your code when perhaps you wanted only immutable sequences. As of Scala 2.13, the default is changed to `immutable.Seq`.

## Maps

Another common data structure is the `Map`, used to hold pairs of keys and values, where the keys must be unique. `Maps` and the common `+map` method reflect a similar concept, associating a key with a value and associating an input element with an output element, respectively.

```
// src/script/scala/progscala3/fp/datastructs/Map.scala

scala> val stateCapitals = Map(
  |   "Alabama" -> "Montgomery",
  |   "Alaska"  -> "Juneau",
  |   "Wyoming" -> "Cheyenne")
val stateCapitals: Map[String, String] =
  Map(Alabama -> Montgomery, Alaska -> Juneau, Wyoming -> Cheyenne)
```

when we map over the values, they are passed to the anonymous function as pairs. The order is undefined for the Map trait, but may be defined by particular implementations, such as **SortedMap**:

```
scala> val lengths = stateCapitals map(kv => (kv._1, kv._2.length))
val lengths: Map[String, Int] = Map(Alabama -> 10, Alaska -> 6,
Wyoming -> 8)
```

Sometimes it's more convenient to pattern match on the key and value:

```
scala> val caps = stateCapitals map { case (k, v) => (k,
v.toUpperCase) }
val caps: Map[String, String] = Map(Alabama -> MONTGOMERY, ...)
```

You can add new key-value pairs, possibly overriding an existing definition, or add multiple key-value pairs. All of these operations return a new Map:

```
scala> val stateCapitals2a = stateCapitals + ("Virginia" ->
"Richmond")
val stateCapitals2a: Map[String, String] = Map(..., Virginia ->
Richmond)

scala> val stateCapitals2b = stateCapitals + ("Alabama" ->
"MONTGOMERY")
val stateCapitals2b: Map[String, String] = Map(Alabama -> MONTGOMERY,
...)

scala> val stateCapitals2c = stateCapitals ++ Seq(
  | "Virginia" -> "Richmond", "Illinois" -> "Springfield")
val stateCapitals2c: Map[String, String] = HashMap(
  Alaska -> Juneau, Illinois -> Springfield, Wyoming -> Cheyenne,
  Virginia -> Richmond, Alabama -> Montgomery)
```

We never called `new Map(...)`, because Map is a trait. Instead, we used the companion object method, `Map.apply`, to construct an instance of a concrete implementation class that is optimal for the data set, usually based on size. In fact, there are concrete map

classes for one, two, three, and four key-value pairs! Why? For very small instances, it's more efficient to have custom implementations than use the hash-based implementations for the general case.

## Sets

Like Maps, **Sets** are unordered collections, so they aren't sequences, per se, but we can still use some sequential methods on them, like `map`. They also require each element to be unique:

```
// src/script/scala/progscala3/fp/datastructs/Set.scala

scala> val states = Set("Alabama", "Alaska", "Wyoming")
scala> val lengths = states map (st => st.length)
val lengths: Set[Int] = Set(7, 6)    // Two names are 7 characters
long

scala> val states2 = states + "Virginia"
val states2: Set[String] = Set(Alabama, Alaska, Wyoming, Virginia)

scala> val states3 = states ++ Seq("New York", "Illinois", "Alaska")
val states3: Set[String] = HashSet(
  Alaska, Alabama, New York, Illinois, Wyoming) // Alaska already
present
```

## Traversing, Mapping, Filtering, Folding, and Reducing

Traversing a collection is a universal operation for working with the contents. Some collections are iterable by design, like `List` and `Vector`, where the ordering is defined. Others, like hash-based `Set` and `Map` types can be converted to iterable collections as needed. `Option` is a collection with zero or one element, implemented by the `None` and `Some` subtypes. Even, **Product**, an abstraction implemented by tuples and case classes, has the notion of iterating through its elements. Your own container types can and should be designed for traversability, when possible. There are also infinite (lazy) sequences

and iterators (essentially “views” on a collection), which can be traversed.

This “protocol” is one of the most reusable, flexible, and powerful concepts in all of programming. Scala programs use this capability extensively. In the sense that all programs boil down to “data in, data out”, traversing is a superpower.

## Traversal

The traversal method for performing only side effects is `foreach`. It is defined in `collection.IterableOnceOps`. Many of the operations we’ll discuss are defined in “mix-in” traits like this, which we won’t discuss further until *later chapters*). Also, we’ll use the generic term “collection” from now on to refer to all the iterable types we just mentioned.

This is the signature for `foreach`:

```
trait IterableOnceOps[A] { // Some details omitted.
  ...
  def foreach[U](f: A => U): Unit = {...}
  ...
}
```

Subclasses of `IterableOnceOps` may redefine the method to exploit local knowledge for better performance.

It actually doesn’t matter what the `U` is, the return type of the function. The output of `foreach` is `Unit`, so it can only perform side effects. This means `foreach` isn’t really consistent with the functional programming principle of writing pure, side effect-free code, but it is useful for writing output and other tasks. Because it takes a function parameter, `foreach` is a *higher-order function*, as are all the operations we’ll discuss.

Performance is  $O(N)$  in the number of elements. Here are examples using it with our `stateCapitals` Map:

```
scala> var str1 = ""
      | stateCapitals.foreach { case (k, v) => str1 += s"${k}: ${v}, "
      }
var str2: String =
  "Alabama: Montgomery, Alaska: Juneau, Wyoming: Cheyenne, "
```

Since nothing is returned, you could say that `foreach` transforms each element into zero elements, or *one to zero* for short.

## Mapping

We have used `map` many times already. Mapping is *one to one*, where for each input element an output element is returned. Hence, an “invariant” is the size of the input and output collections must be equal. For any collection `C[A]` for elements of type `A`, `map` has the following “logical” signature:

```
class C[A]:
  def map[B](f: (A) => B): C[B]
```

The real signature is more involved, in part because of the need to handle subtyping, but the details don’t concern us now. This will be true for all the combinators we explore below, too. I’ll simplify the signatures to focus on the concepts, then return to more of the implementation details in *a later chapter*.

Note the signature for `f`. It must transform an `A` to a `B`. Usually the type `B` is inferred from this function, while `A` is known from the original collection:

```
scala> val in = Seq("one", "two", "three")
val in: Seq[String] = List(one, two, three)

scala> val out1 = in.map(_.length)
val out1: Seq[Int] = List(3, 3, 5)

scala> val out2 = in.map(s => (s, s.length))
val out2: Seq[(String, Int)] = List((one,3), (two,3), (three,5))
```

Another way to think of `map` is that it transforms `Seq[A] => Seq[B]`. This fact is obscured by the object syntax, e.g., `myseq.map(f)`. If instead we had a separate module of functions that take `Seq` instances as parameters, it would look something like this:

```
// src/script/scala/progscala3/fp/combinators/MapF.scala

object MapF:
  def map[A,B](f: (A) => B)(list: Seq[A]): Seq[B] = list map f
```

- ❶ “F” used to avoid conflict with `Map` data structure.
- ❷ A `map` that takes the transforming function as the first parameter list, then the collection. I’m cheating and using `Seq.map` to implement the transformation for simplicity.

Now use it. Note the types of the values returned when using `MapLF.map` and `MapF.map`:

```
scala> val intToString = (i:Int) => s"N=$i"
      | val input = Seq(1, 2, 3, 4)

scala> val ff = MapF.map(intToString)
      | val seq = ff(input)
val ff: Seq[Int] => Seq[String] = Lambda$...
val seq: Seq[String] = List(N=1, N=2, N=3, N=4)
```

Partial application of `MapF` *lifts* a function `Int => String` into a new function `Seq[Int] => Seq[String]`, the type of `ff`. When we call this new function with a collection argument, we get a new collection.

Put another way, partial application of `MapF.map` is a transformer of *functions*. There is a “symmetry” between the idea of mapping over a collection with a function to create a new collection vs. mapping a function to a new form that is able to transform one collection into another.



## Flat Mapping

A generalization of the Map operation is `flatMap`, where we generate zero to many new elements for each element in the original collection. In other words, while `map` is *one to one*, `flatMap` is *one to many*.

Here the logical signature, with `map` for comparison, for some collection `C[A]`:

```
class C[A]:  
  def flatMap[B](f: A => Seq[B]): C[B]  
  def map[B](f: (A) => B): C[B]
```

We pass a function that returns a collection, instead of a single element, and `flatMap` “flattens” those collections into one collection, `C[B]`. It’s not required for `f` to return collections of the same type as `C`.

Consider this example that compares `flatMap` and `map`. First we map over a `Range`. For example integer, we construct a new range from that value until 5:

```
// src/script/scala/progscala3/fp/datastructs/FlatMap.scala  
  
scala> val seq = 0 until 5  
val seq: Range = Range 0 until 5  
  
scala> val seq1 = seq.map(i => i until 5)  
val seq1: IndexedSeq[Range] = Vector(Range 0 until 5, Range 1 until 5, ...)
```

Now use another method, `flatten`:

```
scala> val seq2 = seq1.flatten  
val seq2: IndexedSeq[Int] = Vector(0, 1, 2, 3, 4, 1, 2, 3, 4, 2, 3, 4, 3, 4, 4)
```

(Vector was used by Scala itself, because it has  $O(1)$  append.)

Finally, use `flatMap` instead. Notice that it effectively combines `map` followed by `flatten`:

```
scala> val seq3 = seq.flatMap(i => i until 5)
val seq3: IndexedSeq[Int] = Vector(0, 1, 2, 3, 4, 1, 2, 3, 4, 2, 3, 4, 3, 4, 4)
```

Note that `flatMap` won't flatten elements beyond one level.

`flatMap` may not seem like a particularly useful operation, but it has far greater significance than might be obvious. Consider the following example using `Options` where we simulate validating account information a user might provide in a form:

```
// src/script/scala/progscala3/fp/datastructs/FlatMapValidate.scala
```

```
scala> case class Account(name: String, password: String, age: Int)
```

❶

```
scala> val validName: Account => Option[Account] =
```

❷

```
  a => if a.name.length > 0 then Some(a) else None
```

```
  val validPwd: Account => Option[Account] =
```

```
    a => if a.password.length > 0 then Some(a) else None
```

```
  val validAge: Account => Option[Account] =
```

```
    a => if a.age > 18 then Some(a) else None
```

❸

```
scala> val accounts = Seq(
  |   Account("bucktrends", "1234", 18),
  |   Account("", "1234", 29),
  |   Account("bucktrends", "", 29),
  |   Account("bucktrends", "1234", 29))
```

```
scala> val validated = accounts.map { account =>
```

❹

```
  |   Some(account).flatMap(validName).flatMap(validPwd).flatMap(validAge)
  | }
```

```
val validated: Seq[Option[Account]] =
```

```
  List(None, None, None, Some(Account(bucktrends,1234,29)))
```

- ❶ Define a case class for Account form data.
- ❷ Define separate validation functions for each field in an Account instance. Note that each has exactly the same signature, other than the name.
- ❸ Disallow minors.
- ❹ Map over the accounts, testing each one with the validators.

Note that only one account passes validation. None is returned for the rest.

Using `Options` and `flatMap` this way might seem like overkill. We could just call simpler validation methods that don't return `Option`. But if we made the list of validators more configurable, perhaps defined in an external library, then using this protocol let's us sequence together the separate tests without having to add logic to handle each success or failure ("if this passes, try that, ..."). As shown, the first three instances fail at one validator, while the last one passes all of them.

What's missing here is information about which validator failed and why for each case, which you would want to show to the user in a form. We'll see alternatives in [Chapter 8](#) that fill in this gap, but still leverage the same `flatMap` approach.

Using `flatMap` in this way is extremely common in Scala code. We'll see many more examples.

## Filtering

It is common to traverse a collection and extract a new collection from it with elements that match certain criteria.

For any collection `C[A]`:

```
class C[A]:  
  def filter(f: A => Boolean): C[A]
```

Hence, filtering is *one to zero or one*. For example:

```
scala> val numbers = Map("one" -> 1, "two" -> 2, "three" -> 3)  
  
scala> val tnumbers = numbers filter { case (k, v) =>  
  k.startsWith("t") }  
val tnumbers: Map[String, Int @uncheckedVariance] = Map(two -> 2,  
three -> 3)
```

Most collections that support `filter` have a set of related methods. Note that some of these methods won't return for infinite collections and some might return different results for different invocations unless the collection type is ordered. The descriptions are adapted from the Scaladocs:

```
def drop(n: Int): C[A]
```

Return a new collection without the first `n` elements. The returned collection will be empty if this collection has less than `n` elements.

```
def dropWhile (p: (A) => Boolean): C[A]
```

Drop the longest prefix of elements that satisfy the *predicate* `p`. The new collection returned starts with the first element that *doesn't* satisfy the predicate.

```
def exists (p: (A) => Boolean): Boolean
```

Return true if the predicate holds for at least one of the elements of this collection. Return false, otherwise.

```
def filter (p: (A) => Boolean): C[A]
```

Return a collection with all the elements of this collection that satisfy a predicate. The order of the elements is preserved.

```
def filterNot (p: (A) => Boolean): C[A]
```

The “negation” of `filter`; select all elements that do not satisfy the predicate.

```
def find (p: (A) => Boolean): Option[A]
```

Find the first element of the collection that satisfies the predicate, if any. Return an `Option` containing that first element or `None`, if no element exists satisfying the predicate.

```
def forall (p: (A) => Boolean): Boolean
```

Return `true` if the predicate holds for all elements of the collection. Return `false`, otherwise.

```
def partition (p: (A) => Boolean): (C[A], C[A])
```

Partition the collection in two new collections according to the predicate. Return the pair of new collections where the first one consists of all elements that satisfy the predicate and the second one consists of all elements that don't. The relative order of the elements in the resulting collections is the same as in the original collection.

```
def take (n: Int): C[A]
```

Return a collection with the first `n` elements. If `n` is greater than the size of the collection, return the whole collection.

```
def takeWhile (p: (A) => Boolean): C[A]
```

Take the longest prefix of elements that satisfy the predicate.

Note that concatenating the results of `take` and `drop` yields the original collection. Same for `takeWhile` and `dropWhile`. Also, the same predicate used with `partition` would return the same two collections:

```
// src/script/scala/progscala3/fp/datastructs/FilterOthers.scala
```

```
val seq = 0 until 10
val f = (i: Int) => i < 5

for i <- 0 until 10 do
  val (l1,r1) = (seq.take(i), seq.drop(i))
  val (l2,r2) = (seq.takeWhile(f), seq.dropWhile(f))
  val (l3,r3) = seq.partition(f)
  assert(seq == l1++r1)
  assert(seq == l2++r2)
  assert(seq == l3++r3)
  assert(l2 == l3 && r3 == r3)
```

## Folding and Reducing

Let's discuss folding and reducing together, because they're similar. Both are operations for “shrinking” a collection down to a smaller collection or a single value, so they are *many to one* operations.

Folding starts with an initial “seed” value and processes each element in the context of that value. In contrast, reducing doesn't start with a user-supplied initial value. Rather, it uses one of the elements as the initial value, usually the first or last element:

```
// src/script/scala/progscala3/fp/datastructs/FoldReduce.scala
```

```
scala> val int1 = Seq(1,2,3,4,5,6).reduce(_ + _) ❶
      |
      | val int2 = Seq(1,2,3,4,5,6).fold(0)(_ + _) ❷
val int1: Int = 21
val int2: Int = 21

scala> val int3 = Seq.empty[Int].reduce(_ + _) ❸
java.lang.UnsupportedOperationException: empty.reduceLeft
...

scala> val int4 = Seq(1).reduce(_ + _) ❹

scala> val opt1 = Seq.empty[Int].reduceOption(_ + _) ❺
val opt1: Option[Int] = None
```

```
scala> val opt2 = Seq(1,2,3,4,5,6).reduceOption(_ * _)
val opt2: Option[Int] = Some(720)
```

- ❶ Reduce a sequence of integers by adding them together, returning 21.
- ❷ Do the same calculation with `fold`, which takes a “seed” value, 0 in this case.
- ❸ An attempt to reduce an empty sequence causes an exception, because there needs to be at least one element for the “reduction”.
- ❹ Having one element is sufficient. The returned value is just the single element.
- ❺ A safer way to invoke `reduce` if you aren’t sure if the collection is empty. `Some(value)` is returned if the collection is not empty. Otherwise, `None` is returned.

If you think about it, reducing can only return the closest, common parent type<sup>3</sup> of the elements. If the elements all have the same type, the final output will have that type. In contrast, because folding takes a seed value, it offers more options for the final result. Here are fold examples that implement mapping, flat mapping, and filtering:

```
// src/script/scala/progscala3/fp/datastructs/FoldMap.scala
```

```
scala> val vector = Vector(1, 2, 3, 4, 5, 6)
```

```
scala> val vector2 = vector.foldLeft(Vector.empty[String]) {
```

```
❶
  |   (vector, x) => vector := ("[" + x + "]")
  | }
```

```
val vector2: Vector[String] = Vector([1], [2], [3], [4], [5], [6])
```

```
scala> val vector3 = vector.foldLeft(Vector.empty[Int]) {
```

```
❷
```

```

    | (vector, x) => if x % 2 == 0 then vector else vector :+ x
    | }
val vector3: Vector[Int] = Vector(1, 3, 5)

scala> val vector4a = vector.foldLeft(Vector.empty[Seq[Int]]) {
❸
    | (vector, x) => vector :+ (1 to x)
    | }
val vector4a: Vector[Seq[Int]] =
    Vector(Range 1 to 1, Range 1 to 2, Range 1 to 3, Range 1 to 4, ...)

scala> val vector4 = vector4a.flatten
❹
val vector4: Vector[Int] =
    Vector(1, 1, 2, 1, 2, 3, 1, 2, 3, 4, 1, 2, 3, 4, 5, 1, 2, 3, 4, 5,
    6)

scala> val vector2b = vector.foldRight(Vector.empty[String]) {
❺
    | (x, vector) => ("[" + x + "]") +: vector
    | }
val vector2b: Vector[String] = Vector([1], [2], [3], [4], [5], [6])

scala> vector2 == vector2b
val res0: Boolean = true

```

- ❶ “Map” over a vector, creating strings [n]. While fold doesn’t guarantee a particular traversal order, foldLeft traverses left to right. Note the anonymous function. For foldLeft, the first parameter is the “accumulator”, the new vector we are building, and the second parameter is an element. We return a new vector with the element string appended to it, using :+. This returned vector will be passed in as the new accumulator on the next iteration or returned to vector2.
- ❷ “Filter” a vector, returning just the odd values. Note that for even values, the anonymous function just returns the current accumulator.
- ❸ A “map” that creates a vector of ranges.



- ④ Flattening the previous output vector, thereby implementing `flatMap`.
- ⑤ Traverse from the right. Note the parameters are reversed in the anonymous function and we prepend to the accumulator. Hence `vector2` and `vector2b` are equal.

Folding really is the universal operator, because it can be used to implement all the others, where `map`, `filter`, and `flatMap` implementations are shown here. (Try doing `foreach` yourself.)

Here are the signatures and descriptions for the various fold and reduce operations available on the iterable collections. The descriptions are paraphrased from the Scaladocs. I'll explain the type parameter `A1 >: A` below:

```
def fold[A1 >: A](z: A1)(op: (A1, A1) => A1): A1
```

Fold the elements of this collection using the specified *associative* binary operator `op`. The order in which operations are performed on elements is unspecified and may be nondeterministic. However, for most ordered collections like `Lists`, `fold` is equivalent to `foldLeft`.

```
def foldLeft[B](z: B)(op: (B, A) => B): B
```

Apply `op` to the start value `z` and all elements of this collection, going left to right.

```
def foldRight[B](z: B)(op: (A, B) => B): B
```

Apply `op` to all elements of this collection and a start value, going right to left.

```
def reduce[A1 >: A](op: (A1, A1) => A1): A1
```

Reduce the elements of this collection using the specified *associative* binary operator `op`. The order in which operations are

performed on elements is unspecified and may be nondeterministic. However, for most ordered collections like Lists, reduce is equivalent to reduceLeft. An exception is thrown if the collection is empty.

```
def reduceLeft[A1 >: A](op: (A1, A1) => A1): A1
```

Apply op to all elements of this collection, going left to right. An exception is thrown if the collection is empty.

```
def reduceRight[A1 >: A](op: (A1, A1) => A1): A1
```

Apply op to all elements of this collection going right to left. An exception is thrown if the collection is empty.

```
def optionReduce[A1 >: A](op: (A1, A1) => A1): Option[A1]
```

Like reduce, but return None if the collection is empty or Some(...) if not.

```
def reduceLeftOption[B >: A](op: (B, A) => B): Option[B]
```

Like reduceLeft, but return None if the collection is empty or Some(...) if not.

```
def reduceRightOption[B >: A](op: (A, B) => B): Option[B]
```

Like reduceRight, but return None if the collection is empty or Some(...) if not.

```
def scan[B >: A](z: B)(op: (B, B) => B): C[B]
```

Compute a prefix scan of the elements of the collection. Note that the neutral element z may be applied more than once. (I'll show an example below.)

```
def scanLeft[B >: A](z: B)(op: (B, B) => B): C[B]
```

Produce a collection containing cumulative results of applying the operator op going left to right.

```
def scanRight[B >: A](z: B)(op: (B, B) => B): C[B]
```

Produce a collection containing cumulative results of applying the operator `op` going right to left.

```
def product[B >: A](implicit num: math.Numeric[B]): B
```

Multiply the elements of this collection together, as long as the elements have an implicit conversion to type `Numeric[A]`, which effectively means `Int`, `Long`, `Float`, `Double`, and `BigInt`, etc. We discussed such conversions in “[Constraining Allowed Instances](#)”.

```
def sum[B >: A](implicit num: math.Numeric[B]): B
```

Similar to `product`; add the elements together.

```
def mkString: String
```

Display all elements of this collection in a string. This is a custom implementation of `fold` used for conveniently generating a custom string from the collection. There will be no delimiter between elements in the string.

```
def mkString(sep: String): String
```

Display all elements of this collection in a string using the specified separator (`sep`) string.

```
def mkString(start: String, sep: String, end: String): String
```

Display all elements of this collection in a string using the specified start (prefix), `sep` (separator), and end (suffix) strings.

Where you see the type parameter `A1 >: A`, it means that the final output type `A1` must be a parent type of `A`, although usually they will be the same.

Pay careful attention to the parameters passed to the anonymous functions for various `reduce`, `fold`, and `scan` methods. For the `Left`

methods, e.g., `foldLeft`, the *first* parameter is the accumulator and the collection is traversed left to right. For the Right functions, e.g., `foldRight`, the *second* parameter is the accumulator and the collection is traversed right to left. For the methods like `fold` and `reduce` that aren't left- or right-biased, the traversal order and which parameter is the accumulator are *undefined*, but usually they delegate to the left-biased methods.

The `fold` and `scan` methods can output a completely different type, based on the seed value, while the `reduce` methods always return the same element type or a supertype.

None of these functions will terminate for infinite collections. Also, they might return different results for different runs if the collection is not a sequence (i.e., the elements are not stored in a defined order) or the operation isn't associative.

The `scan` methods are useful for processing successive subsets of a collection. Consider the following example:

```
scala> val ints = Seq(1,2,3,4,5,6)

scala> val plus = ints.scan(0)(_ + _)
val plus: Seq[Int] = List(0, 1, 3, 6, 10, 15, 21)

scala> val mult = ints.scan(1)(_ * _)
val mult: Seq[Int] = List(1, 1, 2, 6, 24, 120, 720)
```

For the `plus` example, first the seed value 0 is emitted, followed by the first element plus the seed,  $1 + 0 = 1$ , followed by the second element plus the previous value,  $1 + 2 = 3$ , and so on. Try rewriting `scan` using `foldLeft`.

Finally, the three `mkString` methods are special-case versions of `fold` for generating Strings from collections. They are also quite handy when the default `toString` for a collection isn't what you want.

## Left Versus Right Folding

There are nonobvious differences going from left to right when folding or reducing. We'll focus on folding, but the same remarks apply to reducing, too.

Recall that `fold` does not guarantee a particular traversal order and the function passed to it must be associative, while `foldLeft` and `foldRight` guarantee traversal order. Consider the following examples:

source,scala]

```
// src/script/scala/progscala3/fp/datastructs/Fold.scala
```

```
scala> val seq6 = Seq(1,2,3,4,5,6)
      | val int1 = seq6.fold(0)(_ + _)
      | val int2 = seq6.foldLeft(0)(_ + _)
      | val int3 = seq6.foldRight(0)(_ + _)
val seq6: Seq[Int] = List(1, 2, 3, 4, 5, 6)
val int1: Int = 21
val int2: Int = 21
val int3: Int = 21
```

All yield the same result, which is hopefully not surprising. It doesn't matter what order we traverse the sequence, as addition is associative and also commutative.

Let's explore examples where order matters. First, recall that for many sequences, `fold` just calls `foldLeft`. So, we'll focus on `foldLeft` and `foldRight`. Second, while we used the same anonymous function above `_ + _`, recall that the parameters passed to this function are actually reversed for `foldLeft` vs. `foldRight`. To spell it out:

```
val int4 = seq6.foldLeft(0)((accum: Int, element: Int) => accum +
element)
val int5 = seq6.foldRight(0)((element: Int, accum: Int) => element +
accum)
```

For addition, the names are meaningless. Suppose instead that we build up a string from the sequence by folding. We'll add parentheses to show the order of evaluation:

```
scala> val left = (accum: String, element: Int) => s"($accum
$element)"
      | val right = (element: Int, accum: String) => s"($accum
$element)"
      | val right2 = (element: Int, accum: String) => s"($element
$accum)"
      |
      | val strLeft = seq6.foldLeft("(0)")(left)
      | val strRight = seq6.foldRight("(0)")(right)
      | val strRight2 = seq6.foldRight("(0)")(right2)
val strLeft: String = ((((((0) 1) 2) 3) 4) 5) 6)
val strRight: String = ((((((0) 6) 5) 4) 3) 2) 1)
val strRight2: String = (1 (2 (3 (4 (5 (6 (0)))))))
```

Note that the bodies of `left` and `right` are the same, while the parameter list is reversed. I wrote them this way so all the parentheses would line up the same way. Clearly the numbers are different. However, `right2` reverses the way the arguments are used in the body, so the parentheses come out very different, but the order of the numbers is the same in `strLeft` and `strRight2`.

It turns out that `foldLeft` has an important advantage over `foldRight`: the left traversals are tail recursive, so they can benefit from Scala's tail-call optimization.

Recall that a tail call must be the last operation in a recursion. Looking at the output for `strRight2`, the outermost string construction `(1 ...)` can't be performed until *all* of the nested strings are constructed, shown as `....`. Hence, right folding is *not* tail recursive and it can't be converted to a loop.

In contrast, for the `reduceLeft` example, we can construct the first substring `((0) 1)`, then the next outer string `((0) 1) 2`, etc. In other words, we can convert this process to a loop, because it *is* tail recursive.

Another way to see this is to implement our own simplified `reduceLeft` and `reduceRight` for `Seqs` using recursion:

```
// src/main/scala/progscala3/fp/datastructs/FoldLeftRight.scala
package progscala3.fp.datastructs

import scala.annotation.tailrec

/**
 * Simplified implementations of foldLeft and foldRight.
 */
object FoldLeftRight:

  def foldLeft[A,B](s: Seq[A])(seed: B)(f: (B,A) => B): B =
    @tailrec
    def fl(accum: B, s2: Seq[A]): B = s2 match
      case head +: tail => fl(f(accum, head), tail)
      case _ => accum
    fl(seed, s)

  def foldRight[A,B](s: Seq[A])(seed: B)(f: (A,B) => B): B =
    s match
      case head +: tail => f(head, foldRight(tail)(seed)(f))
      case _ => seed
```

Using them, we should get the same results as before:

```
scala> import progscala3.fp.datastructs.FoldLeftRight._

scala> val strLeft3 = foldLeft(seq6)("()")(left)
      | val strRight3 = foldRight(seq6)("()")(right)
      | val strRight4 = foldRight(seq6)("()")(right2)
val strLeft3: String = (((((((() 1) 2) 3) 4) 5) 6)
val strRight3: String = (((((((() 6) 5) 4) 3) 2) 1)
val strRight4: String = (1 (2 (3 (4 (5 (6 (0)))))))
```

These implementations are simplified in the sense that they don't attempt to construct the correct subclass of the input `Seq` for the output. For example, if you pass in a `Vector`, you'll get a `List` back instead. The Scala collections handle this correctly.

You should learn these two recursion patterns well enough to always remember the behaviors and trade-offs of left versus right recursion, even though in practice you'll almost always use Scala's built-in functions instead of writing your own.

Because we are processing a `Seq`, we should normally work with the elements left to right. It's true that `Seq.apply(index: Int)` returns the element at position `index` (counting from zero). However, for a linked list, this would require an  $O(N)$  traversal for each call to `apply`, yielding an  $O(N^2)$  algorithm rather than  $O(N)$ , which we want. So, the implementation of `foldRight` "suspends" prefixing the value to the rest of the new `Seq` until the recursive invocation of `foldRight` returns. Hence, `foldRight` is not tail recursive.

For `foldLeft`, we use a nested function `rl` to implement the recursion. It carries along an `accum` parameter that accumulates the new `Seq[B]`. When we no longer match on `head` `+: tail`, we've hit the empty tail `Seq`, at which point we return `accum`, which has the completed `Seq[B]` we'll return. When we make a recursive call to `rl`, it is the last thing we do (the tail call), because we prepend the new element to `accum` before passing its updated value to `rl`. Hence, `foldLeft` is tail recursive.

In contrast, when we hit the end of the input `Seq` in `foldRight`, we return an empty `Seq[B]` and *then* the new elements are prefixed to it as we pop the stack.

However, right recursion has one advantage over left recursion. Consider the case where you have a potentially infinite stream of data coming in. You can't conceivably put all that data into a collection in memory, but perhaps you only need to process the first  $N$  elements, for some  $N$ , and then discard the rest. The library's `LazyList` is designed for this purpose. `LazyList` only evaluates the head and tail on demand. We discussed it briefly near the beginning of this chapter.



This on-demand evaluation is the only way to define an infinite stream, and the assumption is that we'll never ask for all of it! That evaluation could be reading from an input channel, like a socket, a **Kafka** *topic*, or a social media “firehose.”. Or it could be a function that generates a sequence of numbers. For example, `LazyList.from(0)` can generate all the natural numbers.

Back to `foldRight`, how is it useful here? Let's develop an intuition for it by reviewing the output we just generated for `strLeft3` and `strRight4`:

```
val strLeft3: String = (((((((0) 1) 2) 3) 4) 5) 6)
val strRight4: String = (1 (2 (3 (4 (5 (6 (0)))))))
```

Suppose we only care about the first four numbers. Visually, we could grab the prefix string `(1 (2 (3 (4` from `strRight4` and then stop. (We might add right parentheses for esthetics...). We're done! In contrast, assume we could generate a string like `strLeft3` with an infinite sequence. To get the first four numbers, we would have to traverse the infinite left parentheses to reach them.

Let's consider an interesting example of using **LazyList** to define a famous infinite sequence, the Fibonacci sequence.

Recall that a Fibonacci number `fib(n)` is defined as follows for natural numbers:

$$f(n) = \begin{cases} 0 & \text{if } n = 0 \\ 1 & \text{if } n = 1 \\ f(n-1) + f(n-2) & \text{otherwise} \end{cases}$$

Like any good recursion, `n` equals 0 or 1 provides the termination condition we need, in which case `f(n) = n`. Otherwise, `f(n) = f(n-1) + f(n-2)`. We saw a tail recursive implementation in “**Nesting Method Definitions and Recursion**”.

Now consider this definition using `LazyList` and described in its documentation:

```
// src/main/scala/progscala3/fp/datastructs/LazyListFibonacci.scala
package progscala3.fp.datastructs

import scala.math.BigInt
import scala.language.implicitConversions // for Int to BigInt

object Fibonacci:
  val fibs: LazyList[BigInt] =
    BigInt(0) #:: BigInt(1) #:: fibs.zip(fibs.tail).map (n => n._1 +
    n._2)
```

Let's try it:

```
scala> import progscala3.fp.datastructs.Fibonacci

scala> Fibonacci.fibs.take(10).toList
val res0: List[BigInt] = List(0, 1, 1, 2, 3, 5, 8, 13, 21, 34)
```

`LazyList` defines its own lazy version the “cons” operation, `#::`.<sup>4</sup> We construct the first two values of the sequence eagerly for the special case of  $n$  equals 0 and 1, then we define the rest of the stream with a *recursive definition*. It is right-recursive, but we'll only take the first  $n$  elements and discard the rest.

Note that we are both defining `fibs` and defining the tail portion using `fibs` itself: `fibs.zip(fibs.tail).map(...)`. This tail expression pairs up all elements of `fibs` with the successor elements, because you always calculate a Fibonacci number by adding its two predecessors together. For example, we have tuple elements like  $(f(2), f(3))$ ,  $(f(3), f(4))$ , etc. going on to infinity, but we don't actually evaluate these expressions until the user asks for values. Note that the tuples are then mapped to an integer, the sum of their values, which is the next  $f(n)$  value!

The last line uses `take` to evaluate the first ten numbers in the sequence, just like for eager sequences.

This is a clever and powerful recursive definition of a sequence. Make sure you understand what it's doing. It helps to play with the pieces of the expression to understand what each one does and then work out the first several values by hand.

It's important to note that the structure of `fibs` is very similar to our implementation of `FoldLeftRight.foldRight, f(0) + f(1) + tail`. Because it is effectively a right recursion, we can stop evaluating `tail` when we have as many head elements as we want. In contrast, trying to construct a left recursion that is also lazy is not possible, because it would look conceptually like this: `f(0 + f(1 + f(tail)))`. (Compare our implementation of `FoldLeftRight.foldLeft`). Hence, a right recursion lets us work with infinite, lazy streams, truncating them where we want, while a left recursion does not.

#### NOTE

Left recursions are tail recursive. Right recursions can handle infinite, lazy streams with truncation.

## Combinators: Software's Best Component Abstractions

When object-oriented programming went mainstream in the late '80s and early '90s, there was great hope that it would usher in an era of reusable software components, even an industry of component libraries. It didn't work out that way, except in special cases, like the windowing APIs of various platforms.

Why did this era of reusable components not materialize? There are certainly many factors, but the fundamental reason is that appropriate, *generative* source code or binary interoperability protocols never materialized that would be the basis of these components. It's a paradox that the richness of object APIs actually

undermined modularization into reusable components, because they didn't define fundamental protocols for interoperability.

In the larger world, component models that succeeded are all based on very simple foundations. Digital integrated circuits (ICs) plug into buses with  $2^n$  signaling wires, each of which is boolean, either on or off. Upon the foundation of this extremely simple protocol, an industry was born with the most explosive growth of any industry in human history.

HTTP is another successful example of a “component model.” Services interact through a narrow, well-defined interface, involving a handful of message types, a naming standard (URLs), and simple standards for message content.

In both cases, higher-level protocols were built upon simple foundations, but these foundations enabled composition to *generate* more complex structures. In digital circuits, some binary patterns are interpreted as CPU instructions, others as memory addresses, and others as data values. REST, data formats like JSON, and other higher-level standards are built upon the foundation elements of HTTP.

When you look at a city, at first it looks like most buildings are unique, completely customized creations. In fact, standards are behind the unique veneers: electric power, water and sewage, standard conventions for room dimensions, furnishings, and windows. Finally, these buildings are surrounded by standard roads holding cars that have their own hierarchy of standards.

Object-oriented programming never established fundamental, *generative* standards. Within each language community, the fundamental unit of composition is the object (with or without a class “template”). Yet objects are not fundamental enough. Every developer invents a new “standard” for a Customer type. No two teams can agree what fields and behaviors the Customer type should have, because each team needs different data and computations for

scenarios related to customers. Something more fundamental than an object is needed.

Attempts to standardize components across language and process boundaries also came up short. Models like CORBA were far from simple. Most defined a binary (versus text) standard. This made interoperability very brittle, subject to *version hell*. It wasn't the choice of binary versus source standardization that was the problem. It was the complexity of the binary standard that led to failure.

In contrast, think about the examples we studied in this chapter. We started with a small set of collections, Seqs (Lists), Vectors, Maps, etc., that all share a set of operations in common, most of which are defined on the Seq abstraction (trait). Most of our examples used sequences, but that was an arbitrary choice.

Except for the utility method `foreach`, the operations were all pure, higher-order functions. They had no side effects and they took other functions as parameters to do the custom work for filtering or transforming each element from the original collection. Such higher-order functions are related to the mathematical concept of *combinators* from *Combinatory Logic*.

We can sequence together these combinator functions to build up nontrivial behavior with relatively little code. We can separate data from the behavior we need to implement for particular problems. This is in opposition to the normal approach of object-oriented programming, which combines data and behavior. It's more typical to create ad hoc implementations of domain logic inside custom classes. We've seen a more productive way in this chapter. This is why the chapter started with the Alan J. Perlis quote.

Let's finish this discussion with a final example, a simplified payroll calculator:

```
// src/test/scala/progscala3/fp/combinators/PayrollSuite.scala  
package progscala3.fp.combinators
```

```

import munit._

class PayrollSuite extends FunSuite:

  case class Employee (name: String, title: String, annualSalary:
Double,
    taxRate: Double, insurancePremiumsPerWeek: Double)

  val employees = List(
    Employee("Buck Trends", "CEO", 200000, 0.25, 100.0),
    Employee("Cindy Banks", "CFO", 170000, 0.22, 120.0),
    Employee("Joe Coder", "Developer", 130000, 0.20, 120.0))

  val weeklyPayroll = employees map { e =>
    val net = (1.0 - e.taxRate) * (e.annualSalary / 52.0) -
      e.insurancePremiumsPerWeek
    (e, net)
  }

  test("weeklyPayroll computes pay for each employee") {
    val results1 = weeklyPayroll map {
      case (e, net) => (e.name, f"${net}%.2f")
    }
    assert(results1 == List(
      ("Buck Trends", "2784.62"),
      ("Cindy Banks", "2430.00"),
      ("Joe Coder", "1880.00")))
  }

  test("from weeklyPayroll, the totals can be calculated") {
    val report = weeklyPayroll.foldLeft( (0.0, 0.0, 0.0) ) {
      case ((totalSalary, totalNet, totalInsurance), (e, net)) =>
        (totalSalary + e.annualSalary/52.0,
          totalNet + net, totalInsurance + e.insurancePremiumsPerWeek)
    }
    assert(f"${report._1}%.2f" == "9615.38", "total salary")
    assert(f"${report._2}%.2f" == "7094.62", "total net")
    assert(f"${report._3}%.2f" == "340.00", "total insurance")
  }

```

We could have implemented this logic in many ways, but let's consider a few of the design choices.

First, although this section criticized object-oriented programming as a component model, OOP is still quite useful. We defined an `Employee` type to hold the fields for each employee. In a real application, we would load this data from a database.

Instead, what if we just use tuples instead of a custom type? You might try rewriting the code this way and compare the two versions. Using `Employee` and the names it provides for the different fields makes it easier to reason about the code. *Meaningful names* is an old principle of good software design. Although I've emphasized the virtues of fundamental collections, functional programming does not say that custom types are bad. As always, design trade-offs should be carefully considered.

However, `Employee` could be called *anemic*. It is a structure with minimal behavior, only the methods generated by the compiler for all case classes. In classic object-oriented design, we might add a lot of behavior to `Employee` to help with the payroll calculation or other domain logic. I believe the design chosen here provides optimal *separation of concerns*. It's also so concise that the maintenance burden is small if the structure of `Employee` changes and this code has to change.

Note also that the logic was implemented in small code snippets, rather than defining lots of classes spread over many file. Of course, it's a toy example, but hopefully you can appreciate that nontrivial applications don't always require large code bases.

There is a counterargument for using a dedicated type, the overhead of constructing instances. Here, this overhead is unimportant. What if we have *billions* of records? We'll return to this question when we explore *Big Data* in *later chapters*.

## What About Making Copies?

Let's finish this chapter by considering a practical problem. Making copies of functional collections is necessary to preserve immutability, but suppose I have a Vector of 100,000 items and I need a copy with the item at index 8 replaced. It would be terribly inefficient to construct a completely new, 100,000-element copy.

Fortunately, we don't have to pay this penalty, nor must we sacrifice immutability. The secret is to realize that 99,999 elements are not changing. If we can share the parts of the original Vector that aren't changing, while representing the change in some way, then creating the new vector can still be very efficient. This idea is called *structure sharing*.

If other code on a different thread is doing something different with the original vector, it is unaffected by the copy operation, because the original vector is not modified. In this way, a "history" of vectors is preserved, as long as there are references to one or more older versions. No version will be garbage-collected until there are no more references to it.


Because this history is maintained, a data structure that uses structure sharing is called a *persistent data structure*.

Our challenge is to select an implementation data structure that lets us expose Vector semantics (or the semantics of another data structure), while providing efficient operations that exploit structure sharing. Let's sketch the underlying data structure and how the copy operation works. We won't cover all the details in depth. For more information start with the [Wikipedia page on persistent data structures](#).

The tree data structure with a branching factor of 32 is used. The branching factor is the maximum number of child nodes each parent node is allowed to have. We said earlier in this chapter that some Vector search and modification operations are  $O(\log(N))$ , but with 32 as the branching factor, that becomes  $O(\log_{32}(N))$ , effectively a constant for even large values of  $N$ !




Figure 7-2 shows an example for `Vector(1,2,3,4,5)`. We'll use just two or three child nodes for legibility.



*Figure 7-2. A Vector represented as a tree*

When you reference the vector, you're actually referencing the root of this tree, marked by #1. As an exercise, you might work through how operations like accessing a particular element by its index, `map`, `flatMap`, etc., would work on a tree implementation of a Vector.

Now suppose we want to insert 2.5 between 2 and 3. To create a new copy, we don't mutate the original tree, but instead create a new tree. Figure 7-3 shows what happens when we add 2.5 between 2 and 3.



*Figure 7-3. Two states of a Vector, before and after element insertion*

Note that the original tree (#1) remains, but we have created a new root (#2) and new nodes between it and the child holding the new element. A new left subtree was created. With a branching factor of 32, we will have to copy up to 32 child references per level, but this number of copy operations is far less than the number required for all references in the original tree. Special handling is required when you need to add an element to a node that already has 32 elements.

Deletion and other operations work similarly. A good textbook on data structures will describe the standard algorithms for tree operations.

Therefore, it is possible to use large, immutable data structures, if their implementations support an efficient copy operation. There is extra overhead compared to a mutable vector, where you can simply modify an entry in place very quickly. Ironically, that doesn't mean

that object-oriented and other *procedural* programs are necessarily simpler and faster.

Because of the dangers of mutability, it's common for OOP classes to wrap mutable collections they hold in accessor methods. This increases the code footprint, testing burden, etc. Worse, if the collection itself is exposed through a “getter” method, the defensive class author might make a `copy_` of the collection to return, so that the internal copy can't be modified. Because collection implementations in nonfunctional languages often have inefficient copy operations, the net effect can be far less efficient programs than corresponding functional programs and greatly increased complexity to work around these limitations.

Immutable collections can not only be efficient, but they eliminate the need for all this extra code to defend against unwanted mutation.

It's analogous to how garbage collected languages are far easier and safe to use than languages where you have to carefully manage memory yourself. You may pay a performance penalty (but definitely not always), but you gain significant benefits.

There are other kinds of functional data structures that are optimized for efficient copying, optimized for modern hardware, such as minimizing cache misses, and other purposes. Many of these data structures were invented as alternatives to mutable data structures that are commonly discussed in classic textbooks on data structures and algorithms.

## Recap and What's Next

We discussed the basic concepts of FP and argued for their importance for solving modern problems in software development. We saw how the fundamental collections and their common higher-order functions, *combinators*, yield concise, powerful, modular code.

Typical functional programs are built on this foundation. At the end of the day, *all* programs input data, perform transformations on it, then output the results. Much of the “ceremony” in typical programs just obscures this essential purpose.

Since functional programming is still relative new for most people, let’s finish this chapter with some references for more information.

For another, gentle introduction to FP aimed at programmers with OOP backgrounds, see *Functional Programming for the Object-Oriented Programmer* by Brian Marick (Leanpub). If you want to convince your Java developer friends that FP is worth their attention, consider my short (but old) *Introduction to Functional Programming for Java Developers* (O’Reilly).

*Functional Programming in Scala*, by Paul Chiusano and Rúnar Bjarnason (Manning Publications), is an excellent, in-depth introduction to functional programming using Scala.

To practice using combinators, see Phil Gold’s [“Ninety-Nine Scala Problems” webpage](#).

For more on functional data structures, see *Purely Functional Data Structures* by Chris Okasaki and *Pearls of Functional Algorithm Design* by Richard Bird (both by Cambridge University Press), and *Algorithms: A Functional Programming Approach*, by Fethi Rabhi and Guy Lapalme (Addison-Wesley).

Next, we’ll return to `for` comprehensions and use our new knowledge of FP to understand how `for` comprehensions are implemented, how we can implement our own data types to exploit them, and how the combination of `for` comprehensions and combinator methods yield concise, powerful code. Along the way, we’ll deepen our understanding of functional concepts.

We’ll return to more advanced features of functional programming in *a later chapter* and we’ll dive into more of the implementation details of Scala’s collections in *a later chapter*.

- 
- 1 I'm ignoring the tricky fact that comparing floating points numbers for equality is fraught with peril.
  - 2 The reason for the `@uncheckedVariance` annotation for `String` in `res1` is discussed in a later chapter.
  - 3 Also called the *least upper bound* or LUB.
  - 4 Two colons are used because `List` defines a cons operator `::`, for historical reasons. We have ignored it, because it is now the convention in Scala to use `++` for prepending elements to any “eager” sequence, including lists.

# Chapter 8. for Comprehensions in Depth

---

## A NOTE FOR EARLY RELEASE READERS

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

This will be the 8th chapter of the final book. Please note that the GitHub repo will be made active later on.

If you have comments about how we might improve the content and/or examples in this book, or if you notice missing material within this chapter, please reach out to the author at [mcronin@oreilly.com](mailto:mcronin@oreilly.com).

“*Scala for Comprehensions*” described many of the details. At this point, they look like a nice, more flexible version of the venerable *for loop*, but not much more. In fact, lots of sophistication lies below the surface, connected to some of the functional *combinators* we discussed in the previous chapter. You can write concise code with elegant solutions to a number of design problems.

In this chapter, we’ll dive below the surface to understand *for* comprehensions and how they are implemented in Scala. You understand how your own types can exploit them.

We’ll finish with some practical design problems implemented using *for* comprehensions, such as error handling during the execution of a sequence of processing steps.

## Recap: The Elements of for Comprehensions

A for comprehension contains one or more generator expressions, optional guard expressions for filtering, and optional value definitions. The output can be “yielded” to create new collections or a side-effecting block of code can be executed on each pass, such as printing output. The following example demonstrates all these features. It removes blank lines from a text file. This is a full program with an example of how to parse input arguments (although there are libraries available for this purpose), handle help messages, etc.

```
// src/main/scala/progscala3/forcomps/RemoveBlanks.scala
package progscala3.forcomps

object RemoveBlanks:
  def apply(path: String, compress: Boolean, numbers: Boolean):
  Seq[String] =
    for
      (line, i) <-
        scala.io.Source.fromFile(path).getLines().toSeq.zipWithIndex
      if line.matches("""^\s*$""") == false
      line2 = if compress then line.trim.replaceAll("\\s+", " ")
              else line
      numLine = if numbers then "%4d: %s".format(i, line2)
                else line2
    yield numLine

  def main(params: Array[String]): Unit =

    val Args(compress, numbers, paths) = parseParams(params.toSeq,
    Args())
    for
      path <- paths
      seq = s"\n== File: $path\n" +: RemoveBlanks(path, compress,
    numbers)
      line <- seq
    do println(line)

    protected val helpMessage = """
      |usage: RemoveBlanks [-h|--help] [-c|--compress] [-n|--numbers]
      |file ...
    """
```

```
|where:
| -h | --help      Print this message and quit
| -c | --compress  Compress whitespace
| -n | --numbers   Print line numbers
| file ...         One or more files to print without blanks
|"".stripMargin
```

```
protected case class Args(
  compress: Boolean = false,
  numbers: Boolean = false,
  paths: Vector[String] = Vector.empty)

protected def help(messages: Seq[String], exitCode: Int) =
  messages.foreach(println)
  println(helpMessage)
  sys.exit(exitCode)

protected def parseParams(params2: Seq[String], args: Args): Args =
  params2 match
  case ("-h" | "--help") +: tail =>
    println(helpMessage)
    sys.exit(0)
  case ("-c" | "--compress") +: tail =>
    parseParams(tail, args.copy(compress = true))
  case ("-n" | "--number") +: tail =>
    parseParams(tail, args.copy(numbers = true))
  case flag +: tail if flag.startsWith("-") =>
    println(s"ERROR: Unknown option $flag")
    println(helpMessage)
    sys.exit(1)
  case path +: tail =>
    parseParams(tail, args.copy(paths = args.paths :+ path))
  case Nil => args
```

- ❶ Use `scala.io.Source` to open the file and get the lines, where `getLines` returns an `Iterator[String]`, which we must convert to a sequence, because we can't return an `Iterator` from the `for` comprehension and the return type is determined by the initial generator. Using `zipWithIndex` adds a line number.

❷



Filter out blank lines using a regular expression. Note that this will result in line number gaps.

- ③ Define a local variable containing the nonblank line, if whitespace compression is not enabled, or a new string with all whitespace compressed to single spaces.
- ④ Format a string with the line number, if enabled.
- ⑤ The `main` method to process the argument list.
- ⑥ A second for comprehension to process the files. Note that we prepend a line with the file name, which will be printed, along with an optional line number created using `zipWithIndex`. Note that the numbers printed won't be the original file's line numbers.
- ⑦ Convenience class to parse the arguments, including flags to show help, whether or not to compress the whitespace in lines, and whether or not to print line numbers.

Try running it at the `sbt` prompt:

```
> runMain progscala3.forcomps.RemoveBlanks --help
> runMain progscala3.forcomps.RemoveBlanks README.md build.sbt -n -c
```

Try different files and different command line options.

## For Comprehensions: Under the Hood

The for comprehension syntax is actually syntactic sugar provided by the compiler for calling the collection methods `foreach`, `map`, `flatMap`, and `withFilter`.

Having a second way to invoke these methods is often easier to understand for nontrivial sequences, compared to using the API

calls. After a while, you develop an intuition about which approach is best for a given context.

The method `withFilter` is used for filtering elements just like the `filter` method, but it doesn't construct its own output collection. For better efficiency, it works with the other methods to combine filtering with their logic so that one less new collection is generated.

Specifically, `withFilter` restricts the domain of the elements allowed to pass through subsequent combinators like `map`, `flatMap`, `foreach`, and other `withFilter` invocations.

To see what the `for` comprehension sugar encapsulates, let's walk through several informal comparisons first, then we'll discuss the details of the precise mapping. As you look at the examples that follow, ask yourself which syntax is easier to understand in each case, the `for` comprehension or the corresponding method calls.

Consider this example of a simple `for` comprehension and the equivalent use of `foreach` on a collection:

```
// src/script/scala/progscala3/forcomps/ForForeach.scala
```

```
scala> val states = Vector("Alabama", "Alaska", "Virginia", "Wyoming")
```

```
scala> var lower1a = Vector.empty[String]
```

```
scala> var lower1b = Vector.empty[String]
```

```
scala> var lower2 = Vector.empty[String]
```

```
scala> for
```

```
|   s <- states
```

```
| do lower1a = lower1a :+ s.toLowerCase
```

```
|
```

```
| for s <- states do lower1b = lower1b :+ s.toLowerCase
```

```
|
```

```
| states.foreach(s => lower2 = lower2 :+ s.toLowerCase)
```

```
var lower1a: Vector[String] = Vector(alabama, alaska, virginia,  
wyoming)
```

```
var lower1b: Vector[String] = Vector(alabama, alaska, virginia,  
wyoming)
```

```
var lower2: Vector[String] = Vector(alabama, alaska, virginia,
wyoming)
```

When there is just one generator (the `s <- states`) in a for comprehension, it can be written on a single line, as shown for `lower1b`. you can still put the `do` clause on the next line, if you prefer.

A single generator expression with a `do` statement corresponds to an invocation of `foreach` on the collection.

What happens if we use `yield` instead?

```
// src/script/scala/progscala3/forcomps/ForMap.scala
```

```
scala> var upper1a = Vector.empty[String]
scala> var upper1b = Vector.empty[String]
scala> var upper2 = Vector.empty[String]
```

```
scala> val upper1a = for
  |   s <- states
  | yield s.toUpperCase
  |
  | val upper1b = for s <- states yield s.toUpperCase
  |
  | val upper2 = states.map(_.toUpperCase)
```

```
val upper1a: Vector[String] = Vector(ALABAMA, ALASKA, VIRGINIA,
WYOMING)
val upper1b: Vector[String] = Vector(ALABAMA, ALASKA, VIRGINIA,
WYOMING)
val upper2: Vector[String] = Vector(ALABAMA, ALASKA, VIRGINIA,
WYOMING)
```

A single generator expression followed by a `yield` expression corresponds to an invocation of `map`. When `yield` is used to construct a new container, its type is determined by the first generator. This is consistent with how `map` works.

What if we have more than one generator?

```
// src/script/scala/progscala3/forcomps/ForFlatmap.scala
```

```

scala> val results1 = for
  |   s <- states
  |   c <- s
  |   yield s"$c-${c.toUpperCase}"
  |
  |   val results2 = states.
  |     flatMap(s => s.toSeq).
  |     map(c => s"$c-${c.toUpperCase}")
val results1: Vector[String] = Vector(A-A, l-L, a-A, b-B, a-A, m-M,
a-A, ...)
val results2: Vector[String] = Vector(A-A, l-L, a-A, b-B, a-A, m-M,
a-A, ...)

```

The second generator iterates through each character in the string `s`. The contrived `yield` statement returns the character and its uppercase equivalent, separated by a dash.

When there are multiple generators, all but the last are converted to `flatMap` invocations. The last is a `map` invocation. Already, you may find the `for` comprehension easier to understand.

What if we add a guard?

*// src/script/scala/progscala3/forcomps/ForGuard.scala*

```

scala> val results1 = for
  |   s <- states
  |   c <- s
  |   if c.isLower
  |   yield s"$c-${c.toUpperCase}"
  |
  |   val results2 = states.
  |     flatMap(s => s.toSeq).
  |     withFilter(c => c.isLower).
  |     map(c => s"$c-${c.toUpperCase}")
  |
val results1: Vector[String] = Vector(l-L, a-A, b-B, a-A, m-M, a-A,
l-L, ...)
val results2: Vector[String] = Vector(l-L, a-A, b-B, a-A, m-M, a-A,
l-L, ...)

```

Note that the `withFilter` invocation is injected before the final `map` invocation.

Finally, defining a variable works as follows:

```
// src/script/scala/progscala3/forcomps/ForVariable.scala

scala> val results1 = for
  |   s <- states
  |   c <- s
  |   if c.isLower
  |   c2 = s"$c-${c.toUpperCase}"
  | yield c2
  |
  | val results2 = states.           // Same as the previous example.
  |   flatMap(s => s.toSeq).
  |   withFilter(c => c.isLower).
  |   map(c => s"$c-${c.toUpperCase}")
val results1: Vector[String] = Vector(l-L, a-A, b-B, a-A, m-M, a-A,
l-L, ...)
val results2: Vector[String] = Vector(l-L, a-A, b-B, a-A, m-M, a-A,
l-L, ...)
```

## Translation Rules of for Comprehensions

Now that we have an intuitive understanding of how for comprehensions are translated to collection methods, let's define the details more precisely.

First, in a generator expression such as `pat <- expr`, `pat` is a pattern expression. For example, `(x, y) <- Seq((1,2),(3,4))`. Similarly, in a value definition `pat2 = expr`, `pat2` is also interpreted as a pattern. For example, `(x, y) = aPair`.

Because they are interpreted as patterns, the compiler translates the expressions using partial functions. For example, first step in the translation is to convert simple comprehension with a generator, `pat <- expr`. The translation is similar to the following for comprehensions (`yield`) and loops (`do`):

```
// src/script/scala/progscala3/forcomps/ForTranslated.scala
```

```
scala> val seq = Seq(1,2,3)

scala> for i <- seq yield 2*i
val res0: Seq[Int] = List(2, 4, 6)

scala> seq.map { case i => 2*i }
val res1: Seq[Int] = List(2, 4, 6)

scala> var sum1 = 0
scala> for i <- seq do sum1 += 1
var sum1: Int = 3

scala> var sum2 = 0
scala> seq.foreach { case i => sum2 += 1 }
var sum2: Int = 3
```

A conditional is translated to withFilter conceptually as shown next:

```
scala> for
  |   i <- seq
  |   if i%2 != 0
  |   yield 2*i
val res2: Seq[Int] = List(2, 6)

scala> for
  |   i <- seq if i%2 != 0
  |   yield 2*i
val res3: Seq[Int] = List(2, 6)

scala> seq.withFilter {
  |   case i if i%2 != 0 => true
  |   case _ => false
  | }.map { case i => 2*i }
val res4: Seq[Int] = List(2, 6)
```

- ❶ You can write the guard on the same line as the previous generator.

After this, the translations are applied repeatedly until all comprehension expressions have been replaced. Note that some steps generate *new* for comprehensions that subsequent iterations will translate.

First, a for comprehension with two generator and a yield expression:

```
scala> for
  |   i <- seq
  |   j <- (i to 3)
  | yield j
val res5: Seq[Int] = List(1, 2, 3, 2, 3, 3)

scala> seq.flatMap { case i => for j <- (i to 3) yield j } ❶
val res6: Seq[Int] = List(1, 2, 3, 2, 3, 3)

scala> seq.flatMap { case i => (i to 3).map { case j => j } } ❷
val res7: Seq[Int] = List(1, 2, 3, 2, 3, 3)
```

- ❶ One level of translation. Note the nested for ... yield.
- ❷ Completed translation.

A for *loop*, with do, again translating in two steps:

```
scala> var sum3=0
scala> for
  |   i <- seq
  |   j <- (i to 3)
  | do sum3 += j
val sum3: Int = 14

scala> var sum4=0
scala> seq.foreach { case i => for j <- (i to 3) do sum4 += j }
val sum4: Int = 14

scala> var sum5=0
scala> seq.foreach { case i => (i to 3).foreach { case j => sum5 += j } }
val sum5: Int = 14
```

A generator followed by a value definition has a surprisingly complex translation. Here I show complete for ... yield ... expressions:

```
scala> for
|   i <- seq
|   i10 = i*10
|   yield i10
val res8: Seq[Int] = List(10, 20, 30)

scala> for
|   (i, i10) <- for
|     x1 @ i <- seq                                ❶
|     yield
|       val x2 @ i10 = x1*10                        ❷
|       (x1, x2)                                    ❸
|     yield i10                                     ❹
val seq9: Seq[Int] = List(10, 20, 30)
```

- ❶ Recall from **Chapter 4** that `x1 @ i` means assign to variable `x1` the value corresponding to the whole expression on the right-hand side of `@`, which is trivially `i` in this case, but it could be an arbitrary pattern with nested variable bindings to the constituent parts.
- ❷ Assign to `x2` the value of `i10`.
- ❸ Return the tuple.
- ❹ Yield `i10`, which will be equivalent to `x2`.

Here is another example of `x @ pat = expr`:

```
scala> val z @ (x, y) = (1 -> 2)
val z: (Int, Int) = (1,2)
val x: Int = 1
val y: Int = 2
```



This completes the translation rules. Whenever you encounter a for comprehension, you can apply these rules to translate it into method invocations on containers. You won't need to do this often, but sometimes it's a useful skill for debugging problems.

## Options and Other Container Types

We used collections like Lists, Arrays, and Maps for our examples, but any types that implement `foreach`, `map`, `flatMap`, and `withFilter` (or `filter`) can be used in for comprehensions and not just the obvious collection types. In the general case, these are *containers* and eligible for use in for comprehensions.

Let's consider several other container types. We'll see how exploiting for comprehensions can transform your code in unexpected ways.

### Option as a Container

**Option** is a *binary* container. It has an item or it doesn't. It implements the four methods we need.

Here is a simplified version of the `Option` abstract class in the Scala library (full source [here](#)):

```
sealed abstract class Option[+A] { self => ❶
  ...
  def isEmpty: Boolean = this eq None ❷

  final def foreach[U](f: A => U): Unit =
    if (!isEmpty) f(this.get)

  final def map[B](f: A => B): Option[B] =
    if (isEmpty) None else Some(f(this.get))

  final def flatMap[B](f: A => Option[B]): Option[B] =
    if (isEmpty) None else f(this.get)

  final def filter(p: A => Boolean): Option[A] =
    if (isEmpty || p(this.get)) this else None
```

```

final def withFilter(p: A => Boolean): WithFilter = new
WithFilter(p)

class WithFilter(p: A => Boolean) {
  def map[B](f: A => B): Option[B] = self filter p map f
  def flatMap[B](f: A => Option[B]): Option[B] = self filter p
flatMap f
  def foreach[U](f: A => U): Unit = self filter p foreach f
  def withFilter(q: A => Boolean): WithFilter =
    new WithFilter(x => p(x) && q(x))
}

```

- ❶ The `self =>` expression defines an alias for this for the `Option` instance. It is needed inside `WithFilter` below. See *Chapter 2* for more details.
- ❷ Test if this is actually the `None` instance, not value equality.
- ❸ The `WithFilter`, which is used by `withFilter` combined with the other operations to avoid creation of an intermediate collection when filtering.
- ❹ Here's where the `self` reference defined above is used to operate on the enclosing `Option` instance. Using this would refer to the instance of `WithFilter` itself.

The `final` keyword prevents subclasses from overriding the implementation. It might be surprising to see the base class refer to derived classes. Normally, in object-oriented design this would be considered bad. However, with sealed type hierarchies, this file knows all the possible subclasses. Referring to derived classes makes the implementation more concise and efficient, overall, as well as safe.

The crucial feature about these `Option` methods shown is that the function arguments are only applied if the `Option` isn't empty. This

feature allows us to address a common design problem in an elegant way.

Say for example that you want to distribute some tasks around a cluster, then gather the results together. Suppose you want an elegant way to ignore those tasks that return empty results.

Wrap each task return value in an `Option`, where `None` is used for empty results and `Some` wraps a nonempty result. We want an easy way to filter out the `None` results. Here is an example, where we have the returned `Options` in a `Vector`:

```
// src/script/scala/progscala3/forcomps/ForOptionsFilter.scala

scala> val options: Seq[Option[Int]] = Vector(Some(10), None,
Some(20))
val options: Seq[Option[Int]] = Vector(Some(10), None, Some(20))

scala> val results = for
  |   case Some(i) <- options
  |   yield (2 * i)
val results: Seq[Int] = Vector(20, 40)
```

`case Some(i) <- options` pattern matches on each element in `results` and extracts the integers inside the `Some` values. Since a `None` won't match, all of them are removed. We then `yield` the final expression we want.

As an exercise, let's work through the translation rules. Try it yourself before reading on! Here is the first step, where we apply the first rule for converting each `pat <- expr` expression to a `withFilter` expression:

```
scala> val results2 = for
  |   case Some(i) <- options withFilter {
  |     case Some(i) => true
  |     case None => false
  |   }
  |   yield (2 * i)
val results2: Seq[Int] = Vector(20, 40)
```

Finally, we convert the outer `for x <- y yield (z)` expression to a `map` call:

```
scala> val results3 = options withFilter {  
  |   case Some(i) => true  
  |   case None => false  
  | } map {  
  |   case Some(i) => (2 * i)  
  |   case None => -1           // hack  
  | }  
val results3: Seq[Int] = Vector(20, 40)
```

The “hack” is there because we don’t actually need the `case None` clause, because the `withFilter` has already removed all `Nones`. However, the compiler doesn’t understand this, so it warns us we’ll risk a `MathError` without the clause. Try removing this clause and observe the warning you get.

Consider another design problem. Instead of independent tasks where we ignore the empty results and combine the nonempty results, consider the case where we run a sequence of dependent steps, and we want to stop the whole process as soon as we encounter a `None`.

Note that we have a limitation that using `None` means we receive no feedback about why the step returned nothing, such as a failure. We’ll address this limitation later.

We could write tedious conditional logic that tries each case, one at a time, and checks the results, but a `for` comprehension is more concise:

```
// src/script/scala/progscala3/forcomps/ForOptionsSeq.scala  
  
scala> def positiveOption(i: Int): Option[Int] =  
  |   if i > 0 then Some(i) else None  
  
scala> val resultSuccess = for  
  |   i1 <- positiveOption(5)  
  |   i2 <- positiveOption(10 * i1)
```

```

    |   i3 <- positiveOption(25 * i2)
    |   i4 <- positiveOption(2 * i3)
    |   yield (i1 + i2 + i3 + i4)
val resultSuccess: Option[Int] = Some(3805)

scala> val resultFail = for
    |   i1 <- positiveOption(5)
    |   i2 <- positiveOption(-1 * i1)
    |   i3 <- positiveOption(25 * i2)
    |   i4 <- positiveOption(-2 * i3)
    |   yield (i1 + i2 + i3 + i4)
val resultFail: Option[Int] = None

```

- ❶ None is returned. The subsequent generators don't call `positiveOption`, they just pass the `None` through.

At each step, the integer in the `Some` returned by `positiveOption` is extracted and assigned to a variable. Subsequent generators use those values. It appears we assume the “happy path” always works, which is true for the first `for` comprehension. It also works fine for the second `for` comprehension, because once a `None` is returned, the subsequent generators simply propagate the `None` and don't call `positiveOption`.

Let's look at three more container types with similar properties, `Either` and `Try` from the Scala library, and `Validated` from the `Typelevel Cats` library. `Validated` is a more sophisticated tool for sequencing validation steps.

## Either: A Logical Extension to Option

We noted that the use of `Option` has the disadvantage that `None` carries no information that could tell us why no value is available. Did an error occur? What kind? Using `Either` instead is one solution. As the name suggests, `Either` is a container that holds one and only one of two things. In other words, where `Option` handled the case of zero or one items, `Either` handles the case of one item or another.

Either is a parameterized type with two parameters, `Either[+A, +B]`, where the `A` and `B` are the two possible types of the element contained in the `Either`. Recall that `+A` indicates that `Either` is *covariant* in the type parameter `A` and similarly for `+B`. This means that if you need a value of type `Either[Any, Any]` (for example, a method parameter), you can use an instance of type `Either[String, Int]`, because `String` and `Int` are subtypes of `Any`, therefore `Either[String, Int]` is a subtype of `Either[Any, Any]`.

`Either` is also a sealed abstract class with two subclasses defined, `Left` and `Right`. That's how we distinguish between the two possible elements.

The concept of `Either` predates Scala. It has been used for a long time as an alternative to throwing exceptions. By historical convention, the `Left` value is used to hold the error indicator, such as a message string or thrown exception, and the normal return value is returned in a `Right`.

Let's port our `Option` example. It's almost identical:

```
// src/script/scala/progscala3/forcomps/ForEithersGood.scala

scala> def positiveEither(i: Int): Either[String, Int] =
  |   if i > 0 then Right(i) else Left(s"nonpositive number $i")

scala> val result1 = for
  |   i1 <- positiveEither(5)
  |   i2 <- positiveEither(10 * i1)
  |   i3 <- positiveEither(25 * i2)
  |   i4 <- positiveEither(2 * i3)
  | yield (i1 + i2 + i3 + i4)
val result1: Either[String, Int] = Right(3805)

scala> val result2 = for
  |   i1 <- positiveEither(5)
  |   i2 <- positiveEither(-1 * i1)
  |   i3 <- positiveEither(25 * i2)
  |   i4 <- positiveEither(-2 * i3)
  | yield (i1 + i2 + i3 + i4)
val result2: Either[String, Int] = Left(nonpositive number -5)
```

- ❶ A Left is returned here, stopping the process.

Note how Left and Right objects are constructed in positiveEither. Note the types for result1 and result2. In particular, result2 now tells us where the first negative number was encountered, but not the second occurrence of one.

Either isn't limited to this error-handling idiom. It could be used for any scenario where you want to hold one object or another, possibly of different types. However, union types are, such as String | Int, are better for this purpose. Superficially, they appear to serve a similar function, but union types can't be used in this context, because they don't have the combinators like map, flatMap, etc.

That raises some questions, though. Why do Lefts stop the for comprehension and Rights don't? It's because Either isn't really symmetric in the types. Since it is always used for this error-handling idiom, the implementations of Left and Right *bias* towards the right as the "happy path".

Let's look how the combinators and some other methods work for these two types, using result1 and result2:

```
scala> result1    // Reminder of these values:
      | result2
val res6: Either[String, Int] = Right(3805)
val res7: Either[String, Int] = Left(nonpositive number -5)

scala> var r1 = 0
      | result1.foreach(i => r1 = i * 2)
      | var r2 = 0
      | result2.foreach(i => r1 = i * 2)
val r1: Int = 7610
val r2: Int = 0

scala> val r3 = result1.map(_ * 2)
      | val r4 = result2.map(_ * 2)
      |
val r3: Either[String, Int] = Right(7610)
```

❶

```

val r4: Either[String, Int] = Left(nonpositive number -5)

scala> val r5a = result1.flatMap(i => Right(i * 2))
      | val r5b = result1.flatMap(i => Left("hello"))
      | val r5c = result1.flatMap(i => Left[String,Double]("hello"))
      | val r5d: Either[String,Double] = result1.flatMap(i =>
Left("hello"))
      | val r6 = result2.flatMap(i => Right(i * 2))
      |
val r5a: Either[String, Int] = Right(7610)
val r5b: Either[String, Nothing] = Left(hello)
val r5c: Either[String, Double] = Left(hello)
val r5d: Either[String, Double] = Left(hello)
val r6: Either[String, Int] = Left(nonpositive number -5)

```

- ❶ No change is made to r2 after initialization.
- ❷ Note the second type for r5b vs. r5c and r5d. Using Left("hello") alone provides no information about the desired second type, so Nothing is used.

The filter and withFilter methods aren't supported. They are somewhat redundant in this case.

You can infer that the Left method implementations ignore the function and just return their value. Right.map extracts the value, applies the function, then constructs a new Right, while Right.flatMap simply returns the value the function returns.

Finally, here is for comprehension that uses Eithers:

```

// src/script/scala/progscala3/forcomps/ForEithersSeq.scala

scala> val seq: Seq[Either[RuntimeException,Int]] =
      | Vector(Right(10), Left(new RuntimeException("boo!")),
Right(20))
      |
      | val results3 = for
      |   case Right(i) <- seq
      |   yield 2 * i
val results3: Seq[Int] = Vector(20, 40)

```



## Throwing exceptions versus returning Either values

Just as `Either` encourages handling of errors as normal return values, avoiding thrown exceptions is also valuable for uniform handling of errors, including maintaining referential transparency, which thrown exceptions violate. To see this, consider the following contrived example:

```
// src/script/scala/progscala3/forcomps/RefTransparency.scala

scala> def addInts(s1: String, s2: String): Int = s1.toInt + s2.toInt

scala> def addInts2(s1: String, s2: String): Either[String, Int] =
  |   try
  |     Right(s1.toInt + s2.toInt)
  |   catch
  |     case nfe: NumberFormatException => Left("NFE:
"+nfe.getMessage)

scala> val add12a = addInts("1", "2")
  |   val add12b = addInts2("1", "2")
val add12a: Int = 3
val add12b: Either[String, Int] = Right(3)

scala> val add1x  = addInts2("1", "x")
  |   val addx2   = addInts2("x", "2")
  |   val addxy   = addInts2("x", "y")
val add1x: Either[String, Int] = Left(NFE: For input string: "x")
val addx2: Either[String, Int] = Left(NFE: For input string: "x")
val addxy: Either[String, Int] = Left(NFE: For input string: "x")
```

We would like to believe that `addInts` is referentially transparent, so we could replace calls to it with values from a cache of previous invocations, for example. However, `addInts` will throw an exception if we pass a `String` that can't be parsed as an `Int`. Hence, we can't replace the function call with values that can be returned for *all* parameter lists.

Also, the type signature of `addInts` provides no indication that trouble lurks.

Using `Either` as the return type of `addInts2` restores referential transparency and the type signature is explicit about potential errors. It is referentially transparent, because we could replace *all* calls with a value, even for bad string input.

Also, instead of grabbing control of the call stack by throwing the exception, we've *reified* the error by returning the exception as a `Left` value.

So, `Either` lets us maintain control of call stack in the event of a wide class of failures. It also makes the behavior more explicit to users of your APIs, through type signatures.

However, look at the implementation of `addInts2` again. Handling exceptions is quite common, so the `try ... catch ...` boilerplate shown appears a lot in code.

So, for handling exceptions, we should encapsulate this boilerplate with types and use names for these types that express more clearly when we have either a “failure” or a “success.” The `Try` type does just that.

## Try: When There Is No Do

`scala.util.Try` is structurally similar to `Either`. It is a sealed abstract class with two subclasses, `Success` and `Failure`.

`Success` is analogous to the conventional use of `Right`. It holds the normal return value. `Failure` is analogous to `Left`, but `Failure` always holds a `Throwable`, which is why `Try` has one type parameter, instead of two, for the value held by `Success`.

Here are the signatures of these types (omitting some traits that aren't relevant to the discussion):

```
sealed abstract class Try[+T] extends AnyRef {...}
final case class Success[+T](value: T) extends Try[T] {...}
final case class Failure[+T](exception: Throwable) extends Try[T]
{...}
```

Try is clearly asymmetric, unlike Either, where the asymmetry isn't clear from the type signature, it just reflects convention and the convention determined how the combinators were implemented asymmetrically.

Let's see how Try is used, again porting our previous example. First, if you have a list of Try values and just want to discard the Failures, a simple for comprehension does the trick:

```
// src/script/scala/progscala3/forcomps/ForTries.scala

scala> import scala.util.{ Try, Success, Failure }

scala> def positiveTries(i: Int): Try[Int] = Try {
  |   assert (i > 0, s"nonpositive number $i")
  |   i
  | }

scala> val result4 = for
  |   i1 <- positiveTries(5)
  |   i2 <- positiveTries(10 * i1)
  |   i3 <- positiveTries(25 * i2)
  |   i4 <- positiveTries(2 * i3)
  | yield (i1 + i2 + i3 + i4)
val result4: scala.util.Try[Int] = Success(3805)

scala> val result5 = for
  |   i1 <- positiveTries(5)
  |   i2 <- positiveTries(-1 * i1)           // FAIL!
  |   i3 <- positiveTries(25 * i2)
  |   i4 <- positiveTries(-2 * i3)
  | yield (i1 + i2 + i3 + i4)
val result5: scala.util.Try[Int] =
  Failure(java.lang.AssertionError: assertion failed: nonpositive
number -5)
```

Note the concise definition of positiveTries. If the assertion fails, the Try block will return a Failure wrapping the thrown java.lang.AssertionError. Otherwise, the result of the Try

expression is wrapped in a `Success`. A more explicit definition of `positiveTries` showing the boilerplate is the following:

```
def positiveTries2(i: Int): Try[Int] =  
  if (i > 0) Success(i)  
  else Failure(new AssertionError("assertion failed"))
```

The `for` comprehensions look exactly like those for the original `Option` example. With type inference, there is very little boilerplate here, too. You can focus on the “happy path” logic and let `Try` capture errors.

## Cats Validator

While using `Option`, `Either`, or `Try` meets most needs, there is one common scenario where using any of them remains tedious. Consider the case of form validation, where a user submits a form with several fields, all of which need to be validated. Ideally, you would validate all at once and report all errors, rather than doing one at a time, which is not a friendly user experience. Using `Option`, `Either`, or `Try` in a `for` comprehension doesn’t support this need, because processing is short-circuited as soon as a failure occurs. This is where `cats.datatypes.Validated` provides several useful approaches.

We’ll consider one approach here. First, start with some domain specific classes:

```
// src/main/scala/progscala3/forcomps/LoginFormValidation.scala  
  
package progscala3.forcomps  
  
case class ValidLoginForm(userName: String, password: String)  
❶  
  
sealed trait LoginValidation:  
❷  
  def error: String
```

```

case class Empty(name: String) extends LoginValidation:
  val error: String = s"The $name field can't be empty"

case class TooShort(name: String, n: Int) extends LoginValidation:
  val error: String = s"The $name field must have at least $n
characters"

case class BadCharacters(name: String) extends LoginValidation:
  val error: String = s"The $name field has invalid characters"

```

- ❶ A case class with the form fields to validate.
- ❷ A trait used by other case classes that encapsulate each error.

Now we use them in the following code, where the acronym *Nec* stands for “non empty chain”. In this context, that means that a failed validation will have a sequence (“chain”) of one or more error objects.

```

// src/main/scala/progscala3/forcomps/LoginFormValidatorNec.scala
package progscala3.forcomps

import cats.implicits._
import cats.data._
import cats.data.Validated._
import scala.language.implicitConversions

/**
 * Nec variant, where NEC stands for "non empty chain".
 * @see https://typelevel.org/cats/datatypes/validated.html
 */
object LoginFormValidatorNec:

  type V[T] = ValidatedNec[LoginValidation, T]

  ❶
  def nonEmpty(field: String, name: String): V[String] =
    ❷
    if field.length > 0 then field.validNec
    else Empty(name).invalidNec

```

```

def notTooShort(field: String, name: String, n: Int): V[String] =
  if field.length >= n then field.validNec
  else TooShort(name, n).invalidNec

/** For simplicity, just disallow whitespace. */
def goodCharacters(field: String, name: String): V[String] =
  val re = raw".*\s.*".r
  if re.matches(field) == false then field.validNec
  else BadCharacters(name).invalidNec

def apply(
  userName: String, password: String): V[ValidLoginForm] =
  (notEmpty(userName, "user name"),
   notTooShort(userName, "user name", 5),
   goodCharacters(userName, "user name"),
   notEmpty(password, "password"),
   notTooShort(password, "password", 5),
   goodCharacters(password, "password")).mapN {
    case (s1, _, _, s2, _, _) => ValidLoginForm(s1, s2)
  }

/**
 * This method uses the matching clauses shown rather something like
 this:
 *   assert(LoginFormValidatorNec("123 45", "678 90") ==
 *     Invalid(Chain(BadCharacters("user name"),
 BadCharacters("password"))))
 * This is necessary because we use -language:strictEquality, which
 causes
 * these == expressions to fail compilation!
 */
@main def TryLoginFormValidatorNec =
  import LoginFormValidatorNec._
  assert(LoginFormValidatorNec("", "") ==
    Invalid(Chain(
      Empty("user name"), TooShort("user name", 5),
      Empty("password"), TooShort("password", 5))))

  assert(LoginFormValidatorNec("1234", "6789") ==
    Invalid(Chain(
      TooShort("user name", 5),
      TooShort("password", 5))))

  assert(LoginFormValidatorNec("12345", "") ==

```

```

Invalid(Chain(
  Empty("password"), TooShort("password", 5)))

assert(LoginFormValidatorNec("123 45", "678 90") ==
  Invalid(Chain(
    BadCharacters("user name"), BadCharacters("password"))))

assert(LoginFormValidatorNec("12345", "67890") ==
  Valid(ValidLoginForm("12345", "67890")))

```

- ❶ Shorthand type alias. ValidationNec will encapsulate errors or successful results.
- ❷ Several functions to test that fields meet desired criteria. When successful, an appropriate ValidationNec is constructed by calling either of the extension methods on String, validNec or invalidNec.
- ❸ The apply method uses a Cats function mapN for mapping over the N elements of a tuple. It returns a final ValidationNec instance with all the accumulated errors in an Invalid(Chain(...)) or if all validation criteria were met, a Valid(ValidLoginForm(...)) holding the passed-in field values.

For comparison, see also in the example code, `src/main/scala/progscala3/forcomps/LoginFormValidatorSingle.scala`, which handles single failures using `Either`, but following a similar implementation approach.

Without a tool like `Validation`, we would have to manage the chain of errors ourselves.

## Recap and What's Next

`Either`, `Try`, and `Validator` express through types a fuller picture of how the program actually behaves. All three say that a valid value or

values will (hopefully) be returned, but if not, they also encapsulate the failure information needed. Similarly, `Option` encapsulates the presence or absence of a value explicitly in the type signature.

Using these types instead of thrown exceptions keeps control of the call stack, signals to the reader the kinds of errors that might occur, and allows error conditions to be less “exceptional” and more amenable to programmatic handling, just like the happy path scenarios.

Another benefit we haven’t mentioned yet is a benefit for asynchronous (concurrent) code. Because asynchronous code isn’t guaranteed to be running on the same thread as the caller, it might not be possible to catch and handle an exception. However, by returning errors the same way normal results are returned, the caller can more easily intercept and handle the problem. We’ll explore the details in *Chapter 18*.

You probably expected this chapter to be a perfunctory explanation of Scala’s fancy *for loops*. Instead, we broke through the facade to find a surprisingly powerful set of tools. We saw how a set of functions, `map`, `flatMap`, `foreach`, and `withFilter`, plug into `for` comprehensions to provide concise, flexible, yet powerful tools for building nontrivial application logic.

We saw how to use `for` comprehensions to work with collections, but we also saw how useful they are for other container types, specifically `Option`, `Either`, `Try`, and `Cats Validated`.

Now we have finished our exploration of the essential parts of functional programming and their support in Scala. We’ll learn more concepts when we discuss the type system in *Chapters 15 and 16* and explore advanced concepts in *Chapter 17*.

Let’s now turn to Scala’s support for object-oriented programming. We’ve already covered many of the details in passing. Now we’ll complete the picture.



# Chapter 9. Object-Oriented Programming in Scala

---

## A NOTE FOR EARLY RELEASE READERS

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

This will be the 9th chapter of the final book. Please note that the GitHub repo will be made active later on.

If you have comments about how we might improve the content and/or examples in this book, or if you notice missing material within this chapter, please reach out to the author at [mcronin@oreilly.com](mailto:mcronin@oreilly.com).

One reason Scala is a superb object-oriented programming (OOP) language is because Martin Odersky and his collaborators have thought long and hard about how to make OOP best practices as concise as possible. While we already know many of Scala’s features for OOP, now we will explore them more systematically. We’ll see more examples of Scala’s concise syntax and how it enables effective OOP in combination with Functional Programming (FP).

Scala is both a OOP language and an FP language. I’ve waited until now to explore Scala as an OOP language for two reasons.

First, I wanted to emphasize that FP has become an essential skill set for modern problems, a skill set that may be new to you. When

you start with Scala, it's easy to use it as a better OOP language, a "better Java," and neglect the power of its FP side.

Second, a common architectural approach with Scala has been to use FP for *programming in the small* and OOP for *programming in the large*. Using FP for implementing algorithms, manipulating data, and managing state in a principled way is our best way to minimize bugs, the amount of code we write, and the risk of schedule delays. On the other hand, Scala's OO model provides tools for designing composable, reusable, and encapsulated *modules*, which are essential for building larger applications. Hence, Scala gives us the best of both worlds.

I've assumed you already know the basics of OOP from other languages, so many concepts were defined quickly and informally. This chapter starts with a quick review of class and object basics, then fills in the details, such as the mechanics of creating type hierarchies, how constructors work for Scala classes, and runtime-efficient types using *opaque type aliases* (new to Scala 3) and *value classes*. The next chapter will dive into *traits* and then we'll spend a few chapters filling in additional details on Scala's object model and the standard library.

## **Class and Object Basics: Review**

Classes are declared with the keyword `class`, while *singleton* objects are declared with the `object` keyword. For this reason, I have used the term *instance* in this book to refer to objects generically, whether they are class instances or declared objects. In most OOP languages, *instance* and *object* are synonymous.

An instance can refer to itself using the `this` keyword, although it's somewhat rare in Scala code. One reason is that constructor boilerplate is absent in Scala. Consider the following Java code:

```
// src/main/java/progscala3/basicoop/JavaPerson.java
package progscala3.basicoop;

public class JavaPerson {
    private String name;
    private int    age;

    public JavaPerson(String name, int age) {
        this.name = name;
        this.age  = age;
    }

    public void setName(String name) { this.name = name; }
    public String getName()          { return this.name; }

    public void setAge(int age) { this.age = age; }
    public int  getAge()        { return this.age; }
}
```

Other OOP languages are similar. Now compare it with the following equivalent Scala declaration, in which all the boilerplate disappears:

```
class Person(var name: String, var age: Int)
```

Prefixing a constructor parameter with a `var` makes it a mutable *field* of the class, also called an *instance variable* or *attribute* in other languages. Prefixing a constructor parameter with a `val` makes it an immutable field. Using the `case` keyword infers the `val` keyword and also adds additional methods, as we've learned:

```
case class Person(name: String, age: Int)
```

This is just one example of how concise OOP can be in Scala.

You can also declare other `val` and `var` fields inside the type body.

The term *member* refers to a field, method, or type alias in a generic way.

The term *method* refers to a function that is tied to an instance. Its parameter list has an “implied” `this`. Method definitions start with the

def keyword.

Scala allows *overloaded methods*. Two or more methods can have the same name as long as their full *signatures* are unique. The signature includes the enclosing type name, method name, and the types of all the parameters. The parameter names are not part of the signature for typing purposes, but they are significant because you can provide them when calling the method, e.g., `log(message = "Error!")`. Also, different return types alone are not sufficient to distinguish methods.

### TIP

In Scala 2, only the parameters in the *first* parameter list were considered when determining the method signature for the purposes of overloading. In Scala 3, all parameter lists are considered.

However, JVM *type erasure* prevents concrete parameterized types from being considered unique. Both of the following methods will have the same type after erasure:

```
// src/script/scala/progscala3/basicoop/GoodBad.scala
object OBad:
  def m(seq: Seq[Int]): String = seq.mkString("|")
  def m(seq: Seq[String]): String = seq.mkString(",")
```

We discussed workarounds when we first discussed this issue in [“Working Around Type Erasure with Context Bounds”](#).

Member *type aliases* are declared using the `type` keyword. They are used to provide shorter names for complex types and to provide a complementary mechanism to type parameterization (see [“Abstract Types Versus Parameterized Types”](#)).

A field and method can have the same name, but only if the method has a parameter list:

```
scala> trait Good:
  |   def x(suffix: String): String
  |   val x: String

scala>
  |   trait Bad:
  |     def x: String
  |     val x: String
4 |   val x: String
  |       ^
  |       Double definition...
```

## Open Versus Closed Types

Scala encourages us to think carefully about what types should be abstract vs. concrete, what types should be *singletons*, what types should be *mixins*, and what types should be *open* vs. *closed* for extension, meaning allowed to be *subtyped* or not. OOP languages also use the terms *subclassing* or *deriving* from a supertype. I've used subtyping to emphasize the more general type system in Scala's combination of FP and OOP.

Mixins promote *composition over inheritance*, discussed in “[Good Object-Oriented Design: A Digression](#)”. Traits are used to define mixins, while abstract classes or traits are used as base classes in a hierarchy. Here is a “sketch” of a hierarchy of services with logging mixed in. First, define a basic logging trait:

```
// tag::logging[]
// src/main/scala/progscala3/basicoop/Abstract.scala
package progscala3.basicoop

enum LoggingLevel:
  ❶ case INFO, WARN, ERROR

trait Logging:
  import LoggingLevel._
  final def info(message: String): Unit = log(INFO, message)
  final def warn(message: String): Unit = log(WARN, message)
```

```

final def error(message: String): Unit = log(ERROR, message)
final def log(level: LoggingLevel, message: String): Unit =
  write(s"$level: $message")

protected val write: String => Unit
❷

trait ConsoleLogging extends Logging:
  protected val write = println
❸
// end::logging[]

// tag::service[]
open abstract class Service(val name: String) extends Logging:
❹
  import Service._
  final def handle(request: Request): Response =
    info(s"($name) Starting handle for request: $request")
    val result = process(request)
    info(s"($name) Finishing handle with result $result")
    result

  protected def process(request: Request): Response

object Service:
❺
  type Request = Map[String,Any]
  type Response = Either[String,Map[String,Any]]

case class HelloService(override val name: String)
❻
  extends Service(name) with ConsoleLogging:
    import Service._

    protected def process(request: Request): Response =
      request.get("user") match
        case Some(user) => Right(Map("message" -> s"Hello, $user"))
        case None => Left("No user field found!")
// end::service[]

// tag::main[]
@main def HelloServiceMain(name: String, users: String*): Unit =
  val hs = HelloService("hello")
  for
    user <- users

```

```

    request = Map("user" -> user)
  do hs.handle(request) match
    case Left(error) => println(s"ERROR! $error")
    case Right(map) => println(map("message"))

  println("Try an empty map:")
  println(hs.handle(Map.empty))
// end::main[]

```

- ❶ Define a simple logging abstraction with three levels.
- ❷ Do as much as possible in the Logging mixin trait. The protected abstract function value `write` is implemented by subtypes for actually writing to the “log”. Every other method is declared `final` to prevent overriding.
- ❸ ConsoleLogger just uses `println`.
- ❹ TK4a
- ❺ TK5a
- ❻ TK6a

Now, define an abstract base class for services that mix in logging:

```

// tag::logging[]
// src/main/scala/progscala3/basicoop/Abstract.scala
package progscala3.basicoop

enum LoggingLevel:
  ❶
    case INFO, WARN, ERROR

trait Logging:
  import LoggingLevel._
  final def info(message: String): Unit = log(INFO, message)
  final def warn(message: String): Unit = log(WARN, message)
  final def error(message: String): Unit = log(ERROR, message)
  final def log(level: LoggingLevel, message: String): Unit =

```

```

    write(s"$level: $message")

    protected val write: String => Unit
2
trait ConsoleLogging extends Logging:
    protected val write = println
3
// end::logging[]

// tag::service[]
open abstract class Service(val name: String) extends Logging:
4
    import Service._
    final def handle(request: Request): Response =
        info(s"($name) Starting handle for request: $request")
        val result = process(request)
        info(s"($name) Finishing handle with result $result")
        result

    protected def process(request: Request): Response

object Service:
5
    type Request = Map[String,Any]
    type Response = Either[String,Map[String,Any]]

case class HelloService(override val name: String)
6
    extends Service(name) with ConsoleLogging:
        import Service._

    protected def process(request: Request): Response =
        request.get("user") match
            case Some(user) => Right(Map("message" -> s"Hello, $user"))
            case None => Left("No user field found!")
// end::service[]

// tag::main[]
@main def HelloServiceMain(name: String, users: String*): Unit =
    val hs = HelloService("hello")
    for
        user <- users
        request = Map("user" -> user)
    do hs.handle(request) match

```



```

    case Left(error) => println(s"ERROR! $error")
    case Right(map) => println(map("message"))

    println("Try an empty map:")
    println(hs.handle(Map.empty))
  // end::main[]

```

- ❶ A service abstraction. It uses the Logging trait as the supertype, because there isn't another supertype (other than AnyRef), but really this is mixin-composition. Concrete subtypes of Service must implement write themselves or mixin a sub-trait of Logging that does this. The way handle is implemented are discussed below.
- ❷ Define convenient type aliases for Request and Response. The choices are inspired by typical web service APIs, where maps of key-value pairs are often used. Note that errors are handled by returning an Either. The logging of service requests is handled in the final method handle, so users of this trait only need to worry about the specific logic of processing a request, by defining the protected method process.
- ❸ A concrete class that extends Service and mixes in the implementation trait ConsoleLogging. The process method expects to find a key user in the map. The open keyword is discussed below.
- ❹ TK4abstract
- ❺ TK5abstract
- ❻ TK6abstract

Finally, an entry point:

TK

Let's run `HelloServiceMain` in `sbt`:

```
> runMain progscale3.basicoop.HelloServiceMain hello Dean Buck
...
INFO: (hello) Starting handle for request: Map(user -> Dean)
INFO: (hello) Finishing handle with result Right(Map(message ->
Hello, Dean))
Hello, Dean
INFO: (hello) Starting handle for request: Map(user -> Buck)
INFO: (hello) Finishing handle with result Right(Map(message ->
Hello, Buck))
Hello, Buck
Try an empty map:
INFO: (hello) Starting handle for request: Map()
INFO: (hello) Finishing handle with result Left(No user field
found!)
Left(No user field found!)
[success] Total time: 1 s, completed Aug 21, 2020, 1:37:05 PM
```

Let's explore the key ideas in this example.

## Classes Open for Extension

Scala 2 constrained subtyping when a type was declared `final` or an abstract supertype was declared `sealed`, but otherwise, you could create subtypes of concrete types. Scala 3 tightens the default rules by adding the `open` keyword. It says, *it is permissible to define a subtype from this concrete type*. Without this keyword, the compiler now issues a warning when a subtype is created from a concrete type.

In the `HelloService` above, we left open the possibility that someone might want to use `HelloService` as a supertype for another `Service`. As a practical matter, the Scaladoc for such a type should provide details about how to extend the type, including what *not* to do. Whether or not `open` is used is a deliberate design decision.

There are two exceptions to the rule that `open` is now required for extension:

1. Subtypes in the same source file, like how sealed hierarchies work.
2. Use of the `adhocExtensions` language feature.

The `adhocExtensions` language feature is enabled globally by adding the compiler flag, `-language:adhocExtensions`, or in single scopes using `import scala.language.adhocExtensions`.

Because this is a breaking change, it is being introduced gradually. In Scala 3.0, the feature warning is only emitted if you use the compiler flag `-source:3.1`. It will be turned on by default in Scala 3.1.

A type that is neither `open` nor `final` now has similar subtyping behavior as a sealed type. The difference is that you can still subtype a “closed” type by enabling the language feature, while sealed type hierarchies can’t be “re-opened”. An important example of this advantage is when you need a *test double* in a unit test, where you create a subtype to stub out certain members. The test source file imports the language feature to enable this subtyping, but subtyping is disallowed in the rest of the code.

As a rule, I try to use *only* abstract types as supertypes and treat all concrete types as *final*, except for the testing scenario. The main reason for this rule is because it’s difficult to get the semantics and implementations correct for `hashCode`, `equals`, and user-defined members. For example, if `Manager` is a subtype of `Employee` (assuming this is a good design...), when should a `Manager` instance be considered equal to an `Employee` instance? If the common subset of fields are equal, is that good enough? The answer depends on the context and other factors. This is one reason why Scala simply prohibits case classes from being subtypes of other case classes.

A few other points. First, `open` is a soft modifier, meaning it can be used as a normal identifier when it’s not in the modifier position. Hence, you don’t need to rename all your `open` methods! Second, `open` can’t be used with `final` or `sealed`, because that would be a

contradiction. Finally, traits and abstract classes are by definition already open, so the keyword is redundant for them.

### TIP

Because composition is usually more robust than inheritance, use open rarely.

## Overriding Methods? The Template Method Pattern

Notice how I declared and implemented the methods in the `Logging` and `Service` types above.

Just as you should avoid subtyping concrete types, you should avoid overriding *concrete* methods. It is a common source of subtle behavioral bugs. Should the subtype implementation call the supertype method? If so, *when* should it call it, at the beginning or end of the overriding implementation? The correct answers depend on the context. It is too easy to make mistakes. Unfortunately, we are so accustomed to overriding the concrete *toString* method, that we consider it normal practice. It should not be normal.

The example above uses the *Template Method Pattern* (GOF1995) to eliminate the need to override concrete methods. The supertypes `Logging` and `Service` define `final`, concrete methods that are publicly visible. `Service.handle` is a template that calls abstract methods, which are the points of allowed variance for subtypes to define. The `Logger` concrete methods are simple templates. They just call the protected, abstract function `Logger.write`. It's easy to implement this correctly, because it does only one thing; write a formatted string somewhere. `ConsoleLogger.write` writes the string to the console.

Similarly, `Service.handle` was implemented carefully to add logging while correctly handling the result of the computation. The

protected, abstract method `Service.process` is *agnostic* to logging. It focuses on processing the request.

However, we can't completely eliminate overriding concrete methods, like `toString`. Fortunately, Scala requires the `override` keyword, which you should treat as a reminder to be careful. When you need to call a supertype method `foo`, use `super().foo(...)`. See also later content (*SelfTypeDeclarations*) for handling the special case when multiple supertypes implement the same method and you need a way to specify a particular one of them.

### TIP

Avoid overriding concrete methods, except when replacing default implementations like `toString`, mixing in orthogonal behavior, or “stubbing” for tests. Be careful to preserve the *contract* of the method. Otherwise, use the *Template Method Pattern* to provide extensibility without overrides. To prevent overrides, declare methods `final`.

## Reference Versus Value Types

While Scala is now a cross-platform language, its roots as a JVM language are reflected in the top-level types.

For performance reasons, the JVM implements a set of special *primitive types*: `short`, `int`, `long`, `float`, `double`, `boolean`, `char`, and `byte`. When used by themselves, meaning not enclosed in other objects, they can be used without the overhead of allocating space for them on the heap and reading and writing the memory location. For example, the compiler-generated byte code or runtime processing could push these values directly onto stack frames or store them in CPU registers. Arrays of these values require only one heap allocation, for the array itself. The values can be “inlined” in the array. These primitives are called *value* types, because the byte code works with these values directly.

All other types are called *reference types*, because all instances of them are allocated on the heap and variables for these instances refer to the corresponding heap locations.

Scala source code breaks the clear distinction between primitives and reference types to provide a more consistent programming model, but without sacrificing performance where possible.

In Scala, all reference types are subtypes of `scala.AnyRef` on the JVM and `js.Object` in Scala.js. `AnyRef` is a subtype of `Any`, the root of the Scala type hierarchy. For Scala.js, `js.Any` is the equivalent supertype of `js.Object`. Note that Java's root type, `Object`, is actually equivalent to `AnyRef`, not `Any`. It is increasingly common to see documentation refer to `Object` directly. Arguably this is more uniform across Scala for the JVM and Scala.js, but it can be confusing to see `Object` and `AnyRef` used interchangeably. I've used `AnyRef` in this book, but keep in mind that you'll see both in documentation.

The Scala types `Short`, `Int`, `Long`, `Float`, `Double`, `Boolean`, `Char`, `Byte`, and `Unit` are *value types*. They correspond to the JVM *primitives* `short`, `int`, `long`, `float`, `double`, `boolean`, `char`, `byte`, and the `void` keyword, respectively. All value types are subtypes of `scala.AnyVal`, which is also a subtype of `Any`. For Scala.js, the JavaScript primitives are used, including `String`, with a rough correspondence to the `AnyVal` types. In the Java and JavaScript object models, primitives don't have a common supertype.

To avoid confusion, I have used `Any`, `AnyRef`, and `AnyVal` consistently with a bias towards the JVM implementations. See the [Scala.js Type Correspondence](#) guide for more details about Scala.js types. The [Scala Native documentation](#) discusses its handling of Scala types.

`Unit` is an `AnyVal` type, but it involves no storage at all. Loosely speaking, `Unit` is analogous to the `void` keyword in many languages. This is only true in the sense that a method returning `Unit` or `void` doesn't return anything you can use. Otherwise, `Unit` or `void` are

quite different. While `void` is a keyword, `Unit` is a real type with *one* literal value, `()`, which we rarely use that value explicitly. This means that *all* functions and methods in Scala return a value, whereas languages with `void` have a separate idea of *functions* that return a value and *procedures* that don't.

### WHY IS UNIT'S LITERAL VALUE ()?

`Unit` really behaves like a tuple with zero elements, written as `()`.<sup>1</sup> If there are no elements, it contains no useful information. The name *unit* comes from algebra, where adding the unit to any value returns the original value, such as 0 for integers. For multiplication, 1 is the unit value. So, if I “add” `()` to `(1, 2.2)`, I get back `(1, 2.2)`, but if I “add” `(3L)` to `(1, 2.2)`, I get back `(3L, 1, 2.2)` (or `(1, 2.2, 3L)`). We'll explore this idea more precisely in a later chapter.

As an aside, consider a sequence of `AnyVals`? What is the *least upper bound*?

### THE LEAST UPPER BOUND FOR ANYVALS

For the special case where a sequence of numbers contains `Floats` and `Ints`, like `Seq(1, 2.2F, 3)`, the inferred type is `Seq[Float]`. The `Ints` are converted to `Floats`. Similarly if `Ints` and `Doubles` are mixed, you get `Seq[Double]`. *For all other combinations* of `AnyVals`, the inferred type is `Seq[AnyVal]`, even when all the values are `Doubles` and `Floats`.

## Opaque Types and Value Classes

In “*Implicit Conversions*”, we defined some wrapper types for `Dollars` and `Percentages`:

```
//
src/main/scala/progscala3/contexts/accounting/NewImplicitConversions.s
cala
package progscala3.contexts.accounting
import scala.language.implicitConversions

case class Dollars(amount: Double):
  ...
case class Percentage(amount: Double):
  ...
```

Now imagine you are writing a *big data* application that creates millions or more instances of these types. The extra overhead of heap allocations and memory accesses for these wrapper types becomes very expensive. They can be quite stressful for the garbage collector. Fundamentally, these types just wrap Doubles, so we would prefer to keep the efficiency of primitive doubles, without giving up the convenience of object orientation and custom types.

Let's consider *three* potential solutions to this issue. First, we could use type aliases for Dollars and Percentage:

```
// tag::definitions[]
// src/script/scala/progscala3/basicoop/DollarsPercentagesTypes.scala

object Accounting:
  type Dollars = Double
  type Percentage = Double

import Accounting._
case class Salary(gross: Dollars, taxes: Percentage):
  def net: Dollars = gross * (1.0 - taxes)
  override def toString =
    f"Salary(gross = $$$gross%.2f, taxes = ${ (taxes*100.0) }%.2f%%)"
// end::definitions[]

// tag::oops[]
import Accounting._
val gross: Dollars = 10000.0
val taxes: Percentage = 0.1
val salary1 = Salary(gross, taxes)
val net1 = salary1.net
```



```
val salary2 = Salary(taxes, gross)    // Error!!  
val net2 = salary2.net  
// end::oops[]
```

Now let's try it:

```
scala> import Accounting._  
      | val gross: Dollars = 10000.0  
      | val taxes: Percentage = 0.1  
val gross: Accounting.Dollars = 10000.0  
val taxes: Accounting.Percentage = 0.1  
  
scala> val salary1 = Salary(gross, taxes)  
      | val net1 = salary1.net  
val salary1: Salary = Salary(gross = $10000.00, taxes = 10.00%)  
val net1: Accounting.Dollars = 9000.0  
  
scala> val salary2 = Salary(taxes, gross)    // Error, but it  
      | val net2 = salary2.net  
val salary2: Salary = Salary(gross = $0.10, taxes = 1000000.00%)  
val net2: Accounting.Dollars = -999.90000000000001
```

This is a simple solution, but it has some problems that impact users of the API. The type aliases are just new names for the same type, so the compiler doesn't catch the mistake of mixing up the arguments. Hence, type aliases provide *no* additional type safety. Furthermore, we can't define custom methods for `Dollars` and `Percentage`. Attempting to use extension methods will add them to `Double`, not separately for the two types.

## Opaque Type Aliases

Scala 3 introduces *opaque type aliases*, which are declared like regular type aliases, but with the *opaque* keyword. They preserve type safety and have zero runtime overhead beyond the value they wrap, but they provide some of the benefits of using “richer” types. Here is the same example rewritten with opaque type aliases:

```

// tag::definitions[]
// src/script/scala/progscala3/basicoop/DollarsPercentagesOpaque.scala

object Accounting:
  opaque type Dollars = Double
  opaque type Percentage = Double

  object Dollars:
    def apply(amount: Double): Dollars = amount

    extension (d: Dollars):
      def +(d2: Dollars): Dollars = d + d2
      def -(d2: Dollars): Dollars = d - d2
      def *(p: Percentage): Dollars = d*p
      def toDouble: Double = d
      def toString = f"$$$d%.2f"

  object Percentage:
    def apply(amount: Double): Percentage = amount

    extension (p: Percentage)
      def +(p2: Percentage): Percentage = p + p2
      def -(p2: Percentage): Percentage = p - p2
      def *(d: Dollars): Dollars = d*p
      def toDouble: Double = p
      def toString = f"${(p*100.0)}%.2f%"

import Accounting._
case class Salary(gross: Dollars, taxes: Percentage):
  def net: Dollars = gross - (gross * taxes)
// end::definitions[]

// tag::usage[]
val gross = Dollars(10000.0)
val taxes = Percentage(0.1)
val salary1 = Salary(gross, taxes)
val net1 = salary1.net
val salary2 = Salary(taxes, gross) // Won't compile!
val net2 = salary2.net
// end::usage[]

```

- ❶ Like regular type aliases, but prefixed with the opaque keyword.

- ② For each opaque type alias, define an object that *looks like a companion object* for factory methods like `apply`, so they behave similar to user expectations for other types. All “instance” methods for an opaque type are defined as extension methods.
- ③ Our original wrapper types for `Dollars` and `Percentage` had nice `toString` methods and we could have implemented `equals` methods that incorporate accounting rules. We don’t have the option to override concrete methods for opaque type aliases.
- ④ The companion for `Percentage`.
- ⑤ A case class that uses these types.

The opaque keyword is soft. It is only treated as a keyword in a declaration as shown. Otherwise, it can be used as a regular identifier.

Compared to case classes, opaque classes have most of the limitations of regular type aliases. You can’t override concrete methods like `equals` and `toString` for opaque type aliases, nor pattern match on them. You can only pattern match on the underlying type.

You can define an object with the same name for factory methods like `apply`, but they aren’t generated automatically like they are for case classes.

Note that regular `Double` methods, like those for arithmetic and comparisons, are not automatically available for users of these types. We have to define extension methods for the operations we want or call `toDouble` first.

However, this only applies for users of an opaque type outside the scope where the type is defined. This is why they are called

“opaque”. Inside the scope, the type looks like a `Double`. Note how the bodies are implemented using `Double` methods.

Outside the defining scope, an opaque type is considered *abstract*, even though the definition inside the scope is concrete. Note how `Dollars` and `Percentage` are used in the next code snippet below. When we construct `Dollars`, we call the `Dollars` object method `apply`. Everywhere else, like arguments to `Salary`, there’s nothing to require `Dollars` to be concrete, just like using `Seq` everywhere, even though it is not a concrete type:

```
scala> import Accounting._
scala> val gross = Dollars(10000.0)
      | val taxes = Percentage(0.1)
val gross: Accounting.Dollars = 10000.0
val taxes: Accounting.Percentage = 0.1

scala> val salary1 = Salary(gross, taxes)
      | val net1 = salary1.net
val salary1: Salary = Salary(10000.0,0.1)
val net1: Accounting.Dollars = 9000.0

scala> val salary2 = Salary(taxes, gross) // Won't compile!
5 | val salary2 = Salary(taxes, gross) // Won't compile!
  |                      ^^^^^
  |                      Found:    (taxes : Accounting.Percentage)
  |                      Required: Accounting.Dollars
5 | val salary2 = Salary(taxes, gross) // Won't compile!
  |                      ^^^^^
  |                      Found:    (gross : Accounting.Dollars)
  |                      Required: Accounting.Percentage
```

When printing the values, we no longer have the dollar and percentage signs. If we still want them, we would have to implement an ad-hoc print method of some kind and use that instead of relying on `toString`.

However, as desired, we get the type checking for `Dollars` vs. `Percentages` that we want and we don’t pay a runtime penalty for wrapping these types.

The inability to override the equality methods, `equals`, `==`, and `!=`, means the underlying type's methods are used. This can cause surprises:

```
scala> val dollars = Dollars(100.0)
      | val percentage = Percentage(100.0)

scala> dollars == percentage
val res0: Boolean = true
```

However, you can define extension methods for `<=`, `>=`, etc.

### WARNING

When comparing instances of different opaque types that alias the same underlying type, the underlying type's equality operations are used, even if the instances should not be considered comparable! Avoid using these methods. Define other, ad-hoc extension methods to fine-tune the behavior and use them instead.

At compile time, opaque types work like regular types, but the byte code generated only uses the overhead of the aliased type, `Double` in this case. For all three approaches we are discussing here, you aren't limited to aliasing `AnyVal` types, either. Even wrapping `Strings`, for example, means one less heap allocation and fewer memory operations.

Opaque types can also have type parameters. The code examples contain two variations of an example that I won't discuss here. They are adapted from an example in the *Scala Improvement Process* proposal for opaque types, [SIP-35](#), which shows a no-overhead way to “tag” values with metadata. In this case, units like meters vs. feet are the tags. The implementations are a bit advanced, but worth studying to appreciate both the idea of tagging data with metadata in a type-safe way and the way it is implemented with no runtime overhead. See

src/main/scala/progscala3/basicoop/tagging/Tags.scala and Tags2.scala. SIP-35 contains other nontrivial examples, too.<sup>2</sup>

## Opaque Type Aliases and Matchable

In an earlier chapter, we saw that pattern matching is now restricted to subtypes of the trait `Matchable`, which fixes a “hole” when pattern matching on certain type aliases like the new `IArray`. I explained pattern matching on abstract types requires them to be bound by `Matchable` and this solves the issue with `IArray` discussed there. In fact, the library’s `IArray` is an opaque type alias for `Array`. so now I can fill in a few more details. Consider the following example with our own `Array` aliases:

```
scala> // src/script/scala/progscala3/basicoop/MatchableOpaque.scala
|
| object Obj:
|   type Arr[T] = Array[T]
|   opaque type OArr[T] = Array[T]
|
scala> summon[Obj.Arr[Int] <: Matchable]           // Okay
val res2: Array[Int] := Array[Int] = generalized constraint

scala> summon[Obj.OArr[Int] <: Matchable]           // Doesn't work
1 | summon[Obj.OArr[Int] <: Matchable]           // Doesn't work
   |                                     ^
   |                                     Cannot prove that Obj.OArr[Int] <: Matchable.
```

Recall we used `summon` like this in “[Implicit Evidence](#)” to check type relationships. So, why is that our type alias `Arr` is considered a `Matchable`, but not our opaque alias `OArr`? It’s because `OArr` is considered *abstract* outside of `Obj`, as we discussed above. Recall from the discussion of `Matchable` that abstract types must be declared bounded by `Matchable`. In contrast, `Arr` not abstract and it aliases the concrete type `Array`, which subtypes `Matchable`.

If you change the definition of `OArr` to this, `opaque type OArr[T] <: Matchable = Array[T]`, the last `summon` will succeed. Try it!

## Value Classes

Scala 2 and 3 offer *value classes*, our third and final mechanism to eliminate the runtime overhead of wrapper types. They have some advantages and drawbacks compared to opaque type aliases.

Here is an example value class for North American phone numbers (excluding the country code):

```
// src/main/scala/progscala3/basicoop/ValueClassPhoneNumber.scala
package progscala3.basicoop

class USPhoneNumber(val s: String) extends AnyVal: ❶
  override def toString = {
    val digs = digits(s)
    val areaCode = digs.substring(0,3)
    val exchange = digs.substring(3,6)
    val subnumber = digs.substring(6,10) // "subscriber number"
    s"($areaCode) $exchange-$subnumber"
  }

  private def digits(str: String): String = str.replaceAll("""\D""",
    "")
```

- ❶ Note it extends `AnyVal`. For simplicity, validation of the input string is not shown.

Now we have the convenience of a domain-specific type, with customized methods, but instances don't require additional memory management beyond what the `String` requires.

To be a valid value class, the following rules must be followed:

1. The value class has one and only one `val` parameter.
2. The type of the parameter must not be a value class itself.
3. The value class doesn't define secondary constructors (see ["Constructors in Scala"](#)).

4. The value class defines only methods, but no other vals and no vars.
5. The value class defines no nested traits, classes, or objects.
6. The value class can't be subtyped.
7. The value class can only inherit from *universal traits* (see below).
8. The value class must be a top-level type or a member of an object that can be referenced.<sup>3</sup>

The compiler provides good error messages when we break the rules.

At compile time the type is the outer type, `NAPhoneNumber`. The runtime type is the wrapped type, `String`. The wrapped type can be any other type, as long as the rules are followed.

A value class can be a case class, but the many extra methods and the companion object generated are less likely to be used and hence more likely to waste space in the output class file.

A *universal trait* has the following properties:

1. It subtypes `Any` (but not from other universal traits).
2. It defines only methods.
3. It does no initialization of its own.

Here is a refined version of `NAPhoneNumber` that mixes in two universal traits:

```
// src/main/scala/progscala3/basicoop/ValueClassUniversalTraits.scala
package progscala3.basicoop

trait Digitizer extends Any:
  def digits(s: String): String = s.replaceAll("""\D""", "")
```



❶

```
trait Formatter extends Any:
```

❷

```
  def format(  
    areaCode: String, exchange: String, subnumber: String): String  
  =  
    s"($areaCode) $exchange-$subnumber"
```

```
case class USPhoneNumberUT(s: String)  
  extends AnyVal with Digitizer with Formatter:  
  override def toString =  
    val digs = digits(s)  
    val areaCode = digs.substring(0,3)  
    val exchange = digs.substring(3,6)  
    val subnumber = digs.substring(6,10)  
    format(areaCode, exchange, subnumber)
```

❸

- ❶ Digitizer is a trait that contains the digits method we originally had in NPhoneNumber.
- ❷ Formatter formats the phone number the way we want it.
- ❸ Use `Formatter.format`.

Formatter actually solves a design problem. We might like to specify a second parameter to NPhoneNumber for a format string to use in `toString`, because there are many popular format conventions for phone numbers. However, we're only allowed to pass one argument to the NPhoneNumber constructor and it can't have any other fields. We can solve this problem by mixing in a universal trait to do the configuration we want. We could define a different Formatter trait and build a different PhoneNumber value class that uses it.

The biggest drawback of value classes is that some non-obvious circumstances can trigger instantiation of the wrapper type, defeating the purpose of using value classes. One situation involves universal traits. Here's a summary of the circumstances requiring instantiation:

1. When a function expects a universal trait instance and it is passed an instance of a value class that implements the trait. However, if a function expects an instance of the value class itself, instantiation isn't required.
2. An Array or another collection of value class instances.
3. The type of a value class is used as a type parameter.

For example, when the following method is called with a `NAPhoneNumber`, an instance of it will be allocated on the heap:

```
def toDigits(d: Digitizer, str: String) = d.digits(str)
...
val digs = toDigits(NAPhoneNumber("987-654-3210"), "123-Hello!-456")
// Result: digs: String = 123456
```

Similarly, when the following parameterized method is passed a `NAPhoneNumber`:

```
def print[T](t: T) = println(t.toString)
print(NAPhoneNumber("987-654-3210"))
// Result: (987) 654-3210
```

Opaque type aliases work around these scenarios, although they bring their own limitations, like inability to override `toString` and `equals` for them.

### NOTE

To clarify terminology, *value type* refers to the `Short`, `Int`, `Long`, `Float`, `Double`, `Boolean`, `Char`, `Byte`, and `Unit` types. *Value class* refers to user-defined classes that subtype `AnyVal`.

To summarize the three approaches for avoiding wrapper type overhead, regular type aliases are a very simple approach, but don't provide any extra type safety. Both value classes and opaque types

provide better type safety and semantics. Use opaque types when performance is the highest priority, when avoiding *all* extra heap allocation and memory accesses are important. Use value classes if you want types that behave more like “real” types, such as the ability to override `toString` and `equals`, and you can tolerate some situations where heap allocation is necessary.

## Super Types

Throughout the book, I have mostly used the terms *supertype* and *subtype*, both as nouns and verbs. *Subtyping* creates a *subtype* from a *supertype*. Other common OOP terms for subtyping include *derivation*, *extension*, and *inheritance*. Scala documentation describes *type class derivation*, so I used that terminology in “**Type Class Derivation**”. The keyword `open` and *extension* are often used together, as I used them in “**Classes Open for Extension**”.

Supertypes are also called *parent* or *base* types. Subtypes are also called *child* or *derived* types.

Scala only supports *single inheritance*, but most of the benefits of *multiple inheritance* are achieved using mixins. All types have a supertype, except for the root of the Scala class hierarchy, `Any`. When declaring a type that isn’t an alias and when a parent type is omitted, the type implicitly subtypes `AnyRef`. In other words, it is automatically a reference type and instances will be heap-allocated.

The keyword `extends` indicates the supertype class or trait. If other traits are mixed in, the `with` keyword is used.

## Constructors in Scala

Scala distinguishes between the *primary constructor* and zero or more *auxiliary constructors*, also called *secondary constructors*. In Scala, the primary constructor is the entire body of the type. Any

parameters that the constructor requires are listed after the type name.

Here is an example with auxiliary constructors. The type represents US zip codes, which have five digits and an optional extension of four digits:

```
//  
src/script/scala/progscala3/basicoop/people/ZipCodeAuxConstructors.scala  
  
case class ZipCodeAuxCtor(zip: Int, extension: Int = 0):  
  override def toString =  
    if extension != 0 then s"$zip-$extension" else zip.toString  
  
  def this(zip: String, extension: String) =  
    this(zip.toInt, if extension.length == 0 then 0 else  
extension.toInt)  
  def this(zip: String) = this(zip, "")
```

The two auxiliary constructors, named `this`, allow users to provide string arguments. They are converted to integers, where 0 is interpreted as “no extension”. All auxiliary constructors are required to call the primary constructor or another auxiliary constructor as the first expression. The compiler also requires that a constructor called is one that appears *earlier* in the source code. So, we must order secondary constructors carefully in our code.

Forcing all construction to go through the primary constructor eliminates duplication of constructor logic and the risk of inconsistent initialization of instances.

We haven’t discussed auxiliary constructors before now, because it’s rare to use them. It’s far more common to overload object apply methods instead when multiple invocation options are desired:

```
// src/script/scala/progscala3/basicoop/people/ZipCodeApply.scala  
  
case class ZipCodeApply(zip: Int, extension: Int = 0):  
  override def toString =
```

```

    if extension != 0 then s"$zip-$extension" else zip.toString

object ZipCodeApply:
  def apply(zip: String, extension: String): ZipCodeApply =
    apply(zip.toInt, if extension.length == 0 then 0 else
extension.toInt)
  def apply(zip: String): ZipCodeApply = apply(zip, "")

```

This has the nice benefit of keeping the case class very simple. The complexity of invocation is moved to the companion object. There are actually three `apply` methods in the companion object. The compiler generates the `apply` method with two `Int` parameters, matching the constructor. The source code adds two more.

The following invocations all work:

```

ZipCodeApply(12345)
ZipCodeApply(12345, 6789)
ZipCodeApply("12345")
ZipCodeApply("12345", "6789")

```

## Calling Supertype Constructors

The primary constructor in a subtype must invoke one of the supertype constructors:

```

class Person(name: String, age: Int)
class Employee(name: String, age: Int, salary: Float) extends
Person(name, age)
class Manager(name: String, age: Int, salary: Float, minions:
Seq[Employee])
  extends Employee(name, age, salary)

```

## Export Clauses

In the next section, we'll discuss the benefits of composition over inheritance. There are times when you want a type to be composed of other types, but you also want members on those other types to be part of the public interface of the "composed" type. This happens

automatically with public members of inherited types, including mixins. However, if dependencies are passed in as constructor arguments (also known as *dependency injection*) or local instances are created, and you would like some of the members of those instances to be part of the composed type's interface, you have to write “forwarding” methods. To see this, consider these types for authentication that some service might want to provide:

```
// tag::traits[]
// src/script/scala/progscala3/basicoop/Exports.scala

case class UserName(value: String):
  ❶ assert(value.length > 0)
case class Password(value: String):
  assert(value.length > 0)

trait Authenticate:
  final def apply(
    ❷ username: UserName, password: Password): Boolean =
    authenticated = auth(username, password)
    authenticated
  def isAuthenticated: Boolean = authenticated

  private var authenticated = false //
  protected def auth(
    username: UserName, password: Password): Boolean

class DirectoryAuthenticate extends Authenticate:
  protected def auth(
    username: UserName, password: Password): Boolean = true

class ResourceManager(
  ❸ private var resources: Map[String,Any] = Map.empty):
  def getResource(key: String): Option[Any] = resources.get(key)
  def setResource(key: String, value: Any) = resources += (key ->
value)
  def deleteResource(key: String) = resources -= key

import scala.concurrent.Future
```

```

import scala.concurrent.ExecutionContext.Implicits.global

trait AsyncWorker[A,B]:
  ④ def apply(data: A): Future[B]

class AsyncTokenizer extends AsyncWorker[String, Seq[String]]:
  def apply(string: String): Future[Seq[String]] =
    Future(string.split("""\W+""").toSeq)
  // end::traits[]

  // tag::service1[]
  object ServiceWithoutExports:
    private val dirAuthenticate = new DirectoryAuthenticate
    private val manager = new ResourceManager(sys.env)
    ⑤ private val tokenizer = new AsyncTokenizer

    def authenticate(username: Username, password: Password): Boolean =
      dirAuthenticate.authenticate(username, password)
    def isAuthenticated: Boolean = dirAuthenticate.isAuthenticated
    def getResource(key: String): Option[Any] = manager.getResource(key)
    def tokenize(string: String): Future[Seq[String]] =
      tokenizer(string)
  // end::service1[]

  // tag::service2[]
  object Service:
    private val dirAuthenticate = new DirectoryAuthenticate
    private val manager = new ResourceManager(sys.env)
    private val tokenizer = new AsyncTokenizer

    export dirAuthenticate.{apply => authenticate, isAuthenticated}
    export manager.getResource
    export tokenizer.{apply => tokenize}
  // end::service2[]

  // tag::example[]
  Service.isAuthenticated
  Service.authenticate(Username("Buck Trends"), Password("1234"))
  Service.isAuthenticated
  Service.getResource("HOME")
  val tokensF = Service.tokenize("Hello from the World!")
  tokensF.value
  // end::example[]

```

- ❶ Types to encapsulate valid user names and passwords (encryption needed!).
- ❷ An abstraction for authentication with a “stub” implementation that uses a directory service available at some URL.
- ❸ TK3
- ❹ TK4
- ❺ TK5

Now let's compose a service that declares a private field for an instance of `DirectoryAuthenticate`, then defines boilerplate methods to expose the two public methods provided by `dirAuthenticate`:

TK

Scala 3 gives us a way to avoid the boilerplate methods using a new export clause:

TK

Before explaining the export clause, let's see `Service` in action:

```
scala> Service.isAuthenticated
val res0: Boolean = false

scala> Service.authenticate(Username("Buck Trends"),
Password("1234"))
val res1: Boolean = true

scala> Service.isAuthenticated
val res1: Boolean = true
```

Export clauses are similar to import clauses. We gave the `apply` method an alias; using `Service.apply` to `authenticate` would be weird. We can export any members, not just methods. We can



exclude items. For example, `export dirAuthenticate.{apply => _, _}` would export all public members *except* for `apply`.

The following rules describe what members are eligible for exporting:

- The member can't be owned by a supertype of the type with the export clause.
- The member can't override a concrete definition in a supertype, but it can be used to implement an abstract member in a supertype.
- The member is accessible at the export clause.
- The member is not a constructor.
- The member is not the synthetic (compiler-generated) class part of an object.
- If the member is a *given instance* or implicit value, then the export must be tagged with `given`.

All exports are `final`. They can't be overridden in subtypes.

Export clauses can also appear outside types, meaning they are defined at the package level. Hence, one way to provide a very controlled view of what's visible in a package is to declare everything *package private*, then use export clauses to expose only those items you want publicly visible. See *VisibilityRules* for more details on public, protected, private, and more fine-grained visibility controls.

## Good Object-Oriented Design: A Digression

Consider the example above where `Person` was a supertype of `Employee`, which was a supertype of `Manager`. It has several *code smells*.

First, there's a lot boilerplate in the constructor argument lists, like `name: String`, `age: Int` repeated three times. Second, it seems like

all three should be case classes, right?

We can create a regular subtype from a case class or the other way around, but we can't subtype one case class to create another. This is because the auto-generated implementations of `toString`, `equals`, and `hashCode` do not work properly for subtypes, meaning they ignore the possibility that an instance could actually be a subtype of the case class type.

This limitation is by design. It reflects the problematic aspects of subtyping. For example, should `Manager`, `Employee`, and `Person` instances be considered equal if they all have the same name and age? A more flexible interpretation of object equality might say yes, while a more restrictive version would say no. Also, the mathematical definition of equality requires commutative behavior: `somePerson == someEmployee` should return the same result as `someEmployee == somePerson`, but you would never expect the `Employee.equals` method to return true in this case.

The real problem is that we are subtyping the *state* of these instances. We are using subtyping to add additional fields that contribute to the instance state. In contrast, subtyping *behavior* (methods) with the *same* state is easier to implement robustly. It avoids the problems with `equals` and `hashCode` just described, for example.

Of course, these problems with inheritance have been known for a long time. Today, good object-oriented design favors *composition over inheritance*, where we compose units of functionality rather than build class hierarchies, when possible.

*Mixin composition* with `traits` is the primary makes composition straightforward. The code examples mostly use type hierarchies with few levels and mixins to enhance them. When bits of cleanly-separated state and behavior are combined, mixin composition is robust.

When subtyping is used, I recommend the following rules:

1. Use only one level of subtyping from a supertype, if at all possible.
2. Concrete classes are never subtyped, except for two cases:
  - a. Classes that mix in other behaviors defined in traits (see [Chapter 10](#)). Ideally, those behaviors should be *orthogonal*, i.e., not overlapping.
  - b. Test-only versions to promote automated unit testing.
3. When subtyping seems like the right approach, consider partitioning behaviors into traits and mix in those traits instead. Recall our NPhoneNumber design earlier in this chapter.
4. Never build up logical state across supertype-subtype boundaries.
5. Only use case classes for “leaf” nodes in a type hierarchy. That is, don’t subtype case classes.
6. Make your intentions explicit by marking types `open`, `final`, or `sealed`, as appropriate.

By “logical” state in the third to last bullet, I mean the fields and methods, which together define a *state machine* for the logical behavior. There might have some private, implementation-specific state that doesn’t affect this external behavior, but be very careful that the internals don’t leak through the type’s abstraction. For example, a library might include private subtypes for special cases. Types might have private fields to implement caching, auditing, or other *concerns* that aren’t part of the public abstraction.

So, what about our Person hierarchy? What should we do instead? The answer really depends on the context of use. If we’re implementing a Human Resources application, do we need a separate concept of Person or can Employee just be the only type,

declared as a case class? Do we even need *any* types for this at all? If we're processing a result set from a database query, is it sufficient to use tuples or maps to hold the values returned from the query for each unique use case? Can we dispense with the “ceremony” of declaring a type altogether?

Here is an alternative with just a single `Employee` case class that embeds the assumption that non-managers will have an empty set of subordinates:

```
// src/script/scala/progscala3/basicoop/people/Employee.scala

case class Employee(
  name:    String,
  age:     Int,
  title:   String,
  manages: Set[Employee] = Set.empty)

val john = Employee("John Smith", 35, "Accountant")
val jane = Employee("Jane Doe", 28, "Full Stack Developer")
val tom  = Employee("Tom Tired", 22, "Junior Minion")
val minions = Set(john, jane, tom)
val ceo = Employee("John Smith", 60, "CEO", minions)
```

## Fields in Types

We started the chapter with a reminder that the primary constructor parameters become instance fields if they are prefixed with the `val` or `var` keyword. For case classes, `val` is implied. This convention greatly reduces source-code boilerplate, but how does it translate to byte code?

Actually, Scala just does implicitly what Java code does explicitly. There is a private field created internal to the class and the equivalent of “getter” and optional “setter” accessor methods are generated. Consider this simple Scala class:

```
class Name(var value: String)
```

Conceptually, it is equivalent to this code:

```
class Name(s: String):  
  private var _value: String = s ❶  
  
  def value: String = _value      ❷  
  def value_=(newValue: String): Unit = _value = newValue
```

- ❶ Invisible field, declared mutable in this case.
- ❷ The “getter” (reader) and the “setter” (writer) methods.

Note the convention used for the `value_ =` method name. When the compiler sees a method named like this, it will allow client code to drop the `_`, effectively enabling infix notation as if we were setting a bare field in the object:

```
scala> val n = Name("Dean")  
val n: Name = Name@333f6b2d  
  
scala> n.value  
val res0: String = Dean  
  
scala> n.value = "Buck"  
  
scala> n.value  
val res1: String = Buck
```

If we declare a field immutable with the `val` keyword, the field is declared with `val` and the writer method is not synthesized, only the reader method.

You can use these conventions yourself, if you want to expose a “field” implemented with custom logic for reading and writing.

Constructor parameters in non case classes without the `val` or `var` don’t become fields. It’s important to note that the value is still in the

scope of the *entire* class body. Let's fill in a bit more detail that might be required for a real `DirectoryAuthenticate` implementation:

```
class DirectoryAuthenticate(location: URL) extends Authenticate:  
  protected def auth(username: UserName, password: Password):  
    Boolean =  
      directory.auth(username, password)  
  
  protected val directory = connect(location) // we can use location!  
  protected def connect(location: URL) = ???
```

While `connect` refers to `location`, it is not a field. It is just in scope!

Why not always make these parameters fields with `val`? Unless these parameters are really part of the logical state exposed to users, they shouldn't be exposed as fields. Instead, they are effectively private to the class body.

## The Uniform Access Principle

In the `Name` example, it appears that users can read and write the “bare” `value` field without going through accessor methods, but in fact they are calling methods. On the other hand, we could just declare a field in the class body with the default public visibility and then access it as a bare field:

```
class Name2(s: String):  
  var value: String = s
```

This uniformity is called the *Uniform Access Principle*. The user experience is identical. We are free to switch between bare field access and accessor methods as needed. For example, if we want to add some sort of validation on writes or lazily construct the field value on reads, then methods are better. Conversely, bare field access is faster than a method call, although some simple method invocations will be inlined by the compiler or runtime environment anyway.

Therefore, the *Uniform Access Principle* has an important benefit in that it minimizes coupling between user code and implementation details. We can change that implementation without forcing client code changes, although a recompilation is required.

Because of the flexibility provided by uniform access, a common convention is to declare abstract, constant “fields” as methods instead:

```
// src/main/scala/progscala3/basicoop/AbstractFields.scala
package progscala3.basicoop

trait Logger:
  def level: LoggingLevel           ❶
  def log(message: String): Unit

case class ConsoleLogger(level: LoggingLevel) extends Logger: ❷
  def log(message: String): Unit = println(s"$level: $message")
```

- ❶ Declare the logging level as a method with no parentheses, rather than a `val`.
- ❷ Implement `loggingLevel` with a `val` constructor parameter instead of a concrete method.

Implementers have the choice of using a concrete method or using a `val`. Using a `val` is legal here because the “contract” of `loggingLevel` is that it returns some `Int`. If the same value is *always* returned, that satisfies the contract. Conversely, if we declared `loggingLevel` to be a field in `Logger`, then using a concrete method implementation would not be allowed, because the compiler can’t confirm that the method would always return a single value consistently.

I have used this convention in earlier examples. Did you notice it?

## TIP

When declaring an abstract field, consider declaring an abstract method instead. That gives implementers the freedom to use a method or a field.

## Unary Methods

We saw above how to define an assignment method `foo_ =` for field `foo`, which enables the intuitive syntax `myinstance.foo = bar`. How can we implement *unary operators*?

An example is negation. If we implement a complex number class, we want to support the negation of some instance `c` with `-c`:

```
// tag::definitions[]
// src/main/scala/progscala3/basicoop/Complex.scala
package progscala3.basicoop

import scala.annotation.alpha

case class Complex(real: Double, imag: Double) derives Eql:
  ❶
  ❷
  @alpha("negate") def unary_- : Complex =
    Complex(-real, imag)
  ❸
  @alpha("minus") def -(other: Complex) =
    Complex(real - other.real, imag - other.imag)
// end::definitions[]

// tag::entrypoint[]
@main def TryComplex: Unit =
  val c1 = Complex(1.1, 2.2)
  assert(c1.real == 1.1)
  assert(c1.imag == 2.2)
  assert(-c1 == Complex(-1.1, 2.2))
  ❹
  assert(c1.unary_- == Complex(-1.1, 2.2))
  assert(c1 - Complex(0.1, 0.2) == Complex(1.0, 2.0))
// end::entrypoint[]
```

- ❶ The method name is `unary_X`, where `X` is the *prefix* operator character we want to use, `-` in this case. Note the space between



the `-` and the `:` is necessary to tell the compiler that the method name ends with `-` and not with `!`

- ② For comparison, we also implement the minus *infix* operator for subtraction.
- ③ TK3c
- ④ TK4c

For example:

```
scala> import progscala3.basicoop.Complex
scala> val c = Complex(1.1, 2.2)
scala> assert(-c == Complex(-1.1, 2.2))
```

## Recap and What's Next

We filled in more details for object-oriented programming in Scala, including constructors vs. object apply methods, inheritance, and some tips on good object-oriented design.

We also set the stage for diving into traits, Scala's powerful tool for composing behaviors from constituent parts. In the next chapter we'll complete our understanding of traits and how to use them to solve various design problems.

- 
- 1 Scala 3 adds an actual `EmptyTuple` type, which is different than `Unit`.
  - 2 Some of the SIP-35 details are obsolete. The Scala 3 documentation is correct.
  - 3 Because of Scala's richer type system, not all types can be referenced in normal variable and method declarations. (However, all the examples we've seen so far work fine.) In *Chapter 11*, we'll explore new kinds of types and learn the rules for what it means to say that a type can or can't be referenced.

# Chapter 10. Traits

---

## A NOTE FOR EARLY RELEASE READERS

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

This will be the 10th chapter of the final book. Please note that the GitHub repo will be made active later on.

If you have comments about how we might improve the content and/or examples in this book, or if you notice missing material within this chapter, please reach out to the author at [mcronin@oreilly.com](mailto:mcronin@oreilly.com).

Scala traits function as *interfaces*, abstract declarations of members (methods, fields, and types) that together express state and behavior. Traits can also provide concrete definitions (implementations) of some of those declarations.

A trait with concrete definitions works best when those definitions provide state and behavior that are well encapsulated and loosely coupled or even “orthogonal” to the rest of the state and behavior in the types that use the trait. The term *mixin* is used for such traits, because we should be able “mix together” such traits to compose different concrete types.

## Traits as Mixins

We first discussed traits as mixins in “[Traits: Interfaces and “Mixins” in Scala](#)”, where we explored an example that mixes logging into a

service type. Logging is a good example of mixin that can be well encapsulated and orthogonal to the state and behavior of the service type.

Let's revisit and expand on what we learned with a new example. Consider the following code for a button in a graphical user interface (GUI) toolkit, which uses callbacks to notify clients when clicks occur:

```
// src/main/scala/progscala3/traits/ui/ButtonCallbacks.scala
package progscala3.traits.ui

abstract class ButtonWithCallbacks(val label: String,
    val callbacks: Seq[() => Unit] = Nil) extends Widget:
  ❶

  def click(): Unit =
    ❷
    updateUI()
    callbacks.foreach(f => f())

  protected def updateUI(): Unit
  ❸
```

- ❶ The class is abstract, because `updateUI` is abstract. `Widget` is define abstract class `Widget` for now.
- ❷ When the button is clicked, invoke the list of callback functions of type `() => Unit`, which can only perform side effects, like sending a message to a back-end service.
- ❸ Update the UI.

This class has two responsibilities: updating the visual appearance and handling callback behavior, including the management of a list of callbacks and calling them whenever the button is clicked. We should strive for *separation of concerns* in our type definitions and not mix multiple responsibilities. Rather, we should mix concerns when necessary through composition.

So, we would like to separate the button-specific logic from the callback logic, such that each part becomes simpler, more modular, easier to test and modify, and more reusable. The callback logic is a good candidate for a mixin.

Callbacks are a special case of the *Observer Design Pattern*. So, let's create two traits that declare and partially implement the Subject and Observer logic in this pattern, then use them to handle callback behavior. To simplify things, we'll start with a single callback that counts the number of button clicks:

```
// src/main/scala/progscala3/traits/observer/Observer.scala
package progscala3.traits.observer

trait Observer[State]:
  ❶
    def receiveUpdate(state: State): Unit

trait Subject[State]:
  ❷
    private var observers: Vector[Observer[State]] = Vector.empty
  ❸

    def addObserver(observer: Observer[State]): Unit =
  ❹
        observers := observer
  ❺

    def notifyObservers(state: State): Unit =
  ❻
        observers foreach (_.receiveUpdate(state))
```

- ❶ The trait for clients who want to be notified of state changes. They must implement the `receiveUpdate` message.
- ❷ The trait for subjects who will send notifications to observers.
- ❸ A list of observers to notify. It's mutable, so it's not thread-safe!
- ❹ A method to add observers.

- ⑤ Since `observers` is a `var`, this expression is equivalent to `observers = observer +: observers`.
- ⑥ A method to notify observers of state changes.

Often, the most convenient choice for the `State` type parameter is just the type of the class mixing in `Subject`. Hence, when the `notifyObservers` method is called, the instance just passes itself, i.e., `this`.

### TIP

Traits with abstract members don't have to be declared abstract by adding the `abstract` keyword before the `trait` keyword. However, if a class has abstract members, it must be declared abstract.

Now, we can define a simple `Button` type:

```
// src/main/scala/progscala3/traits/ui/Button.scala
package progscala3.traits.ui

abstract class Button(val label: String) extends Widget:

  def click(): Unit = updateUI()

  def updateUI(): Unit
```

`Button` is considerably simpler. It has only one concern, handling clicks. At the moment, it seems trivial now to have a public method `click` that does nothing except delegate to a protected method `updateUI`, but we'll see next why this is still useful.

Now we construct `ObservableButton`, which subclasses `Button` and mixes in `Subject`:

```
// src/main/scala/progscala3/traits/ui/ObservableButton.scala
package progscala3.traits.ui
import progscala3.traits.observer._

abstract class ObservableButton(name: String)
❶
  extends Button(name) with Subject[Button]:
❷

  override def click(): Unit =
❸
    super.click()
    notifyObservers(this)
❹
❺
```

- ❶ A subclass of Button that mixes in observability.
- ❷ Extends Button and mixes in Subject and uses Button as the Subject type parameter, named State in the declaration of Subject.
- ❸ In order to notify observers, we have to override the click method.
- ❹ First, call the parent class click to perform the normal GUI update logic.
- ❺ Notify the observers, passing this as the State. In this case, there isn't any state other than the "event" that a button click occurred.

So, we modified click, but kept it simple for future subclasses to correctly implement updateUI, without having to worry about correct handling of the observer logic. This is why we kept both click and updateUI in Button.

Let's try it. First, let's define an observer to count button clicks:

```
// src/main/scala/progscala3/traits/ui/ButtonCountObserver.scala
package progscala3.traits.ui
import progscala3.traits.observer._

class ButtonCountObserver extends Observer[Button]:
  var count = 0
  def receiveUpdate(state: Button): Unit = count += 1
```

Now try it:

```
// src/script/scala/progscala3/traits/ui/ButtonCountObserver1.scala
import progscala3.traits.ui._
import progscala3.traits.observer._

val button = new ObservableButton("Click Me!"):
  def updateUI(): Unit = println(s"$label clicked")

val bco1 = new ButtonCountObserver
val bco2 = new ButtonCountObserver

button.addObserver(bco1)
button.addObserver(bco2)

(1 to 5) foreach (_ => button.click())

assert(bco1.count == 5)
assert(bco2.count == 5)
```

The script declares an observer type, `ButtonCountObserver`, that counts clicks. Then it creates an anonymous subclass of `ObservableButton` with a definition of `updateUI`. Next it creates and registers two observers with the button, clicks the button five times and then verifies the counts for each observer equals five. You'll also see the string `Click Me! clicked` printed five times.

Suppose we only need one instance of an `ObservableButton`. We don't need to declare a class that mixes in the traits we want. Instead, we can declare a `Button` and mix in the `Subject` trait in one step:

```
// src/script/scala/progscala3/traits/ui/ButtonCountObserver2.scala
import progscala3.traits.ui._
import progscala3.traits.observer._

// tag::button
val button = new Button("Click Me!") with Subject[Button]:
  override def click(): Unit =
    super.click()
    notifyObservers(this)
  def updateUI(): Unit = println(s"$label clicked")
// end::button

val bco1 = new ButtonCountObserver

button.addObserver(bco1)

(1 to 5) foreach (_ => button.click())

assert(bco1.count == 5)
```

## NOTE

When declaring a class that only mixes in traits and doesn't extend another class, you must use the `extends` keyword anyway for the first trait listed and the `with` keyword for the rest of the traits. However, when instantiating a class and mixing in traits with the declaration, use the `with` keyword with all the traits.

In “[Good Object-Oriented Design: A Digression](#)”, I recommended that you avoid method overrides. We didn't have a choice here, but we proceeded carefully. In the `click` override, we added invocation of the subject-observer logic, but didn't modify the UI update logic. Instead, we left in place the protected method `updateUI` to keep that logic separated from observation logic. Our click logic is an example of the *Template Method Pattern* described *earlier*.

## Stackable Traits



There are several, further refinements we can do to improve the reusability of our code and to make it easier to use more than one trait at a time, i.e., to “stack” traits.

First, clicking is not limited to buttons in a GUI. We should make that logic abstract, too. We could put it in `Widget`, the so-far empty parent type of `Button`, but it may not be true that all GUI widgets accept clicks. Instead, let’s introduce another trait, `Clickable`:

```
// src/main/scala/progscala3/traits/ui2/Clickable.scala
package progscala3.traits.ui2
❶

trait Clickable:
❷
  def click(): String = updateUI()
  protected def updateUI(): String
```

- ❶ Use a new package because we’re reimplementing types in `traits.ui`.
- ❷ Essentially just like the previous `Button` definition, except now we return a `String` from `click`.

Having `click` return a `String` is useful for discussing the stacking of method calls.

Here is the refactored button, which uses the trait:

```
// src/main/scala/progscala3/traits/ui2/Button.scala
package progscala3.traits.ui2
import progscala3.traits.ui.Widget

abstract class Button(val label: String) extends Widget with
  Clickable
```

It is still abstract, little more than a convenient name for mixed-together types!

Observation should now be tied to `Clickable` and not `Button`, as it was before. When we refactor the code this way, it becomes clear that we don't really care about observing buttons. We really care about observing *events*, such as clicks. Here is a trait that focuses solely on observing `Clickable`:

```
// src/main/scala/progscala3/traits/ui2/ObservableClicks.scala
package progscala3.traits.ui2
import progscala3.traits.observer._

trait ObservableClicks extends Clickable with Subject[Clickable]:
  abstract override def click(): String =
    ❶    val result = super.click()
        notifyObservers(this)
        result
```

- ❶ Note the `abstract override` keyword combination, discussed next.

The implementation is very similar to the previous `ObservableButton` example. The important difference is the `abstract` keyword. We had `override` before.

Look closely at this method. It calls `super.click()`, but what is `super` in this case? At this point, it could only appear to be `Clickable`, which *declares* but does not *define* the `click` method, or it could be `Subject`, which doesn't have a `click` method. So, `super` can't be bound to a real instance, at least not yet. This is why `abstract` is required here.

In fact, `super` will be bound when this trait is mixed into a concrete instance that defines the `click` method, such as `Button`. The `abstract` keyword tells the compiler (and the reader) that `click` is not yet fully implemented, even though `ObservableClicks.click` has a body.

## NOTE

The `abstract` keyword is only required on a method in a trait when the method *has* a body, but it invokes another method in super that doesn't have a concrete implementation in parents of the trait.

Let's use this trait with `Button` and its concrete `click` method. First we'll define an observer class:

```
// src/main/scala/progscala3/traits/ui2/ClickCountObserver.scala
package progscala3.traits.ui2
import progscala3.traits.observer._

class ClickCountObserver extends Observer[Clickable]:
  var count = 0
  def receiveUpdate(state: Clickable): Unit = count += 1
```

Now use it:

```
// src/script/scala/progscala3/traits/ui2/ClickCountObserver.scala
import progscala3.traits.ui2._
import progscala3.traits.observer._

// No override of "click" in Button required.
val button = new Button("Button") with ObservableClicks:
  def updateUI(): String = s"$label clicked"

val cco = new ClickCountObserver
button addObserver cco

(1 to 5) foreach (_ => assert("Button clicked" == button.click()))
assert(cco.clickCount == 5)
```

Note that we can now declare a `Button` instance and mix in `ObservableClicks` without having to override the `click` method ourselves. We have also gained a reusable mixin for click handling, `Clickable`.

Let's finish our example by adding a second trait, where an observer can veto a click after a certain number have been received:

```
// src/main/scala/progscala3/traits/ui2/VetoableClicks.scala
package progscala3.traits.ui2

trait VetoableClicks(val maxAllowed: Int = 1) extends Clickable:
  ❶ private var count = 0
  ❷

  abstract override def click(): String =
    count += 1
    if count <= maxAllowed then
      super.click()
    else
      s"Max allowed clicks $maxAllowed exceeded. Received $count
      clicks!"
    ❸

  def resetCount(): Unit = count = 0
    ❹
```

- ❶ Also extends Clickable.
- ❷ Use a private variable to avoid collision with any other fields named count.
- ❸ The maximum number of allowed clicks. (A “reset” feature would be useful.)
- ❹ Once the number of clicks exceeds the allowed value (counting from zero), no further clicks are sent to super.

Note that `maxAllowed` is declared a `val` and the comment says it is the “default” value, which implies it can be changed. How can that be if it's a `val`? The answer is that we can override the value in a class or trait that mixes in this trait.

Note that this count should be different from the count used in `ClickCountObserver`. Keeping this one private prevents confusion.

The compiler will flag collisions.

In this use of both traits, we limit the number of handled clicks to 2:

```
//  
src/script/scala/progscala3/traits/ui2/VetoableClickCountObserver.scala  
a  
import progscala3.traits.ui2._  
import progscala3.traits.observer._  
  
val button = new Button("Button!")  
  with ObservableClicks with VetoableClicks(maxAllowed = 2):  
  def updateUI(): String = s"$label clicked"  
  
val cco = new ClickCountObserver  
button addObserver cco  
  
(1 to 5) map (_ => button.click())  
assert(cco.count == 2)
```

The map with calls to click returns the following sequence of strings:

```
Vector("Button! clicked", "Button! clicked",  
  "Max allowed clicks 2 exceeded. Received 3 clicks!",  
  "Max allowed clicks 2 exceeded. Received 4 clicks!",  
  "Max allowed clicks 2 exceeded. Received 5 clicks!")
```

Try this experiment. Switch the order of the traits in the declaration of button to this:

```
val button = new Button("Click Me!")  
  with VetoableClicks(maxAllowed = 2) with ObservableClicks:  
  def updateUI(): String = s"$label clicked"
```

What happens when you run this code now? The strings returned will be the same, but the assertion that `cco.count == 2` will now fail. The count is actually 5.

We have three versions of `click` wrapped like an onion. The question is which version of `click` gets called first when we mix in

VetoableClicks and ObservableClicks? The answer is determined by the declaration order, from *right* to *left*.

This means that ObservableClicks will be notified *before* VetoableClicks has the chance to prevent calling up the chain, `super.click()`. Hence, declaration order *matters*.

An algorithm called *linearization* is used to resolve the priority of traits and classes in the inheritance tree when resolving which overridden method to call for `super.method()`. However, it can get confusing quickly, so avoid complicated mixin structures!

This fine-grained composition through mixin traits is quite powerful, but it can be overused:

- It can be difficult to understand code if many mixin traits are used.
- Resolving method overrides in the correct order (when they are required to use the traits) can get confusing quickly.
- Lots of traits can slow down compile times.

We'll cover the full details of how linearization works for more complex declarations in "[Linearization of an Object's Hierarchy](#)".

## Union and Intersection Types

Now is a good time to revisit *union types* and introduce *intersection types*, both of which are new in Scala 3.

We encountered union types in “When You Can’t Avoid Nulls”, where we discussed using `T | Null` as a type declaration for the value returned by calling a method that will return *either* a value of type `T` *or* `null`.

Here’s another example, where `Int | String` is used as an alternative to `Either[String, Int]` (note the different order) as a conventional way to return *either* a success (`Int`) or a failure described by a `String`.

```
// src/script/scala/progscala3/traits/UnionTypes.scala
scala> def process(i: Int): Int | String =
  |   if (i < 0) then "Negative integer!" else i

  |   val result1: Int | String = process(-1)
  |   val result2: Int | String = process(1)
val result1: Int | String = Negative integer!
val result2: Int | String = 1

scala> val result1b = process(-1)
  |   val result2b = process(1)
val result1b: Any = Negative integer!
val result2b: Any = 1

scala> Seq(process(-1), process(1)).map {
  |   case i: Int => "integer"
  |   case s: String => "string"
  | }
val res1: Seq[String] = List(string, integer)
```

Note the types for the `resultN` values, especially if we leave off explicit type annotations. The pattern match correctly matches on the types.

These types are *unions* in the sense that any `Int` *or* `String` value can be assigned to `result1`, for example. That is, the type is the

union of the set of Int and String values.

I didn't mention it in the previous sections, but perhaps you noticed the type printed for this declaration (removing the package prefixes for clarity):

```
scala> val button = new Button("Button!")
      |   with ObservableClicks with VetoableClicks(maxAllowed = 2):
      |   def updateUI(): String = s"$label clicked"
val button: Button & ObservableClicks & VetoableClicks = ...
```

The type `Button & ObservableClicks & VetoableClicks` is an *intersection type*, which results when we construct instances with mixins. In Scala 2, the returned type signature would use the same `with` and `extends` keywords as the definition, `Button with ObservableClicks with VetoableClicks`. It's perhaps a little confusing that we get back something different. Unfortunately, you can't use `&` instead of `with` or `extends` in the definition. However, we can use `&` in a type annotation:

```
scala> val button2: Button & ObservableClicks & VetoableClicks =
      |   new Button("Button!")
      |   with ObservableClicks with VetoableClicks(maxAllowed = 2):
      |   def updateUI(): String = s"$label clicked"
val button2: Button & ObservableClicks & VetoableClicks = ...
```

These types are *intersections* in the sense that the only allowed values that you can assign to `button2` are those that belong to `Button` *and* `ObservableClicks` *and* `VetoableClicks`, which won't include values of each of those types separately. Note this compilation error (some output omitted again):

[illegible]



```
|Required: Button & ObservableClicks & VetoableClicks  
|)
```

Most of the time, this won't be an issue, but imagine a scenario where you have a var for a button that you want to observe now, but eventually replace with an instance without observation:

```
scala> var button4 = new Button("Button!")  
|   with ObservableClicks with VetoableClicks(maxAllowed = 2):  
|   def updateUI(): String = s"$label clicked"  
  
scala> // later...  
scala> button4 = new Button("New Button!"):  
|   def updateUI(): String = s"$label clicked"  
2 |   def updateUI(): String = s"$label clicked"  
|                                     ^  
| Found:    Button {...}  
| Required: Button & ObservableClicks & VetoableClicks
```

Not convenient. One easy workaround is to use a type annotation, `var button5: Button = ...`. Then the reassignment will work.

But wait, isn't that a type error, based on how I just described intersection types? It's not, because the type restriction declared for `button5` allows any `Button`, but the initial assignment happens to be a more restricted instance, a `Button` that also mixes in other types.

Scala 3 offers a new option, discussed next.

## Super Traits

Note that `ObservableClicks` and `VetoableClicks` are implementation details and not really part of the core domain types of a user interface. An old example from Scala 2 is the way that several common traits are automatically added as parents for case classes and objects. Here's an example:

```
// Scala 2 REPL!  
scala> trait Base
```

```
scala> case object Obj1 extends Base
scala> case object Obj2 extends Base
scala> val condition = true
scala> val x = if (condition) Obj1 else Obj2
x: Product with Serializable with Base = Obj1
```

The inferred type is the parent Base, but with `scala.Product` and `java.lang.Serializable`. Usually, we only care about Base.

Scala 3 lets you add a keyword to the declaration of a trait so that it isn't part of the inferred type:

```
scala>
| class Text(val label: String, value: String) extends Widget
|
| trait Flag
| val t1 = new Text("foo", "bar") with Flag
val t1: Text & Flag = anon$1@28f8a295

scala> super trait SFlag
| val t2 = new Text("foo", "bar") with SFlag
val t2: Text = anon$2@4a2219d
```

Note the different inferred types.

Common traits, like `scala.Product`, `java.lang.Serializable`, and `java.lang.Comparable` are now treated as “super”. In the Scala 3 REPL, the Base example above has the inferred type Base without Product and Serializable.

If you need to cross compile to Scala 2 and 3 for a while, use the annotation `scala.annotation.superTrait` instead of the super keyword.

## Using Commas Instead of With

Scala 3 also allows you to substitute a comma (“,”) instead of with, but *only when declaring a type*:

```

scala> class B extends Button("Button!"),
      |   ObservableClicks, VetoableClicks(maxAllowed = 2):
      |   def updateUI(): String = s"$label clicked"

scala> var button4b: Button = new Button("Button!"),
      |   ObservableClicks, VetoableClicks(maxAllowed = 2):
      |   def updateUI(): String = s"$label clicked"
1 | var button4b: Button = new Button("Button!"), ObservableClicks,
  |   ...
  |   ^
  |   end of statement expected but
  |   ',' found
1 | ...

```

If this substitution worked for all situations where with is used, it would be useful, but it will be confusing to remember when it's allowed and when it isn't.

## Trait Parameters

Scala 3 allows traits to have parameters, whereas in Scala 2, you had to declare fields in the trait body.

Here is a logging abstraction that uses a trait parameter for the level:

```

// src/main/scala/progscala3/traits/Logging.scala
package progscala3.traits.logging

enum LoggingLevel:
  case DEBUG, INFO, WARN, ERROR, FATAL
  ❶

trait Logger(val level: LoggingLevel):
  ❷
  def log(message: String): Unit

trait ConsoleLogger extends Logger:
  ❸
  def log(message: String): Unit = println(s"$level: $message")

```

```
class Service(val name: String, level: LoggingLevel)
  ④ extends ConsoleLogger with Logger(level)
```

- ❶ Define logging levels.
- ❷ The `level` is a trait parameter. This will also be a field in concrete types that use this trait.
- ❸ Use the console.
- ❹ A concrete class or one of its parents must pass a `LoggingLevel` value to `Logger`.

Note that we have to mix in *both* `Logging` and `ConsoleLogger`, even though the latter extends the former, because we must specify the `level` parameter for `Logger` explicitly. `ConsoleLogger` can't be declared as follows, like we might see in a hierarchy of classes:

```
trait ConsoleLogger(level: LoggingLevel) extends Logger(level) //
ERROR
```

Note that the `name` argument for `Service` is declared to be a field (with `val`), but if we try declaring `level` as a field, it will conflict with the definition already provided by `Logger`. (Try it!) However, as shown, we can use the same name for a non-field parameter.

Here is the Scala 2-compatible approach without trait parameters:

```
// src/main/scala/progscala3/traits/LoggingNoParameters.scala
package progscala3.traits.logging

trait LoggerNP:
  ❶ def level: LoggingLevel
    def log(message: String): Unit
```

```

trait ConsoleLoggerNP extends LoggerNP:
  def log(message: String): Unit = println(s"$level: $message")

class ServiceNP(val name: String, val level: LoggingLevel)
  ❷
  extends ConsoleLoggerNP

```

- ❶ “NP” for “no parameters”. Note that the abstract `level` is now a method. It could be an abstract field, but using a method provides more flexibility for the implementer (see “[The Uniform Access Principle](#)”).
- ❷ The implementation of `level` is a `ServiceNP` constructor parameter.

The same `LoggingLevel` enumeration is used here.

It’s convenient that we can now declare parameters for traits, but it comes with limitations. If a hierarchy of classes mix in a parameterized trait, only one can pass arguments to it. Also, if `Service` were declared as a case class, the `level` argument would now be a field, which would conflict with `level` defined by `Logging`. (Try adding `case` and recompile.) We would need to use a different name, adding a redundant field!

However, to be fair, types like this that are composed of other types aren’t good candidates for case classes. For example, if two service instances differ only by the logging level, should they still be considered equivalent?

On balance, I think parameterized traits should be used relatively rarely, for simple scenarios. They fix some scenarios with the order of initialization. Scala 2 had a feature to handle these cases called *early initializers*. This feature was dropped in Scala 3, because trait parameters now address these scenarios. [Chapter 12](#) discusses construction of compound types with mixins and a hierarchy of parents.

The idioms demonstrated in the LoggingNP example are concise and elegant. In particular, replacing trait declarations of fields with methods provides more flexibility for derived types, although with the cost that they have to implement the method and understand exactly how it is used in the parent types.

## Should That Type Be a Class or Trait?

When considering whether a type should be a trait or a class, keep in mind that traits are best for pure interfaces and when used as mixins for “adjunct” behavior. If you find that a particular trait is used most often as a parent of other classes, so that the child classes *behave as* the parent trait, then consider defining the type as a class instead to make this logical relationship more clear. I said *behaves as*, rather than *is a*, because the former is the more precise definition of how inheritance needs to work, based on the *Liskov Substitution Principle*.

## Recap and What's Next

In this chapter, we learned how to use traits to encapsulate and share cross-cutting concerns between classes. We covered when and how to use traits, how to “stack” multiple traits, and the rules for initializing values within traits.

In the next few chapters, we explore Scala’s object system and class hierarchy, with particular attention to the collections. We also revisit construction of complex types, like our stacked traits example above.

# Chapter 11. Variance Behavior and Equality

---

## A NOTE FOR EARLY RELEASE READERS

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

This will be the 11th chapter of the final book. Please note that the GitHub repo will be made active later on.

If you have comments about how we might improve the content and/or examples in this book, or if you notice missing material within this chapter, please reach out to the author at [mcronin@oreilly.com](mailto:mcronin@oreilly.com).

An important concept in object-oriented type systems goes by the name *variance under inheritance*. More generally, we need well-defined rules for when an instance of one type can be *substituted* for an instance of another type. This chapter begins with an exploration of these concepts.

A logical follow-on is the subject of instance equality, which is trickier than it first appears in object-oriented languages.

## Parameterized Types: Variance Under Inheritance

## NOTE

The content of this section is somewhat challenging. It will take some effort to master the arguments, but mastery isn't required to use Scala effectively. Consider skimming this part, but definitely read the equality discussion that follows.

Suppose a method takes a parameter of type `Seq[AnyRef]`. Are you allowed to pass a `Seq[String]` value? In other words, is `Seq[String]` considered *substitutable* for `Seq[AnyRef]`? The *Liskov Substitution Principle* was the first to define formally what this means. In object oriented terms, we usually say that instances of one type are substitutable for instances of another type if the former type is a *subtype* of the latter. We'll discuss the defined hierarchy of types in the Scala library in [Chapter 13](#), but we've already learned that `AnyRef` is a parent type of all *reference* types, like `String`. Therefore, instances of `String` are substitutable where instances of `AnyRef` are required. More concisely, we'll say that `String` is substitutable for `AnyRef`.

So what about parameterized types, such as collections like `Seq[AnyRef]` and `Seq[String]`? The type parameter `A` is declared like this: `Seq[+A]`, where `+A` means that `Seq` is *covariant* in the `A` parameter. Since `String` is substitutable for `AnyRef`, then `Seq[String]` is substitutable for a `Seq[AnyRef]`. *Covariance* means the supertype-subtype relationship of the container (the parameterized type) "goes in the same direction" as the relationship between the type parameters.

We can also have types that are *contravariant*. A declaration `X[-A]` means that `X[String]` is a *supertype* of `X[Any]`. The substitutability goes in the opposite direction of the type parameter values. This behavior is less intuitive, but we'll study an example shortly.

If a parameterized type is neither covariant nor contravariant, it is called *invariant*.



Finally, some parameterized types can mix two or more of these behaviors.

We first encountered these concepts in “Abstract Types Versus Parameterized Types”. Now we’ll explore them in more depth.

NOTE

If you are familiar with Java parameterized types, called generics, you do specify variance behavior when they are defined. Instead, the variance behavior is specified when the type is *used*, i.e., at the *call site*, when variables are declared. Scala makes it the responsibility of the type designer to specify the correct behavior, rather than users.

The three kinds of variance notations and their meanings are summarized in Table 11-1.  $T^{\text{sup}}$  is a *supertype* of  $T$  and  $T_{\text{sub}}$  is a *subtype* of  $T$ . Java syntax is shown for comparison.

Table 11-1. Type variance annotations and their meanings

Scala	Java	Description
+T	? extends T	<i>Covariant</i> (e.g., Seq[T <sub>sub</sub> ] is a subtype of Seq[T]).
-T	? super T	<i>Contravariant</i> (e.g., X[T <sup>sup</sup> ] is a subtype of X[T]).
T	T	<i>Invariant</i> subclassing (e.g., can’t substitute Y[T <sup>sup</sup> ] or Y[T <sub>sub</sub> ] for Y[T]).

Seq is declared Seq[+A], which means that Seq[String] is a subclass of Seq[AnyRef], so Seq is covariant in the type parameter A. When a type like Seq has only one covariant type parameter, you’ll often hear the shorthand expression “Seqs are covariant” and similarly for contravariant types.

Covariant and invariant types are reasonably easy to understand. What about contravariant types?

## Functions Under the Hood

Let's dive into a functions a bit more, then explore how they provide a good example of contravariant behavior.

In Scala 2, functions and tuples were limited to 22 parameters. This sounds like a lot, but in *big data* scenarios, it's easy to have far more values in a record, making it impossible to use tuples for them or to pass the values as separate function parameters. Scala 3 eliminates this restriction.

As in Scala 2, a function literal with two arguments implements the trait `Function2[-T1, -T2, +R]`, where the two input parameters have types `T1` and `T2` and the return type is `R`. Note that input parameter types are contravariant. We'll come back to that.

There are corresponding `FunctionN` types for `N` between 0 and 22, inclusive. Beyond 22, the function is instantiated with `scala.FunctionXXL`, which encodes the parameters in an array.

We've been using anonymous functions, also known as function literals, throughout the book. For example:

```
scala> Seq(1, 2, 3, 4).foldLeft(0)((result,i) => result+i)
val res0: Int = 10
```

The function expression `(result,i) => result+i` is actually *syntactic sugar* that the compiler converts to the following instantiation of an anonymous subclass of `Function2`:

```
scala> val f: (Int,Int) => Int = new Function2[Int,Int,Int]:
    |   def apply(i: Int, j: Int): Int = i + j
    |
val f: (Int, Int) => Int = <function2>
```

```
scala> Seq(1, 2, 3, 4).foldLeft(0)(f)
val res1: Int = 10
```

The literal type signature syntax `(Int,Int) => Int` can also be used.

### TIP

When compiling for the JVM, Java *lambdas* are generated for the byte code, which also means that lambdas can be passed as arguments to Scala methods expecting function arguments. This promotes better interoperability between Java and Scala code.

We learned before that when an instance used like a function, i.e., it is passed an argument list, the compiler searches for an `apply` method to call. The name `apply` is conventional. It originated with the idea of *function application*. For example, once `f` is defined, we call it by *applying* a parameter list to it. Therefore `f(1,2)` is actually `f.apply(1,2)`.

For completeness, tuples with 0 to 22 elements are represented by `TupleN` final case classes, for `N` between 0 and 22, inclusive. For example, `scala.Tuple2[+T1,+T2]`. Beyond 22 elements, `scala.TupleXXL` is used, which also uses an array for the elements. Also, there are some new operations on tuples for concatenation and indexing, analogous to similar operations on sequential collections.

Now let's return to contravariance. The best example of it is the types of the parameters passed to the `FunctionNs` `apply` methods. We'll use `scala.Function1[-T,+R]`. The same arguments apply for the other functions with more parameters.

The last type parameter, `+R`, is the return type. It is *covariant*. The type parameter for the function parameter, `-T`, is *contravariant*. Therefore, functions have mixed variance behavior under inheritance.

Let's look at an example to understand what the variance behavior really means. It will also help us think about what *substitutability* really means, in general:

```
//
src/script/scala/progscala3/objectsystem/variance/FunctionVariance.sca
la

class CSuper                                     ❶
class C      extends CSuper
class CSub   extends C

val f1: C => C = (c: C)      => new C              ❷
val f2: C => C = (c: CSuper) => new CSub
val f3: C => C = (c: CSuper) => new C
val f4: C => C = (c: C)      => new CSub
// Compilation errors!
// val f5: C => C = (c: CSub)  => new CSuper        ❸
// val f6: C => C = (c: CSub)  => new C
// val f7: C => C = (c: C)     => new CSuper
```

- ❶ Define a three-type inheritance hierarchy.
- ❷ Four different, valid assignments for functions of the same type `C => C`. Note that `C` is the middle type in the three-type hierarchy.
- ❸ Three invalid assignments, which would trigger compilation errors.

All valid function instances must be `C => C` (in other words, `Function1[C,C]`) The values we assign must satisfy the constraints of variance under inheritance for functions.

The assignment for `f1` is `(c: C) => new C`, which matches the types exactly. It ignores the parameter `c`.

The assignment for `f2` is `(c: CSuper) => new CSub`. It is valid, because the function parameter `C` is contravariant, so `CSuper` is a

valid substitution, while the return value is covariant, so CSub is a valid replacement for C.

The assignments for f3 and f4 are like f2, but we just use C for the return and the parameter, respectively.

Similarly, f5 through f7 are analogous to f2 through f4, but using invalid substitutions. So, let's just discuss f5, where both type substitutions fail. Using CSub for the parameter type is invalid, because the parameter must be C or a parent type (covariance). Using CSuper for the return type is invalid because C or a subtype is required.

Let's try to understand intuitively why these variance behaviors are correct.

The key insight is to know that the function type signature  $C \Rightarrow C$  is like a *contract*. It says I'm expecting you to give me a function that will accept any valid C instance and it will promise to return any valid C instance.

So, consider the return type covariance first. If you pass me a function of type  $C \Rightarrow \text{CSub}$ , which always returns an instance of the subtype CSub, that satisfies the contract, because a CSub instance is always substitutable for a C instance. I can accept any C instance your function returns, so I can easily accept the fact you only ever return the *subset* of instances of type CSub. In this sense, your function is more "restrictive" than it needs to be for return values, but that's okay.

In contrast, if you pass a function  $\text{CSuper} \Rightarrow C$ , your function is more "permissive". I will only pass C instances to your function, because that's what I said I will do with the  $\text{fN}: C \Rightarrow C$  declarations. However, your function is able to handle the wider set of all CSuper instances, including instances of C and even CSub. So, if you give me a  $\text{CSuper} \Rightarrow C$  function, it will also satisfy the contract.

I said that f5 breaks the contract for both types. Let's consider what would happen if I allowed this substitution.

In this case, the function you give me would only know how to handle CSub instances passed to it, but I said I want to pass C instances to your function. Hence your function would be “surprised” when it has to handle a C instance. Similarly, if your function returns a CSuper, I will be “surprised”, because I only know how to work with C return values.

This is why function parameter must be contravariant and return values must be covariant.

### TIP

When thinking about *variance under inheritance*, you should really ask yourself what *substitutions* are valid in a particular situation.

Variance annotations only make sense on the type parameters for *types*, not for parameterized methods, because the annotations affect the behavior of subtyping. Methods aren't subtyped themselves. For example, signature for the Seq.map method looks conceptually like this:

```
abstract class Seq[+A](...):  
  ...  
  def map[B](f: A => B): Seq[B] = ...
```

There is no variance annotation on B and if you tried to add one, the compiler would throw an error.

## NOTE

The + *variance annotation* means the parameterized type is *covariant* in that type parameter. The - variance annotation means the parameterized type is *contravariant* in that type parameter. No variance annotation means the parameterized type is *invariant* in the type parameter.

Finally, the compiler checks your use of variance annotations for invalid uses. Here's what happens if you attempt to define your own function with the wrong annotations:

```
scala> class MyFunction2[+T1, +T2, -R]:  
      |   def apply(v1:T1, v2:T2): R = ???  
2 |   def apply(v1:T1, v2:T2): R = ???  
    | ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^  
    |contravariant type R occurs in covariant position in type  
    | (v1: T1, v2: T2): R of method apply
```

If we change  $-R$  to  $-R$ , we see the next errors:

```
scala> class MyFunction2[+T1, +T2, +R]:
      |   def apply(v1:T1, v2:T2): R = ???
2 |   def apply(v1:T1, v2:T2): R = ???
      |           ^^^^^
      | covariant type T1 occurs in contravariant position in type T1 of
value v1
...

```

## Variance of Mutable Types

The parameterized types we've discussed so far have been for immutable fields. What about the variance behavior of types for mutable fields? The short answer is that only *invariance* is allowed. Consider the following class definitions:

```
// tag::first[]  
//  
src/script/scala/progscala3/objectsystem/variance/MutableVariance.scala
```

*a*

```
class Invariant[A](var mut: A)
class Covariant[+A](var mut: A)           // ERROR
class Contravariant[-A](var mut: A)       // ERROR
// end::first[]

// tag::second[]
class Invariant[A](val mutInit: A):
  private var _mut: A = mutInit
  def mut_=(a: A): Unit = _mut = a
  def mut: A = _mut
// end::second[]

// tag::third[]
class Covariant[+A](val mutInit: A):
  private var _mut: A = mutInit
  def mut_=(a: A): Unit = _mut = a
  def mut: A = _mut

class Contravariant[-A](val mutInit: A):
  private var _mut: A = mutInit
  def mut_=(a: A): Unit = _mut = a
  def mut: A = _mut
// end::third[]
```

Only the invariant definition compiles. When `mut` is declared as a mutable field, it behaves like a private field with public read *and* write methods, each of which limits the allowed variance. Here are logically equivalent definitions of the second and third classes, this time with REPL error messages:

```
scala> class Covariant[+A](val mutInit: A):
  |   private var _mut: A = mutInit
  |   def mut_=(a: A): Unit = _mut = a
  |   def mut: A = _mut
  |
3 |   def mut_=(a: A): Unit = _mut = a
  |               ^^^^
  |               covariant type A occurs in contravariant position in type A of
value a

scala> class Contravariant[-A](val mutInit: A):
```



```

|     private var _mut: A = mutInit
|     def mut_(a: A): Unit = _mut = a
|     def mut: A = _mut
|
1 | class Contravariant[-A](val mutInit: A):
|                               ^^^^^^^^^^^^^^^^^
|     contravariant type A occurs in covariant position in type A of
value mutInit
4 |     def mut: A = _mut
|       ^^^^^^^^^^^^^^^^^
|     contravariant type A occurs in covariant position in type => A of
method mut

```

Recall from our discussion above about function variance that the types for function parameters have to be contravariant. That's why we get an error for `Covariant.mut_`. We are attempting to use a covariant type `A` in a parameter list, which requires a contravariant type.

Similarly, we learned above that function return types need to be covariant, yet `Contravariant.mut_` attempts to return a value of a contravariant type `A`.

As for functions, valid *substitution* is the requirement. Let's pretend these classes compile and reason about what would happen if we used them. We'll work with the same `CSuper`, `C`, and `CSub` used previously. First, let's check `Invariant`:

```

val obj: Invariant[C] = new Invariant(new C)
val c: C = obj.mut      // Type checks correctly
obj.mut = new C         // Type checks correctly

```

Now try `Covariant`. Recall that we are attempting to prove that `Covariant[CSub]` is substitutable for `Covariant[C]`, since `CSub` is substitutable for `C`:

```

val obj: Covariant[C] = new Covariant(new CSub) // Okay??
val c: C = obj.mut      // Type checks correctly, because CSub <: C
obj.mut = new C         // ERROR, because obj expects CSub!!

```

Recall that  $A <: B$  means A is a subtype of B.

The actual type of `obj` is `Covariant[CSub]`. When we read `obj.mut`, we get a `CSub` instance, but that's okay because `CSub` is a subclass of `C`. However, if we attempt to assign a `C` instance to `obj.mut`, we get an error, because a `C` is not a subtype of `CSub`, it is a supertype. Hence, a `C` value isn't substitutable for a `CSub`. This is consistent with the error message we got when we attempted to compile `Covariant` above.

Finally, let's try `Contravariant`. Recall that we are attempting to prove that `Contravariant[CSuper]` is substitutable for `Covariant[C]`, since `C` is substitutable for `CSuper`:

```
val obj: Contravariant[C] = new Contravariant(new CSuper) // Okay??
val c:C = obj.mut      // ERROR, because c:C, but we return a
                        CSuper!
obj.mut = new C         // Type checks correctly, because C <:
                        CSuper
```

Here, we can assign a `C` to `obj.mut`, because of contravariance, but if we read `obj.mut`, we get an instance of `CSuper`, a superclass of `C`, and therefore not substitutable for a `C`.

Hence, if you think of a mutable field in terms of the corresponding getter and setter methods, it appears in both covariant position when read and contravariant position when written. There is no such thing as a type parameter that is *both* contravariant and covariant, so invariance is the only option for the type of a mutable field.

## Variance in Scala Versus Java

While the variance behavior is defined in the *declaration* for Scala types, in Java it is defined when used, at the *call site*. The *client* of a type defines the variance behavior, defaulting to invariant. Java doesn't allow you to specify variance behavior at the definition site, although you can use expressions that look similar. Those

expressions define *type bounds*, which are written as  $A <: B$  in Scala.

There are two drawbacks of Java's call-site variance specifications. First, it should be the library designer's job to understand the correct variance behavior and encode that behavior in the library itself. Instead, it's the user who bears this burden. This leads to the second drawback. It's easy for a Java user to apply an incorrect annotation that results in unsafe code, like in the scenarios we just discussed.

Another problem in Java's type system is that Arrays are covariant in the type  $T$ , but Arrays are *mutable*. Consider this example:

```
// src/main/java/progscala3/objectsystem/JavaArrays.java
package progscala3.objectsystem;

public class JavaArrays {
    public static void main(String[] args) {
        Integer[] array1 = new Integer[] {
            Integer.valueOf(1), Integer.valueOf(2), Integer.valueOf(3) };
        Number[] array2 = array1;           // Compiles fine, but
        // shouldn't!!
        array2[2] = Double.valueOf(3.14);  // Compiles, but throws a
        // runtime error!
    }
}
```

This file compiles without error. However, when you run it with SBT, hilarity ensues:

```
> runMain progscala3.objectsystem.JavaArrays
[info] Running progscala3.objectsystem.JavaArrays
[error] (runMain-4) java.lang.ArrayStoreException: java.lang.Double
java.lang.ArrayStoreException: java.lang.Double
    at progscala3.objectsystem.JavaArrays.main(JavaArrays.java:10)
...
```

Because Java arrays are covariant, we're allowed by the compiler to assign an `Array[Integer]` instance to an `Array[Number]` reference.

Then the compiler thinks it's OK to assign any `Number` to elements of the array, but in fact, the array “knows” internally that it can only accept `Integer` values, so it throws a runtime exception. This defeats the purpose of static type checking.

Even though Scala wraps Java Arrays, the Scala class `scala.Array` is invariant in the type parameter, so the equivalent Scala program won't compile.

See *Naftalin2006* for more details of Java's generics and arrays, from which this example was adapted.

## Equality of Instances

Implementing a reliable equality test for instances is difficult to do correctly. *Bloch2008* and the *Scaladoc* page for `AnyRef.eq` describe the requirements for a good equality test. *Odersky2009* is a very good article on writing `equals` and `hashCode` methods correctly.

Since these methods are created automatically for case classes, tuples, and enumerations, I *never* write my own `equals` and `hashCode` methods anymore.

### TIP

Any objects that might be tested for equality or used as keys in a `Set` or a `Map` (where `hashCode` is used) *should* be instances of case classes, tuples, or enumerations.

Let's explore equality of instances in Scala, which can be tricky because of inheritance.

## CAUTION

Some of the equality methods have the same names as equality methods in other languages, but the semantics are sometimes different!

## The equals Method

We'll use a case class to demonstrate how the different equality methods work:

```
// tag::definitions[]
// src/script/scala/progscala3/objectsystem/equality/Equality.scala

case class Person(firstName: String, lastName: String, age: Int)

val p1a = Person("Dean", "Wampler", 29)
val p1b = Person("Dean", "Wampler", 29)
val p2  = Person("Buck", "Trends", 30)
// end::definitions[]

// tag::equals1[]
assert((p1a equals p1a) == true)
assert((p1a equals p1b) == true)
assert((p1a equals p2)  == false)
assert((p1a equals null) == false)
// end::equals1[]

// tag::equals2[]
assert((p1a == p1a) == true)
assert((p1a == p1b) == true)
assert((p1a == p2)  == false)
assert((p1a == null) == false)

assert((p1a != p1a) == false)
assert((p1a != p1b) == false)
assert((p1a != p2)  == true)
assert((p1a != null) == true)
// end::equals2[]

// tag::equals3[]
assert((null == p1a) == false)
assert((p1a == null) == false)
```

```

assert((null != p1a) == true)
assert((p1a != null) == true)
// end::equals3[]

// tag::equals4[]
assert((p1a eq p1a) == true)
assert((p1a eq p1b) == false) // But p1a == p1b
assert((p1a eq p2) == false)
assert((p1a eq null) == false)
assert((null eq p1a) == false)
assert((null eq null) == true)
// end::equals4[]

// tag::equals5[]
assert((p1a ne p1a) == false)
assert((p1a ne p1b) == true)
assert((p1a ne p2) == true)
assert((p1a ne null) == true)
assert((null ne p1a) == true)
assert((null ne null) == false) // Compiler warns that it's always
false.
// end::equals5[]

// tag::equals6[]
val a1 = Array(1, 2)
val a2 = Array(1, 2)
assert((a1 equals a1) == true)
assert((a1 equals a2) == false)
// end::equals6[]

// tag::equals7[]
assert((a1 sameElements a1) == true)
assert((a1 sameElements a2) == true)
// end::equals7[]

// tag::equals8[]
val s1 = Seq(1, 2)
val s2 = Seq(1, 2)
assert((s1 == s1) == true)
assert((s1 == s2) == true)
assert((s1 sameElements s2) == true)

val m1 = Map("one" -> 1, "two" -> 2)
val m2 = Map("one" -> 1, "two" -> 2)
assert((m1 == m1) == true)

```

```
assert((m1 == m2) == true)
assert((m1.toSeq sameElements m2.toSeq) == true)
// end::equals8[]
```

The equals method tests for *value* equality. That is, obj1 equals obj2 is true if both obj1 and obj2 have the same value. They do not need to refer to the same instance:

```
// tag::definitions[]
// src/script/scala/progscala3/objectsystem/equality/Equality.scala

case class Person(firstName: String, lastName: String, age: Int)

val p1a = Person("Dean", "Wampler", 29)
val p1b = Person("Dean", "Wampler", 29)
val p2  = Person("Buck", "Trends", 30)
// end::definitions[]

// tag::equals1[]
assert((p1a equals p1a) == true)
assert((p1a equals p1b) == true)
assert((p1a equals p2)  == false)
assert((p1a equals null) == false)
// end::equals1[]

// tag::equals2[]
assert((p1a == p1a) == true)
assert((p1a == p1b) == true)
assert((p1a == p2)  == false)
assert((p1a == null) == false)

assert((p1a != p1a) == false)
assert((p1a != p1b) == false)
assert((p1a != p2)  == true)
assert((p1a != null) == true)
// end::equals2[]

// tag::equals3[]
assert((null == p1a) == false)
assert((p1a == null) == false)
assert((null != p1a) == true)
assert((p1a != null) == true)
// end::equals3[]
```

```

// tag::equals4[]
assert((p1a eq p1a) == true)
assert((p1a eq p1b) == false) // But p1a == p1b
assert((p1a eq p2) == false)
assert((p1a eq null) == false)
assert((null eq p1a) == false)
assert((null eq null) == true)
// end::equals4[]

// tag::equals5[]
assert((p1a ne p1a) == false)
assert((p1a ne p1b) == true)
assert((p1a ne p2) == true)
assert((p1a ne null) == true)
assert((null ne p1a) == true)
assert((null ne null) == false) // Compiler warns that it's always
false.
// end::equals5[]

// tag::equals6[]
val a1 = Array(1, 2)
val a2 = Array(1, 2)
assert((a1 equals a1) == true)
assert((a1 equals a2) == false)
// end::equals6[]

// tag::equals7[]
assert((a1 sameElements a1) == true)
assert((a1 sameElements a2) == true)
// end::equals7[]

// tag::equals8[]
val s1 = Seq(1, 2)
val s2 = Seq(1, 2)
assert((s1 == s1) == true)
assert((s1 == s2) == true)
assert((s1 sameElements s2) == true)

val m1 = Map("one" -> 1, "two" -> 2)
val m2 = Map("one" -> 1, "two" -> 2)
assert((m1 == m1) == true)
assert((m1 == m2) == true)
assert((m1.toSeq sameElements m2.toSeq) == true)
// end::equals8[]

```



Hence, `equals` behaves like the `equals` method in Java and the `eq?` method in Ruby, for example.

## The `==` and `!=` Methods

Whereas `==` is an operator in many languages, it is a method in Scala that delegates to `equals`. Similarly, `!=` call `==` and returns the opposite value:

```
// tag::definitions[]
// src/script/scala/progscala3/objectsystem/equality/Equality.scala

case class Person(firstName: String, lastName: String, age: Int)

val p1a = Person("Dean", "Wampler", 29)
val p1b = Person("Dean", "Wampler", 29)
val p2  = Person("Buck", "Trends", 30)
// end::definitions[]

// tag::equals1[]
assert((p1a equals p1a) == true)
assert((p1a equals p1b) == true)
assert((p1a equals p2)  == false)
assert((p1a equals null) == false)
// end::equals1[]

// tag::equals2[]
assert((p1a == p1a)    == true)
assert((p1a == p1b)    == true)
assert((p1a == p2)     == false)
assert((p1a == null)   == false)

assert((p1a != p1a)    == false)
assert((p1a != p1b)    == false)
assert((p1a != p2)     == true)
assert((p1a != null)   == true)
// end::equals2[]

// tag::equals3[]
assert((null == p1a) == false)
assert((p1a == null) == false)
assert((null != p1a) == true)
```

```

assert((p1a != null) == true)
// end::equals3[]

// tag::equals4[]
assert((p1a eq p1a) == true)
assert((p1a eq p1b) == false) // But p1a == p1b
assert((p1a eq p2) == false)
assert((p1a eq null) == false)
assert((null eq p1a) == false)
assert((null eq null) == true)
// end::equals4[]

// tag::equals5[]
assert((p1a ne p1a) == false)
assert((p1a ne p1b) == true)
assert((p1a ne p2) == true)
assert((p1a ne null) == true)
assert((null ne p1a) == true)
assert((null ne null) == false) // Compiler warns that it's always
false.
// end::equals5[]

// tag::equals6[]
val a1 = Array(1, 2)
val a2 = Array(1, 2)
assert((a1 equals a1) == true)
assert((a1 equals a2) == false)
// end::equals6[]

// tag::equals7[]
assert((a1 sameElements a1) == true)
assert((a1 sameElements a2) == true)
// end::equals7[]

// tag::equals8[]
val s1 = Seq(1, 2)
val s2 = Seq(1, 2)
assert((s1 == s1) == true)
assert((s1 == s2) == true)
assert((s1 sameElements s2) == true)

val m1 = Map("one" -> 1, "two" -> 2)
val m2 = Map("one" -> 1, "two" -> 2)
assert((m1 == m1) == true)
assert((m1 == m2) == true)

```

```
assert((m1.toSeq sameElements m2.toSeq) == true)
// end::equals8[]
```

Comparing to null behaves as you might expect:

```
// tag::definitions[]
// src/script/scala/progscala3/objectsystem/equality/Equality.scala
```

```
case class Person(firstName: String, lastName: String, age: Int)
```

```
val p1a = Person("Dean", "Wampler", 29)
val p1b = Person("Dean", "Wampler", 29)
val p2   = Person("Buck", "Trends", 30)
// end::definitions[]
```

```
// tag::equals1[]
assert((p1a equals p1a) == true)
assert((p1a equals p1b) == true)
assert((p1a equals p2)  == false)
assert((p1a equals null) == false)
// end::equals1[]
```

```
// tag::equals2[]
assert((p1a == p1a) == true)
assert((p1a == p1b) == true)
assert((p1a == p2)  == false)
assert((p1a == null) == false)
```

```
assert((p1a != p1a) == false)
assert((p1a != p1b) == false)
assert((p1a != p2)  == true)
assert((p1a != null) == true)
// end::equals2[]
```

```
// tag::equals3[]
assert((null == p1a) == false)
assert((p1a == null) == false)
assert((null != p1a) == true)
assert((p1a != null) == true)
// end::equals3[]
```

```
// tag::equals4[]
assert((p1a eq p1a) == true)
assert((p1a eq p1b) == false) // But p1a == p1b
```

```

assert((p1a eq p2)      == false)
assert((p1a eq null)    == false)
assert((null eq p1a)    == false)
assert((null eq null)   == true)
// end::equals4[]

// tag::equals5[]
assert((p1a ne p1a)     == false)
assert((p1a ne p1b)     == true)
assert((p1a ne p2)      == true)
assert((p1a ne null)    == true)
assert((null ne p1a)    == true)
assert((null ne null)   == false) // Compiler warns that it's always
false.
// end::equals5[]

// tag::equals6[]
val a1 = Array(1, 2)
val a2 = Array(1, 2)
assert((a1 equals a1) == true)
assert((a1 equals a2) == false)
// end::equals6[]

// tag::equals7[]
assert((a1 sameElements a1) == true)
assert((a1 sameElements a2) == true)
// end::equals7[]

// tag::equals8[]
val s1 = Seq(1, 2)
val s2 = Seq(1, 2)
assert((s1 == s1) == true)
assert((s1 == s2) == true)
assert((s1 sameElements s2) == true)

val m1 = Map("one" -> 1, "two" -> 2)
val m2 = Map("one" -> 1, "two" -> 2)
assert((m1 == m1) == true)
assert((m1 == m2) == true)
assert((m1.toSeq sameElements m2.toSeq) == true)
// end::equals8[]

```

## The eq and ne Methods

The eq method tests for *reference* equality. That is, obj1 eq obj2 is true if and only if both obj1 and obj2 point to the same location in memory. These methods are only defined for AnyRef:

```
// tag::definitions[]
// src/script/scala/progscala3/objectsystem/equality/Equality.scala

case class Person(firstName: String, lastName: String, age: Int)

val p1a = Person("Dean", "Wampler", 29)
val p1b = Person("Dean", "Wampler", 29)
val p2  = Person("Buck", "Trends", 30)
// end::definitions[]

// tag::equals1[]
assert((p1a equals p1a) == true)
assert((p1a equals p1b) == true)
assert((p1a equals p2)  == false)
assert((p1a equals null) == false)
// end::equals1[]

// tag::equals2[]
assert((p1a == p1a) == true)
assert((p1a == p1b) == true)
assert((p1a == p2)  == false)
assert((p1a == null) == false)

assert((p1a != p1a) == false)
assert((p1a != p1b) == false)
assert((p1a != p2)  == true)
assert((p1a != null) == true)
// end::equals2[]

// tag::equals3[]
assert((null == p1a) == false)
assert((p1a == null) == false)
assert((null != p1a) == true)
assert((p1a != null) == true)
// end::equals3[]

// tag::equals4[]
assert((p1a eq p1a) == true)
assert((p1a eq p1b) == false) // But p1a == p1b
assert((p1a eq p2)  == false)
```

```

assert((p1a eq null) == false)
assert((null eq p1a) == false)
assert((null eq null) == true)
// end::equals4[]

// tag::equals5[]
assert((p1a ne p1a) == false)
assert((p1a ne p1b) == true)
assert((p1a ne p2) == true)
assert((p1a ne null) == true)
assert((null ne p1a) == true)
assert((null ne null) == false) // Compiler warns that it's always
false.
// end::equals5[]

// tag::equals6[]
val a1 = Array(1, 2)
val a2 = Array(1, 2)
assert((a1 equals a1) == true)
assert((a1 equals a2) == false)
// end::equals6[]

// tag::equals7[]
assert((a1 sameElements a1) == true)
assert((a1 sameElements a2) == true)
// end::equals7[]

// tag::equals8[]
val s1 = Seq(1, 2)
val s2 = Seq(1, 2)
assert((s1 == s1) == true)
assert((s1 == s2) == true)
assert((s1 sameElements s2) == true)

val m1 = Map("one" -> 1, "two" -> 2)
val m2 = Map("one" -> 1, "two" -> 2)
assert((m1 == m1) == true)
assert((m1 == m2) == true)
assert((m1.toSeq sameElements m2.toSeq) == true)
// end::equals8[]

```

## NOTE

In many languages, such as Java, C++, and C#, == behaves like eq instead of equals.

The ne method is the negation of eq. It is equivalent to !(obj1 eq obj2).

## Array Equality and the sameElements Method

Because Arrays are defined by Java, equals does reference comparison, like eq, rather than value comparison:

```
// tag::definitions[]
// src/script/scala/progscala3/objectsystem/equality/Equality.scala

case class Person(firstName: String, lastName: String, age: Int)

val p1a = Person("Dean", "Wampler", 29)
val p1b = Person("Dean", "Wampler", 29)
val p2  = Person("Buck", "Trends", 30)
// end::definitions[]

// tag::equals1[]
assert((p1a equals p1a) == true)
assert((p1a equals p1b) == true)
assert((p1a equals p2)  == false)
assert((p1a equals null) == false)
// end::equals1[]

// tag::equals2[]
assert((p1a == p1a) == true)
assert((p1a == p1b) == true)
assert((p1a == p2)  == false)
assert((p1a == null) == false)

assert((p1a != p1a) == false)
assert((p1a != p1b) == false)
assert((p1a != p2)  == true)
assert((p1a != null) == true)
```

```

// end::equals2[]

// tag::equals3[]
assert((null == p1a) == false)
assert((p1a == null) == false)
assert((null != p1a) == true)
assert((p1a != null) == true)
// end::equals3[]

// tag::equals4[]
assert((p1a eq p1a) == true)
assert((p1a eq p1b) == false) // But p1a == p1b
assert((p1a eq p2) == false)
assert((p1a eq null) == false)
assert((null eq p1a) == false)
assert((null eq null) == true)
// end::equals4[]

// tag::equals5[]
assert((p1a ne p1a) == false)
assert((p1a ne p1b) == true)
assert((p1a ne p2) == true)
assert((p1a ne null) == true)
assert((null ne p1a) == true)
assert((null ne null) == false) // Compiler warns that it's always
false.
// end::equals5[]

// tag::equals6[]
val a1 = Array(1, 2)
val a2 = Array(1, 2)
assert((a1 equals a1) == true)
assert((a1 equals a2) == false)
// end::equals6[]

// tag::equals7[]
assert((a1 sameElements a1) == true)
assert((a1 sameElements a2) == true)
// end::equals7[]

// tag::equals8[]
val s1 = Seq(1, 2)
val s2 = Seq(1, 2)
assert((s1 == s1) == true)
assert((s1 == s2) == true)

```



```

assert((s1 sameElements s2) == true)

val m1 = Map("one" -> 1, "two" -> 2)
val m2 = Map("one" -> 1, "two" -> 2)
assert((m1 == m1) == true)
assert((m1 == m2) == true)
assert((m1.toSeq sameElements m2.toSeq) == true)
// end::equals8[]

```

Instead, you have to use the sameElements method:

```

// tag::definitions[]
// src/script/scala/progscala3/objectsystem/equality/Equality.scala

case class Person(firstName: String, lastName: String, age: Int)

val p1a = Person("Dean", "Wampler", 29)
val p1b = Person("Dean", "Wampler", 29)
val p2  = Person("Buck", "Trends", 30)
// end::definitions[]

// tag::equals1[]
assert((p1a equals p1a) == true)
assert((p1a equals p1b) == true)
assert((p1a equals p2)  == false)
assert((p1a equals null) == false)
// end::equals1[]

// tag::equals2[]
assert((p1a == p1a) == true)
assert((p1a == p1b) == true)
assert((p1a == p2)  == false)
assert((p1a == null) == false)

assert((p1a != p1a) == false)
assert((p1a != p1b) == false)
assert((p1a != p2)  == true)
assert((p1a != null) == true)
// end::equals2[]

// tag::equals3[]
assert((null == p1a) == false)
assert((p1a == null) == false)
assert((null != p1a) == true)

```

```

assert((p1a != null) == true)
// end::equals3[]

// tag::equals4[]
assert((p1a eq p1a) == true)
assert((p1a eq p1b) == false) // But p1a == p1b
assert((p1a eq p2) == false)
assert((p1a eq null) == false)
assert((null eq p1a) == false)
assert((null eq null) == true)
// end::equals4[]

// tag::equals5[]
assert((p1a ne p1a) == false)
assert((p1a ne p1b) == true)
assert((p1a ne p2) == true)
assert((p1a ne null) == true)
assert((null ne p1a) == true)
assert((null ne null) == false) // Compiler warns that it's always
false.
// end::equals5[]

// tag::equals6[]
val a1 = Array(1, 2)
val a2 = Array(1, 2)
assert((a1 equals a1) == true)
assert((a1 equals a2) == false)
// end::equals6[]

// tag::equals7[]
assert((a1 sameElements a1) == true)
assert((a1 sameElements a2) == true)
// end::equals7[]

// tag::equals8[]
val s1 = Seq(1, 2)
val s2 = Seq(1, 2)
assert((s1 == s1) == true)
assert((s1 == s2) == true)
assert((s1 sameElements s2) == true)

val m1 = Map("one" -> 1, "two" -> 2)
val m2 = Map("one" -> 1, "two" -> 2)
assert((m1 == m1) == true)
assert((m1 == m2) == true)

```

```
assert((m1.toSeq sameElements m2.toSeq) == true)
// end::equals8[]
```

Because Arrays are mutable and have this behavior when comparing them, consider when it's better to use another collection instead. However, an argument for *not* using alternatives is when you really need the performance benefits of arrays. For example, an array of Doubles is a single block of memory that greatly reduces the references scattered over the heap, which are typical of most collections. This can greatly improve performance when fetching data into the CPU cache.

In contrast, Seqs, Maps, and most other collections work as you would expect:

```
// tag::definitions[]
// src/script/scala/progscala3/objectsystem/equality/Equality.scala
```

```
case class Person(firstName: String, lastName: String, age: Int)
```

```
val p1a = Person("Dean", "Wampler", 29)
val p1b = Person("Dean", "Wampler", 29)
val p2  = Person("Buck", "Trends", 30)
// end::definitions[]
```

```
// tag::equals1[]
assert((p1a equals p1a) == true)
assert((p1a equals p1b) == true)
assert((p1a equals p2)  == false)
assert((p1a equals null) == false)
// end::equals1[]
```

```
// tag::equals2[]
assert((p1a == p1a) == true)
assert((p1a == p1b) == true)
assert((p1a == p2)  == false)
assert((p1a == null) == false)
```

```
assert((p1a != p1a) == false)
assert((p1a != p1b) == false)
assert((p1a != p2)  == true)
assert((p1a != null) == true)
```

```

// end::equals2[]

// tag::equals3[]
assert((null == p1a) == false)
assert((p1a == null) == false)
assert((null != p1a) == true)
assert((p1a != null) == true)
// end::equals3[]

// tag::equals4[]
assert((p1a eq p1a) == true)
assert((p1a eq p1b) == false) // But p1a == p1b
assert((p1a eq p2) == false)
assert((p1a eq null) == false)
assert((null eq p1a) == false)
assert((null eq null) == true)
// end::equals4[]

// tag::equals5[]
assert((p1a ne p1a) == false)
assert((p1a ne p1b) == true)
assert((p1a ne p2) == true)
assert((p1a ne null) == true)
assert((null ne p1a) == true)
assert((null ne null) == false) // Compiler warns that it's always
false.
// end::equals5[]

// tag::equals6[]
val a1 = Array(1, 2)
val a2 = Array(1, 2)
assert((a1 equals a1) == true)
assert((a1 equals a2) == false)
// end::equals6[]

// tag::equals7[]
assert((a1 sameElements a1) == true)
assert((a1 sameElements a2) == true)
// end::equals7[]

// tag::equals8[]
val s1 = Seq(1, 2)
val s2 = Seq(1, 2)
assert((s1 == s1) == true)
assert((s1 == s2) == true)

```

```

assert((s1 sameElements s2) == true)

val m1 = Map("one" -> 1, "two" -> 2)
val m2 = Map("one" -> 1, "two" -> 2)
assert((m1 == m1) == true)
assert((m1 == m2) == true)
assert((m1.toSeq sameElements m2.toSeq) == true)
// end::equals8[]

```

## Equality and Inheritance

We learned previously that case classes can't inherit from other case classes, in part because equals and hashCode would break. The generated versions don't account for derivation, where additional fields might be added. Consider this example:

```

// tag::definitions[]
//
src/script/scala/progscala3/objectsystem/equality/InheritanceEquality.
scala

class Employee(val name: String, val annualSalary: Double)
class Manager(name: String, annualSalary: Double, val minions:
Seq[Employee])
  extends Employee(name, annualSalary)

val e1  = new Employee("Buck Trends", 50000.0)
val e1b = new Employee("Buck Trends", 50000.0)
val e2  = new Employee("Jane Doe", 50000.0)
val m1  = new Manager("Jane Doe", 50000.0, Seq(e1, e2))
val all = Seq(e1, e1b, e2, m1)
// end::definitions[]

// tag::equals1[]
assert((e1 == e1) == true)
assert((e1 == e1b) == false) // Different references, so == returns
false.
assert((e1 == e2) == false)
assert((e2 == m1) == false)
// end::equals1[]

// tag::equals2[]

```

```
val same = all.filter(e => e2 == e)
// end::equals2[]
```

Note that `e2` and `m1` have the same name and `annualSalary`.

I'll ignore whether or not this is a good use of inheritance. I didn't use case classes to avoid the decision of which one of the two would be a good case class. I didn't define equality methods either, which means that reference equality, the default for objects on the JVM, will be used:

```
// tag::definitions[]
//
src/script/scala/progscala3/objectsystem/equality/InheritanceEquality.
scala

class Employee(val name: String, val annualSalary: Double)
class Manager(name: String, annualSalary: Double, val minions:
Seq[Employee])
  extends Employee(name, annualSalary)

val e1  = new Employee("Buck Trends", 50000.0)
val e1b = new Employee("Buck Trends", 50000.0)
val e2  = new Employee("Jane Doe", 50000.0)
val m1  = new Manager("Jane Doe", 50000.0, Seq(e1, e2))
val all = Seq(e1, e1b, e2, m1)
// end::definitions[]

// tag::equals1[]
assert((e1 == e1) == true)
assert((e1 == e1b) == false) // Different references, so == returns
false.
assert((e1 == e2) == false)
assert((e2 == m1) == false)
// end::equals1[]

// tag::equals2[]
val same = all.filter(e => e2 == e)
// end::equals2[]
```

Suppose we decided to add `equals` methods? In particular, what should happen with `e2 == m1`, since they have the same name and

annualSalary fields? For example, if I filter the list of all employees for those that are equal to e2:

```
// tag::definitions[]
//
src/script/scala/progscala3/objectsystem/equality/InheritanceEquality.
scala

class Employee(val name: String, val annualSalary: Double)
class Manager(name: String, annualSalary: Double, val minions:
Seq[Employee])
  extends Employee(name, annualSalary)

val e1  = new Employee("Buck Trends", 50000.0)
val e1b = new Employee("Buck Trends", 50000.0)
val e2  = new Employee("Jane Doe", 50000.0)
val m1  = new Manager("Jane Doe", 50000.0, Seq(e1, e2))
val all = Seq(e1, e1b, e2, m1)
// end::definitions[]

// tag::equals1[]
assert((e1 == e1) == true)
assert((e1 == e1b) == false) // Different references, so == returns
false.
assert((e1 == e2) == false)
assert((e2 == m1) == false)
// end::equals1[]

// tag::equals2[]
val same = all.filter(e => e2 == e)
// end::equals2[]
```

I might be tempted to argue that I'm just comparing e2 to each Employee, so I only care about the name and annualSalary. When two objects have the same values for these two fields, I consider them equal, so I should consider e2 == m1 equal. Certainly the Employee.equals method I might implement will only compare those two fields.

What if I write m1 == e2 instead? Am I now thinking about whether two *managers* are equal? Now the Manager.equals method I might

write will compare all three fields and return false for this comparison. In mathematics, equality is usually considered *symmetric*, where `a == b` means `b == a`.

The point is that equality might be contextual. If we really only care about employee fields, then maybe it's okay if `e2 == m1` returns true, but most of the time this creates ambiguities we should avoid. It reflects a poor design.

The resources mentioned at the beginning of this section all point out that a good `equals` method obeys the following rules for an *equivalence relation* (using `==`):

- It is *reflexive*: For any `x` of type `Any`, `x == x`.
- It is *symmetric*: For any `x` and `y` of type `Any`, `x == y` returns the same value as `y == x`.
- It is *transitive*: For any `x`, `y`, and `z` of type `Any`, if `x == y` and `y == z`, then `x == z`.

### TIP

Never use equality with instances in a type hierarchy where fields are spread between parent and child types. Only allow equality when comparing instances of the same concrete type. Obey the rules of an *equivalence relation*.

## Multiversal Equality

The only reason I could even discuss the possibilities in the previous section is because Scala has historically followed Java's model of equality checking, which can be called *universal equality*. This means the compiler allows us to compare any values to any other values. I used `Any` as the type in the list of rules for equivalence relations, following the documentation for `Any.equals`.

In Scala, this translates to the following declarations for `equals`:



```

abstract class Any:
  def equals(other: Any): Boolean
  ...
abstract class AnyRef:
  def equals(other: AnyRef): Boolean
  ...

```

For implementation reasons, AnyRef can't use Any as the type for other, but otherwise, Any.equals defines how *all* equals methods work.

Scala 3 introduces *multiversal equality*, a mechanism for limiting what types can be compared to provide more rigorous type safety while also supporting backwards compatibility. Instead of a “universe”, we have a “multiverse”, where equality checking is only allowed within each part of the “multiverse”.

Multiversal equality is implemented as a language feature and the `scala.Eql` type class, which we first encountered in “[Type Class Derivation](#)”. We used an example enum `Tree[T]` derives `Eql` to constrain what comparisons can be done with `==` and `!=`.

The derives `Eql` clause generates the following given instance:

```

given Eql[Tree[T], Tree[T]] = Eql.derived

```

We are also limited to types `T` that have their own given instance, `given Eql[T, T] = Eql.derived`.

For backwards compatibility, this language feature must be turned on either with the compiler flag `-language:strictEquality` or the import statement `import scala.language.strictEquality`.

Scala 3 automatically derives `Eql` for the primitive types `Byte`, `Short`, `Char`, `Int`, `Long`, `Float`, `Double`, `Boolean`, `Unit`, `java.lang.Number`, `java.lang.Boolean`, `java.lang.Character`, `Seq`, and `Set`.

Additional given instances are defined as necessary to permit the following comparisons:

- Primitive numeric types can be compared with each other and with subtypes of `java.lang.Number`.
- `Boolean` can be compared with `java.lang.Boolean`.
- `Char` can be compared with `java.lang.Character`.
- Two arbitrary subtypes of `Seq` can be compared with each other if their element types can be compared. The two sequence types need not be the same. (This is a compromise with pragmatism...)
- Similarly, two arbitrary subtypes of `Set` can be compared with each other if their element types can be compared. The two set types need not be the same.
- Any subtype of `AnyRef` can be compared with `null`.

All these comparisons are symmetric, of course.

How does this feature improve type safety? Consider this contrived example:

```
// src/script/scala/progscala3/objectsystem/equality/EqlBug.scala

case class X(name: String)
❶

def findMarkers[T](seq: Seq[T]): Seq[T] =
❷
  seq.filter(_ == X("marker"))

findMarkers(Seq(X("one"), X("two"), X("marker"), X("three")))
❸

// Refactoring
case class Y(name: String)
❹
findMarkers(Seq(Y("one"), Y("two"), Y("marker"), Y("three")))
❺
```

- ❶ Define some domain class.

- ② Define a method to local all the “special” instances of this domain class.
- ③ Use it, returning `List(X(marker))`.
- ④ Introduce a new domain class as part of a refactoring.
- ⑤ Reuse the same `findMarkers` code as before, knowing that the type parameter `T` will correctly change to `Y`. But now it returns `Nil`!

The flaw is that `findMarkers` still uses the old value for the “marker”. Because of universal equality, the compiler happily allows us to do the comparison `X("marker") == Y(...)`, which always returns false and hence the filtering will always return an empty collection!

Enabling multiversal equality forces us to use a better design:

```
// src/script/scala/progscala3/objectsystem/equality/EqlBugFix.scala

import scala.language.strictEquality
❶

case class X(name: String) derives Eql
❷

def findMarkers[T](marker: T, seq: Seq[T])(using Eql[T,T]): Seq[T] =
❸
    seq.filter(_ == marker)

findMarkers(X("marker"), Seq(X("one"), X("two"), X("marker"),
X("three")))

// Refactoring
case class Y(name: String) derives Eql
findMarkers(Y("marker"), Seq(Y("one"), Y("two"), Y("marker"),
Y("three")))
```

- ❶ Enable the language feature.

- ② Derive from EqL.
- ③ Pass in a type-compatible marker for filtering. Note the using clause; it requires all T values to derive from EqL. Without this clause, you'll get a compilation error for `_ == marker`.

The output will be the same as before, except now the last line will return `List(Y(marker))`, which is what we desire.

More details on how Scala implements multiversal equality can be found in this [Scala documentation](#).

## Recap and What's Next

We discussed the nontrivial behavior of type parameterization when inheritance is involved. We explored the comparison methods for Scala types and how careful design is required to correctly implement *equivalence relations*. Finally, we explored a new Scala 3 feature, *multiversal equality*, which aims to make equality checking more type safe than in the past.

Next we'll continue our discussion of the object system by examining the behavior of field initialization during construction and member overriding and resolution in a type with multiple parent types.

# Chapter 12. Instance Initialization and Method Resolution

---

## A NOTE FOR EARLY RELEASE READERS

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

This will be the 12th chapter of the final book. Please note that the GitHub repo will be made active later on.

If you have comments about how we might improve the content and/or examples in this book, or if you notice missing material within this chapter, please reach out to the author at [mcronin@oreilly.com](mailto:mcronin@oreilly.com).

A type has a *directed acyclic graph* (DAG) of dependencies with its parent types. When the fields in a type are spread over this DAG, initialization order becomes important to prevent access before initialization. When one or more types in the DAG define and override the same method, then method resolution rules need to be understood. This chapter explores these concepts. As we’ll see, they are closely related, governed by a concept called *linearization*.

## Linearization of an Object’s Hierarchy

Because of single inheritance, if we ignored mixed-in traits, the inheritance hierarchy would be linear, one ancestor after another.

When traits are considered, each of which may be derived from other traits and classes, the inheritance hierarchy forms a directed, acyclic graph.

The term *linearization* refers to the algorithm used to “flatten” this graph to determine the order of construction of the types in the tree, as well as the ordering of method invocations, including binding of super to invoke a parent-type method. Think of the linearized traversal path flowing left to right with the super types to the left and the subtypes to the right. Construction follows this *left to right* ordering, meaning super types are constructed before sub types. In contrast, when resolving which overloaded method to call in a hierarchy, the traversal goes *right to left*, where nearest super class definitions of the method take precedence over definitions higher in the graph.

We saw an example of these behaviors in “**Stackable Traits**”. The `VetoableClickCountObserver` example required us to declare the two mixin traits in the order `Button(...)` with `ObservableClicks` with `VetoableClicks(maxAllowed = 2)`. All three types defined a `click` method. We wanted to invoke the `VetoableClicks.click()` method *first*, which would only call the super-type `click` methods if we had not already exceeded the max number of allowed clicks. If not, it would call `ObservableClicks.click()`, which would first call `+Button.click()`, *then* notify observers. Switching the order of declaration to `+VetoableClicks` with `ObservableClicks` would mean that `Button.click()` is *still* only called up to an allowed number of invocations, but *all* invocations would trigger observation, because `ObservableClicks.click()` would be called before `Vetoable.clicks()`.

Consider the following, deliberately-complicated example that demonstrates linearization:

```
//  
src/script/scala/progscala3/objectsystem/linearization/Linearization.s  
cala
```

```

trait Base:
  var str = "Base"
  def m(): String = "Base"

trait T1 extends Base:
  str = str + " T1"
  override def m(): String = "T1 " + super.m()

trait T2 extends Base:
  str = str + " T2"
  override def m(): String = "T2 " + super.m()

trait T3 extends Base:
  str = str + " T3"
  override def m(): String = "T3 " + super.m()

class C2 extends T2:
  str = str + " C2"
  override def m(): String = "C2 " + super.m()

class C3A extends C2 with T1 with T2 with T3:
  str = str + " C3A"
  override def m(): String = "C3A " + super.m()

class C3B extends C2 with T3 with T2 with T1:
  str = str + " C3B"
  override def m(): String = "C3B " + super.m()

val c3a = new C3A
val c3b = new C3B
val c3c = new C2 with T1 with T2 with T3
val c3d = new C2 with T3 with T2 with T1

```

```

// Construction precedence:
assert(c3a.str == "Base T2 C2 T1 T3 C3A")
assert(c3b.str == "Base T2 C2 T3 T1 C3B")
assert(c3c.str == "Base T2 C2 T1 T3")
assert(c3d.str == "Base T2 C2 T3 T1")

```

```

// Method invocation precedence:
assert(c3a.m() == "C3A T3 T1 C2 T2 Base")
assert(c3b.m() == "C3B T1 T3 C2 T2 Base")
assert(c3c.m() == "T3 T1 C2 T2 Base")
assert(c3d.m() == "T1 T3 C2 T2 Base")

```

1  
2

3

- ❶ Used to track construction ordering
- ❷ Used to track method invocation ordering
- ❸ Compare these strings in the two assertion blocks, for example, `c3a.str` vs. `c3a.m()`.

Figure 12-1 diagrams the type hierarchy for C3A. The diagrams for C3B and the anonymous types for c3c and c3d would be identical.

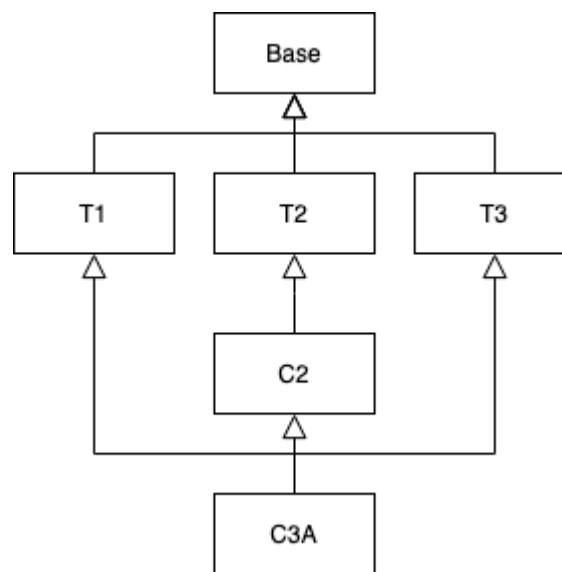


Figure 12-1. Type hierarchy example

Note that we don't need to draw a line from C3A to T2, even though it explicitly mixes it in, because it also gets mixed in through C2.

The first four assertions show how the constructors are invoked, essentially *left to right*.

Consider the case of `c3a`. Before we can invoke the C3A constructor body, we have to construct the parents. The first parent is C2, but it is derived from T2, which is derived from Base, so the construction order starts out Base, then T2, then C2. Next we move to T1, which depends on Base, we've already constructed Base, so we can



immediately construct T1. The same logic holds for T2 and T3. Finally, the C3A constructor body is invoked.

C3B just flips the order of the three traits, but note that T2 is always the first trait constructed! This is because C2 goes first, but it requires Base then T2 first, etc., as we described for C3A.

Finally, c3c is very similar to c3a, because it just constructs an anonymous type instead of the *named* type C3A, but with the exact same parents. The same argument applies for c3d and c3b.

Keep in mind that the construction of Any and AnyRef happen before Base, but of course they don't modify Base.str.

The last four assertions show how method invocation precedence is implemented for overridden methods. It traverses the linearization in the opposite direction, *right to left*. You can see this when you compare c3a.str to c3a.m(), etc.

Working through c3a.m(), the C3A.m() method is called first. Then, to determine what to call for super.m() inside C3A.m(), we go right to left. T3 is right-most type, so T3.m() is called next. You might think that the call to super.m() inside T3.m() should invoke Base.m(), but the precedence rules are evaluated globally for the instance, not by the static declaration of the trait. Hence, we go back to the declaration of C3A and see that T2 is next, *but* T2 is also a parent dependency of C2, so we have to wait until C2.m() can be called. We keep going and arrive at T1; we can call its m() immediately, then we reach C2. C2 overrides m(), so we call it, then we call the super().m() call in C2.m(), which resolves to T2.m(). The super().m() call in T2.m() resolves to Base.m().

Here is the actual algorithm for calculating the linearization:

## LINEARIZATION ALGORITHM

1. Put the actual type of the instance as the first element.
2. Starting with the *rightmost* parent type, work *left*, computing the linearization of each type, appending its linearization to the cumulative linearization. (Ignore AnyRef and Any for now.)
3. Working from *left to right*, remove any type if it appears again to the *right* of the current position.
4. Append AnyRef and Any.

For value classes, replace AnyRef with AnyVal.

It's worth looking at the linearization example just discussed and convince yourself that the rules achieve the result we described.

### TIP

Overly complex type hierarchies can result in method lookup “surprises.” If you have to work through this algorithm to figure out what's going on in your code, try simplifying it!

## Overriding Members of Classes and Traits

Types can declare *concrete* or *abstract* members: fields, methods, and type aliases. All abstract members must be defined by a derived class or trait before an instance can be created.

In Scala, when overriding a concrete member, the `override` keyword is required. This keyword is not allowed for member definitions that don't override a parent type's concrete definition.

Requiring the `override` keyword has several benefits:

- It catches misspelled members that were intended to be overrides. The compiler will throw an error that the member doesn't override anything.
- It catches a subtle bug that can occur if a new member is added to a base class where the member's name collides with a preexisting member in a derived class, one which is unknown to the base class developer. That is, the derived-class member was never intended to override a base-class member. Because the derived-class member won't have the `override` keyword, the compiler will throw an error when the new base-class member is introduced.
- Having to add the keyword reminds you to consider carefully how to handle the parent type's definition.

## Consider Making Concrete Members Final

Should you *ever* override a concrete member? The correct answer is *almost never*.

The relationship between parent and child types is a *contract* and care is required to ensure that a child class does not break the implemented behaviors specified by the parent type. This is especially true when different people or teams maintain the parent and child types. It is very easy to break this contract.

Let's first consider method overrides. Should an override of the `foo` method call `super.foo`? If so, then *when* should it be called within the child's implementation? Of course, there is no right answer.

A far more robust *contract* is the *Template Method Pattern* described in *an earlier section*. In this pattern, the parent class provides a concrete and `final` implementation of a method, defining the outline of the required behavior. The method calls `abstract` protected methods at the points where polymorphic behavior is needed. Then,

subclasses implement the protected, abstract members, but never override the final concrete method.

Here is an example, a sketch of a payroll calculator for a company based in the US:

```
//
src/script/scala/progscala3/objectsystem/overrides/PayrollTemplateMethod.scala

case class Address(city: String, state: String, zip: String)
case class Employee(name: String, salary: Double, address: Address)

abstract class Payroll:
  final def netPay(employee: Employee): Double =
    ❶
    val fedTaxes    = calcFedTaxes(employee.salary)
    val stateTaxes = calcStateTaxes(employee.salary, employee.address)
    employee.salary - fedTaxes - stateTaxes

    protected def calcFedTaxes(salary: Double): Double
    ❷
    protected def calcStateTaxes(salary: Double, address: Address):
    Double

object Payroll2020 extends Payroll:
  val stateRate = Map(
    "YY" -> 0.03,
    "ZZ" -> 0.0)

  def calcFedTaxes(salary: Double): Double = salary * 0.25
  ❸
  def calcStateTaxes(salary: Double, address: Address): Double =
    salary * stateRate(address.state)

val tom  = Employee("Tom Jones", 100000.0, Address("MyTown", "YY",
"98765"))
val jane = Employee("Jane Doe", 110000.0, Address("BigCity", "ZZ",
"67890"))

assert(Payroll2020.netPay(tom) == 72000.0)
assert(Payroll2020.netPay(jane) == 82500.0)
```

- ❶ The concrete template method, declared `final` so that no child class can override and potentially break it.
- ❷ The protected methods that are “hooks” for polymorphic behavior.
- ❸ A concrete implementation.

The `netPay` method defines the “protocol” for calculating payroll, and delegates to abstract methods for details that change year over year. Note that `override` is not needed anywhere.

These days, when I see the `override` keyword in code, I see it as a potential *design smell*. Someone is overriding concrete behavior and subtle bugs might be lurking, because it’s easy to forget when or if to call the parent method being overridden and to ensure that the “contract” of the parent method is faithfully preserved.

I can think of two exceptions to this rule. First, some parent-class implementations of a method do nothing useful, such as the default `toString`, `equals`, and `hashCode` methods. Unfortunately, overriding such methods is so common that the practice has made us too comfortable with overriding concrete methods, in general.

We encountered a good example of the second kind of exception in “**Stackable Traits**”, where we needed to mix in non-overlapping behaviors. In the `ObservableClicks` example there, we override the `click` method to add notification of observers. We took special care to invoke the super-class method first, then invoke the notification logic, and finally return the value that was returned from the super-class method invocation. Mixing in logging calls is another common example.

Another, special case example is when you override a method in a *test double*. One example is replacing a random number generator with a deterministic replacement so test results are deterministic.

Another example is replacing a call to an external service with fixed values.

In those situations, it's critical to understand the parent-class behavior and treat it as an *invariant* that you preserve in the subclass override. The behavior is orthogonal, like logging method calls, it's easier to do the right thing. The observer example required care, because we wanted to notify observers *after* a state change, which meant we needed to call the parent-class method *first*.

### TIP

Avoid overriding concrete members, except when replacing default implementations like `toString`, mixing in orthogonal behavior, or testing scenarios. Be careful to preserve the *contract* of the method you are overriding.

You can prevent overrides with `final`, but doing so precludes those valid situations when overriding is necessary. Consider when and if overriding might be necessary. Otherwise, use `final`.

What about overriding concrete fields? Sometimes a field will be assigned a default value in a parent type, for convenience, with the expectation that derived types will provide a more useful value.

If you are designing such parent types, resist the urge to define a default value unless it is likely to be used almost all the time. Consider using the following idiom.

```
trait Alarm:
  val panicLevel: Int = DEFAULT_PANIC_LEVEL

object Alarm:
  val DEFAULT_PANIC_LEVEL = "PANIC!!"
```

The use of `DEFAULT` (or `default`) indicates to the user that overriding this value is encouraged.

## TIP

Consider using `final` for almost all concrete definitions, to prevent undesired overrides.

The disadvantages of this recommendation include the hassle of adding this keyword everywhere and the fact that it prevents useful overrides in unit tests when you want to *stub* out behavior, like a database call. However, the latter concern can be addressed with a refactoring where all such test “seams” are encapsulated in abstract, protected methods (along with the template method pattern), so it’s easy to implement these overrides.

## When Should Types Be Final?

A type can also be declared `final`, preventing subtyping. This is a good way to prevent *any* overrides, but also precludes this mechanism for special cases, like stubbing behavior for unit testing.

Most of the time, you’ll use sealed hierarchies or enums to prevent unexpected subtyping. Recall our Message hierarchy was sealed in “A Sample Application”, because we defined all the message types we cared about it in one place. In contrast, the Shape hierarchy wasn’t sealed to allow users to add new shapes. However, a user might decide to implement a Polygon class that handles all polygon shapes of three or more vertices. That would be a good type to make `final`.

If you want to make a *concrete* class open for subtyping, use the open keyword discussed in “Classes Open for Extension”.

## NOTE

Some of the types in the Scala library are final, including hooks for JDK classes like `String` and all the “value” types derived from `AnyVal` (see [Chapter 13](#)).

## Defining Abstract Methods and Overriding Concrete Methods

For members that aren’t final, let’s examine the rules for overriding them and the behaviors that result, starting with methods.

We’ll extend the `Widget` trait we introduced in “[Traits as Mixins](#)” with an abstract method `draw`, to support “rendering” the widget to a display, web page, etc. We’ll also override `toString()` using an ad hoc format convention. Instead of a trait, we’ll declare `Widget` an abstract class, since it functions as a logical parent for all UI widgets (but this change isn’t required):

```
// src/main/scala/progscala3/objectsystem/ui/Widget.scala
package progscala3.objectsystem.ui

abstract class Widget:
  def draw(): String
  override def toString(): String = "(widget)"
```

The `draw` method has no body, so it is abstract. Therefore, `Widget` has to be declared abstract. To simplify the example, `draw` will return a string for the instance, rather than render the shape in an actual GUI.

Because `AnyRef` defines `toString`, the `override` keyword is required for `Widget.toString`.



## NOTE

Drawing is actually a *cross-cutting concern*. The state of a `Widget` is one thing; how it is rendered on different platforms—“fat” clients, web pages, mobile devices, etc.—is a separate issue. So, drawing is a very good candidate for a trait, especially if you want your GUI abstractions to be portable. However, to keep things simple, we will handle drawing in the `Widget` hierarchy itself.

Here is the revised `Button` class, with `draw` and `toString` methods. It also mixes in the `Clickable` trait we introduced in “[Stackable Traits](#)”:

```
// src/main/scala/progscala3/objectsystem/ui/Button.scala
package progscala3.objectsystem.ui
import progscala3.traits.ui2.Clickable

open class Button(val label: String) extends Widget with Clickable:

  ❶ def draw(): String = s"Drawing: $this"

  ❷ protected def updateUI(): String = s"$this clicked; updating UI"

  ❸ override def toString(): String =
    s"(button: label=$label, ${super.toString()})"
```

- ❶ A simple “drawing” hack for demonstration purposes.
- ❷ Method declared in `Clickable`
- ❸ Override `toString` again, but incorporate the parent `toString`.

The `super` keyword is analogous to `this`, but it binds to the parent type, which is the aggregation of the parent class and any mixed-in traits. Since `Button` extends `Widget` and `Clickable`, which `toString` is invoked by `super.toString()`? The “closest” parent type `toString` is invoked, as determined by the linearization process discussed

below (“[Linearization of an Object’s Hierarchy](#)”). In this case, because Clickable doesn’t define toString, Widget.toString is called.

Button isn’t a case class, because we will subclass it for other button types and we want to avoid the previously-discussed issues with case class inheritance. So instead, we declare a regular class and mark it open.

Here is a simple script that exercises Button:

```
// src/script/scala/progscala3/objectsystem/ui/Button.scala

scala> import progscala3.objectsystem.ui.Button
scala> val b = new Button("Submit")
scala> b.draw()
val res0: String = Drawing: (button: label=Submit, (widget))
```

Notice how Widget.toString is used by Button.toString in an ad-hoc way.

## Initializing Abstract Fields

Initializing abstract fields in parent types requires attention to initialization order. Consider this example that uses an undefined field before it is properly initialized:

```
// tag::include[]
//
src/script/scala/progscala3/objectsystem/init/BadFieldInitOrder.scala

trait T1:
  val denominator: Int
  val inverse = 1.0/denominator ❶

val obj1 = new T1: ❷
  val denominator = 10

println(s"obj1: denominator = ${obj1.denominator}, inverse =
${obj1.inverse}")
```

```
// end::include[]

assert(obj1.denominator == 10)
assert(obj1.inverse.isPosInfinity)
```

- ❶ What is denominator when inverse is initialized?
- ❷ Construct an instance of an anonymous class where T1 is the parent.

This script prints obj1: denominator = 10, inverse = Infinity.

So, denominator was 0, the default value for an Int, when inverse was calculated in T1. Afterwards, denominator was initialized.

Specifically, the trait body (constructor) was executed before the anonymous class' body. A divide-by-zero exception wasn't thrown, but the compiler recognized the value is infinite.

Scala 3 adds a new experimental compiler flag `-Ycheck-init` that will issue a warning for common initialization problems like this. Pasting the previous example in a REPL with `-Ycheck-init` enabled prints the following:

```
...
9 |   val denominator = 10
  |   ^
  | Access non-initialized field denominator. Calling trace:
  | -> val inverse = 1.0/denominator [ rs$line$1:6 ]
```

The code examples don't use this flag because it rejects some safe code, such as the `LazyList` examples in “[Left Versus Right Folding](#)”. This may work once the feature is no longer experimental. Also, checking can be disabled for specific sections of code using the `@unchecked` annotation.

Scala provides three solutions to this problem, the first two of which work for Scala 2.

The first solution is *lazy values*, which we discussed in “*lazy val*”:

```
// tag::include[]  
// src/script/scala/progscala3/objectsystem/init/LazyValInit.scala  
  
trait T2:  
  val denominator: Int  
  lazy val inverse = 1.0/denominator  
  
val obj2 = new T2:  
  val denominator = 10  
  
println(s"obj2: denominator = ${obj2.denominator}, inverse =  
${obj2.inverse}")  
// end::include[]  
  
assert(obj2.denominator == 10)  
assert(obj2.inverse == 0.1)
```

❶

❶ Added the keyword `lazy`.

This time, the print statement is `obj2: denominator = 10, inverse = 0.1`. Hence, `inverse` is initialized to a valid value, `0.1`, after `denominator` is initialized.

However, `lazy` only helps if the `inverse` isn't used too soon. For example, if you put a print statement in the body of `T2` to see if `inverse` is correct, you'll force evaluation too soon!

### TIP

If a `val` is `lazy`, make sure all uses of the `val` are also as `lazy` as possible.

The second solution is to define `inverse` as a method, which also delays evaluation, but only as long as someone doesn't ask for it:

```
// tag::include[]
// src/script/scala/progscala3/objectsystem/init/DefValInit.scala

trait T3:
  val denominator: Int
  def inverse = 1.0/denominator

val obj3 = new T3:
  val denominator = 10

println(s"obj3: denominator = ${obj3.denominator}, inverse =
${obj3.inverse}")
// end::include[]

assert(obj3.denominator == 10)
assert(obj3.inverse == 0.1)
```

- ❶ Use a method to define inverse.

While a lazy val is only evaluated once, a method is evaluated every time you invoke it. Recall that a lazy val also has some internal logic to check if initialization has already happened, which adds a bit of overhead on each invocation.

Finally, Scala 3 introduced parameters for traits, which is arguably the most robust solution to this issue and to other initialization ordering problems that can occur:

```
// tag::include[]
//
// src/script/scala/progscala3/objectsystem/init/TraitParamValInit.scala

trait T4(val denominator: Int):
  ❶ val inverse = 1.0/denominator
  ❷

val obj4 = new T4(10) {}
  ❸

println(s"obj4: denominator = ${obj4.denominator}, inverse =
${obj4.inverse}")
// end::include[]
```

```
assert(obj4.denominator == 10)
assert(obj4.inverse == 0.1)
```

- ❶ Pass a parameter to the trait, just as you do for classes.
- ❷ Use a regular, “eager” value for inverse.
- ❸ We have to provide a body, even though it is empty, when instantiating a trait.

The output is the same as before.

To recap, the bad initialization showed us that the parent type constructors are evaluated before the child type, so `inverse` prematurely referenced `denominator`. We can either ensure that `denominator` is initialized first using a trait parameter or we can delay evaluation of `inverse` by making it either lazy or a method.

## Overriding Concrete Fields

The same order of construction rules apply for concrete fields and accessing them. Here is an example with both a `val` that is overridden in a child class and a `var` that is reassigned in the child class:

```
//
src/script/scala/progscala3/objectsystem/overrides/ClassFields.scala

trait T5:
  val name = "T5"
  var count = 0

class ClassT5 extends T5:
  override val name = "ClassT5"
  count = 1

val c = new ClassT5()
```

```
assert(c.name == "ClassT5")
assert(c.count == 1)
```

Just as for methods, the `override` keyword is required for the concrete `val` field name when it is overridden, as opposed to defined for the first time. `ClassT5` doesn't override the definition of the `var` field `count`. It just changes the assignment.

### TIP

Use caution when overriding concrete fields, for the same reasons you should use caution when overriding concrete methods.

## Abstract and Concrete Type Aliases

We introduced abstract type declarations in “[Abstract Types Versus Parameterized Types](#)”. Do abstract and concrete types have similar initialization order issues that fields experience?

```
// src/script/scala/progscala3/objectsystem/init/TypeInitOrder.scala
```

```
trait TT1:
  type TA
  type TB = Seq[TA]
  ❶ val seed: TA
  ❷ val seq: TB = Seq.fill(5)(seed)

class TT2 extends TT1:
  ❸ type TA = Int
  val seed: TA = 5

val obj = new TT2
println(s"obj: seq = ${obj.seq}")
```

❶ This works fine!

- ② Use TB to create a value
- ③ TA is initialized here, but TB still works.

This example doesn't quite work, but not because of the type aliases; the seed is prematurely initialized to 0 so the sequence has five zeros instead of five fives. Otherwise, the type aliases work fine. You can even move the definition of seed before TA inside TT2 without changing the behavior.

Unlike fields and methods, it is not possible to override a concrete type definition, even with the `override` keyword.

## Recap and What's Next

We walked through the details of Scala's linearization algorithm for type construction and method lookup resolution. We explored the fine points of overriding members in derived types.

In the next chapter, we'll learn about Scala's collections library.



# Chapter 13. The Scala Type Hierarchy

---

## A NOTE FOR EARLY RELEASE READERS

With Early Release ebooks, you get books in their earliest form—the author’s raw and unedited content as they write—so you can take advantage of these technologies long before the official release of these titles.

This will be the 13th chapter of the final book. Please note that the GitHub repo will be made active later on.

If you have comments about how we might improve the content and/or examples in this book, or if you notice missing material within this chapter, please reach out to the author at [mcronin@oreilly.com](mailto:mcronin@oreilly.com).

We’ve learned a lot about Scala’s implementation of object-oriented programming. We’ve already seen many of the types available in Scala’s library. In this chapter, I’ll fill in details about the object-oriented type hierarchy and I’ll discuss some useful types in the standard library we haven’t discussed yet. We’ll focus on the collections part of the library in *Collections*.

**Figure 13-1** shows the large-scale structure of the inheritance hierarchy for Scala types. Unless otherwise noted, all the types we’ll discuss are in the top-level `scala` package.

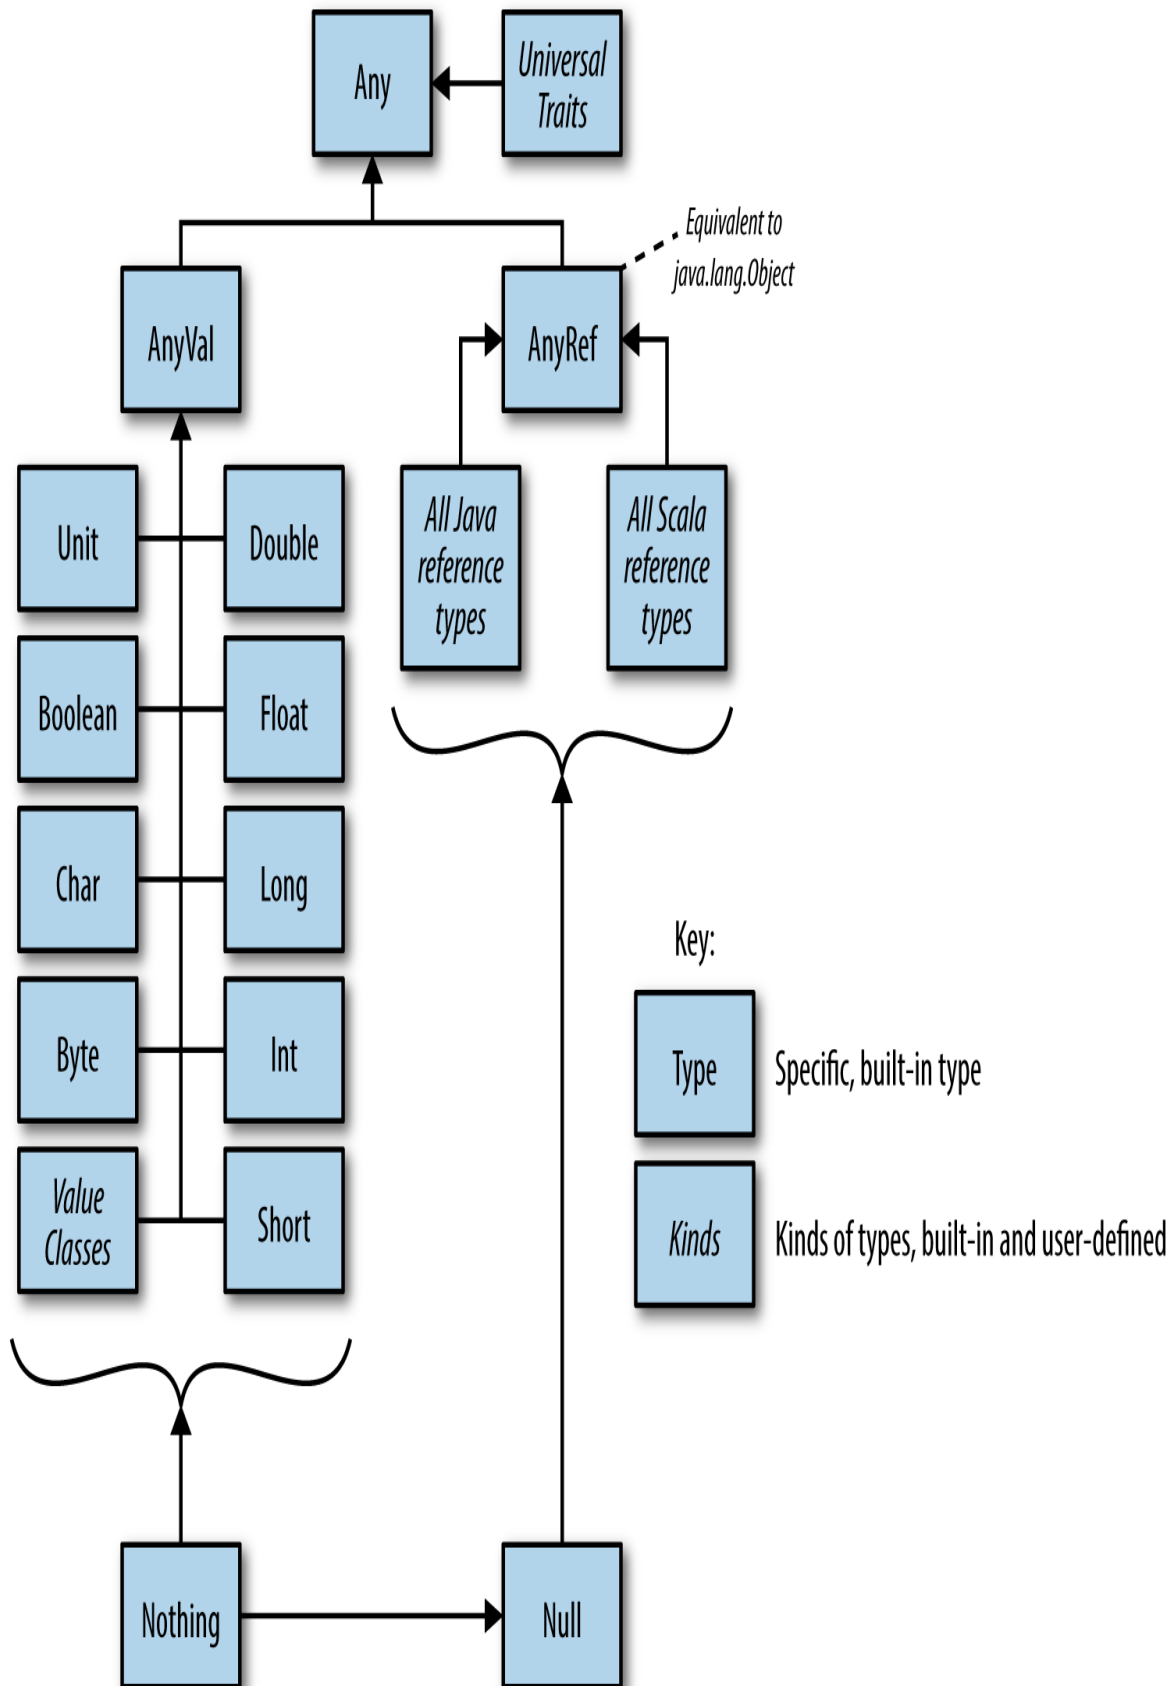


Figure 13-1. Scala's type hierarchy

At the root of the type hierarchy is `Any`. It has no parents and three children:

- `AnyVal`, the parent of value types and value classes
- `AnyRef`, the parent of all *reference* types
- *Universal Traits*, which we discussed in “Value Classes”

`AnyVal` has nine concrete subtypes, called the *value types*. They don't require heap allocation of instances. Instead, the values fit in CPU registers and they are pushed onto the call stack, when needed. Seven of them are numeric value types: `Byte`, `Char`, `Short`, `Int`, `Long`, `Float`, and `Double`. The remaining two are nonnumeric, `Unit` and `Boolean`.

*Value classes* also extend `AnyVal` (see “Value Classes”) and `Nothing` is a subtype of `AnyVal`, which we'll discuss below.

In contrast, all the other types are *reference types*, because instances of them are allocated in the heap and managed by reference. They are derived from `AnyRef`, which is the analog of `java.lang.Object`, when compiling for the JVM. Java's object model doesn't have a parent type of `Object` that encapsulates both primitives and reference types, because primitive types have special treatment.

Let's discuss a few of the widely used types.

## Much Ado About Nothing (and Null)

`Nothing` and `Null` are two unusual types at the bottom of the type system. Specifically, `Nothing` is a subtype of all other types, while `Null` is a subtype of all reference types.

`Null` is the familiar concept from most programming languages, although other languages may not define a `Null` type. It exists in Scala's type hierarchy because of the need to interoperate with Java and JavaScript libraries that support `null` values. However, we learned in “[Option, Some, and None: Avoiding nulls](#)” that avoiding `nulls` is recommended to prevent common software bugs.

`Null` has the following definition:

```
abstract final class Null extends AnyRef
```

How can it be both `final` and `abstract`? This declaration disallows subtyping and creating your own instances, but the runtime environment provides one instance, `null`, our old nemesis.

`Null` is defined as a subtype of `AnyRef`, but it is also treated by the compiler as a subtype of all `AnyRef` types. This is the type system's formal way of allowing you to assign `null` to references for any reference type. On the other hand, because `Null` is not a subtype of `AnyVal`, it is not possible to assign `null` in place of an `Int`, for example. Hence, Scala's `null` behaves exactly like Java's `null`.

In contrast, `Nothing` has no analog in Java, but it fills a hole that exists in Java's type system. `Nothing` has the following definition:

```
abstract final class Nothing extends Any
```

In contrast to `Null`, `Nothing` is a subclass of *all* other types, value types as well as reference types. In other words, `Nothing` subclasses *everything*, as weird as that sounds.

Unlike `Null`, `Nothing` has no instances. We say the type is *uninhabited*. It is still useful, because it provides two capabilities in the type system that contribute to robust, type-safe design.

The first capability is best illustrated with the `List[+A]` class. We now understand that `List` is *covariant* in `A`, so `List[String]` is a subtype of `List[AnyRef]`, because `String` is a subtype of `AnyRef`. Therefore,

a `List[String]` instance can be assigned to a variable of type `List[Any]`.

Scala declares a type for the special case of an empty list, `Nil`. In Java, `Nil` would have to be a parameterized class like `List`, but this is unfortunate, because by definition, `Nil` never holds any elements, so a `Nil[String]` and a `Nil[AnyRef]` would be different, but without distinction.

Scala solves this problem by having `Nothing`. `Nil` is actually declared like this:

```
package scala.collection.immutable
case object Nil extends List[Nothing] with Product with Serializable
```

We'll discuss `Product` in the next section. `Serializable` is the familiar Java “marker” interface for objects that can be serialized using Java's built-in mechanism for this purpose.

Note that `Nil` is an object and it extends `List[Nothing]`. There is only one instance of it needed, because it carries no “state” (elements). Because `List` is covariant in the type parameter, `Nil` is a subtype of `List[A]` for *all* types `A`. Therefore, we don't need separate `Nil[A]` instances. One will do.

`Nothing` and `Null` are called *bottom* types, because they reside at the bottom of the type hierarchy, so they are subtypes of all or a subset of the rest of the types.

The other use for `Nothing` is to represent expressions that terminate the program, such as by throwing an exception. An example is the special `???` method. It can be called in a temporary method definition so the method is concrete, allowing an enclosing, concrete type to compile:

```
object Foo:
  def bar(str: String): String = ???
```

However, if `Foo.bar` is called, an exception is thrown. Here is the definition of `???` inside `scala.Predef`:

```
def ??? : Nothing = throw new NotImplementedError
```

Because `???` “returns” `Nothing`, it can be called by any other method, no matter what type it returns.

Use `Nothing` as the return type of any method that always throws an exception, like `???` and `sys.error`, or terminates the application, like `sys.exit`.

This means that a method can declare that it returns a “normal” type, yet choose to call `sys.exit` if necessary, and still type check. Consider this example that processes command-line arguments, but exits if an unrecognized argument is provided:<sup>1</sup>

```
// src/main/scala/progscala3/objectsystem/CommandArgs.scala
package progscala3.objectsystem

object CommandArgs:

  val help = """
    |Usage: progscala3.objectsystem.CommandArgs arguments
    |Where the allowed arguments are:
    |  -h | --help           Show help
    |  -i | --in | --input path Path for input (required)
    |  -o | --on | --output path Path for output (required)
    |""".stripMargin

  def quit(status: Int = 0, message: String = ""): Nothing =
    ❶ if message.length > 0 then println("ERROR: "+message)
      println(help)
      sys.exit(status)

  case class Args(inputPath: Option[String], outputPath:
    Option[String])

  def parseArgList(params: Array[String]): Args =
    ❷ def pa(params2: Seq[String], args: Args): Args = params2 match
```

```

3      case Nil => args
4      case ("-h" | "--help") +: Nil => quit()
5      case ("-i" | "--in" | "--input") +: path +: tail =>
        pa(tail, args.copy(inputPath = Some(path)))
6      case ("-o" | "--out" | "--output") +: path +: tail =>
        pa(tail, args.copy(outputPath = Some(path)))
7      case _ => quit(1, s"Unrecognized argument ${params2.head}")

    val argz = pa(params.toList, Args(None, None))

8    if argz.inputPath == None || argz.outputPath == None then
        quit(1, "Must specify input and output paths.")
    argz

def main(params: Array[String]): Unit =
    val argz = parseArgList(params)
    println(argz)

```

- ❶ Print an optional message, followed by the help message, then exit with the specified error status. Following Unix conventions, 0 is used for normal exits and nonzero values are used for abnormal termination. Note that `quit` returns `Nothing`.
- ❷ Parse the argument list, returning an instance of `Args`, which holds the specified input and output paths determined from the argument list.
- ❸ A nested, recursively invoked function to process the argument list. We use the idiom of passing an `Args` instance to accumulate new settings, by making a copy of it.
- ❹ End of input, so return the accumulated `Args`.
- ❺ Process options for help, input, and output. Note the three ways the input and output flags can be specified. Note the recursive invocations of `pa`.
- ❻ Handle the error of an unrecognized argument.

- ⑦ Call `pa` to process the arguments with a seed value for `Args`.
- ⑧ Verify that the input and output arguments were provided.

I find this example of pattern matching particularly elegant and concise. It is fully type safe, even though an error or help request triggers termination.

Try it in `sbt` with various combinations of arguments: `runMain progscala3.objectsystem.CommandArgs ...`

## The `scala` Package

The top-level `scala` package defines many of the types we've discussed in this book, including most of the ones in [Figure 13-1](#). This package is automatically in scope when compiling, so it's unnecessary to import it.

[Table 13-1](#) describes the most interesting packages under `scala`:



*Table 13-1. Interesting Packages Under `scala`*

Name	Description
<code>annotation</code>	Where most annotations are defined.
<code>collection</code>	All the collection types. See <i>Collections</i> .
<code>concurrent</code>	Tools for concurrent programming. See <i>Chapter 18</i> .
<code>io</code>	A limited set of tools for input and output, such as <code>io.Source</code> .
<code>math</code>	Mathematical operations and definitions, including numeric comparisons.
<code>reflect</code>	Introspection on types. See <i>Metaprogramming</i>
<code>sys</code>	Access to underlying system resources, like properties and environment variables, along with operations like terminating the program.
<code>util</code>	Miscellaneous useful types like <code>Either</code> , <code>Left</code> , <code>Right</code> , <code>Try</code> , <code>Success</code> , <code>Failure</code> , <code>Using</code> .

The `scala` package defines aliases to popular types in other packages, so they are easily accessible. Examples include `scala.math` types like `Numeric`, `Ordering`, and `Ordered`, `java.lang` exceptions like `Throwable` and `NullPointerException`, and some of the collections types, like `Iterable`, `List`, `Seq`, and `Vector`.

Some annotations are also defined here, mostly those that signal information to the compiler, such as `@inline` that directs the compiler

to be aggressive about inlining the definition in byte code.

Finally, there is a class of types for case classes, tuples, and functions.

## Products, Case Classes, Tuples, and Functions

Case classes mix in the `scala.Product` trait, which provides a few generic methods for working with the fields of type and case class instances:

```
scala> case class Person(name: String, age: Int)
scala> val p: Product = Person("Dean", 29)

scala> p.productArity
val res0: Int = 2

scala> p.productIterator.foreach(println)
Dean
29

scala> (p.productElement(0), p.productElementName(0))
val res1: (Any, String) = (Dean,name)

scala> (p.productElement(1), p.productElementName(1))
val res8: (Any, String) = (29,age)

scala> p.productElement(2)
java.lang.IndexOutOfBoundsException: 2
...
```

```
scala> val tup = ("Wampler", 39, "hello")
val tup: (String, Int, String) = (Wampler,39,hello)

scala> (tup.productArity, tup.productElement(2),
tup.productElementName(2))
val res9: (Int, Any, String) = (3,hello,_3)
```

Note the element name used for tuple elements, corresponding to the methods used to extract those elements.

While having generic ways of accessing fields can be useful, its value is limited by the fact that `Any` is used for the fields' types, not their actual types.

There are also subtypes of `Product` for specific arities, up to 22. They are parent types of the corresponding `TupleN` types. For example, `Product3[+T1,+T2,+T3]` is defined for three-element products. These types add methods for selecting particular elements with the correct type information preserved. For example, `Product3[+T1,+T2,+T3]` adds these methods:

```
package scala
trait Product3[+T1, +T2, +T3] extends Product {
  abstract def _1: T1
  abstract def _2: T2
  abstract def _3: T3
  ...
}
```

For tuples and functions with 22 elements or less (a somewhat arbitrary value), there are corresponding `TupleN` and `FunctionN` types, for example `Tuple3[+T1,+T2,+T3]` and `Function3[-T1,-T2,-T3,+R]`. Note that `Function3` has a fourth type parameter for the return type. Recall from “[Functions Under the Hood](#)” that the types for arguments must be *contravariant*.

In Scala 2, tuples and functions were limited to 22 elements, but Scala 3 removes this limitation. For more than 22 elements, the compiler generates an instance of `scala.TupleXXL` and `scala.FunctionXXL`, respectively. In this case, an immutable array type, `scala.opaques.IArray`, is used to store the elements for tuples and the arguments for functions.

While it may seem that 22 elements should be more than enough, especially when considering human comprehension, think about the case of a SQL query. It's common to want to put each record returned into a tuple or case class, but the number of columns can easily exceed 22.

# The Tuple Trait

Tuples are also derived from a new Scala 3 trait called **Tuple**, which adds several useful operations that makes it easy to use tuples like collections or convert to collections. First, you can construct new tuples by prepending an element or concatenation:

```
// src/script/scala/progscala3/basicoop/Tuple.scala
scala> val t1 = (1, "two", 3.3)
      | val t2 = (4.4F, (5L, 6.6))
val t1: (Int, String, Double) = (1,two,3.3)
val t2: (Float, (Long, Double)) = (4.4,(5,6.6))

scala> val t01 = 0L *: t1           // Prepend an element
      | val t12 = t1 ++ t2         // Concatenate tuples
val t01: (Long, Int, String, Double) = (0,1,two,3.3)
val t12: Int *: String *: Double *: t2.type = (1,two,3.3,4.4,(5,6.6))
```

You can drop or take elements or split a tuple:

```
scala> val t12d3 = t12.drop(3)      // Drop leading 3 elements
      | val t12d4 = t12.drop(4)    // Drop leading 4 elements
      | val t12t3 = t12.take(3)
      | val t12s3 = t12.splitAt(3)  // Like take and drop
combined
val t12d3: (Float, (Long, Double)) = (4.4,(5,6.6))
val t12d4: (Long, Double) *: scala.Tuple$package.EmptyTuple.type =
((5,6.6),)
val t12t3: Int *: String *: Double *: EmptyTuple = (1,two,3.3)
val t12s3: (scala.Tuple.Take[Int *: String *: Double *: t2.type, 3],
  scala.Tuple.Drop[Int *: String *: Double *: t2.type, 3]
) = ((1,two,3.3),(4.4,(5,6.6)))
```

You can convert to a few collection types and use the Tuple companion object to convert them to a tuple:

```
scala> val a  = t1.toArray          // Convert to collections
      | val ia = t1.toIArray
      | val l  = t1.toList
val a: Array[Object] = Array(1, two, 3.3)
val ia: opaques.IArray[Object] = Array(1, two, 3.3)
```

```

val l: List[Tuple.Union[t1.type]] = List(1, two, 3.3)

scala> val ta = Tuple.fromArray(a)      // Convert to collections
      | val tia = Tuple.fromIArray(ia)
      | // val tl = Tuple.fromList(l)  // Doesn't exist
val ta: Tuple = (1,two,3.3)
val tia: Tuple = (1,two,3.3)

scala> case class Person(name: String, age: Int)
      | val tp = Tuple.fromProduct(Person("Dean", 29))
val tp: Tuple = (Dean,29)

```

The last example extracted the elements of a case class instance, an example of a Product, into a tuple. Finally, you can zip tuples:

```

scala> val z1 = t1.zip(t2)              // Note that t1's "3.3" is
      | dropped
      | val z2 = t1.zip((4.4,5,6.6))    // Two, three-element tuples
      | zipped
val z1: (Int, Float) *: (String, (Long, Double)) *: EmptyTuple =
  ((1,4.4),(two,(5,6.6)))
val z2: (Int, Double) *: (String, Int) *: (Double, Double) *:
  EmptyTuple =
  ((1,4.4),(two,5),(3.3,6.6))

```

Many other definitions can be found in the `scala` package, but we'll spend the rest of this chapter discussing the definitions in `Predef`.

## The Predef Object

Some of the widely used definitions are in `scala.Predef`. Like the `scala` package, it is automatically in scope when you compile code, so you don't need to import it.

Let's summarize the contents of `Predef`. Some definitions we've seen before. A few definitions we'll wait to discuss until *our chapter on Metaprogramming*.

## Implicit Conversions

First, Predef defines many implicit conversions to wrap AnyVal and Java types in order to add extension methods. The conversions have existed in Scala for a long time, so Scala 2 implicit conversions are used instead of Scala 3 extension methods. I provide examples here, but I won't list all the methods. See the Predef documentation for all the details.

First are conversions from `scala.Array` for specific AnyVal types and all AnyRef types. They can be converted to corresponding wrappers of type `scala.collection.mutable.ArraySeq`, providing all the methods from sequential collections for arrays.

```
implicit def wrapByteArray(
  xs: Array[Boolean]): scala.collection.mutable.ArraySeq.ofBoolean
...
implicit def wrapRefArray[T <: AnyRef](
  xs: Array[T]): scala.collection.mutable.ArraySeq.ofRef[T]
```

Having separate types for each of the AnyVal type exploits the fact that Java arrays of primitives are more efficient than arrays of boxed elements.

There are similar conversions to `scala.collection.ArrayOps`, which are similar to ArraySeqs, but the methods return Array instances instead of ArraySeq instances. The latter are better when calling a chain of ArraySeq transformations.

Similarly, for Strings, which are effectively *character arrays*, there are `WrappedString` and `StringOps` conversions:

```
implicit def wrapString(s: String): WrappedString
implicit def augmentString(x: String): StringOps
```

## NOTE

Having pairs of similar wrapper types, like `ArraySeq/ArrayOps` and `WrappedString/StringOps`, is confusing, but fortunately the implicit conversions are invoked automatically, selecting the correct wrapper type for the method you need.

Several conversions add methods to `AnyVal` types. For example:

```
implicit def booleanWrapper(b: Boolean): RichBoolean
implicit def byteWrapper(b: Byte): RichByte
...
```

The `Rich*` types add methods like comparison methods, such as `<=` and `compare`.

Why have two separate types for bytes, for example? Why not put all the methods in `Byte` itself? The reason is that the extra methods would force boxing of the value and allocation on the heap, due to implementation requirements for byte code. Recall that `AnyVal` instances are not heap-allocated, but represented as the corresponding Java primitives. So, having separate `Rich*` types avoids the heap allocation except for those times when the extra methods are needed.

There are methods for converting between Java's boxed types for primitives and Scala's `AnyVal` types, which make Java interoperability easier:<sup>2</sup>

```
implicit def boolean2Boolean(x: Boolean): java.lang.Boolean
...
implicit def Boolean2boolean(x: java.lang.Boolean): Boolean
...
```

Finally, recall `ArrowAssoc` from “**Extension Methods**”. It is defined in `Predef`. For pattern matching on tuples, there is a definition in `Predef` `val ->: Tuple2.type` that supports using the same syntax:

```
scala> val x -> y = (1, 2)
val x: Int = 1
val y: Int = 2

scala> (1, 2) match:
|   case x -> y => println(s"$x and $y")
|   case _ => println("Error!!")
1 and 2
```

## Type Definitions

Predef defines several types and type aliases for convenience.

To encourage the use of immutable collections, Predef defines aliases for the most popular, immutable collection types:

```
type Map[A, +B]      = collection.immutable.Map[A, B]
type Set[A]          = collection.immutable.Set[A]
type Function[-A, +B] = (A) => B
```

Several convenient aliases point to Java types:

```
type Class[T] = java.lang.Class[T]
type String = java.lang.String
```

## Condition Checking Methods

Sometimes you want to assert a condition is true, perhaps to “fail fast” and especially during testing. Predef defines a number of methods that assist in this goal. The following methods come in pairs. One method takes a Boolean value. If false, an exception is thrown. The second version of the method takes the boolean value and an error string to include in the exception’s message. All the methods actually behave the same, but the names convey different meanings, as shown in this example for a factorial method:

```
scala> val n = 5
scala> assert(n > 0, s"Must be a positive number. $n given.")
...
```



```
scala> def factorial(n: Long): Long = { // Should really return
    BigDecimal.
    |   require(n >= 0, s"factorial($n): Must pass a positive
number!")
    |   if n == 1 then n
    |   else n*factorial(n-1)
    | } ensuring(_ > 0, "BUG!!")
def factorial(n: Long): Long

scala> factorial(-1)
java.lang.IllegalArgumentException: requirement failed:
factorial(-1):
  Must pass a positive number!
  ...

scala> factorial(5)
val res1: Long = 120
```

Many languages provide some form of assert. The `require` and `assume` (not shown) methods behave identically, but `require` is meant to convey that an input failed to meet the *requirements*, while `assume` verifies that *assumptions* are true. There is also `assertFail()`, which behaves like `assert(false)`, and a corresponding `assertFail(message)`.

Notice how `ensuring` is used to perform a final verification on the result and then return it if the assertion passes. Another Predef implicit conversion class, `Ensuring[A]`, is invoked on the value returned from the block. It provides four variants of an `ensuring` method:

```
def ensuring(cond: (A) => Boolean, msg: => Any): A
def ensuring(cond: (A) => Boolean): A
def ensuring(cond: Boolean, msg: => Any): A
def ensuring(cond: Boolean): A
```

If the condition `cond` is true, the `A` value is returned. Otherwise an exception is thrown.

I'll discuss an approach to using these methods for writing robust code in *Chapter 23*.

In Scala 2, if you wanted to turn the assertions off in production, passing `-Xelide-below 2001` to the compiler would suppress compilation of them (except for the two `require` methods), because they were annotated with `@elidable(ASSERTION)`, where `@elidable.ASSERTION` is 2000. However, this feature is not implemented in Scala 3.

## Input and Output Methods

We've enjoyed the convenience of writing `println("foo")`. `Predef` gives us four variants for writing strings to `stdout`:

```
def print(x: Any): Unit    // Print x as a String to stdout; no line
                             feed
def printf(format: String, xs: Any*): Unit  // Printf-formatted
                             string
def println(x: Any): Unit  // Print x as a String to stdout, with a
                             line feed
def println(): Unit        // Print a blank line
```

All delegate to the corresponding `scala.Console` methods. For the `printf` format syntax, see `java.util.Formatter`.

## Miscellaneous Methods

Finally, in addition to the `???` placeholder method, there are a few more methods defined in `Predef` you'll find useful.

When using some of the *combinator* methods we discussed in *Chapter 7*, sometimes you'll apply a combinator, but no actual change is required for the inputs. Pass `identity` to the combinator for the function, which simply returns the input value.

When you use a context bound, like `foo[A : Comparable](...)`, but you need a reference to the given `Comparable`, use

`implicitly[Comparable]` to bind the object, Then you can invoke methods on it. For Scala 3, there is a `summon` method, which does the same thing.

In any context, if there is one unique value for a type that you want to get, use `valueOf[T]` to fetch it. For example:

```
scala> valueOf[EmptyTuple]
val res0: EmptyTuple.type = ()
```

Finally, the `locally` method works around some rare parsing ambiguities. See its documentation for details.

## Recap and What's Next

We introduced the Scala type hierarchy and explored several definitions in the top-level package `scala` and the object `scala.Predef`. I encourage you to browse the library documentation at [scala-lang.org/api](https://scala-lang.org/api) to learn more.

In the next chapter, we'll learn about Scala's collections library.

- 
- 1 Or use the `scopt` library.
  - 2 We'll see corresponding conversions between collections in `Collections`.

# Index

---

## Symbols

\$ (dollar sign) placeholder, [Reserved Words](#)

\$ command, [Using SBT](#)

& (ampersand), [Reserved Words](#)

' (quotation marks, single)

in symbol literals, [Symbol Literals](#)

( ) (parentheses)

omitting in method invocations, [Methods with Empty Parameter Lists-Methods with Empty Parameter Lists](#)

\* (asterisk)

as wildcard character, [Importing Types and Their Members](#)

+ (plus sign)

in type parameters, [Abstract Types Versus Parameterized Types](#)

- (minus sign)

in type parameters, [Abstract Types Versus Parameterized Types](#)

/ (slash), [A Taste of Scala](#)

: (colon)

as separator, [Reserved Words](#)

in method definitions, [A Taste of Scala](#)

; (semicolons)

as expression separators, [Semicolons](#)

< (left arrow)

in for comprehensions, [Expanded Scope and Value Definitions](#)

<> (angle brackets), [Abstract Types Versus Parameterized Types](#)

in method names, [A Taste of Scala](#)

= (equals sign)

assignment, [Reserved Words](#)

in for comprehensions, [Expanded Scope and Value Definitions](#)

in method definitions, [A Taste of Scala](#)

== (equals sign), [Conditional Operators](#)

[] (square brackets), [Abstract Types Versus Parameterized Types](#)

with parameterized types, [A Taste of Scala](#)

\_ (underscore)

as wildcard character, [Importing Types and Their Members](#)

## A

abstract keyword, [Reserved Words](#), [Traits as Mixins](#), [Stackable Traits](#)

abstract types, [Why Scala?](#), [Abstract Types Versus Parameterized Types](#)-[Abstract Types Versus Parameterized Types](#)

anonymous functions, [Anonymous Functions](#), [Lambdas](#), and [Closures-Methods as Functions](#)

application resource manager example, [Call by Name](#), [Call by Value](#)-[Call by Name](#), [Call by Value](#)

as keyword, [Reserved Words](#)

## B

back-quote literals, [Allowed Characters in Identifiers](#)

Boolean literals, [Boolean Literals](#)

bottom types, [Much Ado About Nothing \(and Null\)](#)

by-name parameters, [Call by Name](#), [Call by Value](#)-[Call by Name](#), [Call by Value](#)

by-value parameters, [Call by Name](#), [Call by Value](#)-[Call by Name](#), [Call by Value](#)

## C

call-by-name parameters, [Call by Name](#), [Call by Value](#)-[Call by Name](#), [Call by Value](#)

call-by-value parameters, [Call by Name](#), [Call by Value](#)-[Call by Name](#), [Call by Value](#)

case classes, [A Sample Application-A Sample Application](#), [Matching on Case Classes and Enums-unapplySeq Method](#), [Products](#), [Case Classes](#), [Tuples](#), and [Functions-Products](#), [Case Classes](#), [Tuples](#), and [Functions](#)

case clauses, [Values](#), [Variables](#), and [Types in Matches](#), [Guards in Case Clauses](#)

case keyword, [Reserved Words](#)

catch clauses, [Using try, catch, and finally Clauses](#)

catch keyword, [Reserved Words](#)

Cats, [Cats Validator](#)

character literals, [Character Literals](#)

characters allowed in identifiers, [Allowed Characters in Identifiers](#)

class, [A Taste of Scala](#)

    abstract, [A Sample Application](#)

class definitions, [A Taste of Scala](#)

class keyword, [A Taste of Scala](#), [Reserved Words](#)

classes

    abstract, [A Sample Application](#)

    defining traits as, [Should That Type Be a Class or Trait?-Should That Type Be a Class or Trait?](#)

    fields in, [Fields in Types-Unary Methods](#)

    in object-oriented programming, [Class and Object Basics: Review-Class and Object Basics: Review](#)

    overrides in (see overrides)

closures, [Anonymous Functions, Lambdas, and Closures-Anonymous Functions, Lambdas, and Closures](#)

code

    organizing in files and namespaces, [Organizing Code in Files and Namespaces-Organizing Code in Files and Namespaces](#)

combinators, [Variables That Aren't, Functional Data Structures, Combinators: Software's Best Component Abstractions-Combinators: Software's Best Component Abstractions](#)

command-line tools

running, [Running the Scala Command-Line Tools-Running the Scala Command-Line Tools](#)

companion object, [A Sample Application-A Sample Application, unapply Method](#)

comprehensions, [Scala for Comprehensions, Pattern Matching in Other Contexts](#)

(see also for comprehensions; loops)

conditional operators, [Conditional Operators](#)

console command, [Using SBT](#)

constructors, [Constructors in Scala-Constructors in Scala](#)

containers (see for comprehensions)

contravariance, [Parameterized Types: Variance Under Inheritance-Functions Under the Hood](#)

contravariant typing, [Abstract Types Versus Parameterized Types](#)

covariance, [Parameterized Types: Variance Under Inheritance](#)

covariant specialization, [Abstract Types Versus Parameterized Types](#)

covariant typing, [Abstract Types Versus Parameterized Types](#)

currying, [Currying and Uncurrying Functions-Partial Functions vs. Functions Returning Options](#)

## D

declarations, [Method Declarations](#)

(see also method declarations)



package-level, [Organizing Code in Files and Namespaces-Importing Types and Their Members](#)

variable, [Variable Declarations-Variable Declarations](#)

deep matching, [Matching on Case Classes and Enums-unapplySeq Method](#)

def, [A Taste of Scala](#)

def keyword, [A Taste of Scala](#), [Reserved Words](#)

definitions, method, [Method Declarations](#)  
(see also [method declarations](#))

dictionary (see [maps](#))

do keyword, [Reserved Words](#)

domain-specific languages (see [DSLs](#))

DSLs (domain-specific languages), [Why Scala?](#)

## E

eager evaluation, [Variables That Aren't](#)

either container, [Either: A Logical Extension to Option-Either: A Logical Extension to Option](#)

else statements, [Reserved Words](#)

enumerations, [Enumerations and Algebraic Data Types-Enumerations and Algebraic Data Types](#)

equality of objects, [Equality of Instances-Array Equality and the sameElements Method](#)

`==` and `!=` methods, [The `==` and `!=` Methods](#)

array equality and same objects method, [Array Equality and the sameElements Method](#)

eq and ne methods, [The eq and ne Methods](#)

equals method, [The equals Method](#)

equals sign

in method definitions, [A Taste of Scala](#)

(see also = (equals sign))

exception handling, [Using try, catch, and finally Clauses](#)

exhaustive matches, [Sealed Hierarchies and Exhaustive Matches-Sealed Hierarchies and Exhaustive Matches](#)

expression, [A Taste of Scala](#)

extends keyword, [Traits as Mixins](#)

extension keyword, [Reserved Words](#)

extension methods, [Value Classes-Value Classes](#)

extracting (see pattern matching)

extractors, [unapply Method-unapplySeq Method](#)

## F

false keyword, [Reserved Words](#)

family polymorphism, [Abstract Types Versus Parameterized Types](#)

fields, in classes, [Fields in Types-Unary Methods](#)

file organization, [Organizing Code in Files and Namespaces-Organizing Code in Files and Namespaces](#)

filtering, [Guards: Filtering Values, Filtering-Filtering](#)

in for expressions, [For Comprehensions: Under the Hood](#)

final declarations, [Reserved Words](#)

final keyword, [Option as a Container](#)

finally clauses, [Using try, catch, and finally Clauses](#)

finally keyword, [Reserved Words](#)

flat mapping, [Flat Mapping-Flat Mapping](#)

folding, [Folding and Reducing-Folding and Reducing](#), [Left Versus Right Folding-Left Versus Right Folding](#)

for comprehensions, [Reserved Words](#), [Scala for Comprehensions](#),  
[for Comprehensions in Depth-Recap and What's Next](#)

collection methods, [For Comprehensions: Under the Hood-For Comprehensions: Under the Hood](#)

container types, [Options and Other Container Types](#)

    Cats validator, [Cats Validator](#)

    either, [Either: A Logical Extension to Option-Either: A Logical Extension to Option](#)

    option, [Option as a Container-Option as a Container](#)

    try, [Try: When There Is No Do-Try: When There Is No Do](#)

elements of, [Recap: The Elements of for Comprehensions](#)

expanded scope, [Expanded Scope and Value Definitions](#)

filters in, [Guards: Filtering Values](#)

for loops, [for Loops](#)

generator expressions, [Generator Expressions](#), [For Comprehensions: Under the Hood](#), [Translation Rules of for](#)

Comprehensions-Translation Rules of for Comprehensions

pattern matching and, Pattern Matching in Other Contexts

throwing exceptionsin, Throwing exceptions versus returning

Either values-Throwing exceptions versus returning Either values

translation rules, Translation Rules of for Comprehensions-  
Translation Rules of for Comprehensions

yielding, Yielding New Values

for loops, for Loops

formal parameters, Anonymous Functions, Lambdas, and Closures

free variables, Anonymous Functions, Lambdas, and Closures

function application, Functions Under the Hood

function literals, Function Literals, Functional Programming in Scala

defining, A Taste of Scala

function, defined, Anonymous Functions, Lambdas, and Closures

functional programming (FP), Functional Programming in Scala-  
Recap and What's Next

anonymous functions, lambdas, and closures, Anonymous  
Functions, Lambdas, and Closures-Methods as Functions

combinators, Combinators: Software's Best Component  
Abstractions-Combinators: Software's Best Component  
Abstractions

currying, Currying and Uncurrying Functions-Partial Functions  
vs. Functions Returning Options

filtering, Filtering-Filtering

flat mapping, Flat Mapping-Flat Mapping

folding and reducing, Folding and Reducing-Folding and Reducing

functional data structures, Functional Data Structures

maps, Maps

sequences, Sequences-Sequences

sets, Sets

functions in mathematics, Functions in Mathematics-Functions in Mathematics

immutable variables in, Variables That Aren't-Variables That Aren't

left versus right foldings, Left Versus Right Folding-Left Versus Right Folding

mapping, Mapping-Mapping

methods as functions, Methods as Functions

mixed paradigm in Scala, Why Scala?

overview, What Is Functional Programming?-Variables That Aren't

partially applied versus partial functions, Partially Applied Functions Versus Partial Functions-Partially Applied Functions Versus Partial Functions

pure functions, Purity Inside Versus Outside

recursion in, Recursion-Recursion

structure sharing, What About Making Copies?-What About Making Copies?

tail calls, [Tail Calls and Tail-Call Optimization](#)

transformations on functions, [Currying and Uncurrying](#)  
[Functions-Partial Functions vs. Functions Returning Options](#)

traversal, [Traversing, Mapping, Filtering, Folding, and Reducing-Traversal](#)

uses, [Functional Programming in Scala](#)

functions

methods as, [Methods as Functions](#)

partial versus partially applied, [Partially Applied Functions Versus Partial Functions-Partially Applied Functions Versus Partial Functions](#)

functions versus methods, [A Taste of Scala](#)

functions, higher-order, [Functions in Mathematics](#)

futures, [A Taste of Futures-A Taste of Futures](#)

## G

generator expressions, [Generator Expressions, For Comprehensions: Under the Hood, Translation Rules of for Comprehensions-Translation Rules of for Comprehensions](#)

generics, [Why Scala?, Abstract Types Versus Parameterized Types](#)

given keyword, [Reserved Words](#)

guards, [Guards: Filtering Values, lazy val](#)

in case clauses, [Guards in Case Clauses](#)

## H

Hash (see maps)

higher-order functions, [A Sample Application](#), [Functions in Mathematics](#)

## I

identifiers

- characters allowed in, [A Taste of Scala](#)

- rules summary for, [Allowed Characters in Identifiers-Allowed Characters in Identifiers](#)

if expressions, [Scala Conditional Expressions](#)

if statements, [Reserved Words](#), [Scala Conditional Expressions](#)

immutable variables, [Variable Declarations](#), [Variables That Aren't-Variables That Aren't](#)

implicit keyword, [Methods with Multiple Parameter Lists-A Taste of Futures](#), [Reserved Words](#)

implicits

- implicit conversions, [Reserved Words](#), [Implicit Conversions-Implicit Conversions](#)

import keyword, [Reserved Words](#)

import statements, [Importing Types and Their Members-Package Imports and Package Objects](#)

- package objects, [Package Imports and Package Objects-Package Imports and Package Objects](#)

infix notation, [Infix Operator Notation](#)

inheritance, [Super Types-Super Types](#), [Good Object-Oriented Design: A Digression](#)

instances, [Class and Object Basics: Review-Class and Object Basics: Review](#)

interfaces (see traits)

interpolated strings, [A Sample Application](#), [Interpolated Strings-Interpolated Strings](#)

invariance, [Parameterized Types: Variance Under Inheritance](#)

invariant typing, [Abstract Types Versus Parameterized Types](#)

## J

Java, [Why Scala?](#)

variance in, versus Scala, [Variance in Scala Versus Java-Variance in Scala Versus Java](#)

JVM (Java Virtual Machine), [Why Scala?](#)

## L

lambdas, [Anonymous Functions](#), [Lambdas](#), and [Closures](#)

lazy evaluation, [Reserved Words](#), [Variables That Aren't](#)

lazy values, [lazy val](#), [Initializing Abstract Fields](#)

LazyList, [Left Versus Right Folding](#)

least-upper bound values, [Scala Conditional Expressions](#)

lifting, [Methods as Functions](#), [Partial Functions vs. Functions Returning Options](#)



linearization, [Stackable Traits](#), [Linearization of an Object's Hierarchy](#)-[Linearization of an Object's Hierarchy](#)

lists

linked, [Sequences](#)

literal values, [Literal Values-Tuple Literals](#)

Boolean literals, [Boolean Literals](#)

character literals, [Character Literals](#)

function literals, [Function Literals](#)

string literals, [String Literals-String Literals](#)

symbol literals, [Symbol Literals](#)

literals

back-quote, [Allowed Characters in Identifiers](#)

partial functions, [Partial Functions-Partial Functions](#)

ranges, [Ranges-Ranges](#)

local type inference, [A Taste of Scala](#)

loops

conditional operators, [Conditional Operators](#)

do-while, [Scala do-while Loops](#)

for loops or comprehensions, [Scala for Comprehensions-Using try, catch, and finally Clauses](#)

while, [Scala while Loops](#)

## M

mapping, [Mapping-Mapping](#)

maps, [Maps](#)

match keyword, [Reserved Words](#)

MatchError, [A Sample Application](#), [Partial Functions](#)

matches (see pattern matching)

members, [Class and Object Basics: Review](#)

method declarations, [Method Declarations](#)

- default and named parameters, [Method Default and Named Parameters-Method Default and Named Parameters](#)

- futures, [A Taste of Futures-A Taste of Futures](#)

- multiple parameter lists, [Methods with Multiple Parameter Lists-Methods with Multiple Parameter Lists](#)

method definitions, [A Taste of Scala](#), [Nesting Method Definitions and Recursion-Nesting Method Definitions and Recursion](#)

methods

- with empty parameter lists, [Methods with Empty Parameter Lists-Methods with Empty Parameter Lists](#)

- as functions, [Methods as Functions](#)

- versus functions, [A Taste of Scala](#)

- in object-oriented programming (OOP), [Class and Object Basics: Review](#)

mixin composition, [Why Scala?](#)

mixins

- defined, [Traits](#)

traits as, [Traits as Mixins-Traits as Mixins](#)

mutable types, [Variance of Mutable Types-Variance of Mutable Types](#)

mutable variables, [Variable Declarations-Variable Declarations](#)

## N

named parameters, [Method Default and Named Parameters-Method Default and Named Parameters](#)

namespaces, [Organizing Code in Files and Namespaces](#)

nested classes, [Why Scala?](#)

nested methods, [Nesting Method Definitions and Recursion-Nesting Method Definitions and Recursion](#)

new keyword, [Reserved Words](#)

no parameter methods, [Methods with Empty Parameter Lists-Methods with Empty Parameter Lists](#)

none class, [Option, Some, and None: Avoiding nulls-Option, Some, and None: Avoiding nulls](#)

nothing, [Much Ado About Nothing \(and Null\)-Much Ado About Nothing \(and Null\)](#)

null keyword, [Reserved Words](#)

nulls, [Option, Some, and None: Avoiding nulls-Option, Some, and None: Avoiding nulls, Much Ado About Nothing \(and Null\)-Much Ado About Nothing \(and Null\)](#)

## O

object, [A Taste of Scala](#)

object definitions, [A Taste of Scala](#)

object keyword, [A Taste of Scala](#), [Reserved Words](#)

object system (Scala)

equality tests, [Equality of Instances-Array Equality and the sameElements Method](#)

hierarchy linearization, [Linearization of an Object's Hierarchy-Linearization of an Object's Hierarchy](#)

nothing and null, [Much Ado About Nothing \(and Null\)-Much Ado About Nothing \(and Null\)](#)

overrides, [Overriding Members of Classes and Traits](#)

overriding

abstract and concrete methods, [Defining Abstract Methods and Overriding Concrete Methods-Defining Abstract Methods and Overriding Concrete Methods](#)

abstract type aliases, [Abstract and Concrete Type Aliases](#)

concrete members, [Consider Making Concrete Members Final-Consider Making Concrete Members Final](#)

final declarations, [When Should Types Be Final?-When Should Types Be Final?](#)

predef object, [The Predef Object-Miscellaneous Methods](#)

products, case classes, and tuples, [Products, Case Classes, Tuples, and Functions-Products, Case Classes, Tuples, and Functions](#)

type hierarchy, [The Scala Type Hierarchy](#)

variance under inheritance, [Parameterized Types: Variance Under Inheritance-Variance in Scala Versus Java](#)

object-oriented design, [Good Object-Oriented Design: A Digression-Good Object-Oriented Design: A Digression](#), [Traits as Mixins](#)

object-oriented programming (OOP), [Functional Programming in Scala](#), [Combinators: Software's Best Component Abstractions](#), [Object-Oriented Programming in Scala-Recap and What's Next](#)

class and object basics, [Class and Object Basics: Review-Class and Object Basics: Review](#)

constructors, [Constructors in Scala-Constructors in Scala](#)

fields in classes, [Fields in Types-Unary Methods](#)

unary methods, [Unary Methods](#)

uniform access principle, [The Uniform Access Principle-The Uniform Access Principle](#)

good design in, [Good Object-Oriented Design: A Digression-Good Object-Oriented Design: A Digression](#), [Traits as Mixins](#)

methods versus functions in, [A Taste of Scala](#)

mixed paradigm in Scala, [Why Scala?](#)

parent class, [Calling Supertype Constructors-Calling Supertype Constructors](#)

reference versus value types, [Reference Versus Value Types-Reference Versus Value Types](#)

supertype, [Super Types-Super Types](#)

value classes, [Value Classes-Value Classes](#)

objects, [Class and Object Basics: Review-Class and Object Basics: Review](#)

Observer Design Pattern, [Traits as Mixins](#)

Odersky, Martin, [Why Scala?](#)

opaque keyword, [Reserved Words](#)

open keyword, [Reserved Words](#)

operator names, [Operator Overloading?-Syntactic Sugar](#)

operator notation, [Infix Operator Notation](#)

operator precedence rules, [Precedence Rules-Left vs. Right Associative Methods](#)

option class, [Option, Some, and None: Avoiding nulls-Option, Some, and None: Avoiding nulls](#)

option container, [Option as a Container-Option as a Container](#)

override keyword, [Reserved Words](#)

overrides, [Overriding Members of Classes and Traits](#)

- abstract and concrete methods, [Defining Abstract Methods and Overriding Concrete Methods-Defining Abstract Methods and Overriding Concrete Methods](#)

- abstract types aliases, [Abstract and Concrete Type Aliases](#)

- concrete members, [Consider Making Concrete Members Final](#)

## P

package declarations, [Organizing Code in Files and Namespaces-Importing Types and Their Members](#)

package keyword, [Reserved Words](#)

package objects, [Package Imports and Package Objects](#)-[Package Imports and Package Objects](#)

parameter lists, [Methods with Multiple Parameter Lists](#)-[Methods with Multiple Parameter Lists](#)

parameterized types, [A Taste of Scala](#), [Abstract Types Versus Parameterized Types](#)-[Abstract Types Versus Parameterized Types](#), [Parameterized Types: Variance Under Inheritance](#)

parametric polymorphism, [Abstract Types Versus Parameterized Types](#)

parent class, [Calling Supertype Constructors](#)-[Calling Supertype Constructors](#)

partial functions, [Partial Functions](#)-[Partial Functions](#), [Partially Applied Functions Versus Partial Functions](#)

`PartialFunction`, [Partial Functions](#)

partially applied function, defined, [Partially Applied Functions Versus Partial Functions](#)

pattern matching, [A Sample Application](#), [Partial Functions](#), [Pattern Matching-Recap and What's Next](#)

- on case classes, [Matching on Case Classes and Enums](#)-[unapplySeq Method](#)

- in for comprehensions, [Expanded Scope and Value Definitions](#)

- in for comprehensions, [Pattern Matching in Other Contexts](#)

- identifiers in, [Allowed Characters in Identifiers](#)

- on regular expressions, [Matching on Regular Expressions](#)-[Matching on Regular Expressions](#), [Pattern Matching in Other Contexts](#)

and sealed hierarchies, [Sealed Hierarchies and Exhaustive Matches-Sealed Hierarchies and Exhaustive Matches](#)

on sequences, [Matching on Sequences-Matching on Sequences](#)

for sequences, [Pattern Matching in Other Contexts](#)

on tuples, [Matching on Tuples, Matching on Case Classes and Enums](#)

type matching, [More on Type Matching](#)

values, variables, and types, [Values, Variables, and Types in Matches](#)

versus subtype polymorphism, [A Sample Application](#)

persistent data structures, [What About Making Copies?](#)

pre-initialized fields, [Initializing Abstract Fields](#)

precedence, operator, [Precedence Rules-Left vs. Right Associative Methods](#)

Predef object, [The Predef Object-Miscellaneous Methods](#)

condition checking methods, [Condition Checking Methods](#)

implicit conversions, [Implicit Conversions-Implicit Conversions](#)

input and output methods, [Input and Output Methods](#)

miscellaneous methods, [Miscellaneous Methods](#)

type definitions, [Type Definitions](#)

primitive data types, [Variable Declarations](#)

primitives, [Reference Versus Value Types-Reference Versus Value Types](#)



private keyword, [Reserved Words](#)

products, [Products](#), [Case Classes](#), [Tuples](#), and [Functions-Products](#),  
[Case Classes](#), [Tuples](#), and [Functions](#)

protected keyword, [Reserved Words](#)

pure functions, [A Sample Application](#), [Functions in Mathematics](#),  
[Variables That Aren't-Variables That Aren't](#), [Purity Inside Versus Outside](#)

## R

ranges, [Ranges-Ranges](#)

recursion, [Nesting Method Definitions and Recursion-Nesting Method Definitions and Recursion](#)

in functional programming, [Recursion-Recursion](#)

in sequence matching, [Matching on Sequences-Matching on Sequences](#)

tail recursion, [Left Versus Right Folding-Left Versus Right Folding](#)

tail-call self-recursion, [Tail Calls and Tail-Call Optimization](#)

reduce function, [Functional Programming in Scala](#)

reducing, [Folding and Reducing-Folding and Reducing](#)

reference equality, [The eq and ne Methods](#)

reference types, [Reference Versus Value Types-Reference Versus Value Types](#)

referential transparency, [Functions in Mathematics-Functions in Mathematics](#)

regexes, [Matching on Regular Expressions](#)

regular expressions, [Matching on Regular Expressions-Matching on Regular Expressions](#)

REPL (read, evaluate, print, loop), [Running the Scala Command-Line Tools](#)

paste mode, [Semicolons](#)

requires keyword, [Reserved Words](#)

reserved keywords, [Reserved Words-Reserved Words](#)

return keyword, [Reserved Words](#)

return type for methods, [Nesting Method Definitions and Recursion](#)

required explicit declarations of, [Inferring Type Information](#)

## S

SBT (standard build tool), [Using SBT-Using SBT](#)

Scala

installing, [Installing Scala](#)

introduction to, [Zero to Sixty: Introducing Scala-The Seductions of Scala](#)

reserved keywords list, [Reserved Words-Reserved Words](#)

type hierarchy, [The Scala Type Hierarchy-The Scala Type Hierarchy](#)

scala command, [Running the Scala Command-Line Tools](#)

Scaladocs, [More Tips](#), [A Taste of Scala](#), [Abstract Types Versus Parameterized Types](#)

Schönfinkel, Moses, [Currying and Uncurrying Functions](#)

sealed class hierarchies, [Sealed Class Hierarchies and Enumerations](#), [Sealed Hierarchies and Exhaustive Matches](#)-[Sealed Hierarchies and Exhaustive Matches](#)

sealed keyword, [Reserved Words](#), [Sealed Class Hierarchies and Enumerations](#), [Option as a Container](#)

self expression, [Option as a Container](#)

separation of concerns, [Traits as Mixins](#)

sequences, [Matching on Sequences](#)-[Matching on Sequences](#), [Sequences](#)-[Sequences](#)

    pattern matching for, [Pattern Matching in Other Contexts](#)

sets, [Sets](#)

Singleton, [A Taste of Scala](#)

Singleton pattern, [A Taste of Scala](#)

some class, [Option](#), [Some](#), and [None](#): [Avoiding nulls](#)-[Option](#), [Some](#), and [None](#): [Avoiding nulls](#)

stackable traits, [Stackable Traits](#)-[Stackable Traits](#)

static typing, [Why Scala?](#)

statics, [A Taste of Scala](#)

strict evaluation, [Variables That Aren't](#)

string interpolation, [Interpolated Strings](#)-[Interpolated Strings](#)

string literals, [String Literals](#)-[String Literals](#)

string parameter, [A Taste of Scala](#)

structure sharing, [What About Making Copies?-What About Making Copies?](#)

subtype polymorphism versus pattern matching, [A Sample Application](#)

subtypes, [Value Classes](#)

super keyword, [Reserved Words](#)

supertype, [Super Types-Super Types](#)

symbol literals, [Symbol Literals](#)

syntactic sugar, [Syntactic Sugar](#), [Functions Under the Hood](#)

## T

tail calls, [Tail Calls and Tail-Call Optimization](#), [Trampoline for Tail Calls](#)

tail recursion, [Left Versus Right Folding-Left Versus Right Folding](#)

Template Method Pattern, [Traits as Mixins](#), [Consider Making Concrete Members Final](#)

then keyword, [Reserved Words](#)

this keyword, [Reserved Words](#)

throw keyword, [Reserved Words](#)

throwing exceptions, [Throwing exceptions versus returning Either values-Throwing exceptions versus returning Either values](#), [Much Ado About Nothing \(and Null\)-Much Ado About Nothing \(and Null\)](#)

trait keyword, [Reserved Words](#)

traits, [Why Scala?](#), [Traits: Interfaces and “Mixins” in Scala-Traits: Interfaces and “Mixins” in Scala](#), [Traits-Recap and What’s Next](#)

defining as classes, [Should That Type Be a Class or Trait?](#)  
[Should That Type Be a Class or Trait?](#)

as mixins, [Traits as Mixins](#)-[Traits as Mixins](#)

overrides in (see [overrides](#))

stackable, [Stackable Traits](#)-[Stackable Traits](#)

trampolines, [Trampoline for Tail Calls](#)

traversal, [Traversing, Mapping, Filtering, Folding, and Reducing](#)-  
[Traversal](#)

true keyword, [Reserved Words](#)

try clauses, [Using try, catch, and finally](#) Clauses

try container, [Try: When There Is No Do-Try: When There Is No Do](#)

try keyword, [Reserved Words](#)

tuples, [Matching on Tuples](#)

    matching on, [Matching on Case Classes and Enums](#)

type annotations, [Running the Scala Command-Line Tools](#)

type annotations, required explicit, [Inferring Type Information](#)

type hierarchy, [The Scala Type Hierarchy](#)-[The Scala Type Hierarchy](#)

type inference, [Why Scala?](#), [Inferring Type Information](#)-[Variadic  
Argument Lists](#)

type keyword, [Reserved Words](#)

type matching, [Values, Variables, and Types in Matches](#), [More on  
Type Matching](#)

type system (Scala), [Why Scala?](#)

types, [Class and Object Basics: Review](#)

abstract versus parameterized, [Abstract Types Versus Parameterized Types-Abstract Types Versus Parameterized Types](#)

importing, [Importing Types and Their Members-Importing Types and Their Members](#)

## U

unapply method, [unapply Method](#)

unapplySeq method, [unapplySeq Method-unapplySeq Method](#)

unary methods, [Unary Methods](#)

Uniform Access Principle, [The Uniform Access Principle-The Uniform Access Principle](#)

universal traits, [Value Classes-Value Classes](#)

using keyword, [Reserved Words](#)

## V

val keyword, [Variable Declarations-Variable Declarations, Reserved Words](#)

value classes, [Value Classes-Value Classes](#)

value matches, [Values, Variables, and Types in Matches](#)

value types, [Reference Versus Value Types](#)

var keyword, [Variable Declarations-Variable Declarations, Reserved Words](#)

variable declarations, [Variable Declarations-Variable Declarations](#)

variable matches, [Values, Variables, and Types in Matches](#)

variables, [Variables That Aren't](#)

(see also immutable variables)

expanded scope in for expressions, [Expanded Scope and Value Definitions](#)

variance annotations, [Functions Under the Hood-Functions Under the Hood](#)

variance under inheritance, [Parameterized Types: Variance Under Inheritance-Variance in Scala Versus Java](#)

mutable types, [Variance of Mutable Types-Variance of Mutable Types](#)

Scala versus Java, [Variance in Scala Versus Java-Variance in Scala Versus Java](#)

## W

while loops, [Reserved Words](#)

with keyword, [Reserved Words](#)

## Y

yield keyword, [Yielding New Values](#)

## 1. Foreword

## 2. Preface

- a. Welcome to Programming Scala, Third Edition
- b. Welcome to Programming Scala, Second Edition
- c. Welcome to Programming Scala, First Edition
- d. Conventions Used in This Book
- e. Using Code Examples
  - i. Getting the Code Examples
- f. O'Reilly Safari
- g. How to Contact Us
- h. Acknowledgments for the Third Edition
  - i. Acknowledgments for the Second Edition
  - j. Acknowledgments for the First Edition

## 3. 1. Zero to Sixty: Introducing Scala

- a. Why Scala?
  - i. The Seductions of Scala
- b. Why Scala 3?
  - i. Migrating to Scala 3
- c. Installing Scala
  - i. Coursier
  - ii. Java JDK
  - iii. SBT



#### iv. Scala

##### d. Building the Code Examples

##### e. More Tips

###### i. Using SBT

###### ii. Running the Scala Command-Line Tools

##### f. A Taste of Scala

##### g. A Sample Application

##### h. Recap and What's Next

#### 4. 2. Type Less, Do More

##### a. New Scala 3 Syntax

##### b. Semicolons

##### c. Variable Declarations

##### d. Ranges

##### e. Partial Functions

##### f. Infix Operator Notation

##### g. Method Declarations

###### i. Method Default and Named Parameters

###### ii. Methods with Multiple Parameter Lists

###### iii. Nesting Method Definitions and Recursion

##### h. Inferring Type Information

###### i. Variadic Argument Lists

###### j. Reserved Words

## k. Literal Values

- i. Numeric Literals
- ii. Boolean Literals
- iii. Character Literals
- iv. String Literals
- v. Symbol Literals
- vi. Function Literals
- vii. Tuple Literals

## l. Option, Some, and None: Avoiding nulls

- i. When You Can't Avoid Nulls

## m. Sealed Class Hierarchies and Enumerations

## n. Organizing Code in Files and Namespaces

## o. Importing Types and Their Members

- i. Package Imports and Package Objects

## p. Abstract Types Versus Parameterized Types

## q. Recap and What's Next

# 5. 3. Rounding Out the Basics

## a. Operator Overloading?

## b. Allowed Characters in Identifiers

- i. Syntactic Sugar

## c. Methods with Empty Parameter Lists

## d. Precedence Rules

- i. Left vs. Right Associative Methods
  - e. Enumerations and Algebraic Data Types
  - f. Interpolated Strings
  - g. Scala Conditional Expressions
  - h. Conditional Operators
  - i. Scala for Comprehensions
    - i. for Loops
    - ii. Generator Expressions
    - iii. Guards: Filtering Values
    - iv. Yielding New Values
    - v. Expanded Scope and Value Definitions
  - j. Scala while Loops
    - i. Scala do-while Loops
  - k. Using try, catch, and finally Clauses
  - l. Call by Name, Call by Value
  - m. lazy val
  - n. Traits: Interfaces and “Mixins” in Scala
  - o. Recap and What’s Next
- 6. 4. Pattern Matching
  - a. Values, Variables, and Types in Matches
  - b. Matching on Sequences
  - c. Matching on Tuples

- i. Parameter Untupling
  - d. Guards in Case Clauses
  - e. Matching on Case Classes and Enums
  - f. Matching on Regular Expressions
  - g. More on Type Matching
  - h. Sealed Hierarchies and Exhaustive Matches
  - i. Chaining Match Expressions
  - j. Pattern Matching in Other Contexts
  - i. Problems in Pattern Bindings
  - k. Extractors
  - i. unapply Method
    - ii. Alternatives to Option Return Values
    - iii. unapplySeq Method
    - iv. Implementing unapplySeq
  - l. Concluding Remarks on Pattern Matching
  - m. Recap and What's Next
- 7. 5. Abstracting over Context: Type Classes and Extension Methods
- a. Four Changes
  - b. Extension Methods
  - c. Type Classes
  - i. Scala 3 Type Classes
    - ii. Scala 2 Type Classes

- d. Implicit Conversions
    - i. Rules for Implicit Conversion Resolution
  - e. Type Class Derivation
    - i. Givens and Imports
  - f. Resolution Rules for Givens and Extension Methods
    - i. Build Your Own String Interpolator
    - ii. The Expression Problem
  - g. Wise Use of Type Extensions
  - h. Recap and What's Next
8. 6. Abstracting over Context: Using Clauses
- a. Using Clauses
  - b. Context Bounds
    - i. By-Name Context Parameters
  - c. Other Context Parameters
  - d. Passing Context Functions
  - e. Constraining Allowed Instances
  - f. Implicit Evidence
  - g. Working Around Type Erasure with Context Bounds
  - h. Rules for Using Clauses
    - i. Improving Error Messages
    - j. Recap and What's Next
9. 7. Functional Programming in Scala

- a. What Is Functional Programming?
  - i. Functions in Mathematics
  - ii. Variables That Aren't
- b. Functional Programming in Scala
  - i. Anonymous Functions, Lambdas, and Closures
  - ii. Purity Inside Versus Outside
- c. Recursion
- d. Tail Calls and Tail-Call Optimization
  - i. Trampoline for Tail Calls
- e. Partially Applied Functions Versus Partial Functions
- f. Currying and Uncurrying Functions
- g. Tupled and Untupled Functions
- h. Partial Functions vs. Functions Returning Options
- i. Functional Data Structures
  - i. Sequences
  - ii. Maps
  - iii. Sets
- j. Traversing, Mapping, Filtering, Folding, and Reducing
  - i. Traversal
  - ii. Mapping
  - iii. Flat Mapping

- iv. Filtering
    - v. Folding and Reducing
    - vi. Left Versus Right Folding
  - k. Combinators: Software's Best Component Abstractions
  - l. What About Making Copies?
  - m. Recap and What's Next
- 10. 8. for Comprehensions in Depth
  - a. Recap: The Elements of for Comprehensions
  - b. For Comprehensions: Under the Hood
  - c. Translation Rules of for Comprehensions
  - d. Options and Other Container Types
    - i. Option as a Container
    - ii. Either: A Logical Extension to Option
    - iii. Try: When There Is No Do
    - iv. Cats Validator
  - e. Recap and What's Next
- 11. 9. Object-Oriented Programming in Scala
  - a. Class and Object Basics: Review
  - b. Open Versus Closed Types
    - i. Classes Open for Extension
    - ii. Overriding Methods? The Template Method Pattern

- c. Reference Versus Value Types
- d. Opaque Types and Value Classes
  - i. Opaque Type Aliases
  - ii. Value Classes
- e. Super Types
- f. Constructors in Scala
  - i. Calling Supertype Constructors
- g. Export Clauses
- h. Good Object-Oriented Design: A Digression
- i. Fields in Types
  - i. The Uniform Access Principle
  - ii. Unary Methods
- j. Recap and What's Next

## 12. 10. Traits

- a. Traits as Mixins
- b. Stackable Traits
- c. Union and Intersection Types
- d. Super Traits
- e. Using Commas Instead of With
- f. Trait Parameters
- g. Should That Type Be a Class or Trait?
- h. Recap and What's Next



## 13. 11. Variance Behavior and Equality

### a. Parameterized Types: Variance Under Inheritance

- i. Functions Under the Hood
- ii. Variance of Mutable Types
- iii. Variance in Scala Versus Java

### b. Equality of Instances

- i. The equals Method
- ii. The == and != Methods
- iii. The eq and ne Methods
- iv. Array Equality and the sameElements Method

### c. Equality and Inheritance

### d. Multiversal Equality

### e. Recap and What's Next

## 14. 12. Instance Initialization and Method Resolution

### a. Linearization of an Object's Hierarchy

### b. Overriding Members of Classes and Traits

- i. Consider Making Concrete Members Final
- ii. When Should Types Be Final?
- iii. Defining Abstract Methods and Overriding Concrete Methods
- iv. Initializing Abstract Fields
- v. Abstract and Concrete Type Aliases

- c. Recap and What's Next

- 15. 13. The Scala Type Hierarchy

- a. Much Ado About Nothing (and Null)

- b. The scala Package

- c. Products, Case Classes, Tuples, and Functions

- i. The Tuple Trait

- d. The Predef Object

- i. Implicit Conversions

- ii. Type Definitions

- iii. Condition Checking Methods

- iv. Input and Output Methods

- v. Miscellaneous Methods

- e. Recap and What's Next

- 16. Index